



ExtremeWare Software User Guide

Software Version 6.1

Extreme Networks, Inc.

3585 Monroe Street

Santa Clara, California 95051

(888) 257-3000

<http://www.extremenetworks.com>

©2000 Extreme Networks, Inc. All rights reserved. Extreme Networks and BlackDiamond are registered trademarks of Extreme Networks, Inc. in the United States and certain other jurisdictions. ExtremeWare, ExtremeWare Vista, ExtremeWorks, ExtremeAssist, ExtremeAssist1, ExtremeAssist2, PartnerAssist, Extreme Standby Router Protocol, ESRP, SmartTraps, Alpine, Summit, Summit1, Summit4, Summit4/FX, Summit71, Summit24, Summit48, Summit Virtual Chassis, SummitLink, SummitGbX, SummitRPS and the Extreme Networks logo are trademarks of Extreme Networks, Inc., which may be registered or pending registration in certain jurisdictions. The Extreme Turbodrive logo is a service mark of Extreme Networks, which may be registered or pending registration in certain jurisdictions. Specifications are subject to change without notice.

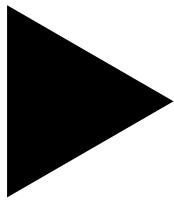
NetWare and Novell are registered trademarks of Novell, Inc. Merit is a registered trademark of Merit Network, Inc. Solaris is a trademark of Sun Microsystems, Inc. F5, BIG/ip, and 3DNS are registered trademarks of F5 Networks, Inc. see/IT is a trademark of F5 Networks, Inc.

Data Fellows™ "Data Fellows", the triangle symbol, and Data Fellows product names and symbols/logos are trademarks of Data Fellows.

F-SECURE® F-Secure SSH is a registered trademark of Data Fellows.



All other registered trademarks, trademarks and service marks are property of their respective owners.



Contents

PREFACE

Introduction	xix
Terminology	xx
Conventions	xx
Related Publications	xxi

1 EXTREMEWARE OVERVIEW

Summary of Features	1-1
Virtual LANs (VLANs)	1-3
Spanning Tree Protocol	1-3
Quality of Service	1-3
Unicast Routing	1-4
IP Multicast Routing	1-4
Load Sharing	1-4
“i” Chipset Products	1-5
“i” Chipset Feature Differences	1-5
Software Licensing	1-6
Router Licensing	1-6
Basic Functionality	1-6
Full L3 Functionality	1-6
Product Support	1-7
Verifying the Router License	1-7
Obtaining a Router License	1-7
Security Licensing	1-7
Obtaining a Security License	1-8

Security Features Under License Control	1-8
Software Factory Defaults	1-8

2 ACCESSING THE SWITCH

Understanding the Command Syntax	2-1
Syntax Helper	2-2
Command Completion with Syntax Helper	2-2
Abbreviated Syntax	2-2
Command Shortcuts	2-2
BlackDiamond and Alpine Switch Numerical Ranges	2-3
Summit Switch Numerical Ranges	2-4
Names	2-4
Symbols	2-4
Line-Editing Keys	2-5
Command History	2-6
Common Commands	2-6
Configuring Management Access	2-9
User Account	2-10
Administrator Account	2-10
Prompt Text	2-10
Default Accounts	2-11
Changing the Default Password	2-11
Creating a Management Account	2-12
Viewing Accounts	2-12
Deleting an Account	2-12
Domain Name Service Client Services	2-13
Checking Basic Connectivity	2-13
Ping	2-14
Traceroute	2-15

3 MANAGING THE SWITCH

Overview	3-1
Using the Console Interface	3-2
Using the 10/100 UTP Management Port	3-2
Using Telnet	3-3
Connecting to Another Host Using Telnet	3-3
Configuring Switch IP Parameters	3-3

Using a BOOTP Server	3-4
Manually Configuring the IP Settings	3-4
Disconnecting a Telnet Session	3-6
Controlling Telnet Access	3-6
Using Secure Shell 2 (SSH2)	3-7
Enabling SSH2	3-7
Using ExtremeWare Vista	3-8
Controlling Web Access	3-9
Using SNMP	3-9
Accessing Switch Agents	3-10
Supported MIBs	3-10
Configuring SNMP Settings	3-10
Displaying SNMP Settings	3-13
Authenticating Users	3-13
RADIUS Client	3-13
Per-Command Authentication Using RADIUS	3-14
Configuring RADIUS Client	3-14
RADIUS RFC 2138 Attributes	3-16
RADIUS Server Configuration Example (Merit)	3-16
RADIUS Per-Command Configuration Example	3-17
Configuring TACACS+	3-20
Using the Simple Network Time Protocol	3-21
Configuring and Using SNTP	3-22
SNTP Configuration Commands	3-25
SNTP Example	3-25

4 CONFIGURING BLACKDIAMOND AND ALPINE SWITCH SLOTS AND PORTS

Configuring a Slot	4-1
BlackDiamond and Alpine Switch Port Configuration	4-2
Enabling and Disabling BlackDiamond and Alpine Switch Ports	4-3
Configuring BlackDiamond and Alpine Switch Port Speed and Duplex Setting	4-3
Turning Off Autonegotiation for a Gigabit Ethernet Port	4-4
BlackDiamond and Alpine Switch Port Commands	4-4
Jumbo Frames	4-7
Enabling Jumbo Frames	4-7
Load Sharing on the BlackDiamond and Alpine Switch	4-7

Load-Sharing Algorithms	4-8
Configuring BlackDiamond and Alpine Switch Load Sharing	4-9
Load-Sharing Example	4-11
Verifying the Load-Sharing Configuration	4-11
BlackDiamond and Alpine Switch Port-Mirroring	4-11
Port-Mirroring Commands	4-12
BlackDiamond Switch Port-Mirroring Example	4-12
Extreme Discovery Protocol	4-13
EDP Commands	4-13

5 CONFIGURING SUMMIT SWITCH PORTS

Enabling and Disabling Summit Switch Ports	5-1
Configuring Summit Switch Port Speed and Duplex Setting	5-2
Turning Off Autonegotiation for a Gigabit Ethernet Port	5-3
Summit Switch Port Commands	5-3
Jumbo Frames	5-5
Enabling Jumbo Frames	5-5
Load Sharing on the Summit Switch	5-6
Load Sharing Algorithms	5-7
Configuring Summit Switch Load Sharing	5-7
Load-Sharing Example	5-10
Verifying the Load Sharing Configuration	5-10
Summit Switch Port-Mirroring	5-10
Port-Mirroring Commands	5-11
Summit Switch Port-Mirroring Example	5-12
Extreme Discovery Protocol	5-12
EDP Commands	5-12
Smart Redundancy	5-13

6 VIRTUAL LANs (VLANs)

Overview of Virtual LANs	6-1
Benefits	6-2
Types of VLANs	6-2
Port-Based VLANs	6-2
Spanning Switches with Port-Based VLANs	6-3
Tagged VLANs	6-6
Uses of Tagged VLANs	6-6

Assigning a VLAN Tag	6-6
Mixing Port-Based and Tagged VLANs	6-9
Protocol-Based VLANs	6-9
Predefined Protocol Filters	6-10
Defining Protocol Filters	6-11
Deleting a Protocol Filter	6-12
Precedence of Tagged Packets Over Protocol Filters	6-12
VLAN Names	6-12
Default VLAN	6-12
Renaming a VLAN	6-13
Configuring VLANs on the Switch	6-13
VLAN Configuration Commands	6-14
VLAN Configuration Examples	6-15
Displaying VLAN Settings	6-16
VLAN Tunneling (vMANs)	6-17
Generic VLAN Registration Protocol	6-18
GVRP and Spanning Tree Domains	6-20
GVRP Commands	6-20
MAC-Based VLANs	6-21
MAC-Based VLAN Guidelines	6-22
MAC-Based VLAN Limitations	6-22
MAC-Based VLAN Commands	6-23
MAC-Based VLAN Example	6-23
Timed Configuration Download for MAC-Based VLANs	6-24
Example	6-24

7 FORWARDING DATABASE (FDB)

Overview of the FDB	7-1
FDB Contents	7-1
FDB Entry Types	7-2
How FDB Entries Get Added	7-3
Associating a QoS Profile with an FDB Entry	7-3
Configuring FDB Entries	7-3
FDB Configuration Examples	7-4
Displaying FDB Entries	7-5

8 SPANNING TREE PROTOCOL (STP)

Overview of the Spanning Tree Protocol	8-1
Spanning Tree Domains	8-2
STPD Status for GVRP-Added Ports	8-2
Defaults	8-3
STP Configurations	8-3
Configuring STP on the Switch	8-6
STP Configuration Example	8-8
Displaying STP Settings	8-8
Disabling and Resetting STP	8-9

9 QUALITY OF SERVICE (QoS)

Overview of Policy-Based Quality of Service	9-2
Applications and Types of QoS	9-3
Voice Applications	9-3
Video Applications	9-3
Critical Database Applications	9-4
Web Browsing Applications	9-4
File Server Applications	9-4
Assigning QoS Attributes	9-5
QoS Profiles	9-6
Configuring a QoS Profile	9-8
Traffic Groupings and Creating a QoS Policy	9-8
IP-Based Traffic Groupings	9-10
MAC-Based Traffic Groupings	9-10
Permanent MAC addresses	9-10
Dynamic MAC Addresses	9-11
Blackhole MAC Address	9-11
Broadcast/Unknown Rate Limiting MAC Address	9-11
Verifying MAC-Based QoS Settings	9-12
Explicit Class of Service (802.1p and DiffServ) Traffic Groupings	9-12
Configuring 802.1p Priority	9-12
Observing 802.1p Information	9-13
802.1p Commands	9-14
Changing the Default 802.1p Mapping	9-14
Replacing 802.1p Priority Information	9-14

Configuring DiffServ	9-15
Observing DiffServ Information	9-17
Changing DiffServ Code point assignments in the QoS Profile	9-17
Replacing DiffServ Code Points	9-18
DiffServ Example	9-19
Physical and Logical Groupings	9-20
Source port	9-20
VLAN	9-20
Verifying Physical and Logical Groupings	9-21
Verifying Configuration and Performance	9-21
QoS Monitor	9-21
Real-Time Performance Monitoring	9-22
Background Performance Monitoring	9-22
Displaying QoS Profile Information	9-23
Modifying a QoS Policy	9-23
Intra-Subnet QoS	9-24
Bi-Directional Rate Shaping	9-25
Bi-Directional Rate Shaping Properties	9-25
Bi-Directional Rate Shaping Limitations	9-26
Bi-Directional Rate Shaping Commands	9-26
Dynamic Link Context System	9-27
DLCS Guidelines	9-28
DLCS Limitations	9-28
DLCS Commands	9-29

10 EXTREME STANDBY ROUTER PROTOCOL

Overview	10-1
ESRP-Aware Switches	10-2
ESRP Basics	10-2
Determining the ESRP Master	10-3
ESRP Tracking	10-4
ESRP VLAN Tracking	10-4
ESRP Route Table Tracking	10-4
ESRP Ping Tracking	10-4
ESRP Election Algorithms	10-5
Master Switch Behavior	10-5
Standby Switch Behavior	10-6

Electing the Master Switch	10-6
Failover Time	10-6
Grouping Blocks of 10/100 Ports	10-7
ESRP Options	10-9
ESRP Host Attach	10-9
ESRP Domains	10-10
ESRP Groups	10-11
Linking ESRP Switches	10-12
Configuring ESRP and Multinetting	10-12
ESRP and Spanning Tree	10-12
ESRP and VLAN aggregation	10-13
ESRP Commands	10-14
ESRP Examples	10-16
Single VLAN Using Layer 2 and Layer 3 Redundancy	10-16
Multiple VLANs Using Layer 2 Redundancy	10-18
Displaying ESRP Information	10-20

11 IP UNICAST ROUTING

Overview of IP Unicast Routing	11-2
Router Interfaces	11-2
Populating the Routing Table	11-3
Dynamic Routes	11-4
Static Routes	11-4
Multiple Routes	11-4
IP Route Sharing	11-5
Proxy ARP	11-5
ARP-Incapable Devices	11-5
Proxy ARP Between Subnets	11-6
Relative Route Priorities	11-6
IP Multinetting	11-7
IP Multinetting Operation	11-8
IP Multinetting Examples	11-9
Configuring IP Unicast Routing	11-10
Verifying the IP Unicast Routing Configuration	11-11
VLAN Aggregation	11-11
VLAN Aggregation Properties	11-13
VLAN Aggregation Limitations	11-13

Isolation Option for Communication Between Sub-VLANs	11-14
VLAN Aggregation Commands	11-14
VLAN Aggregation Example	11-15
Verifying the VLAN Aggregation Configuration	11-15
Configuring DHCP/BOOTP Relay	11-16
Verifying the DHCP/BOOTP Relay Configuration	11-16
UDP-Forwarding	11-16
Configuring UDP-Forwarding	11-17
UPD-Forwarding Example	11-17
ICMP Packet Processing	11-18
UDP-Forwarding Commands	11-18
IP Commands	11-19
Routing Configuration Example	11-25
Displaying Router Settings	11-27
Resetting and Disabling Router Settings	11-28

12 INTERIOR GATEWAY ROUTING PROTOCOLS

Overview	12-2
RIP Versus OSPF	12-2
Overview of RIP	12-3
Routing Table	12-3
Split Horizon	12-4
Poison Reverse	12-4
Triggered Updates	12-4
Route Advertisement of VLANs	12-4
RIP Version 1 Versus RIP Version 2	12-5
Overview of OSPF	12-5
Link-State Database	12-5
Areas	12-6
Area 0	12-6
Stub Areas	12-7
Not-So-Stubby-Areas (NSSA)	12-7
Normal Area	12-8
Virtual Links	12-8
Route Re-distribution	12-10
Configuring Route Re-Distribution	12-11
Re-Distributing Routes into OSPF	12-12

Previous Release Issues with OSPF Re-Distribution	12-12
Re-Distributing Routes into RIP	12-13
OSPF Timers and Authentication	12-13
Configuring RIP	12-14
RIP Configuration Example	12-17
Displaying RIP Settings	12-19
Resetting and Disabling RIP	12-20
Configuring OSPF	12-21
OSPF Configuration Example	12-25
Configuration for ABR1	12-27
Configuration for IR1	12-27
Displaying OSPF Settings	12-28
Resetting and Disabling OSPF Settings	12-28

13 EXTERIOR GATEWAY ROUTING PROTOCOLS

Overview	13-2
BGP Attributes	13-2
BGP Communities	13-3
BGP Features	13-3
Route Reflectors	13-3
Route Confederations	13-4
Route Confederation Example	13-4
Route Aggregation	13-8
Using Route Aggregation	13-8
IGP Synchronization	13-9
Using The Loopback Interface	13-9
OSPF to BGP Route Re-Distribution	13-9
Configuring BGP	13-10
Displaying BGP Settings	13-15
Resetting and Disabling BGP	13-15

14 IP MULTICAST ROUTING

Overview	14-2
DVMRP Overview	14-2
PIM Overview	14-2
PIM Dense Mode	14-3
PIM Sparse Mode (PIM-SM)	14-3

IGMP Overview	14-3
IGMP Snooping	14-4
Configuring IP Multicasting Routing	14-4
Configuration Examples	14-9
PIM-DM Configuration Example	14-10
Configuration for IR1	14-11
Configuration for ABR1	14-13
Displaying IP Multicast Routing Settings	14-13
Deleting and Resetting IP Multicast Settings	14-14

15 IPX ROUTING

Overview of IPX	15-1
Router Interfaces	15-1
IPX Routing Performance	15-3
IPX Encapsulation Types	15-3
Populating the Routing Table	15-4
Dynamic Routes	15-4
Static Routes	15-4
IPX/RIP Routing	15-4
GNS Support	15-5
Routing SAP Advertisements	15-5
Configuring IPX	15-6
Verifying IPX Router Configuration	15-6
Protocol-Based VLANs for IPX	15-7
IPX Commands	15-7
IPX Configuration Example	15-11
Displaying IPX Settings	15-13
Resetting and Disabling IPX	15-14

16 ACCESS POLICIES

Overview of Access Policies	16-1
IP Access Lists	16-2
Routing Access Policies	16-2
Route Maps	16-2
Using IP Access Lists	16-2
How IP Access Lists Work	16-3
Precedence Numbers	16-3

Specifying a Default Rule	16-3
The permit-established Keyword	16-4
Adding and Deleting Access List Entries	16-4
Maximum Entries	16-5
Access Lists for ICMP	16-5
Verifying Access List Configurations	16-6
Access List Commands	16-6
IP Access List Examples	16-11
Using the Permit-Established Keyword	16-11
Example 2: Filter ICMP Packets	16-14
Using Routing Access Policies	16-15
Creating an Access Profile	16-16
Configuring an Access Profile Mode	16-16
Adding an Access Profile Entry	16-17
Specifying Subnet Masks	16-17
Sequence Numbering	16-17
Permit and Deny Entries	16-18
Autonomous System Expressions	16-18
Deleting an Access Profile Entry	16-18
Applying Access Profiles	16-18
Routing Access Policies for RIP	16-19
Examples	16-19
Routing Access Policies for OSPF	16-21
Example	16-22
Routing Access Policies for DVMRP	16-23
Example	16-23
Routing Access Policies for PIM	16-24
Example	16-24
Routing Access Policies for BGP	16-25
Making Changes to a Routing Access Policy	16-25
Removing a Routing Access Policy	16-26
Routing Access Policy Commands	16-26
Using Route Maps	16-29
Creating a Route Map	16-30
Add Entries to the Route Map	16-30
Add Statements to the Route Map Entries	16-30
Route Map Operation	16-32
Route Map Example	16-32

Changes to Route Maps 16-34
Route Maps in BGP 16-34
Route Map Commands 16-35

17 SERVER LOAD BALANCING (SLB)

Overview 17-2
SLB Components 17-2
 Nodes 17-3
 Pools 17-3
 Virtual Servers 17-3
 Using Standard or Wildcard Virtual Servers 17-4
Forwarding Modes 17-5
 Transparent Mode 17-5
 Translational Mode 17-8
 Port Translation Mode 17-10
 GoGo Mode 17-11
VIP Network Advertisement 17-12
Balancing Methods 17-13
 Round-Robin 17-13
 Ratio 17-13
 Ratio Weight 17-14
 Least Connections 17-14
 Priority 17-14
Basic SLB Commands 17-15
Advanced SLB Application Example 17-18
Health Checking 17-22
 Ping-Check 17-23
 Ping-Check Commands 17-23
 TCP-Port-Check 17-23
 TCP-Port-Check Commands 17-23
 Service-Check 17-24
 Service-Check Commands 17-25
External Health Checking 17-25
 Maintenance Mode 17-25
Persistence 17-26
 Client Persistence 17-26
 Sticky Persistence 17-26

Server Load Balancing with ESRP	17-27
Configuring the Switches for SLB and ESRP	17-29
Notes regarding configuration for SLB with ESRP	17-30
Web-Server Configuration	17-30
Using High Availability System Features	17-31
Redundant SLB	17-31
Using Ping-Check	17-32
Configuring Active-Active Operation	17-32
Sample Active-Active Configuration	17-33
Using Manual Fail-Back	17-35
Using SLB High Availability	17-36
Configuring Clients	17-37
Configuring Switches for SLB H/A	17-37
Notes Regarding SLB/HA	17-38
Web-Server configuration	17-39
3DNS Support	17-40
Advanced SLB Commands	17-40
Web Cache Redirection	17-46
Flow Redirection	17-46
Flow Redirection Commands	17-47
Flow Redirection Example	17-47

18 STATUS MONITORING AND STATISTICS

Status Monitoring	18-1
Slot Diagnostics	18-3
Port Statistics	18-4
Port Errors	18-5
Port Monitoring Display Keys	18-6
Setting the System Recovery Level	18-7
Logging	18-7
Local Logging	18-9
Real-Time Display	18-9
Remote Logging	18-9
Logging Configuration Changes	18-10
Logging Commands	18-11
RMON	18-12
About RMON	18-12

RMON Features of the Switch	18-13
Statistics	18-13
History	18-13
Alarms	18-13
Events	18-14
Configuring RMON	18-14
Event Actions	18-15

19 USING EXTREMEWARE VISTA

Enabling and Disabling Web Access	19-2
Setting Up Your Browser	19-2
Accessing ExtremeWare Vista	19-3
Navigating ExtremeWare Vista	19-4
Task Frame	19-4
Content Frame	19-4
Browser Controls	19-5
Status Messages	19-5
Standalone Buttons	19-5
Saving Changes	19-6
Filtering Information	19-6
Do a GET When Configuring a VLAN	19-7
Sending Screen Output to Extreme Networks	19-7

20 SOFTWARE UPGRADE AND BOOT OPTIONS

Downloading a New Image	20-1
Rebooting the Switch	20-2
Saving Configuration Changes	20-3
Returning to Factory Defaults	20-3
Using TFTP to Upload the Configuration	20-4
Using TFTP to Download the Configuration	20-5
Downloading a Complete Configuration	20-5
Downloading an Incremental Configuration	20-5
Scheduled Incremental Configuration Download	20-6
Remember to Save	20-6
Synchronizing MSMs	20-7
Upgrading and Accessing BootROM	20-7
Upgrading BootROM	20-7

Accessing the BootROM menu	20-7
Boot Option Commands	20-8

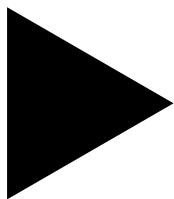
A SUPPORTED STANDARDS

B TROUBLESHOOTING

LEDs	B-1
Using the Command-Line Interface	B-3
Port Configuration	B-5
VLANs	B-6
STP	B-7
Debug Tracing	B-8
TOP Command	B-8
Contacting Extreme Technical Support	B-8

INDEX

INDEX OF COMMANDS



Preface

This Preface provides an overview of this guide, describes guide conventions, and lists other publications that may be useful.

INTRODUCTION

This guide provides the required information to configure ExtremeWare™ software running on a BlackDiamond™, Alpine™, or Summit™ switch.

This guide is intended for use by network administrators who are responsible for installing and setting up network equipment. It assumes a basic working knowledge of the following:

- Local area networks (LANs)
- Ethernet concepts
- Ethernet switching and bridging concepts
- Routing concepts
- Internet Protocol (IP) concepts
- Routing Information Protocol (RIP) and Open Shortest Path First (OSPF) concepts
- Border Gateway Protocol (BGP-4) concepts
- IP Multicast concepts
- Distance Vector Multicast Routing Protocol (DVMRP) concepts
- Protocol Independent Multicast (PIM) concepts

- Internet Packet Exchange (IPX) concepts
- Server Load Balancing (SLB) concepts
- Simple Network Management Protocol (SNMP)

 *If the information in the “Release Notes” shipped with your switch differs from the information in this guide, follow the “Release Notes.”*

TERMINOLOGY

When features, functionality, or operation is specific to the Summit, Alpine, or BlackDiamond switch family, the family name is used. Explanations about features and operations that are the same across all switch product families simply refer to the product as the “switch.”

CONVENTIONS

[Table 1](#) and [Table 2](#) list conventions that are used throughout this guide.

Table 1: Notice Icons

Icon	Notice Type	Alerts you to...
	Note	Important features or instructions.
	Caution	Risk of personal injury, system damage, or loss of data.
	Warning	Risk of severe personal injury.

Table 2: Text Conventions

Convention	Description
Screen displays	This typeface indicates command syntax, or represents information as it appears on the screen.
Screen displays bold	This typeface indicates how you would type a particular command.
The words “enter” and “type”	When you see the word “enter” in this guide, you must type something, and then press the Return or Enter key. Do not press the Return or Enter key when an instruction simply says “type.”
[Key] names	Key names are written with brackets, such as [Return] or [Esc]. If you must press two or more keys simultaneously, the key names are linked with a plus sign (+). Example: Press [Ctrl]+[Alt]+[Del].
Words in <i>italicized</i> type	Italics emphasize a point or denote new terms at the place where they are defined in the text.

RELATED PUBLICATIONS

The following is a list of related publications:

- *ExtremeWare Quick Reference Guide*
- ExtremeWare “Release Notes”
- *BlackDiamond 6800 Hardware Installation Guide*
- *BlackDiamond 3800 Hardware Installation Guide*
- *Summit Hardware Installation Guide*

Documentation for Extreme Networks products is available on the World Wide Web at the following location:

- <http://www.extremenetworks.com/>



ExtremeWare Overview

This chapter covers the following topics:

- [Summary of Features on page 1-1](#)
- [“i” Chipset Products on page 1-5](#)
- [Software Licensing on page 1-6](#)
- [Software Factory Defaults on page 1-8](#)

ExtremeWare is the full-featured software operating system that is designed to run on the BlackDiamond, Alpine, and Summit families of Gigabit Ethernet switches.

SUMMARY OF FEATURES

The features of ExtremeWare include the following:

- Virtual local area networks (VLANs) including support for IEEE 802.1Q and IEEE 802.1p
- VLAN aggregation
- Spanning Tree Protocol (STP) (IEEE 802.1D) with multiple STP domains
- Policy-Based Quality of Service (PB-QoS)
- Wire-speed Internet Protocol (IP) routing
- IP Multinetting
- DHCP/BOOTP Relay

- Extreme Standby Router Protocol (ESRP)
- Routing Information Protocol (RIP) version 1 and RIP version 2
- Open Shortest Path First (OSPF) routing protocol
- Border Gateway Protocol (BGP) version 4
- Wire-speed IP multicast routing support
- Diffserv support
- Access-policy support for routing protocols
- Access list support for packet filtering
- IGMP snooping to control IP multicast traffic
- Distance Vector Multicast Routing Protocol (DVMRP)
- Protocol Independent Multicast-Dense Mode (PIM-DM)
- Protocol Independent Multicast-Sparse Mode (PIM-SM)
- Wire-speed IPX, IPX/RIP, and IPX/SAP support
- Server Load Balancing (SLB) support
- Load sharing on multiple ports, across all blades (BlackDiamond only)
- RADIUS client and per-command authentication support
- TACACS+ support
- Console command-line interface (CLI) connection
- Telnet CLI connection
- SSH2 connection
- ExtremeWare Vista Web-based management interface
- Simple Network Management Protocol (SNMP) support
- Remote Monitoring (RMON)
- Traffic mirroring for all ports, across all blades (BlackDiamond only)



For more information on BlackDiamond 6808 switch components, refer to the BlackDiamond 6800 Switch Hardware Installation Guide. For more information on Alpine 3800 switch components, refer to the Alpine 3800 Switch Hardware Installation Guide. For more information on Summit switch components, refer to the Summit Hardware Installation Guide.

VIRTUAL LANs (VLANs)

ExtremeWare has a VLAN feature that enables you to construct your broadcast domains without being restricted by physical connections. A VLAN is a group of location- and topology-independent devices that communicate as if they were on the same physical local area network (LAN).

Implementing VLANs on your network has the following three advantages:

- They help to control broadcast traffic. If a device in VLAN *Marketing* transmits a broadcast frame, only VLAN *Marketing* devices receive the frame.
- They provide extra security. Devices in VLAN *Marketing* can only communicate with devices on VLAN *Sales* using routing services.
- They ease the change and movement of devices on networks.



For more information on VLANs, refer to [Chapter 6](#).

SPANNING TREE PROTOCOL

The switch supports the IEEE 802.1D Spanning Tree Protocol (STP), which is a bridge-based mechanism for providing fault tolerance on networks. STP enables you to implement parallel paths for network traffic, and ensure the following:

- Redundant paths are disabled when the main paths are operational.
- Redundant paths are enabled if the main traffic paths fail.

A single spanning tree can span multiple VLANs.



For more information on STP, refer to [Chapter 8](#).

QUALITY OF SERVICE

ExtremeWare has Policy-Based Quality of Service (QoS) features that enable you to specify service levels for different traffic groups. By default, all traffic is assigned the “normal” QoS policy profile. If needed, you can create other QoS policies and apply them to different traffic types so that they have different guaranteed minimum bandwidth, maximum bandwidth, and priority.



For more information on Quality of Service, refer to [Chapter 9](#).

UNICAST ROUTING

The switch can route IP or IPX traffic between the VLANs that are configured as virtual router interfaces. Both dynamic and static IP routes are maintained in the routing table. The following routing protocols are supported:

- RIP version 1
- RIP version 2
- OSPF
- IPX/RIP
- BGP version 4



For more information on IP unicast routing, refer to [Chapter 11](#). For more information on IPX/RIP, refer to [Chapter 15](#).

IP MULTICAST ROUTING

The switch can use IP multicasting to allow a single IP host to transmit a packet to a group of IP hosts. ExtremeWare supports multicast routes that are learned by way of the Distance Vector Multicast Routing Protocol (DVMRP) or the Protocol Independent Multicast (dense mode or sparse mode).



For more information on IP multicast routing, refer to [Chapter 14](#).

LOAD SHARING

Load sharing allows you to increase bandwidth and resiliency by using a group of ports to carry traffic in parallel between systems. The sharing algorithm allows the switch to use multiple ports as a single logical port. For example, VLANs see the load-sharing group as a single virtual port. The algorithm also guarantees packet sequencing between clients.



For information on load sharing, refer to [Chapter 4](#) and [Chapter 5](#).

“i” CHIPSET PRODUCTS

Summit switches and BlackDiamond 6800 switch modules that use naming conventions ending with an “i” have additional capabilities that are documented throughout this User Guide. For the most current list of products supporting the “i” chipset, consult your Release Notes.

Unless otherwise specified, a feature requiring the “i” chipset requires the use of the BlackDiamond MSM64i and an “i” chipset-based I/O module, such as the G8Xi.

“i” CHIPSET FEATURE DIFFERENCES

The following list summarizes the feature areas specific to the “i” chipset products:

- QoS and Access Policies – Complete use of IP access lists (products without the “i” chipset are capable of a subset of this functionality); support for IP DiffServ; and support for eight QoS queues per port, instead of four.
- Bridging/Switching – Support for jumbo frames; support for address- and round-robin-based load-sharing algorithms; ports belonging to a load-sharing group do not need to be contiguous.
- Routing – Wire-speed IPX routing
- BGP-4 – Requires the use of the “i” chipset, but requires only the MSM64i on the BlackDiamond.
- Server Load Balancing – Requires the use of the “i” chipset.
- Web cache redirection – Requires the use of the “i” chipset.
- ESRP – No port blocking restrictions, use of the additional tracking ESRP feature.
- Load sharing – No contiguous port restrictions.

SOFTWARE LICENSING

Some Extreme Networks products have capabilities that are enabled by using a license key. Keys are typically unique to the switch, and are not transferable. Keys are stored in NVRAM and, once entered, persist through reboots, software upgrades, and reconfigurations. The following sections describe the features that are associated with license keys.

ROUTER LICENSING

Some switches support software licensing for different levels of router functionality. In ExtremeWare version 6.0 and above, routing protocol support is separated into two sets: Basic and Full L3. Basic is a subset of Full L3.

BASIC FUNCTIONALITY

Basic functionality requires *no license key*. All Extreme switches have Basic layer 3 functionality, without the requirement of a license key. Basic functionality includes all switching functions, and also includes all available layer 3 QoS, access list, and ESRP functions. Layer 3 routing functions include support for the following:

- IP routing using RIP version 1 and/or RIP version 2
- IP routing between directly attached VLANs
- IP routing using static routes

FULL L3 FUNCTIONALITY

On switches that support router licensing, the Full L3 license enables support of additional routing protocols and functions, including the following:

- IP routing using OSPF
- IP multicast routing using DVMRP
- IP multicast routing using PIM (Dense Mode or Sparse Mode)
- IPX routing (direct, static, and dynamic using IPX/RIP and IPX/SAP)
- IP routing using BGP
- Server load balancing
- Web cache redirection

PRODUCT SUPPORT

ExtremeWare version 6.0 and above supports router licensing on the Summit24 switch, Summit48 switch, and Summit7i switch. The BlackDiamond 6808 switch supports all documented router functions, without the need for additional router licensing. Consult the Release Notes for the most current set of products that require router licensing support.

VERIFYING THE ROUTER LICENSE

To verify the router license, use the `show switch` command.

OBTAINING A ROUTER LICENSE

You can order the desired functionality from the factory, using the appropriate model of the desired product. If you order licensing from the factory, the switch arrives packaged with a certificate that contains the unique license key(s), and instructions for enabling the correct functionality on the switch. The certificate is typically packaged with the switch documentation. Once the license key is entered, it should not be necessary to enter the information again. However, we recommend keeping the certificate for your records.

You may upgrade the router licensing of an existing product by purchasing a voucher for the desired product and functionality. Please contact your supplier to purchase a voucher.

Once received, the voucher contains information and instructions on obtaining a license key for the switch using the Extreme Networks Support website at:

<http://www.extremenetworks.com/extreme/support/upgrade.htm>

or by phoning Extreme Networks Technical Support at:

- (800) 998-2408
- (408) 579-2826

SECURITY LICENSING

Certain additional ExtremeWare security features, such as the use of Secure Shell (SSH2) encryption, may be under United States export restriction control. Extreme Networks ships these security features in a disabled state. You may obtain information on enabling these features at no charge from Extreme Networks.

OBTAINING A SECURITY LICENSE

To obtain information on enabling features that require export restriction, access the Extreme Networks Support website at:

<http://www.extremenetworks.com/go/security.htm>

Fill out a contact form to indicate compliance or non-compliance with the export restrictions. If you are in compliance, you will be given information that will allow you to enable security features.

SECURITY FEATURES UNDER LICENSE CONTROL

ExtremeWare version 6.0 and above supports the SSH2 protocol. SSH2 allows the encryption of Telnet session data. The encryption methods used are under U.S. export restriction control.

SOFTWARE FACTORY DEFAULTS

Table 1-1 shows factory defaults for global ExtremeWare features.

Table 1-1: ExtremeWare Global Factory Defaults

Item	Default Setting
Serial or Telnet user account	<i>admin</i> with no password and <i>user</i> with no password
Web network management	Enabled
Telnet	Enabled
SSH2	Disabled
SNMP	Enabled
SNMP read community string	<i>public</i>
SNMP write community string	<i>private</i>
RMON	Disabled
BOOTP	Enabled on the default VLAN (<i>default</i>)
QoS	All traffic is part of the default queue
QoS monitoring	Automatic roving
802.1p priority	Recognition enabled
802.3x flow control	Enabled on Gigabit Ethernet ports

Table 1-1: ExtremeWare Global Factory Defaults (continued)

Item	Default Setting
Virtual LANs	Three VLANs pre-defined. VLAN named <i>default</i> contains all ports and belongs to the STPD named <i>s0</i> . VLAN <i>mgmt</i> exists only on switches that have an Ethernet management port, and contains only that port. The Ethernet management port is DTE only, and is not capable of switching or routing. VLAN <i>MacVlanDiscover</i> is used only when using the MAC VLAN feature.
802.1Q tagging	All packets are untagged on the default VLAN (<i>default</i>)
Spanning Tree Protocol	Disabled for the switch; enabled for each port in the STPD
Forwarding database aging period	300 seconds (5 minutes)
IP Routing	Disabled
RIP	Disabled
OSPF	Disabled
IP multicast routing	Disabled
IGMP	Enabled
IGMP snooping	Enabled
DVMRP	Disabled
GVRP	Disabled
PIM-DM	Disabled
IPX routing	Disabled
NTP	Disabled
DNS	Disabled
Port mirroring	Disabled



For default settings of individual ExtremeWare features, refer to individual chapters in this guide.

Accessing the Switch

This chapter covers the following topics:

- [Understanding the Command Syntax on page 2-1](#)
- [Line-Editing Keys on page 2-5](#)
- [Command History on page 2-6](#)
- [Common Commands on page 2-6](#)
- [Configuring Management Access on page 2-9](#)
- [Domain Name Service Client Services on page 2-13](#)
- [Checking Basic Connectivity on page 2-13](#)

UNDERSTANDING THE COMMAND SYNTAX

This section describes the steps to take when entering a command. Refer to the sections that follow for detailed information on using the command-line interface.

When entering a command at the prompt, ensure that you have the appropriate privilege level. Most configuration commands require you to have the administrator privilege level. To use the command-line interface (CLI), follow these steps:

- 1 Enter the command name.

If the command does not include a parameter or values, skip to Step 3. If the command requires more information, continue to Step 2.

- 2 If the command includes a parameter, enter the parameter name and values.

- 3 The value part of the command specifies how you want the parameter to be set. Values include numerics, strings, or addresses, depending on the parameter.

- 4 After entering the complete command, press [Return].



If an asterisk () appears in front of the command-line prompt, it indicates that you have outstanding configuration changes that have not been saved. For more information on saving configuration changes, refer to [Chapter 20](#).*

SYNTAX HELPER

The CLI has a built-in syntax helper. If you are unsure of the complete syntax for a particular command, enter as much of the command as possible and press [Return]. The syntax helper provides a list of options for the remainder of the command.

The syntax helper also provides assistance if you have entered an incorrect command.

COMMAND COMPLETION WITH SYNTAX HELPER

ExtremeWare provides command completion by way of the [Tab] key. If you enter a partial command, pressing the [Tab] key posts a list of available options, and places the cursor at the end of the command.

ABBREVIATED SYNTAX

Abbreviated syntax is the most unambiguous, shortest allowable abbreviation of a command or parameter. Typically, this is the first three letters of the command.



When using abbreviated syntax, you must enter enough characters to make the command unambiguous and distinguishable to the switch.

COMMAND SHORTCUTS

All named components of the switch configuration must have a unique name. Components are named using the `create` command. When you enter a command to configure a named component, you do not need to use the keyword of the component. For example, to create a VLAN, you must enter a unique VLAN name:

```
create vlan engineering
```

Once you have created the VLAN with a unique name, you can then eliminate the keyword `vlan` from all other commands that require the name to be entered. For example, instead of entering the BlackDiamond switch command

```
config vlan engineering delete port 1:3,4:6
```

you could enter the following shortcut:

```
config engineering delete port 1:3,4:6
```

Similarly, on the Summit switch, instead of entering the command

```
config vlan engineering delete port 1-3,6
```

you could enter the following shortcut:

```
config engineering delete port 1-3,6
```

BLACKDIAMOND AND ALPINE SWITCH NUMERICAL RANGES

Commands that require you to enter one or more port numbers on a BlackDiamond and Alpine switch use the parameter `<portlist>` in the syntax. A `<portlist>` can be one port on a particular slot. For example,

```
port 3:1
```

A `<portlist>` can be a range of numbers. For example,

```
port 3:1-3:3
```

You can add additional slot and port numbers to the list, separated by a comma:

```
port 3:1,4:8,6:10
```

You can specify all ports on a particular slot. For example,

```
port 3:*
```

indicates all ports on slot 3.

You can specify a range of slots and ports. For example,

```
port 2:3-4:5
```

indicates slot 2, port 3 through slot 4, port 5.

SUMMIT SWITCH NUMERICAL RANGES

Commands that require you to enter one or more port numbers on a Summit switch use the parameter <portlist> in the syntax. A portlist can be a range of numbers, for example:

```
port 1-3
```

You can add additional port numbers to the list, separated by a comma:

```
port 1-3,6,8
```

NAMES

All named components of the switch configuration must have a unique name. Names must begin with an alphabetical character and are delimited by whitespace, unless enclosed in quotation marks.

SYMBOLS

You may see a variety of symbols shown as part of the command syntax. These symbols explain how to enter the command, and you do not type them as part of the command itself. [Table 2-1](#) summarizes command syntax symbols.

Table 2-1: Command Syntax Symbols

Symbol	Description
angle brackets < >	Enclose a variable or value. You must specify the variable or value. For example, in the syntax <code>config vlan <name> ipaddress <ip_address></code> you must supply a VLAN name for <name> and an address for <ip_address> when entering the command. Do not type the angle brackets.
square brackets []	Enclose a required value or list of required arguments. One or more values or arguments can be specified. For example, in the syntax <code>use image [primary secondary]</code> you must specify either the primary or secondary image when entering the command. Do not type the square brackets.

Table 2-1: Command Syntax Symbols (continued)

Symbol	Description
vertical bar	Separates mutually exclusive items in a list, one of which must be entered. For example, in the syntax <code>config snmp community [read-only read-write] <string></code> you must specify either the read or write community string in the command. Do not type the vertical bar.
braces { }	Enclose an optional value or a list of optional arguments. One or more values or arguments can be specified. For example, in the syntax <code>reboot {<date> <time> cancel}</code> you can specify either a particular date and time combination, or the keyword <code>cancel</code> to cancel a previously scheduled reboot. If you do not specify an argument, the command will prompt, asking if you want to reboot the switch now. Do not type the braces.

LINE-EDITING KEYS

[Table 2-2](#) describes the line-editing keys available using the CLI.

Table 2-2: Line-Editing Keys

Key(s)	Description
Backspace	Deletes character to left of cursor and shifts remainder of line to left.
Delete or [Ctrl] + D	Deletes character under cursor and shifts remainder of line to left.
[Ctrl] + K	Deletes characters from under cursor to end of line.
Insert	Toggles on and off. When toggled on, inserts text and shifts previous text to right.
Left Arrow	Moves cursor to left.
Right Arrow	Moves cursor to right.
Home or [Ctrl] + A	Moves cursor to first character in line.
End or [Ctrl] + E	Moves cursor to last character in line.
[Ctrl] + L	Clears screen and moves cursor to beginning of line.
[Ctrl] + P or Up Arrow	Displays previous command in command history buffer and places cursor at end of command.
[Ctrl] + N or Down Arrow	Displays next command in command history buffer and places cursor at end of command.

Table 2-2: Line-Editing Keys (continued)

Key(s)	Description
[Ctrl] + U	Clears all characters typed from cursor to beginning of line.
[Ctrl] + W	Deletes previous word.

COMMAND HISTORY

ExtremeWare “remembers” the last 49 commands you entered. You can display a list of these commands by using the following command:

```
history
```

COMMON COMMANDS

[Table 2-3](#) describes common commands used to manage the switch. Commands specific to a particular feature are described in the other chapters of this guide.

Table 2-3: Common Commands

Command	Description
clear session <number>	Terminates a Telnet session from the switch.
config account <username> {<password>}	Configures a user account password. Passwords must have a minimum of 1 character and can have a maximum of 32 characters. User names and passwords are case-sensitive.
config banner	Configures the banner string. You can enter up to 24 rows of 79-column text that is displayed before the login prompt of each session. Press [Return] at the beginning of a line to terminate the command and apply the banner. To clear the banner, press [Return] at the beginning of the first line.
config ports <portlist> auto off {speed [10 100 1000]} duplex [half full]	Manually configures the port speed and duplex setting of one or more ports on a switch.

Table 2-3: Common Commands (continued)

Command	Description
config slot <slot> module [f32t f32f f48t g4x g6x g8x g12x]	Configures a slot for a particular I/O module card.
config ssh2 key {pregenerated}	Generates the SSH2 host key.
config sys-recovery-level [none critical all]	Configures a recovery option for instances where an exception occurs in ExtremeWare. Specify one of the following: <ul style="list-style-type: none"> ■ none — No recovery mode. ■ critical — ExtremeWare logs an error to the syslog, and reboots the system after a critical task exceptions ■ all — ExtremeWare logs an error to the syslog, and reboots the system after any exception.
config time <date> <time>	Configures the system date and time. The format is as follows: mm/dd/YYYY hh:mm:ss The time uses a 24-hour clock format. You cannot set the year past 2036.
config timezone <gmt_offset> {autodst nosautodst}	Configures the time zone information to the configured offset from GMT time. The format of <code>gmt_offset</code> is +/- minutes from GMT time. Specify: <ul style="list-style-type: none"> ■ autodst — Enables automatic Daylight Savings Time change. ■ nosautodst — Disables automatic Daylight Savings Time change. The default setting is <code>autodst</code> .
config vlan <name> ipaddress <ip_address> <mask>	Configures an IP address and subnet mask for a VLAN.
create account [admin user] <username> <password>	Creates a user account. This command is available to admin-level users and to users with RADIUS command authorization. The <code>username</code> is between 1 and 32 characters, the <code>password</code> is between 0 and 16 characters.
create vlan <name>	Creates a VLAN.
delete account <username>	Deletes a user account.

Table 2-3: Common Commands (continued)

Command	Description
delete vlan <name>	Deletes a VLAN.
disable bootp vlan [<name> all]	Disables BOOTP for one or more VLANs.
disable cli-config-logging	Disables logging of CLI commands to the Syslog.
disable clipaging	Disables pausing of the screen display when a show command output reaches the end of the page.
disable idletimeout	Disables the timer that disconnects all sessions. Once disabled, console sessions remain open until the switch is rebooted or you logoff. Telnet sessions remain open until you close the Telnet client.
disable port <portlist>	Disables a port on the switch.
disable ssh2	Disables SSH2 Telnet access to the switch.
disable telnet	Disables Telnet access to the switch.
disable web	Disables Web access to the switch.
enable bootp vlan [<name> all]	Enables BOOTP for one or more VLANs.
enable cli-config-logging	Enables the logging of CLI configuration commands to the Syslog for auditing purposes. The default setting is enabled.
enable clipaging	Enables pausing of the screen display when show command output reaches the end of the page. The default setting is enabled.
enable idletimeout	Enables a timer that disconnects all sessions (both Telnet and console) after 20 minutes of inactivity. The default setting is disabled.
enable license [full_L3 service-provider security] <license_key>	Enables a particular software feature license. Specify <license_key> as an integer. The command unconfig switch all does not clear licensing information. This license cannot be disabled once it is enabled on the switch.

Table 2-3: Common Commands (continued)

Command	Description
enable ssh2 {access-profile [<access_profile> none]} {port <tcp_port_number>}	Enables SSH2 Telnet sessions. By default, SSH2 is enabled with no access profile, and uses TCP port number 22. To cancel a previously configured access-profile, use the <code>none</code> option.
enable telnet {access-profile [<access_profile> none]} {port <tcp_port_number>}	Enables Telnet access to the switch. By default, Telnet is enabled with no access profile, and uses TCP port number 23. To cancel a previously configured access-profile, use the <code>none</code> option.
enable web {access-profile [<access_profile> none]} {port <tcp_port_number>}	Enables ExtremeWare Vista Web access to the switch. By default, Web access is enabled with no access profile, using TCP port number 80. Use the <code>none</code> option to cancel a previously configured access-profile. You must reboot the switch for this command to take effect.
history	Displays the previous 49 commands entered on the switch.
show banner	Displays the user-configured banner.
unconfig switch {all}	Resets all switch parameters (with the exception of defined user accounts, and date and time information) to the factory defaults. If you specify the keyword <code>all</code> , the switch erases the currently selected configuration image in flash memory and reboots. As a result, all parameters are reset to default settings.

CONFIGURING MANAGEMENT ACCESS

ExtremeWare supports the following two levels of management:

- User
- Administrator

In addition to the management levels, you can optionally use an external RADIUS server to provide CLI command authorization checking for each command. For more information on RADIUS, refer to “[RADIUS Client](#),” in [Chapter 3](#).

USER ACCOUNT

A user-level account has viewing access to all manageable parameters, with the exception of the following:

- User account database
- SNMP community strings

A user-level account can use the ping command to test device reachability, and change the password assigned to the account name. If you have logged on with user capabilities, the command-line prompt ends with a (>) sign. For example:

```
Summit1:2>
```

ADMINISTRATOR ACCOUNT

An administrator-level account can view and change all switch parameters. It can also add and delete users, and change the password associated with any account name. The administrator can disconnect a management session that has been established by way of a Telnet connection. If this happens, the user logged on by way of the Telnet connection is notified that the session has been terminated.

If you have logged on with administrator capabilities, the command-line prompt ends with a (#) sign. For example:

```
Summit1:18#
```

PROMPT TEXT

The prompt text is taken from the SNMP sysname setting. The number that follows the colon indicates the sequential line/command number.

If an asterisk (*) appears in front of the command-line prompt, it indicates that you have outstanding configuration changes that have not been saved. For example:

```
*Summit1:19#
```

DEFAULT ACCOUNTS

By default, the switch is configured with two accounts, as shown in [Table 2-4](#).

Table 2-4: Default Accounts

Account Name	Access Level
admin	This user can access and change all manageable parameters. The admin account cannot be deleted.
user	This user can view (but not change) all manageable parameters, with the following exceptions: <ul style="list-style-type: none"> ■ This user cannot view the user account database. ■ This user cannot view the SNMP community strings.

CHANGING THE DEFAULT PASSWORD

Default accounts do not have passwords assigned to them. Passwords must have a minimum of 4 characters and can have a maximum of 12 characters.



User names and passwords are case-sensitive.

To add a password to the default admin account, follow these steps:

- 1 Log in to the switch using the name *admin*.
- 2 At the password prompt, press [Return].
- 3 Add a default admin password by entering the following:
config account admin
- 4 Enter the new password at the prompt.
- 5 Re-enter the new password at the prompt.

To add a password to the default user account, follow these steps:

- 1 Log in to the switch using the name *admin*.
- 2 At the password prompt, press [Return], or enter the password that you have configured for the *admin* account.
- 3 Add a default user password by entering the following:
config account user
- 4 Enter the new password at the prompt.

- 5 Re-enter the new password at the prompt.

 *If you forget your password while logged out of the command-line interface, contact your local technical support representative, who will advise on your next course of action.*

CREATING A MANAGEMENT ACCOUNT

The switch can have a total of 16 management accounts. You can use the default names (*admin* and *user*), or you can create new names and passwords for the accounts. Passwords can have a minimum of 0 characters and can have a maximum of 31 characters.

To create a new account, follow these steps:

- 1 Log in to the switch as *admin*.
- 2 At the password prompt, press [Return], or enter the password that you have configured for the *admin* account.
- 3 Add a new user by using the following command:
`create account [admin | user] <username>`
- 4 Enter the password at the prompt.
- 5 Re-enter the password at the prompt.

VIEWING ACCOUNTS

To view the accounts that have been created, you must have administrator privileges. Use the following command to see the accounts:

```
show accounts
```

DELETING AN ACCOUNT

To delete a account, you must have administrator privileges. Use the following command to delete an account:

```
delete account <username>
```



The account name admin cannot be deleted.

DOMAIN NAME SERVICE CLIENT SERVICES

The Domain Name Service (DNS) client in ExtremeWare augments the following commands to allow them to accept either IP addresses or host names:

- telnet
- download [bootrom | configuration | image]
- upload configuration
- ping
- traceroute

In addition, the nslookup utility can be used to return the IP address of a hostname.

[Table 2-5](#) describes the commands used to configure DNS.

Table 2-5: DNS Commands

Command	Description
config dns-client add <ipaddress>	Adds a DNS name server(s) to the available server list for the DNS client. Up to three name servers can be configured.
config dns-client default-domain <domain_name>	Configures the domain that the DNS client uses if a fully qualified domain name is not entered. For example, if the default domain is configured to be <code>foo.com</code> , executing <code>ping bar</code> searches for <code>bar.foo.com</code> .
config dns-client delete <ipaddress>	Removes a DNS server.
nslookup <hostname>	Displays the IP address of the requested host.
show dns-client	Displays the DNS configuration.

CHECKING BASIC CONNECTIVITY

The switch offers the following commands for checking basic connectivity:

- ping
- traceroute

PING

The ping command enables you to send Internet Control Message Protocol (ICMP) echo messages to a remote IP device. The ping command is available for both the user and administrator privilege level.

The ping command syntax is

```
ping {continuous} {size <start_size> {- <end_size>}} [<ip_address> |  
<hostname>] {from <src_address> | with record-route | from  
<src_ipaddress> with record-route}
```

Options for the ping command are described in [Table 2-6](#).

Table 2-6: Ping Command Parameters

Parameter	Description
continuous	Specifies ICMP echo messages to be sent continuously. This option can be interrupted by pressing any key.
size	Specifies the size of the ICMP request. If both the start_size and end_size are specified, transmits ICMP requests using 1 byte increments, per packet. If no end_size is specified, packets of start_size are sent.
<ipaddress>	Specifies the IP address of the host.
<hostname>	Specifies the name of the host. To use the hostname, you must first configure DNS.
from	Uses the specified source address in the ICMP packet. If not specified, the address of the transmitting interface is used.
with record-route	Decodes the list of recorded routes and displays them when the ICMP echo reply is received.

If a ping request fails, the switch continues to send ping messages until interrupted. Press any key to interrupt a ping request.

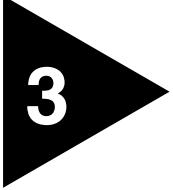
TRACEROUTE

The traceroute command enables you to trace the routed path between the switch and a destination endstation. The traceroute command syntax is

```
traceroute [<ip_address> | <hostname>] {from <src_ipaddress>} {ttl <TTL>} {port <port>}
```

where:

- `ip_address` is the IP address of the destination endstation.
- `hostname` is the hostname of the destination endstation. To use the hostname, you must first configure DNS.
- `from` uses the specified source address in the ICMP packet. If not specified, the address of the transmitting interface is used.
- `ttl` configures the switch to trace up to the time-to-live number of the switch.
- `port` uses the specified UDP port number.



3

Managing the Switch

This chapter covers the following topics:

- [Overview on page 3-1](#)
- [Using the Console Interface on page 3-2](#)
- [Using Telnet on page 3-3](#)
- [Using Secure Shell 2 \(SSH2\) on page 3-7](#)
- [Using ExtremeWare Vista on page 3-8](#)
- [Using SNMP on page 3-9](#)
- [Authenticating Users on page 3-13](#)
- [Using the Simple Network Time Protocol on page 3-21](#)

OVERVIEW

Using ExtremeWare, you can manage the switch using the following methods:

- Access the CLI by connecting a terminal (or workstation with terminal-emulation software) to the console port.
- Access the switch remotely using TCP/IP through one of the switch ports or through the dedicated 10/100 unshielded twisted pair (UTP) Ethernet management port (on switches that are so equipped). Remote access includes the following:
 - Telnet using the CLI interface
 - SSH2 using the CLI interface

- ExtremeWare Vista Web access using a standard Web browser
- SNMP access using ExtremeWare Enterprise Manager or another SNMP manager

The switch supports up to the following number of concurrent user sessions:

- One console session
 - Two console sessions are available on a BlackDiamond switch that has two Management Switch Fabric Modules (MSMs) installed.
- Eight Telnet sessions
- Eight SSH2 sessions
- One Web session

USING THE CONSOLE INTERFACE

The CLI built into the switch is accessible by way of the 9-pin, RS-232 port labeled *console*, located on the back of the Summit switch, or on the front of either of the BlackDiamond switch MSMs.



For more information on the console port pinouts, refer to the BlackDiamond Hardware Installation Guide or the Summit Hardware Installation Guide.

Once the connection is established, you will see the switch prompt and you may log in.

USING THE 10/100 UTP MANAGEMENT PORT

Some Extreme switch models provide a dedicated 10/100 UTP management port. This port provides dedicated remote access to the switch using TCP/IP. It supports the following management methods:

- Telnet using the CLI interface
- ExtremeWare Vista Web access using a standard Web browser
- SNMP access using ExtremeWare Enterprise Manager or another SNMP manager

The management port is a DTE port, and is not capable of supporting switching or routing functions. The TCP/IP configuration for the management port is done using the same syntax as used for VLAN configuration. The VLAN *mgmt* comes preconfigured with only the 10/100 UTP management port as a member.

You can configure the IP address, subnet mask, and default router for the VLAN *mgmt*, using the following commands:

```
config vlan mgmt ipaddress <ip_address>/<subnet_mask>
config iproute add default <gateway>
```

USING TELNET

Any workstation with a Telnet facility should be able to communicate with the switch over a TCP/IP network.

Up to eight active Telnet sessions can access the switch concurrently. If *idle timeouts* are enabled, the Telnet connection will time out after 20 minutes of inactivity. If a connection to a Telnet session is lost inadvertently, the switch terminates the session within two hours.

Before you can start a Telnet session, you must set up the IP parameters described in the section “[Configuring Switch IP Parameters](#),” later in this chapter. Telnet is enabled by default.

To open the Telnet session, you must specify the IP address of the device that you want to manage. Check the user manual supplied with the Telnet facility if you are unsure of how to do this.

Once the connection is established, you will see the switch prompt and you may log in.

CONNECTING TO ANOTHER HOST USING TELNET

You can Telnet from the current CLI session to another host using the following command:

```
telnet [<ipaddress> | <hostname>] {<port_number>}
```

If the TCP port number is not specified, the Telnet session defaults to port 23. Only VT100 emulation is supported.

CONFIGURING SWITCH IP PARAMETERS

To manage the switch by way of a Telnet connection or by using an SNMP Network Manager, you must first configure the switch IP parameters.

USING A BOOTP SERVER

If you are using IP and you have a Bootstrap Protocol (BOOTP) server set up correctly on your network, you must add the following information to the BOOTP server:

- Switch Media Access Control (MAC) address, found on the rear label of the switch
- IP address
- Subnet address mask (optional)

Once this is done, the IP address and subnet mask for the switch will be downloaded automatically. You can then start managing the switch without further configuration.

You can enable BOOTP on a per-VLAN basis by using the following command:

```
enable bootp vlan [<name> | all]
```

By default, BOOTP is enabled on the *default* VLAN.

If you configure the switch to use BOOTP, the switch IP address is not retained through a power cycle, even if the configuration has been saved. To retain the IP address through a power cycle, you must configure the IP address of the VLAN using the command-line interface, Telnet, or Web interface.

All VLANs within a switch that are configured to use BOOTP to get their IP address use the same MAC address. Therefore, if you are using BOOTP relay through a router, the BOOTP server must be capable of differentiating its relay based on the gateway portion of the BOOTP packet.

 For more information on DHCP/BOOTP relay, refer to [Chapter 11](#).

MANUALLY CONFIGURING THE IP SETTINGS

If you are using IP without a BOOTP server, you must enter the IP parameters for the switch in order for the SNMP Network Manager, Telnet software, or Web interface to communicate with the device. To assign IP parameters to the switch, you must do the following:

- Log in to the switch with administrator privileges.
- Assign an IP address and subnet mask to a VLAN.

The switch comes configured with a default VLAN named *default*. To use Telnet or an SNMP Network Manager, you must have at least one VLAN on the switch, and it

must be assigned an IP address and subnet mask. IP addresses are always assigned to a VLAN. The switch can be assigned multiple IP addresses.

 For information on creating and configuring VLANs, refer to [Chapter 6](#).

To manually configure the IP settings, perform the following steps:

- 1 Connect a terminal or workstation running terminal-emulation software to the console port.
- 2 At your terminal, press [Return] one or more times until you see the login prompt.
- 3 At the login prompt, enter your user name and password. Note that they are both case-sensitive. Ensure that you have entered a user name and password with administrator privileges.
 - If you are logging in for the first time, use the default user name *admin* to log in with administrator privileges. For example:

```
login: admin
```

Administrator capabilities enable you to access all switch functions. The default user names have no passwords assigned.

 - If you have been assigned a user name and password with administrator privileges, enter them at the login prompt.
- 4 At the password prompt, enter the password and press [Return].

When you have successfully logged in to the switch, the command-line prompt displays the name of the switch in its prompt.

- 5 Assign an IP address and subnetwork mask for the default VLAN by using the following command:

```
config vlan <name> ipaddress <ipaddress> {<subnet_mask>}
```

For example:

```
config vlan default ipaddress 123.45.67.8 255.255.255.0
```

Your changes take effect immediately.

 As a general rule, when configuring any IP addresses for the switch, you can express a subnet mask by using dotted decimal notation, or by using classless inter-domain routing notation (CIDR). CIDR uses a forward slash plus the number of bits in the subnet mask. Using CIDR notation, the command identical to the one above would be:

```
config vlan default ipaddress 123.45.67.8 / 24
```

- 6 Configure the default route for the switch using the following command:

```
config iproute add default <gateway> {<metric>}
```

For example:

```
config iproute add default 123.45.67.1
```

- 7 Save your configuration changes so that they will be in effect after the next switch reboot, by typing

```
save
```

- 8 When you are finished using the facility, log out of the switch by typing

```
logout or quit
```

DISCONNECTING A TELNET SESSION

An administrator-level account can disconnect a Telnet management session. If this happens, the user logged in by way of the Telnet connection is notified that the session has been terminated.

To terminate a Telnet session, follow these steps:

- 1 Log in to the switch with administrator privileges.
- 2 Determine the session number of the session you want to terminate by using the following command:

```
show session
```

- 3 Terminate the session by using the following command:

```
clear session <session_number>
```

CONTROLLING TELNET ACCESS

By default, Telnet services are enabled on the switch. Telnet access can be restricted by the use of an access profile. An access profile permits or denies a named list of IP addresses and subnet masks. To configure Telnet to use an access profile, use the following command:

```
enable telnet {access-profile [<access_profile> | none]} {port <tcp_port_number>}
```

Use the `none` option to remove a previously configured access profile.

To display the status of Telnet, use the following command:

```
show management
```

You can choose to disable Telnet by using the following command:

```
disable telnet
```

To re-enable Telnet on the switch, at the console port use the following command at the console port:

```
enable telnet
```

You must be logged in as an administrator to enable or disable Telnet.



For more information on Access Profiles, see [Chapter 16](#).

USING SECURE SHELL 2 (SSH2)

Secure Shell 2 (SSH2) is a feature of ExtremeWare that allows you to encrypt Telnet session data between the switch and a network administrator using SSH2 client software. The ExtremeWare SSH2 switch application is based on the Data FellowsTM SSH2 server implementation. It is highly recommended that you use the F-Secure[®] SSH client products from Data Fellows corporation. These applications are available for most operating systems. For more information, refer to the Data Fellows website at:

<http://www.datafellows.com>.



*SSH2 is compatible with the Data Fellows SSH2 client version 2.0.12 or above.
SSH2 is not compatible with SSH1.*

ENABLING SSH2

Because SSH2 is currently under U.S. export restrictions, before enabling SSH2, you must first obtain a security license from Extreme Networks. The procedure for obtaining a security license key is described in [Chapter 1](#).

To enable SSH2, use the following command:

```
enable ssh2 {access-profile [<access_profile> | none]} {port <tcp_port_number>}
```

An authentication key must be generated for each SSH2 session. This can be done automatically by the switch or by the client application. To have the key generated by the switch, use the following command:

```
config ssh2 key {pregenerated}
```

If you do not select automatic key generation, you are prompted to enter the key when you enable SSH2.

You can specify a list of pre-defined clients that are allowed SSH2 access to the switch. To do this, you must create an access profile that contains a list of allowed IP addresses. For more information on creating access profiles, refer to [Chapter 16](#).

You can also specify a TCP port number to be used for SSH2 communication. By default the TCP port number is 22.

The supported cipher is 3DES-CBC. The supported key exchange is DSA.

For additional information on the SSH protocol refer to [FIPS-186] Federal Information Processing Standards Publication (FIPSPUB) 186, Digital Signature Standard, 18 May 1994. This can be download from: <ftp://ftp.cs.hut.fi/pub/ssh>. General technical information is also available from <http://www.ssh.fi>.

After you obtain the SSH2 key value, copy the key to the SSH2 client application. Also, ensure that the client is configured for any non-default access list or TCP port information that you have configured on the switch. Once these tasks are accomplished, you may form an SSH2-encrypted session with the switch.

USING EXTREMEWARE VISTA

ExtremeWare Vista is device-management software running in the switch that enables you to access the switch over a TCP/IP network using a standard Web browser. Any properly configured standard Web browser that supports frames (such as Netscape Navigator 3.0 or above, or Microsoft Internet Explorer 3.0 or above) can manage the switch over a TCP/IP network.



For more information on assigning an IP address, refer to the section, "Configuring Switch IP Parameters," on page 3-3.

The default home page of the switch can be accessed using the following command:

```
http://<ipaddress>
```

When you access the home page of the switch, you are presented with the Logon screen.



For more information on using ExtremeWare Vista, refer to [Chapter 19](#).

CONTROLLING WEB ACCESS

By default, Web access is enabled on the switch. Use of ExtremeWare Vista Web access can be restricted through the use of an access profile. An access profile permits or denies a named list of IP addresses and subnet masks. To configure Vista Web access to use an access profile, use the following command:

```
enable web {access-profile <access-profile> | none} {port  
<tcp_port_number>}
```

Use the none option to remove a previously configured access profile.

To display the status of Web access, use the following command:

```
show management
```

To disable ExtremeWare Vista, use the following command:

```
disable web
```

To re-enable Web access, use the following command:

```
enable web {access-profile <access-profile> | none} {port  
<tcp_port_number>}
```

When you disable or enable ExtremeWare Vista, you must reboot the switch for the changes to take effect. Apply an access profile only when ExtremeWare Vista is enabled.

USING SNMP

Any Network Manager running the Simple Network Management Protocol (SNMP) can manage the switch, provided the Management Information Base (MIB) is installed correctly on the management station. Each Network Manager provides its own user interface to the management facilities.

The following sections describe how to get started if you want to use an SNMP manager. It assumes you are already familiar with SNMP management. If not, refer to the following publication:

The Simple Book
by Marshall T. Rose
ISBN 0-13-8121611-9
Published by Prentice Hall

ACCESSING SWITCH AGENTS

To have access to the SNMP agent residing in the switch, at least one VLAN must have an IP address assigned to it.

SUPPORTED MIBs

In addition to private MIBs, the switch supports the standard MIBs listed in [Appendix A](#).

CONFIGURING SNMP SETTINGS

The following SNMP parameters can be configured on the switch:

- **Authorized trap receivers** — An authorized trap receiver can be one or more network management stations on your network. The switch sends SNMP traps to all trap receivers. You can have a maximum of 16 trap receivers configured for each switch. Entries in this list can also be created, modified, and deleted using the RMON2 trapDestTable MIB variable, as described in RFC 2021.
- **SNMP read access** — The ability to read SNMP information can be restricted through the use of an access profile. An access profile permits or denies a named list of IP addresses and subnet masks.

To configure SNMP read access to use an access profile, use the command:

```
config snmp access-profile readonly [<access_profile> | none]
```

Use the none option to remove a previously configured access profile.

- **SNMP read/write access** — The ability to read and write SNMP information can be restricted through the use of an access profile. An access profile permits or denies a named list of IP addresses and subnet masks.

To configure SNMP read/write access to use an access profile, use the command:

```
config snmp access-profile readwrite [<access_profile> | none]
```

Use the `none` option to remove a previously configured access-profile.

- **Community strings** — The community strings allow a simple method of authentication between the switch and the remote Network Manager. There are two types of community strings on the switch. Read community strings provide read-only access to the switch. The default read-only community string is *public*. Read-write community strings provide read and write access to the switch. The default read-write community string is *private*. A total of eight community strings can be configured on the switch. The community string for all authorized trap receivers must be configured on the switch for the trap receiver to receive switch-generated traps. SNMP community strings can contain up to 127 characters.
- **System contact** (optional) — The system contact is a text field that enables you to enter the name of the person(s) responsible for managing the switch.
- **System name** — The system name is the name that you have assigned to this switch. The default name is the model name of the switch (for example, Summit1 switch).
- **System location** (optional) — Using the system location field, you can enter an optional location for this switch.

[Table 3-1](#) describes SNMP configuration commands.

Table 3-1: SNMP Configuration Commands

Command	Description
<code>config snmp access-profile readonly [<access_profile> none]</code>	Assigns an access profile that limits which stations have read-only access to the switch.
<code>config snmp access-profile readwrite [<access_profile> none]</code>	Assigns an access profile that limits which stations have read-write access to the switch.
<code>config snmp add trapreceiver <ipaddress> community <string></code>	Adds the IP address of a specified trap receiver. The IP address can be a unicast, multicast, or broadcast address. A maximum of 16 trap receivers is allowed.
<code>config snmp community [read-only read-write] <string></code>	Adds an SNMP read or read/write community string. The default read-only community string is <i>public</i> . The default read-write community string is <i>private</i> . Each community string can have a maximum of 127 characters, and can be enclosed by double quotation marks.

Table 3-1: SNMP Configuration Commands (continued)

Command	Description
config snmp delete trapreceiver [<ip_address> community <string> all]	Deletes the IP address of a specified trap receiver or all authorized trap receivers.
config snmp syscontact <string>	Configures the name of the system contact. A maximum of 255 characters is allowed.
config snmp syslocation <string>	Configures the location of the switch. A maximum of 255 characters is allowed.
config snmp sysname <string>	Configures the name of the switch. A maximum of 32 characters is allowed. The default sysname is the model name of the device (for example, Summit1). The sysname appears in the switch prompt.
disable snmp access	Disables SNMP on the switch. Disabling SNMP access does not affect the SNMP configuration (for example, community strings).
disable snmp traps	Prevents SNMP traps from being sent from the switch. Does not clear the SNMP trap receivers that have been configured.
enable snmp access	Turns on SNMP support for the switch.
enable snmp traps	Turns on SNMP trap support.
unconfig management	Restores default values to all SNMP-related entries.

DISPLAYING SNMP SETTINGS

To display the SNMP settings configured on the switch, use the following command:

```
show management
```

This command displays the following information:

- Enable/disable state for Telnet, SSH2, SNMP, and Web access, along with access profile information
- SNMP community strings
- Authorized SNMP station list
- SNMP trap receiver list
- RMON polling configuration
- Login statistics

AUTHENTICATING USERS

ExtremeWare provides two methods to authenticate users who login to the switch:

- Radius client
- TACACS+

RADIUS CLIENT

Remote Authentication Dial In User Service (RADIUS, RFC 2138) is a mechanism for authenticating and centrally administrating access to network nodes. The ExtremeWare RADIUS client implementation allows authentication for Telnet, Vista, or console access to the switch.



You cannot configure RADUIS and TACACS+ at the same time.

You can define a primary and secondary RADIUS server for the switch to contact. When a user attempts to login using Telnet, http, or the console, the request is relayed to the primary RADIUS server, and then to the secondary RADIUS server, if the primary does not respond. If the RADIUS client is enabled, but access to the RADIUS primary and secondary server fails, the switch uses its local database for authentication.

The privileges assigned to the user (admin versus non-admin) at the RADIUS server take precedence over the configuration in the local switch database.

PER-COMMAND AUTHENTICATION USING RADIUS

The RADIUS implementation can be used to perform per-command authentication. Per-command authentication allows you to define several levels of user capabilities by controlling the permitted command sets based on the RADIUS username and password. You do not need to configure any additional switch parameters to take advantage of this capability. The RADIUS server implementation automatically negotiates the per-command authentication capability with the switch. For examples on per-command RADIUS configurations, refer to the next section.

CONFIGURING RADIUS CLIENT

You can define primary and secondary server communication information, and for each RADIUS server, the RADIUS port number to use when talking to the RADIUS server. The default port value is 1645. The client IP address is the IP address used by the RADIUS server for communicating back to the switch.

RADIUS commands are described in [Table 3-2](#).

Table 3-2: RADIUS Commands

Command	Description
<pre>config radius [primary secondary] server [<ipaddress> <hostname>] {<udp_port>} client-ip <ipaddress></pre>	<p>Configures the primary and secondary RADIUS server. Specify the following:</p> <ul style="list-style-type: none"> ■ [primary secondary] — Configure either the primary or secondary RADIUS server. ■ [<ipaddress> <hostname>] — The IP address or hostname of the server being configured. ■ <udp_port> — The UDP port to use to contact the RADUIS server. The default UDP port setting is 1645. ■ client-ip <ipaddress> — The IP address used by the switch to identify itself when communicating with the RADIUS server. <p>The RADIUS server defined by this command is used for user name authentication and CLI command authentication.</p>
<pre>config radius [primary secondary] shared-secret <string></pre> <pre>config radius-accounting [primary secondary] server [<ipaddress> <hostname>] {<udp_port>} client-ip <ipaddress></pre>	<p>Configures the authentication string used to communicate with the RADIUS server.</p> <p>Configures the RADIUS accounting server. Specify the following:</p> <ul style="list-style-type: none"> ■ [primary secondary] — Configure either the primary or secondary RADIUS server. ■ [<ipaddress> <hostname>] — The IP address or hostname of the server being configured. ■ <udp_port> — The UDP port to use to contact the RADUIS server. The default UDP port setting is 1646. ■ client-ip <ipaddress> — The IP address used by the switch to identify itself when communicating with the RADIUS server.
	<p>The accounting server and the RADIUS authentication server can be the same.</p>

Table 3-2: RADIUS Commands (continued)

Command	Description
config radius-accounting [primary secondary] shared-secret <string>	Configures the authentication string used to communicate with the RADIUS accounting server.
disable radius	Disables the RADIUS client.
disable radius-accounting	Disables RADIUS accounting.
enable radius	Enables the RADIUS client. When enabled, all Web and CLI logins are sent to the RADIUS servers for authentication. When used with a RADIUS server that supports ExtremeWare CLI authorization, each CLI command is sent to the RADIUS server for authentication before it is executed.
enable radius-accounting	Enables RADIUS accounting. The RADIUS client must also be enabled.
show radius	Displays the current RADIUS and RADIUS accounting client configuration and statistics.

RADIUS RFC 2138 ATTRIBUTES

The RADIUS RFC 2138 optional attributes supported are as follows:

- User-Name
- User-Password
- Service-Type
- Login-IP-Host

RADIUS SERVER CONFIGURATION EXAMPLE (MERIT)

Many implementations of RADIUS server use the publicly available “Merit® AAA” server application, available on the World Wide Web at:

<http://www.merit.edu/aaa>

Included below are excerpts from relevant portions of a sample Merit RADIUS server implementation. The example shows excerpts from the client and user configuration

files. The client configuration file (ClientCfg.txt) defines the authorized source machine, source name, and access level. The user configuration file (users) defines username, password, and service type information.

ClientCfg.txt

#Client Name	Key	[type]	[version]	[prefix]
#-----	-----	-----	-----	-----
#10.1.2.3:256	test	type = nas	v2	pfx
#pm1	%^\$%#*(&!(*&)+	type=nas		pm1.
#pm2	:)-:-(;^):-}!	type nas		pm2.
#merit.edu/homeless	hmoemreilte.ses			
#homeless	testing	type proxy	v1	
#xyz.merit.edu	moretesting	type=Ascend:NAS	v1	
#anyoldthing:1234	whoknows?	type=NAS+RAD RFC+ACCT RFC		
10.202.1.3	andrew-linux	type=nas		
10.203.1.41	eric	type=nas		
10.203.1.42	eric	type=nas		
10.0.52.14	samf	type=nas		

users

```

user      Password = ""
        Filter-Id = "unlim"
admin    Password = "", Service-Type = Administrative
        Filter-Id = "unlim"

eric     Password = "", Service-Type = Administrative
        Filter-Id = "unlim"

albert   Password = "password", Service-Type = Administrative
        Filter-Id = "unlim"

samuel   Password = "password", Service-Type = Administrative
        Filter-Id = "unlim"

```

RADIUS PER-COMMAND CONFIGURATION EXAMPLE

Building on the example configuration above, you can use RADIUS to perform per-command authentication to differentiate user capabilities. To do so, use the Extreme-modified RADIUS Merit software that is available from the Extreme Networks web server at <http://www.extremenetworks.com/extreme/support/otherapps.htm> or by contacting Extreme Networks technical support. The software is available in

compiled format for Solaris™ or Linux™ operating systems, as well as in source code format. For all clients that use RADIUS per-command authentication, you must add the following type to the client file:

```
type:extreme:nas + RAD_RFC + ACCT_RFC
```

Within the `users` configuration file, additional keywords are available for `Profile-Name` and `Extreme-CLI-Authorization`. To use per-command authentication, enable the CLI authorization function and indicate a profile name for that user. If authorization is enabled without specifying a valid profile, the user is unable to perform any commands.

Next, define the desired profiles in an ASCII configuration file called `profiles`. This file contains named profiles of exact or partial strings of CLI commands. A named profile is linked with a user through the `users` file. A profile with the `permit` on keywords allows use of only the listed commands. A profile with the `deny` keyword allows use of all commands *except* the listed commands.

CLI commands may be defined easily in a hierachal manner by using an asterisk (*) to indicate any possible subsequent entry. The parser performs exact string matches on other text to validate commands. Commands are separated by a comma (,) or newline.

Looking at the following example content in `profiles` for the profile named `PROFILE1`, which uses the `deny` keyword, the following attributes are associated with the user of this profile:

- Cannot use any command starting with `enable`.
- Cannot issue the `disable ipforwarding` command.
- Cannot issue a `show switch` command.
- Can perform all other commands.

We know from the `users` file that this applies to the users `albert` and `lulu`. We also know that `eric` is able to log in, but is unable to perform any commands, because he has no valid profile assigned.

In `PROFILE2`, a user associated with this profile can use any `enable` command, the `clear counter` command and the `show management` command, but can perform no other functions on the switch. We also know from the `users` file that `gerald` has these capabilities.

The following lists the contents of the file `users` with support for per-command authentication:

```

user      Password = ""
          Filter-Id = "unlim"

admin    Password = "", Service-Type = Administrative
          Filter-Id = "unlim"

eric     Password = "", Service-Type = Administrative, Profile-Name = ""
          Filter-Id = "unlim"
          Extreme:Extreme-CLI-Authorization = Enabled

albert   Password = "", Service-Type = Administrative, Profile-Name =
"Profile1"
          Filter-Id = "unlim"
          Extreme:Extreme-CLI-Authorization = Enabled

lulu     Password = "", Service-Type = Administrative, Profile-Name =
"Profile1"
          Filter-Id = "unlim"
          Extreme:Extreme-CLI-Authorization = Enabled

gerald   Password = "", Service-Type = Administrative, Profile-Name
"Profile2"
          Filter-Id = "unlim"
          Extreme:Extreme-CLI-Authorization = Enabled

```

Contents of the file "profiles":

```

PROFILE1 deny
{
enable  *, disable ipforwarding
show switch
}

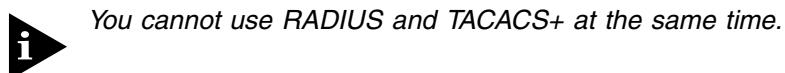
PROFILE2
{
enable  *, clear counters
show management
}

PROFILE3 deny
{
create vlan *, configure iproute *, disable *, show fdb
delete *, configure rip add
}

```

CONFIGURING TACACS+

Terminal Access Controller Access Control System Plus (TACACS+) is a mechanism for providing authentication, authorization, and accounting on a centralized server, similar in function to the RADIUS client. The ExtremeWare version of TACACS+ is used to authenticate prospective users who are attempting to administer the switch. TACACS+ is used to communicate between the switch and an authentication database.



You can configure two TACACS+ servers, specifying the primary server address, secondary server address, and UDP port number to be used for TACACS+ sessions.

Table 3-3 describes the commands that are used to configure TACACS+.

Table 3-3: TACACS+ Commands

Command	Description
config tacacs [primary secondary] server [<ipaddress> <hostname>] {<udp_port>} client-ip <ipaddress>	Configure the server information for a TACACS+ server. Specify the following: <ul style="list-style-type: none"> ■ primary secondary — Specifies primary or secondary server configuration. To remove a server, use the address 0.0.0.0. ■ <ipaddress> <hostname> — Specifies the TACACS+ server. ■ <udp_port> — Optionally specifies the UDP port to be used. ■ client-ip — Specifies the IP address used by the switch to identify itself when communicating with the TACACS+ server.
config tacacs [primary secondary] shared-secret {encrypted} <string>	Configures the shared secret string used to communicate with the TACACS+ server.
config tacacs-accounting [primary secondary] server [<ipaddress> <hostname>] {<udp_port>} client-ip <ipaddress>	Configures the TACACS+ accounting server. You can use the same server for accounting and authentication.
config tacacs-accounting [primary secondary] shared-secret {encrypted} <string>	Configures the shared secret string used to communicate with the TACACS+ accounting server.
disable tacacs	Disables TACACS+.

Table 3-3: TACACS+ Commands (continued)

Command	Description
disable tacacs-accounting	Disables TACACS+ accounting.
disable tacacs-authorization	Disables CLI command authorization.
enable tacacs	Enables TACACS+. Once enabled, all WEB and CLI logins are sent to one of the two TACACS+ server for login name authentication and accounting.
enable tacacs-accounting	Enables TACACS+ accounting. If accounting is use, the TACACS+ client must also be enabled.
enable tacacs-authorization	Enables CLI command authorization. When enabled, each command is transmitted to the remote TACACS+ server for authorization before the command is executed.
show tacacs	Displays the current TACACS+ configuration and statistics.
show tacacs-accounting	Displays the current TACACS+ accounting client configuration and statistics.
unconfig tacacs {server [primary secondary]}	Unconfigures the TACACS+ client configuration.
unconfig tacacs-accounting {server [primary secondary]}	Unconfigures the TACACS+ accounting client configuration.

USING THE SIMPLE NETWORK TIME PROTOCOL

ExtremeWare supports the client portion of the Simple Network Time Protocol (SNTP) Version 3 based on RFC1769. SNTP can be used by the switch to update and synchronize its internal clock from a Network Time Protocol (NTP) server. When enabled, the switch sends out a periodic query to the indicated NTP server, or the switch listens to broadcast NTP updates. In addition, the switch supports the configured setting for Greenwich Mean time (GMT) offset and the use of Daylight Savings Time. These features have been tested for year 2000 compliance.

CONFIGURING AND USING SNTP

To use SNTP, follow these steps:

- 1 Identify the host(s) that are configured as NTP server(s). Additionally, identify the preferred method for obtaining NTP updates. The options are for the NTP server to send out broadcasts, or for switches using NTP to query the NTP server(s) directly. A combination of both methods is possible. You must identify the method that should be used for the switch being configured.
- 2 Configure the Greenwich Mean Time (GMT) offset and Daylight Savings Time preference. The command syntax to configure GMT offset and usage of Daylight Savings is as follows:

```
config timezone <GMT_offset> {autodst | noautodst}
```

The `GMT_OFFSET` is in +/- minutes from the GMT time. Automatic Daylight Savings Time (DST) changes can be enabled or disabled. The default setting is enabled.

- 3 Enable the SNTP client using the following command:

```
enable sntp-client
```

Once enabled, the switch sends out a periodic query to the NTP servers defined later (if configured) or listens to broadcast NTP updates from the network. The network time information is automatically saved into the on-board real-time clock.

- 4 If you would like this switch to use a directed query to the NTP server, configure the switch to use the NTP server(s). If the switch listens to NTP broadcasts, skip this step. To configure the switch to use a directed query, use the following command:

```
config sntp-client [primary | secondary] server [<ip_address> | <hostname>]
```

NTP queries are first sent to the primary server. If the primary server does not respond within 1 second, or if it is not synchronized, the switch queries the secondary server (if one is configured). If the switch cannot obtain the time, it restarts the query process. Otherwise, the switch waits for the `sntp-client update interval` before querying again.

- 5 Optionally, the interval for which the SNTP client updates the real-time clock of the switch can be changed using the following command:

```
config sntp-client update-interval <seconds>
```

The default `sntp-client update-interval` value is 64 seconds.

- 6 You can verify the configuration using the following commands:

- `show sntp-client`

This command provides configuration and statistics associated with SNTP and its connectivity to the NTP server.

- `show switch`

This command indicates the GMT offset, Daylight Savings Time, and the current local time.

NTP updates are distributed using GMT time. To properly display the local time in logs and other timestamp information, the switch should be configured with the appropriate offset to GMT based on geographical location. [Table 3-4](#) describes GMT offsets.

Table 3-4: Greenwich Mean Time Offsets

GMT Offset in Hours	GMT Offset in Minutes	Common Time Zone References	Cities
+0:00	+0	GMT - Greenwich Mean UT or UTC - Universal (Coordinated) WET - Western European	London, England; Dublin, Ireland; Edinburgh, Scotland; Lisbon, Portugal; Reykjavik, Iceland; Casablanca, Morocco
-1:00	-60	WAT - West Africa	Azores, Cape Verde Islands
-2:00	-120	AT - Azores	
-3:00	-180		Brasilia, Brazil; Buenos Aires, Argentina; Georgetown, Guyana;
-4:00	-240	AST - Atlantic Standard	Caracas; La Paz
-5:00	-300	EST - Eastern Standard	Bogota, Columbia; Lima, Peru; New York, NY, Trevor City, MI USA
-6:00	-360	CST - Central Standard	Mexico City, Mexico
-7:00	-420	MST - Mountain Standard	Saskatchewan, Canada
-8:00	-480	PST - Pacific Standard	Los Angeles, CA, Cupertino, CA, Seattle, WA USA
-9:00	-540	YST - Yukon Standard	
-10:00	-600	AHST - Alaska-Hawaii Standard CAT - Central Alaska HST - Hawaii Standard	
-11:00	-660	NT - Nome	

Table 3-4: Greenwich Mean Time Offsets (continued)

GMT Offset in Hours	GMT Offset in Minutes	Common Time Zone References	Cities
-12:00	-720	IDLW - International Date Line West	
+1:00	+60	CET - Central European FWT - French Winter MET - Middle European MEWT - Middle European Winter SWT - Swedish Winter	Paris, France; Berlin, Germany; Amsterdam, The Netherlands; Brussels, Belgium; Vienna, Austria; Madrid, Spain; Rome, Italy; Bern, Switzerland; Stockholm, Sweden; Oslo, Norway
+2:00	+120	EET - Eastern European, Russia Zone 1	Athens, Greece; Helsinki, Finland; Istanbul, Turkey; Jerusalem, Israel; Harare, Zimbabwe
+3:00	+180	BT - Baghdad, Russia Zone 2	Kuwait; Nairobi, Kenya; Riyadh, Saudi Arabia; Moscow, Russia; Tehran, Iran
+4:00	+240	ZP4 - Russia Zone 3	Abu Dhabi, UAE; Muscat; Tblisi; Volgograd; Kabul
+5:00	+300	ZP5 - Russia Zone 4	
+5:30	+330	IST – India Standard Time	New Delhi, Pune, Allahabad, India
+6:00	+360	ZP6 - Russia Zone 5	
+7:00	+420	WAST - West Australian Standard	
+8:00	+480	CCT - China Coast, Russia Zone 7	
+9:00	+540	JST - Japan Standard, Russia Zone 8	
+10:00	+600	EAST - East Australian Standard GST - Guam Standard Russia Zone 9	
+11:00	+660		
+12:00	+720	IDLE - International Date Line East NZST - New Zealand Standard NZT - New Zealand	Wellington, New Zealand; Fiji, Marshall Islands

SNTP CONFIGURATION COMMANDS

[Table 3-5](#) describes SNTP configuration commands.

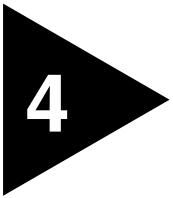
Table 3-5: SNTP Configuration Commands

Command	Description
config sntp-client [primary secondary] server [<ipaddress> <host_name>]	Configures an NTP server for the switch to obtain time information. Queries are first sent to the primary server. If the primary server does not respond within 1 second, or if it is not synchronized, the switch queries the second server.
config sntp-client update-interval <seconds>	Configures the interval between polling for time information from SNTP servers. The default setting is 64 seconds.
disable sntp-client	Disables SNTP client functions.
enable sntp-client	Enables Simple Network Time Protocol (SNTP) client functions.
show sntp-client	Displays configuration and statistics for the SNTP client.

SNTP EXAMPLE

In this example, the switch queries a specific NTP server and a backup NTP server. The switch is located in Cupertino, CA, and an update occurs every 20 minutes. The commands to configure the switch are as follows:

```
config timezone -480 autodst
config sntp-client update interval 1200
enable sntp-client
config sntp-client primary server 10.0.0.1.1
config sntp-client secondary server 10.0.0.1.2
```

4

Configuring BlackDiamond and Alpine Switch Slots and Ports

This chapter covers the following topics:

- [Configuring a Slot on page 4-1](#)
- [BlackDiamond and Alpine Switch Port Configuration on page 4-2](#)
- [Jumbo Frames on page 4-7](#)
- [Load Sharing on the BlackDiamond and Alpine Switch on page 4-7](#)
- [BlackDiamond and Alpine Switch Port-Mirroring on page 4-11](#)

For information on configuring ports on the Summit switch, refer to [Chapter 5](#).

CONFIGURING A SLOT

If a slot has not been configured for a particular type of I/O module, then any type of module is accepted in that slot, and a default port and VLAN configuration is automatically generated.

Once any port on the module is configured (for example, a VLAN association, a VLAN tag configuration, or port parameters), all the port information and the module type for that slot must be saved to non-volatile storage. Otherwise, if the BlackDiamond or Alpine switch is rebooted or the module is removed from the slot, the port, VLAN, and module configuration information is not saved.



For information on saving the configuration, refer to [Chapter 20](#).

You can configure the BlackDiamond or Alpine switch with the type of I/O module that is installed in each I/O slot. To do this, use the following command:

```
config slot <slot> module [f32t | f32f | f48t | g4x | g6x | g8x | g12x]
```

You can also pre-configure the slot before inserting the module card. This allows you to begin configuring the module and ports before installing the card in the chassis.

If a slot is configured for one type of module, and a different type of module is inserted, the inserted module is put into a mismatch state, and is not brought online. To use the new module type in a slot, the slot configuration must be cleared or configured for the new module type. To clear the slot of a previously assigned I/O module type, use the following command:

```
clear slot <slot>
```

All configuration information related to the slot and the ports on the module is erased. If a module is present when you issue this command, the module is reset to default settings.

To display information about a particular slot, use the following command:

```
show slot {<slot>}
```

Information displayed includes the following:

- Card type, serial number, part number
- Current state (power down, operational, diagnostic, mismatch)
- Port information

If no slot is specified, information for all slots is displayed.

BLACKDIAMOND AND ALPINE SWITCH PORT CONFIGURATION

On the BlackDiamond or Alpine switch, the port number is a combination of the slot number and the port number. The nomenclature for the port number is as follows:

slot:port

For example, if a G4X I/O module (having a total of four ports) is installed in slot 2 of the BlackDiamond 6808 chassis, the following ports are valid:

- 2:1
- 2:2
- 2:3
- 2:4

You can also use wildcard combinations (*) to specify multiple BlackDiamond or Alpine slot and port combinations. The following wildcard combinations are allowed:

- slot:* — Specifies all ports on a particular I/O module.
- slot:x-slot:y — Specifies a contiguous series of ports on a particular I/O module.
- slota:x-slotb:y — Specifies a contiguous series of ports that begin on one I/O module and end on another I/O module.

ENABLING AND DISABLING BLACKDIAMOND AND ALPINE SWITCH PORTS

By default, all ports are enabled. To enable or disable one or more ports, use the following command:

```
[enable | disable] ports <portlist>
```

For example, to disable slot 7, ports 3, 5, and 12 through 15 on the BlackDiamond switch, enter the following:

```
disable port 7:3,7:5,7:12-7:15
```

Even though a port is disabled, the link remains enabled for diagnostic purposes.

CONFIGURING BLACKDIAMOND AND ALPINE SWITCH PORT SPEED AND DUPLEX SETTING

By default, the switch is configured to use autonegotiation to determine the port speed and duplex setting for each port. You can select to manually configure the duplex setting and the speed of 10/100 Mbps ports, and you can manually configure the duplex setting on Gigabit Ethernet ports.

10BASE-T and 100BASE-TX ports can connect to either 10BASE-T or 100BASE-T networks. By default, the ports autonegotiate port speed. You can also configure each port for a particular speed (either 10 Mbps or 100 Mbps).

Gigabit Ethernet ports are statically set to 1 Gbps, and their speed cannot be modified.

To configure port speed and duplex setting, use the following command:

```
config ports <portlist> auto off {speed [10 | 100 | 1000]} duplex [half | full]
```

To configure the system to autonegotiate, use the following command:

```
config ports <portlist> auto on
```

Flow control is supported only on Gigabit Ethernet ports. It is enabled or disabled as part of autonegotiation.

TURNING OFF AUTONEGOTIATION FOR A GIGABIT ETHERNET PORT

In certain interoperability situations, you may need to turn autonegotiation off on a Gigabit Ethernet port. Even though a Gigabit Ethernet port runs only at full duplex, you must specify the duplex setting.

The following example turns autonegotiation off for port 1 on a G4X or G6X module located in slot 1:

```
config ports 1:1 auto off duplex full
```

BLACKDIAMOND AND ALPINE SWITCH PORT COMMANDS

[Table 4-1](#) describes the BlackDiamond and Alpine switch port commands.

Table 4-1: BlackDiamond and Alpine Switch Port Commands

Command	Description
clear slot <slot>	Clears a slot of a previously assigned module type.
config jumbo-frame size <jumbo_frame_mtu>	Configures the jumbo frame size. The range is between 1523 and 9216, including 4 bytes of CRC. The default setting is 9216.

Table 4-1: BlackDiamond and Alpine Switch Port Commands (continued)

Command	Description
config ports <portlist> auto off {speed [10 100 1000]} duplex [half full]	Changes the configuration of a group of ports. Specify the following: <ul style="list-style-type: none">■ auto off — The port will not autonegotiate the settings.■ speed — The speed of the port (for 10/100 Mbps or 100/1000 Mbps ports only).■ duplex — The duplex setting (half- or full-duplex).
config ports <portlist> auto on	Enables autonegotiation for the particular port type; 802.3u for 10/100 Mbps ports or 802.3z for Gigabit Ethernet ports.
config ports <portlist> display-string <string>	Configures a user-defined string for a port. The string is displayed in certain <code>show</code> commands (for example, <code>show port all info</code>). The string can be up to 16 characters.
config ports <portlist> qosprofile <qosname>	Configures one or more ports to use a particular QoS profile.
config slot <slot> module [f32t f32f f48t g4x g6x g8x g12x]	Configures a slot for a particular I/O module card.
disable jumbo-frame ports [<portlist> all]	Disables jumbo frame support on a port.
disable learning ports <portlist>	Disables MAC address learning on one or more ports for security purposes. If MAC address learning is disabled, only broadcast traffic, Extreme Discovery Protocol (EDP) traffic, and packets destined to a permanent MAC address matching that port number are forwarded to the port. The default setting is enabled.
disable ports <portlist>	Disables a port. Even when disabled, the link is available for diagnostic purposes.
disable sharing <port>	Disables a load-sharing group of ports.
enable jumbo-frame ports [<portlist> all]	Enables reception and transmission of jumbo frames. A jumbo frame is dropped if it is received on a port with jumbo frames disabled, or if the jumbo frame needs to be forwarded out of a port that has jumbo frames disabled.
enable learning ports <portlist>	Enables MAC address learning on one or more ports. The default setting is enabled.
enable ports <portlist>	Enables a port.

Table 4-1: BlackDiamond and Alpine Switch Port Commands (continued)

Command	Description
enable sharing <port> grouping <portlist> {port-based address-based round-robin}	<p>Defines a load-sharing group of ports. The ports specified in <portlist> are grouped to the master port. Optional load-sharing algorithms include:</p> <ul style="list-style-type: none"> ■ port-based — Uses the ingress port as criteria for egress port selection. ■ address-based — Uses addressing information as criteria for egress port selection. ■ round-robin — Forwards packets to all egress ports in a round-robin fashion. <p>If not specified, port-based load-sharing is used.</p>
restart ports <portlist>	Resets autonegotiation for one or more ports by resetting the physical link.
show ports {<portlist>} collisions	Displays real-time collision statistics.
show ports {<portlist>} configuration	Displays the port configuration.
show ports {<portlist>} info {detail}	Displays detailed system-related information.
show ports {<portlist>} packet	Displays a histogram of packet statistics.
show ports {<portlist>} qosmonitor	Displays real-time QoS statistics.
show ports {<portlist>} rxerrors	Displays real-time receive error statistics.
show ports {<portlist>} stats	Displays real-time port statistics.
show ports {<portlist>} txerrors	Displays real-time transmit error statistics.
show ports {<portlist>} utilization	Displays real-time port utilization information. Use the [Spacebar] to toggle between packet, byte, and bandwidth utilization information.
show slot {<slot>}	<p>Displays slot-specific information, including the following:</p> <ul style="list-style-type: none"> ■ Card type, serial number, part number ■ Current state ■ Port information <p>If not slot is specified, information for all slots is displayed.</p>
unconfig ports <portlist> display-string <string>	Clears the user-defined display string from a port.
unconfig slot <slot>	Clears a slot of a previously assigned module type.

JUMBO FRAMES

Jumbo frames are Ethernet frames that are larger than 1523 bytes, including four bytes used for the cyclic redundancy check (CRC). Extreme products that use the “i” chipset support switching and routing of jumbo frames at wire-speed on all ports.

Jumbo frames are used between endstations that support larger frame sizes for more efficient transfers of bulk data. Both endstations involved in the transfer must be capable of supporting jumbo frames. The switch does not perform IP fragmentation, or participate in MTU negotiation on behalf of devices that do not support jumbo frames.

ENABLING JUMBO FRAMES

To enable jumbo frame support, you must configure the maximum MTU size of a jumbo frame that will be allowed by the switch. To set the maximum MTU size, use the following command

```
config jumbo-frame size <jumbo_frame_mtu>
```

The `jumbo_frame_mtu` range is 1523 to 9216. The value describes the maximum size “on the wire,” and includes 4 bytes of CRC plus another 4 bytes if 802.1Q tagging is being used.

Next, enable support on the physical ports that will carry jumbo frames, using the following command:

```
enable jumbo-frame ports [<portlist | all>]
```



Some network interface cards (NICs) have a configured maximum MTU size that does not include the additional 4 bytes of CRC. Ensure that the NIC maximum MTU size is at or below the maximum MTU size configured on the switch.

Frames that are larger than the MTU size configured on the switch are dropped at the ingress port.

LOAD SHARING ON THE BLACKDIAMOND AND ALPINE SWITCH

Load sharing with BlackDiamond or Alpine switches allows you to increase bandwidth and resiliency by using a group of ports to carry traffic in parallel between switches. The sharing algorithm allows the switch to use multiple ports as a single logical port.

For example, VLANs see the load-sharing group as a single logical port. Most load-sharing algorithms guarantee packet sequencing between clients.

If a port in a load-sharing group fails, traffic is redistributed to the remaining ports in the load-sharing group. If the failed port becomes active again, traffic is redistributed to include that port.

-  *Load sharing must be enabled on both ends of the link, or a network loop may result. The load-sharing algorithms do not need to be the same on both ends.*

This feature is supported between Extreme Networks switches only, but may be compatible with third-party “trunking” or link-aggregation algorithms. Check with an Extreme Networks technical representative for more information.

LOAD-SHARING ALGORITHMS

Load sharing algorithms allow you to select the distribution technique used by the load-sharing group to determine the output port selection. Algorithm selection is not intended for use in predictive traffic engineering. You can configure one of three load-sharing algorithms on the switch, as follows:

- Port-based — Uses the ingress port to determine which physical port in the load-sharing group is used to forward traffic out of the switch.
- Address-based — Uses addressing information to determine which physical port in the load-sharing group to use for forwarding traffic out of the switch. Addressing information is based on the packet protocol, as follows:
 - IP packets — Uses the source and destination MAC and IP addresses, and the TCP port number.
 - IPX packets — Uses the source and destination MAC address, and IPX network identifiers.
 - All other packets — Uses the source and destination MAC address.
- Round-robin — When the switch receives a stream of packets, it forwards one packet out of each physical port in the load-sharing group using a round-robin scheme.

-  *Using the round-robin algorithm, packet sequencing between clients is not guaranteed.*

If you do not explicitly select an algorithm, the port-based scheme is used. However, the address-based algorithm has a more even distribution and is the recommended choice.

The address-based and round-robin load-sharing algorithms are supported by BlackDiamond switch modules that use the “*i*” chipset and all Alpine 3800 switch modules. The modules end with an “*i*” in their model designation (for example, G12SX*i*), and require the use of the MSM64*i*. For more information on “*i*” chipset products, refer to [Chapter 1](#).

CONFIGURING BLACKDIAMOND AND ALPINE SWITCH LOAD SHARING

To set up the BlackDiamond or Alpine switch to load share among ports, you must create a load-sharing group of ports. Load-sharing groups are defined according to the following rules:

- The first port in the load-sharing group is configured to be the “master” logical port. This is the reference port used in configuration commands. It can be thought of as the logical port representing the entire port group.

On I/O modules that do not have the “*i*” chipset, the following additional rules apply:

- Ports in a load-sharing group must be contiguous.
- Ports on the I/O module are divided into groups of two or four.

Follow the outlined boxes in [Table 4-2](#), [Table 4-3](#), and [Table 4-4](#) to determine the valid port combinations.

[Table 4-2](#), [Table 4-3](#), and [Table 4-4](#) show the possible load-sharing port group combinations for the G4X module, the G6X module, and the F32T and F32F modules, respectively.

Table 4-2: Port Combinations for the G4X Module

Load-Sharing Group	1	2	3	4
4-port groups	x	x	x	x
2-port groups	x	x	x	x

Table 4-3: Port Combinations for the G6X Module

Load-Sharing Group	1	2	3	4	5	6
4-port groups		x	x	x	x	
2-port groups	x	x	x	x	x	x

Table 4-4: Port Combinations for the F32T and F32F Modules

Load-Sharing Group	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6
4-port groups	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x
2-port groups	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x

Load-Sharing Group	1	1	1	2	2	2	2	2	2	2	2	2	2	3	3	3
4-port groups	7	8	9	0	1	2	3	4	5	6	7	8	9	0	1	2
2-port groups	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x

BlackDiamond switch modules that use the “*i*” chipset do not require contiguous ports in the load-sharing group. On “*i*” chipset BlackDiamond switch modules, the following rules apply:

- One group can contain up to 8 ports.
- The ports in the group must be on the same I/O module.
- The ports in the group do not need to be contiguous.

To define a load-sharing group, you assign a group of ports to a single, logical port number. To enable or disable a load-sharing group, use the following commands:

```
enable sharing <port> grouping <portlist> {port-based | address-based | round-robin}
disable sharing <port>
```

LOAD-SHARING EXAMPLE

The following example defines a load-sharing group on slot 3 that contains ports 9 through 12, and uses the first port in the group as the master logical port 9:

```
enable sharing 3:9 grouping 3:9-3:12
```

In this example, logical port 3:9 represents physical ports 3:9 through 3:12.

When using load sharing, you should always reference the master logical port of the load-sharing group (port 3:9 in the previous example) when configuring or viewing VLANs. VLANs configured to use other ports in the load-sharing group will have those ports deleted from the VLAN when load sharing becomes enabled.



Do not disable a port that is part of a load-sharing group. Disabling the port prevents it from forwarding traffic, but still allows the link to initialize. As a result, a partner switch does not receive a valid indication that the port is not in a forwarding state, and the partner switch will continue to forward packets.

VERIFYING THE LOAD-SHARING CONFIGURATION

The screen output resulting from the `show ports configuration` command lists the ports that are involved in load sharing and the master logical port identity.

BLACKDIAMOND AND ALPINE SWITCH PORT-MIRRORING

Port-mirroring configures the BlackDiamond or Alpine switch to copy all traffic associated with one or more ports. The monitor port can be connected to a network analyzer or RMON probe for packet analysis. The system uses a traffic filter that copies a group of traffic to the monitor port. The traffic filter can be defined based on one of the following criteria:

- **Physical port** — All data that traverses the port, regardless of VLAN configuration, is copied to the monitor port.
- **VLAN** — All data to and from a particular VLAN, regardless of the physical port configuration, is copied to the monitor port.
- **Virtual port** — All data specific to a VLAN on a specific port is copied to the monitor port.

Up to eight mirroring filters and one monitor port can be configured. Once a port is specified as a monitor port, it cannot be used for any other function.

-  *Frames that contain errors are not mirrored.*
-  *On switches that do not support the "i" chipset, mirrored frames that are transmitted from the switch do not contain 802.1Q VLAN tagging information.*

PORt-MIRRORING COMMANDS

BlackDiamond and Alpine switch port-mirroring commands are described in [Table 4-5](#).

Table 4-5: BlackDiamond and Alpine Switch Port-Mirroring Configuration Commands

Command	Description
config mirroring add [vlan <name> port <port> vlan <name> port <portlist>]	Adds a single mirroring filter definition. Up to eight mirroring definitions can be added. You can mirror traffic from a VLAN, a physical port, or a specific VLAN/port combination.
config mirroring delete [vlan <name> port <slot:port> vlan <name> port <portlist> all]	Deletes a particular mirroring filter definition, or all mirroring filter definitions.
disable mirroring	Disables port-mirroring.
enable mirroring to <port>	Dedicates a port to be the mirror output port.
show mirroring	Displays the port-mirroring configuration.

BLACKDIAMOND SWITCH PORT-MIRRORING EXAMPLE

The following example selects port 3 on slot 7 as the mirror port, and sends all traffic coming into or out of the BlackDiamond switch on slot 7, port 1 to the mirror port:

```
enable mirroring port 7:3
config mirroring add port 7:1
```

The following example sends all traffic coming into or out of the system on slot 8, port 1 and the VLAN *default* to the mirror port:

```
enable mirroring port 8:4
config mirroring add port 8:1 vlan default
```

EXTREME DISCOVERY PROTOCOL

The Extreme Discovery Protocol (EDP) is used to gather information about neighbor Extreme Networks switches. EDP is used by the switches to exchange topology information. EDP is also used by the Extreme Standby Router Protocol (ESRP), described in [Chapter 10](#). Information communicated using EDP includes the following:

- Switch MAC address (switch ID)
- Switch software version information
- Switch IP address
- Switch VLAN-IP information
- Switch port number

EDP COMMANDS

[Table 4-6](#) lists EDP commands.

Table 4-6: EDP Commands

Command	Description
disable edp ports <portlist>	Disables the EDP on one or more ports.
enable edp ports <portlist>	Enables the generation and processing of EDP messages on one or more ports. The default setting is enabled.
show edp	Displays EDP information.



5

Configuring Summit Switch Ports

This chapter covers the following topics:

- [Enabling and Disabling Summit Switch Ports on page 5-1](#)
- [Configuring Summit Switch Port Speed and Duplex Setting on page 5-2](#)
- [Summit Switch Port Commands on page 5-3](#)
- [Jumbo Frames on page 5-5](#)
- [Load Sharing on the Summit Switch on page 5-6](#)
- [Summit Switch Port-Mirroring on page 5-10](#)
- [Smart Redundancy on page 5-13](#)

For information on how to configure ports on the BlackDiamond or Alpine switch, refer to [Chapter 4](#).

ENABLING AND DISABLING SUMMIT SWITCH PORTS

By default, all ports are enabled. To enable or disable one or more ports, use the following command:

```
[enable | disable] ports <portlist>
```

For example, to disable ports 3, 5, and 12 through 15 on the Summit switch, enter the following:

```
disable port 3,5,12-15
```

Even though a port is disabled, the link remains enabled for diagnostic purposes.

CONFIGURING SUMMIT SWITCH PORT SPEED AND DUPLEX SETTING

By default, the Summit switch is configured to use autonegotiation to determine the port speed and duplex setting for each port. You can select to manually configure the duplex setting and the speed of 10/100 Mbps ports, and you can manually configure the duplex setting on Gigabit Ethernet ports.

Fast Ethernet ports can connect to either 10Base-T or 100Base-T networks. By default, the ports autonegotiate port speed. You can also configure each port for a particular speed (either 10 Mbps or 100 Mbps).

Gigabit Ethernet ports are statically set to 1 Gbps, and their speed cannot be modified.

All ports on the Summit switch can be configured for half-duplex or full-duplex operation. By default, the ports autonegotiate the duplex setting.

To configure port speed and duplex setting, use the following command:

```
config ports <portlist> auto off {speed [10 | 100 | 1000]} duplex [half | full]
```

To configure the switch to autonegotiate, use the following command:

```
config ports <portlist> auto on
```

Flow control is supported only on Gigabit Ethernet ports. It is enabled or disabled as part of autonegotiation. If autonegotiation is set to off, flow control is disabled. When autonegotiation is turned on, flow control is enabled.

TURNING OFF AUTONEGOTIATION FOR A GIGABIT ETHERNET PORT

In certain interoperability situations, it is necessary to turn autonegotiation off on a Gigabit Ethernet port. Even though a Gigabit Ethernet port runs only at full duplex and gigabit speeds, the command that turns off autonegotiation must still include the duplex setting.

The following example turns autonegotiation off for port 4 (a Gigabit Ethernet port):

```
config ports 4 auto off duplex full
```

SUMMIT SWITCH PORT COMMANDS

[Table 5-1](#) describes the Summit switch port commands.

Table 5-1: Summit Switch Port Commands

Command	Description
config jumbo-frame size <jumbo_frame_mtu>	Configures the jumbo frame size. The range is between 1523 and 9216. The default setting is 9216.
config ports <portlist> auto off {speed [10 100 1000]} duplex [half full]	<p>Changes the configuration of a group of ports. Specify the following:</p> <ul style="list-style-type: none"> ■ auto off — The port will not autonegotiate the settings. ■ speed — The speed of the port (for 10/100 Mbps or 100/1000 Mbps ports only). ■ duplex — The duplex setting (half- or full-duplex).
config ports <portlist> auto on	Enables autonegotiation for the particular port type; 802.3u for 10/100 Mbps ports or 802.3z for Gigabit Ethernet ports.
config ports <portlist> display-string <string>	Configures a user-defined string for a port. The string is displayed in certain <code>show</code> commands (for example, <code>show port all info</code>). The string can be up to 16 characters.
config ports <portlist> qosprofile <qosname>	Configures one or more ports to use a particular QoS profile.
disable jumbo-frame ports [<portlist> all]	Disables jumbo frame support on a port.

Table 5-1: Summit Switch Port Commands (continued)

Command	Description
disable learning ports <portlist>	Disables MAC address learning on one or more ports for security purposes. If MAC address learning is disabled, only broadcast traffic, EDP traffic, and packets destined to a permanent MAC address matching that port number, are forwarded to the port. The default setting is enabled.
disable ports <portlist>	Disables a port. Even when disabled, the link is available for diagnostic purposes.
disable sharing <port>	Disables a load-sharing group of ports.
disable smartredundancy <portlist>	Disables the smart redundancy feature. If the feature is disabled, the switch changes to the active link only when the current active link becomes inoperable.
enable jumbo-frame ports [<portlist> all]	Enables reception and transmission of jumbo frames. A jumbo frame is dropped if it is received on a port with jumbo frames disabled, or if the jumbo frame needs to be forwarded out of a port that has jumbo frames disabled.
enable learning ports <portlist>	Enables MAC address learning on one or more ports. The default setting is enabled.
enable ports <portlist>	Enables a port.
enable sharing <port> grouping <portlist> {port-based address-based round-robin}	Defines a load-sharing group of ports. The ports specified in <portlist> are grouped to the master port. Optional load-sharing algorithms include: <ul style="list-style-type: none"> ■ port-based — Uses the ingress port as criteria for egress port selection. ■ address-based — Uses addressing information as criteria for egress port selection. ■ round-robin — Forwards packets to all egress ports in a round-robin fashion. If not specified, port-based load-sharing is used.
enable smartredundancy <portlist>	Enables the smart redundancy feature on the redundant Gigabit Ethernet port. When the Smart Redundancy feature is enabled, the switch always uses the primary link when the primary link is available. The default setting is enabled.

Table 5-1: Summit Switch Port Commands (continued)

Command	Description
restart ports <portlist>	Resets autonegotiation for one or more ports by resetting the physical link.
show ports {<portlist>} collisions	Displays real-time collision statistics.
show ports {<portlist>} configuration	Displays the port configuration.
show ports {<portlist>} info {detail}	Displays detailed system-related information.
show ports {<portlist>} packet	Displays a histogram of packet statistics.
show ports {<portlist>} qosmonitor	Displays real-time QoS statistics.
show ports {<portlist>} rxerrors	Displays real-time receive error statistics.
show ports {<portlist>} stats	Displays real-time port statistics.
show ports {<portlist>} txerrors	Displays real-time transmit error statistics.
show ports {<portlist>} utilization	Displays real-time port utilization information. Use the [Spacebar] to toggle between packet, byte, and bandwidth utilization information.
unconfig ports <portlist> display-string <string>	Clears the user-defined display string from a port.

JUMBO FRAMES

Jumbo frames are Ethernet frames that are larger than 1523 bytes, including four bytes used for the cyclic redundancy check (CRC). Extreme products that use the “i” chipset support switching and routing of jumbo frames at wire-speed on all ports.

Jumbo frames are used between endstations that support larger frame sizes for more efficient transfers of bulk data. Both endstations involved in the transfer must be capable of supporting jumbo frames. The switch does not perform IP fragmentation, or participate in MTU negotiation on behalf of devices that do not support jumbo frames.

ENABLING JUMBO FRAMES

To enable jumbo frame support, you must configure the maximum MTU size of a jumbo frame that will be allowed by the switch. To set the maximum MTU size, use the following command

```
config jumbo-frame size <jumbo_frame_mtu>
```

The jumbo_frame_mtu range is 1523 to 9216. The value describes the maximum size “on the wire,” and includes 4 bytes of CRC plus another 4 bytes if 802.1Q tagging is being used.

Next, enable support on the physical ports that will carry jumbo frames, using the following command:

```
enable jumbo-frame ports [<portlist | all]
```

 *Some network interface cards (NICs) have a configured maximum MTU size that does not include the additional 4 bytes of CRC. Ensure that the NIC maximum MTU size is at or below the maximum MTU size configured on the switch. Frames that are larger than the MTU size configured on the switch are dropped at the ingress port.*

LOAD SHARING ON THE SUMMIT SWITCH

Load sharing with Summit switches allows you to increase bandwidth and resilience between switches by using a group of ports to carry traffic in parallel between switches. The sharing algorithm allows the switch to use multiple ports as a single logical port. For example, VLANs see the load-sharing group as a single logical port. Most load-sharing algorithms also guarantee packet sequencing between clients.

If a port in a load-sharing group fails, traffic is redistributed to the remaining ports in the load-sharing group. If the failed port becomes active again, traffic is redistributed to include that port.

 *Load sharing must be enabled on both ends of the link, or a network loop will result.*

This feature is supported between Extreme Networks switches only, but may be compatible with third-party “trunking” or sharing algorithms. Check with an Extreme Networks technical representative for more information.

LOAD SHARING ALGORITHMS

Load sharing algorithms allow you to select the distribution technique used by the load-sharing group to determine the output port selection. Algorithm selection is not intended for use in predictive traffic engineering.

You can configure one of three load-sharing algorithms on the switch, as follows:

- Port-based — Uses the ingress port to determine which physical port in the load-sharing group is used to forward traffic out of the switch.
- Address-based — Uses addressing information to determine which physical port in the load-sharing group to use for forwarding traffic out of the switch. Addressing information is based on the packet protocol, as follows:
 - IP packets — Uses the source and destination MAC and IP addresses, and the TCP port number.
 - IPX packets — Uses the source and destination MAC address, and IPX network identifiers.
 - All other packets — Uses the source and destination MAC address.
- Round-robin — When the switch receives a stream of packets, it forwards one packet out of each physical port in the load-sharing group using a round-robin scheme.



Using the round-robin algorithm, packet sequencing between clients is not guaranteed.

If you do not explicitly select an algorithm, the port-based scheme is used. However, the address-based algorithm has a more even distribution and is the recommended choice.

The address-based and round-robin load sharing algorithms are supported by Summit switch products that use the “*i*” chipset, such as the Summit7i switch. For more information on the “*i*” chipset products, refer to [Chapter 1](#).

CONFIGURING SUMMIT SWITCH LOAD SHARING

To set up the Summit switch to load share among ports, you must create a load-sharing group of ports. The first port in the load-sharing group is configured to be the “master” logical port. This is the reference port used in configuration commands. It can be thought of as the logical port representing the entire port group.

On switches that do not have the “*i*” chipset, the following additional rules apply:

- Ports in a load-sharing group must be contiguous.
- Ports on the switch are divided into groups of two or four.
- Address-based and round-robin load sharing algorithms do not apply.

Follow the outlined boxes in [Table 5-2](#) through [Table 5-7](#) to determine the valid port combinations.

[Table 5-2](#), [Table 5-3](#), [Table 5-4](#), [Table 5-5](#), [Table 5-6](#), and [Table 5-7](#) show the possible load-sharing port group combinations for the Summit1, Summit2, Summit3, Summit24, Summit4 and Summit4/FX, and Summit48 switches, respectively.

Table 5-2: Port Combinations for the Summit1 Switch

Load-Sharing Group	1	2	3	4	5	6	7	8
4-port groups				x	x	x	x	
2-port groups	*	x	x	x	x	x	x	*

* In addition, ports 1 and 8 can be combined into a two-port load sharing group on the Summit1.

Table 5-3: Port Combinations for the Summit2 Switch

Load-Sharing Group	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18
4-port groups	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	
2-port groups	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x

Table 5-4: Port Combinations for the Summit3 Switch

Load-Sharing Group	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5
4-port groups	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x
2-port groups	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x

Table 5-5: Port Combinations for the Summit4 Switch and Summit4/FX Switch

Load-Sharing Group	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7	8	9	1	2	2	2
4-port groups	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x
2-port groups	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x

Table 5-6: Port Combinations for the Summit24 Switch

Load-Sharing Group	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7	8	9	0	1	2	2	2	2
4-port groups	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x
2-port groups	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x

Table 5-7: Port Combinations for the Summit48 Switch

Load-Sharing Group	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7	8	9	0	1	2	2	2	2
4-port groups	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x
2-port groups	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x

Load-Sharing Group	2	2	2	2	2	3	3	3	3	3	3	3	3	3	3	3	3	3	4	4	4	4	4	4	4	4
	5	6	7	8	9	0	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7	8	9	
4-port groups	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	
2-port groups	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	

Load-Sharing Group	4	5
	9	0
4-port groups		
2-port groups	x	x

On all other Summit switch models, the following rules apply:

- A group can contain up to 8 ports.
- The ports in a group do not need to be contiguous.

To define a load-sharing group, you assign a group of ports to a single, logical port number. To enable or disable a load-sharing group, use the following commands:

```
enable sharing <port> grouping <portlist> {port-based | address-based |  
round-robin}  
disable sharing <port>
```

LOAD-SHARING EXAMPLE

The following example defines a load-sharing group that contains ports 9 through 12, and uses the first port in the group as the master logical port 9:

```
enable sharing 9 grouping 9-12
```

In this example, logical port 9 represents physical ports 9 through 12.

When using load sharing, you should always reference the master logical port of the load-sharing group (port 9 in the previous example) when configuring or viewing VLANs. VLANs configured to use other ports in the load-sharing group will have those ports deleted from the VLAN when load sharing becomes enabled.



Do not disable a port that is part of a load-sharing group. Disabling the port prevents it from forwarding traffic, but still allows the link to initialize. As a result, a partner switch does not receive a valid indication that the port is not in a forwarding state, and the partner switch will continue to forward packets.

VERIFYING THE LOAD SHARING CONFIGURATION

The screen output resulting from the `show ports configuration` command indicates the ports are involved in load sharing and the master logical port identity.

SUMMIT SWITCH PORT-MIRRORING

Port-mirroring configures the switch to copy all traffic associated with one or more ports to a monitor port on the switch. The monitor port can be connected to a network analyzer or RMON probe for packet analysis. The switch uses a traffic filter that copies a group of traffic to the monitor port.

The traffic filter can be defined based on one of the following criteria:

- **Physical port** — All data that traverses the port, regardless of VLAN configuration, is copied to the monitor port.
- **VLAN** — All data to and from a particular VLAN, regardless of the physical port configuration, is copied to the monitor port.
- **Virtual port** — All data specific to a VLAN on a specific port is copied to the monitor port.

Up to eight mirroring filters and one monitor port can be configured on the switch. Once a port is specified as a monitor port, it cannot be used for any other function.



Frames that contain errors are not mirrored.



On switches that do not support the "i" chipset, mirrored frames that are transmitted from the switch do not contain 802.1Q VLAN tagging information.

PORT-MIRRORING COMMANDS

Summit switch port-mirroring commands are described in [Table 5-8](#).

Table 5-8: Summit Switch Port-Mirroring Configuration Commands

Command	Description
config mirroring add [vlan <name> port <port> vlan <name> port <port>]	Adds a single mirroring filter definition. Up to eight mirroring definitions can be added. You can mirror traffic from a VLAN, a physical port, or a specific VLAN/port combination.
config mirroring delete [vlan <name> port <port> vlan <name> port <port> all]	Deletes a particular mirroring filter definition, or all mirroring filter definitions.
disable mirroring	Disables port-mirroring.
enable mirroring to <port>	Dedicates a port to be the mirror output port.
show mirroring	Displays the port-mirroring configuration.

SUMMIT SWITCH PORT-MIRRORING EXAMPLE

The following example selects port 3 as the mirror port, and sends all traffic coming into or out of the switch on port 1 to the mirror port:

```
enable mirroring to port 3
config mirroring add port 1
```

The following example sends all traffic coming into or out of the switch on port 1 and the VLAN *default* to the mirror port:

```
config mirroring add port 1 vlan default
```

EXTREME DISCOVERY PROTOCOL

The Extreme Discovery Protocol (EDP) is used to gather information about neighbor Extreme Networks switches. EDP is used by the switches to exchange topology information. EDP is also used by the Extreme Standby Router Protocol (ESRP), described in [Chapter 10](#). Information communicated using EDP includes the following:

- Switch MAC address (switch ID)
- Switch software version information
- Switch IP address
- Switch VLAN-IP information
- Switch port number

EDP COMMANDS

[Table 5-9](#) lists EDP commands.

Table 5-9: EDP Commands

Command	Description
disable edp ports <portlist>	Disables the EDP on one or more ports.
enable edp ports <portlist>	Enables the generation and processing of EDP messages on one or more ports. The default setting is enabled.
show edp	Displays EDP information.

SMART REDUNDANCY

Smart redundancy defines the behavior of switches equipped with redundancy Gigabit Ethernet ports (for example, the Summit 24 and Summit48). When the switch becomes operational, it attempts to establish connectivity on the primary link. If this fails, the redundancy port is attempted. When connectivity is established (or re-established) on the primary link, the primary link is used.

On the Summit48, which has two redundant Gigabit Ethernet ports, an additional capability is provided when both ports are configured for load-sharing. Using this configuration, the same failover actions apply simultaneously to both failover ports. For example, if the primary connection on port 49 fails, the redundant port on both port 49 and 50 are activated simultaneously. When connectivity is re-established on the primary links of both ports 49 and 50, the primary links are used.

[Figure 5-1](#) shows a load-shared link with smart redundancy failover.

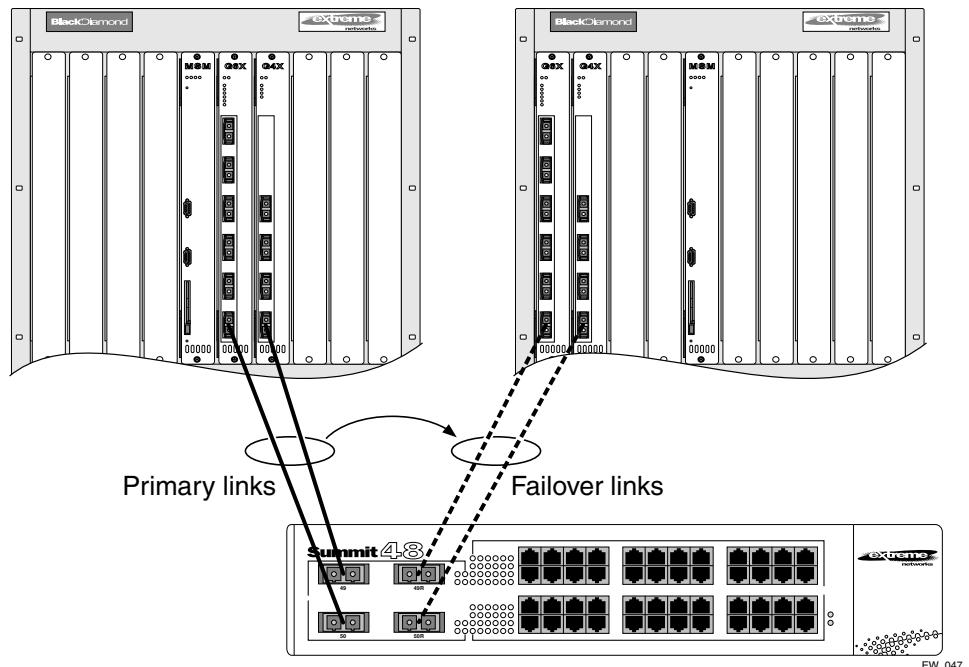
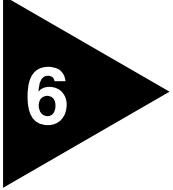


Figure 5-1: Load-shared link with smart redundancy failover

If smart redundancy is disabled, both the primary and redundant ports are dual-homed to active Gigabit Ethernet ports. It is not possible to predict which port will become active, and the first port to initialize becomes the primary. Enabling smart redundancy allows you to predict port failover and fail-back behavior.



6

Virtual LANs (VLANs)

This chapter covers the following topics:

- [Overview of Virtual LANs on page 6-1](#)
- [Types of VLANs on page 6-2](#)
- [VLAN Names on page 6-12](#)
- [Configuring VLANs on the Switch on page 6-13](#)
- [Displaying VLAN Settings on page 6-16](#)
- [VLAN Tunneling \(vMANs\) on page 6-17](#)
- [Generic VLAN Registration Protocol on page 6-18](#)
- [MAC-Based VLANs on page 6-21](#)

Setting up Virtual Local Area Networks (VLANs) on the switch eases many time-consuming tasks of network administration while increasing efficiency in network operations.

OVERVIEW OF VIRTUAL LANS

The term “VLAN” is used to refer to a collection of devices that communicate as if they were on the same physical LAN. Any set of ports (including all ports on the switch) is considered a VLAN. LAN segments are not restricted by the hardware that physically connects them. The segments are defined by flexible user groups you create with the command-line interface.

BENEFITS

Implementing VLANs on your networks has the following advantages:

- **VLANs help to control traffic.**

With traditional networks, congestion can be caused by broadcast traffic that is directed to all network devices, regardless of whether they require it. VLANs increase the efficiency of your network because each VLAN can be set up to contain only those devices that must communicate with each other.

- **VLANs provide extra security.**

Devices within each VLAN can only communicate with member devices in the same VLAN. If a device in VLAN *Marketing* must communicate with devices in VLAN *Sales*, the traffic must cross a routing device.

- **VLANs ease the change and movement of devices.**

With traditional networks, network administrators spend much of their time dealing with moves and changes. If users move to a different subnetwork, the addresses of each endstation must be updated manually.

TYPES OF VLANs

VLANs can be created according to the following criteria:

- Physical port
- 802.1Q tag
- Ethernet, LLC SAP, or LLC/SNAP Ethernet protocol type
- MAC address
- A combination of these criteria

PORT-BASED VLANs

In a port-based VLAN, a VLAN name is given to a group of one or more ports on the switch. A port can be a member of only one port-based VLAN.

For example, on the Summit7i switch in [Figure 6-1](#), ports 9 through 14 are part of VLAN *Marketing*; ports 25 through 29 are part of VLAN *Sales*; and ports 21 through 24 and 30 through 32 are in VLAN *Finance*.

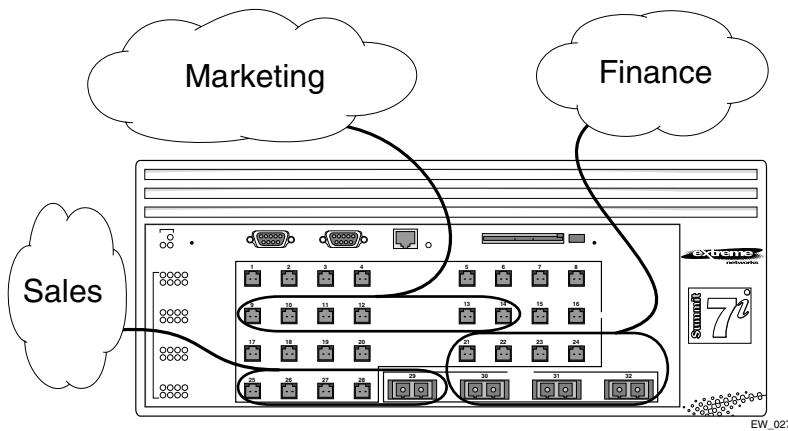


Figure 6-1: Example of a port-based VLAN on the Summit7i switch

For the members of the different IP VLANs to communicate, the traffic must be routed by the switch, even if they are physically part of the same I/O module. This means that each VLAN must be configured as a router interface with a unique IP address.

SPANNING SWITCHES WITH PORT-BASED VLANs

To create a port-based VLAN that spans two switches, you must do two things:

- 1 Assign the port on each switch to the VLAN.
- 2 Cable the two switches together using one port on each switch per VLAN.

[Figure 6-2](#) illustrates a single VLAN that spans a BlackDiamond switch and a Summit7i switch. All ports on the BlackDiamond switch belong to VLAN *Sales*. Ports 1 through 29 on the Summit 7i switch also belong to VLAN *Sales*. The two switches are connected using slot 8, port 4 on System 1 (the BlackDiamond switch), and port 29 on System 2 (the Summit7i switch).

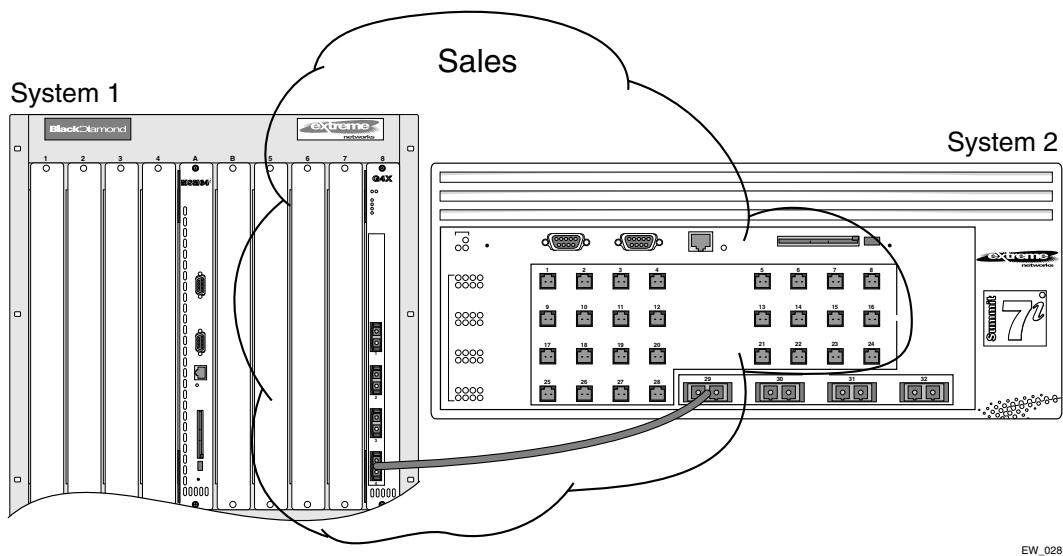


Figure 6-2: Single port-based VLAN spanning two switches

To create multiple VLANs that span two switches in a port-based VLAN, a port on System 1 must be cabled to a port on System 2 for each VLAN you want to have span across the switches. At least one port on each switch must be a member of the corresponding VLANs, as well.

[Figure 6-3](#) illustrates two VLANs spanning two switches. On System 1, ports 25 through 29 are part of VLAN *Accounting*; ports 21 through 24 and ports 30 through 32 are part of VLAN *Engineering*. On System 2, all ports on slot 3 are part of VLAN *Accounting*; all ports on slot 7 are part of VLAN *Engineering*.

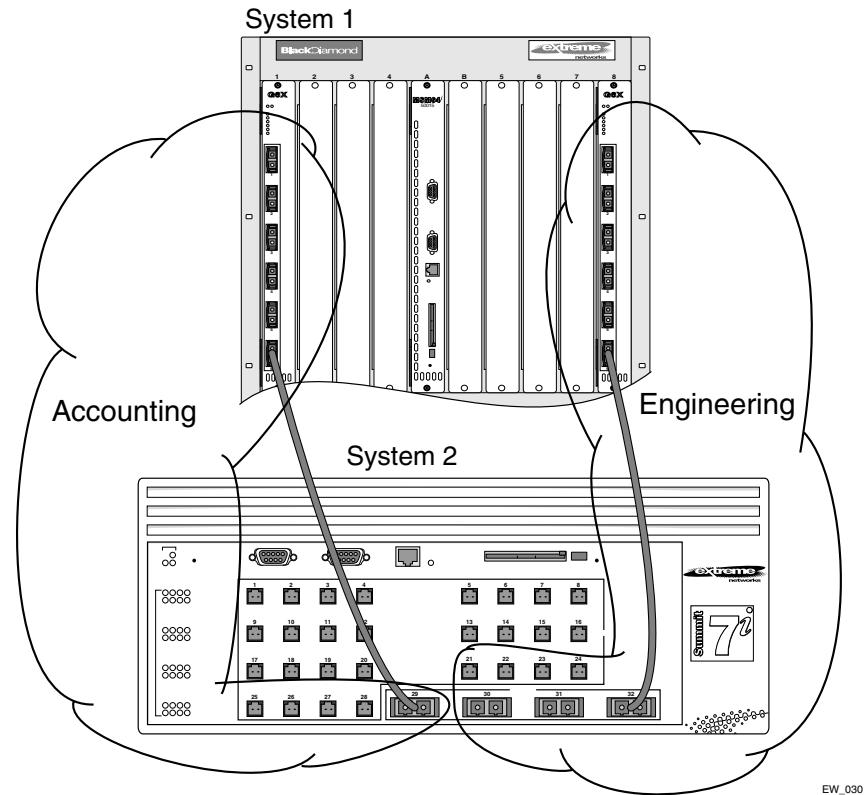


Figure 6-3: Two port-based VLANs spanning two BlackDiamond switches

VLAN *Accounting* spans System 1 and System 2 by way of a connection between System 1, port 29 and System 2, slot 3, port 1. VLAN *Engineering* spans System 1 and System 2 by way of a connection between System 1, port 32, and System 2, slot 7, port 1.

Using this configuration, you can create multiple VLANs that span multiple switches, in a daisy-chained fashion. Each switch must have a dedicated port for each VLAN. Each dedicated port must be connected to a port that is a member of its VLAN on the next switch.

TAGGED VLANS

Tagging is a process that inserts a marker (called a *tag*) into the Ethernet frame. The tag contains the identification number of a specific VLAN, called the *VLANid*.

 *The use of 802.1Q tagged packets may lead to the appearance of packets slightly bigger than the current IEEE 802.3/Ethernet maximum of 1,518 bytes. This may affect packet error counters in other devices, and may also lead to connectivity problems if non-802.1Q bridges or routers are placed in the path.*

USES OF TAGGED VLANs

Tagging is most commonly used to create VLANs that span switches. The switch-to-switch connections are typically called *trunks*. Using tags, multiple VLANs can span multiple switches using one or more trunks. In a port-based VLAN, each VLAN requires its own pair of trunk ports, as shown in [Figure 6-3](#). Using tags, multiple VLANs can span two switches with a single trunk.

Another benefit of tagged VLANs is the ability to have a port be a member of multiple VLANs. This is particularly useful if you have a device (such as a server) that must belong to multiple VLANs. The device must have a NIC that supports 802.1Q tagging.

A single port can be a member of only one port-based VLAN. All additional VLAN membership for the port must be accompanied by tags. In addition to configuring the VLAN tag for the port, the server must have a *Network Interface Card (NIC)* that supports 802.1Q tagging.

ASSIGNING A VLAN TAG

Each VLAN may be assigned an 802.1Q VLAN tag. As ports are added to a VLAN with an 802.1Q tag defined, you decide whether each port will use tagging for that VLAN. The default mode of the switch is to have all ports assigned to the VLAN named *default* with an 802.1Q VLAN tag (VLANid) of 1 assigned.

Not all ports in the VLAN must be tagged. As traffic from a port is forwarded out of the switch, the switch determines (in real time) if each destination port should use tagged or untagged packet formats for that VLAN. The switch adds and strips tags, as required, by the port configuration for that VLAN.

 *Packets arriving tagged with a VLANid that is not configured on a port will be discarded.*

Figure 6-4 illustrates the physical view of a network that uses tagged and untagged traffic.

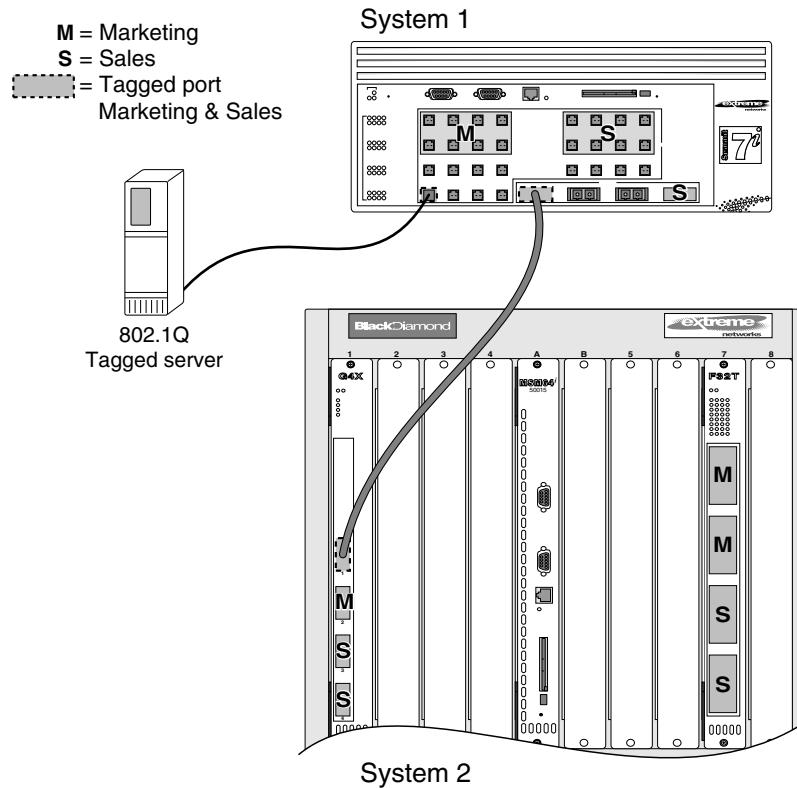


Figure 6-4: Physical diagram of tagged and untagged traffic

[Figure 6-5](#) shows a logical diagram of the same network.

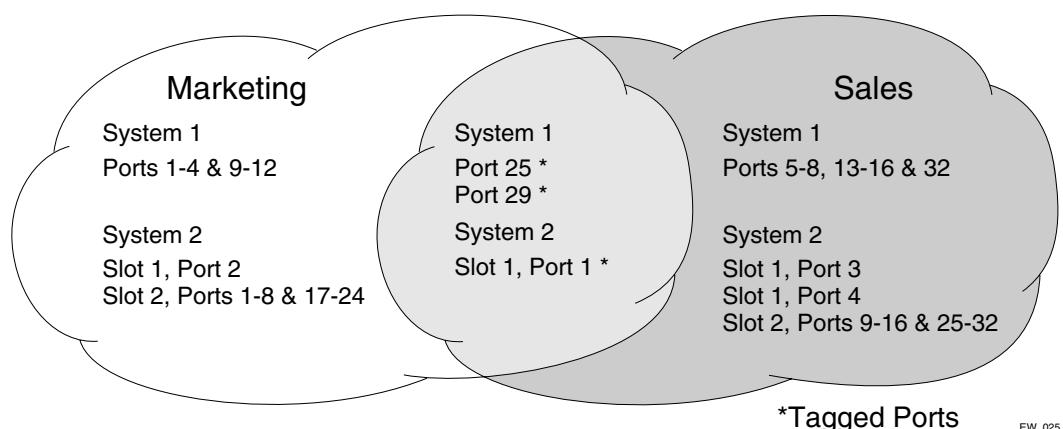


Figure 6-5: Logical diagram of tagged and untagged traffic

In [Figure 6-4](#) and [Figure 6-5](#):

- The trunk port on each switch carries traffic for both VLAN *Marketing* and VLAN *Sales*.
- The trunk port on each switch is tagged.
- The server connected to port 25 on System 1 has a NIC that supports 802.1Q tagging.
- The server connected to port 25 on System 1 is a member of both VLAN *Marketing* and VLAN *Sales*.
- All other stations use untagged traffic.

As data passes out of the switch, the switch determines if the destination port requires the frames to be tagged or untagged. All traffic coming from and going to the server is tagged. Traffic coming from and going to the trunk ports is tagged. The traffic that comes from and goes to the other stations on this network is not tagged.

MIXING PORT-BASED AND TAGGED VLANS

You can configure the switch using a combination of port-based and tagged VLANs. A given port can be a member of multiple VLANs, with the stipulation that only one of its VLANs uses untagged traffic. In other words, a port can simultaneously be a member of one port-based VLAN and multiple tag-based VLANs.



For the purposes of VLAN classification, packets arriving on a port with an 802.1Q tag containing a VLANid of zero are treated as untagged.

PROTOCOL-BASED VLANS

Protocol-based VLANs enable you to define a packet filter that the switch uses as the matching criteria to determine if a particular packet belongs to a particular VLAN.

Protocol-based VLANs are most often used in situations where network segments contain hosts running multiple protocols. For example, in [Figure 6-6](#), the hosts are running both the IP and NetBIOS protocols.

The IP traffic has been divided into two IP subnets, 192.207.35.0 and 192.207.36.0. The subnets are internally routed by the switch. The subnets are assigned different VLAN names, *Finance* and *Personnel*, respectively. The remainder of the traffic belongs to the VLAN named *MyCompany*. All ports are members of the VLAN *MyCompany*.

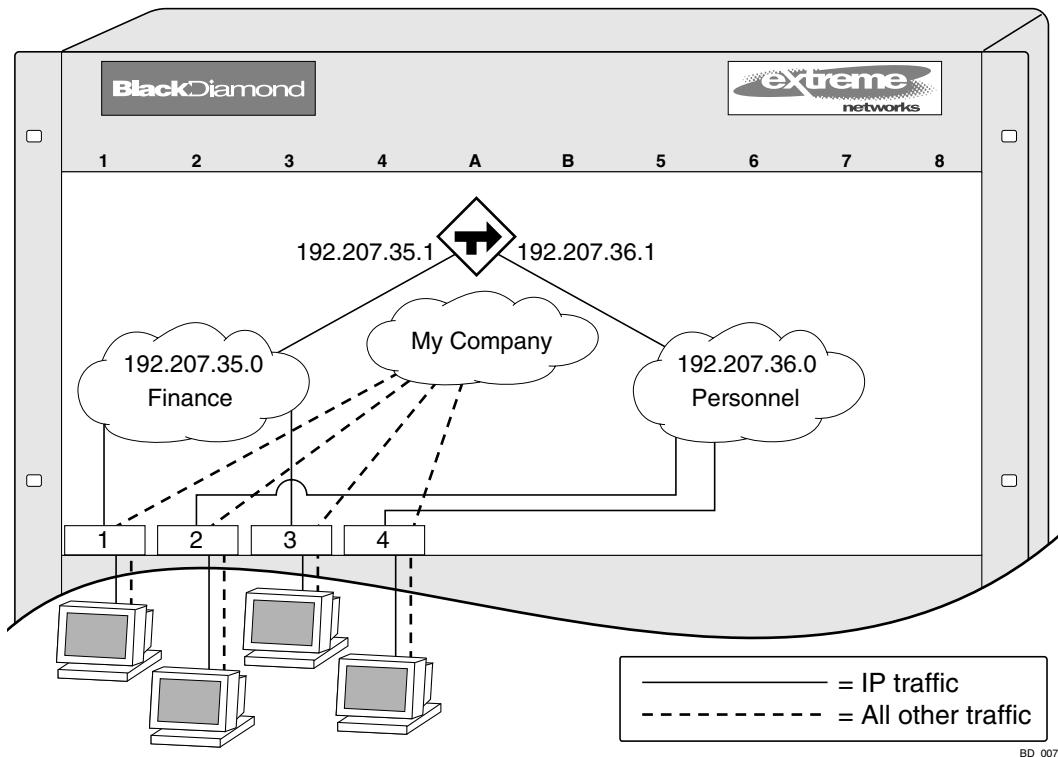


Figure 6-6: Protocol-based VLANs

PREDEFINED PROTOCOL FILTERS

The following protocol filters are predefined on the switch:

- IP
- IPX
- NetBIOS
- DECNet
- IPX_8022
- IPX_SNAP
- AppleTalk

DEFINING PROTOCOL FILTERS

If necessary, you can define a customized protocol filter based on EtherType, Logical Link Control (LLC), and/or Subnetwork Access Protocol (SNAP). Up to six protocols may be part of a protocol filter. To define a protocol filter, do the following:

- 1 Create a protocol using the following command:

```
create protocol <protocol_name>
```

For example:

```
create protocol fred
```

The protocol name can have a maximum of 32 characters.

- 2 Configure the protocol using the following command:

```
config protocol <protocol_name> add <protocol_type> <hex_value>
```

Supported protocol types include:

- **etype** — EtherType

The values for **etype** are four-digit hexadecimal numbers taken from a list maintained by the IEEE. This list can be found at the following URL:

<http://standards.ieee.org/regauth/ethertype/index.html>

- **llc** — LLC Service Advertising Protocol (SAP)

The values for **llc** are four-digit hexadecimal numbers that are created by concatenating a two-digit LLC Destination SAP (DSAP) and a two-digit LLC Source SAP (SSAP).

- **snap** — Ethertype inside an IEEE SNAP packet encapsulation.

The values for **snap** are the same as the values for **etype**, described previously.

For example:

```
config protocol fred add llc feff
```

```
config protocol fred add snap 9999
```

A maximum of fifteen protocol filters, each containing a maximum of six protocols, can be defined. On products that use the Inferno chip set, all fifteen protocol filters can be active and configured for use. On all other platforms, no more than seven protocols can be active and configured for use.



For more information on SNAP for Ethernet protocol types, see TR 11802-5:1997 (ISO/IEC) [ANSI/IEEE std. 802.1H, 1997 Edition].

DELETING A PROTOCOL FILTER

If a protocol filter is deleted from a VLAN, the VLAN is assigned a protocol filter of none. You can continue to configure the VLAN. However, no traffic is forwarded to the VLAN until a protocol is assigned to it.

PRECEDENCE OF TAGGED PACKETS OVER PROTOCOL FILTERS

If a VLAN is configured to accept tagged packets on a particular port, incoming packets that match the tag configuration take precedence over any protocol filters associated with the VLAN.

VLAN NAMES

Each VLAN is given a name that can be up to 32 characters. VLAN names can use standard alphanumeric characters. The following characters are not permitted in a VLAN name:

- Space
- Comma
- Quotation mark

VLAN names must begin with an alphabetical letter. Quotation marks can be used to enclose a VLAN name that does not begin with an alphabetical character, or that contains a space, comma, or other special character.

VLAN names are locally significant. That is, VLAN names used on one switch are only meaningful to that switch. If another switch is connected to it, the VLAN names have no significance to the other switch.



You should use VLAN names consistently across your entire network.

DEFAULT VLAN

The switch ships with one default VLAN that has the following properties:

- The VLAN name is *default*.
- It contains all the ports on a new or initialized switch.
- The default VLAN is untagged on all ports. It has an internal VLANid of 1.

RENAMING A VLAN

To rename an existing VLAN, use the following command:

```
config vlan <old_name> name <new_name>
```

The following rules apply to renaming VLANs:

- Once you change the name of the default VLAN, it cannot be changed back to *default*.
- You cannot create a new VLAN named *default*.
- You cannot change the VLAN name *MacVlanDiscover*. Although the switch accepts a name change, once it is rebooted, the original name is recreated.

CONFIGURING VLANs ON THE SWITCH

This section describes the commands associated with setting up VLANs on the switch. Configuring a VLAN involves the following steps:

- 1 Create and name the VLAN.
- 2 Assign an IP address and mask (if applicable) to the VLAN, if needed.



Each IP address and mask assigned to a VLAN must represent a unique IP subnet. You cannot configure the same IP subnet on different VLANs.

- 3 Assign a VLANid, if any ports in this VLAN will use a tag.
- 4 Assign one or more ports to the VLAN.

As you add each port to the VLAN, decide if the port will use an 802.1Q tag.

VLAN CONFIGURATION COMMANDS

[Table 6-1](#) describes the commands used to configure a VLAN.

Table 6-1: VLAN Configuration Commands

Command	Description
config dot1q ethertype <ethertype>	Configures an IEEE 802.1Q Ethertype. Use this command only if you have another switch that supports 802.1Q, but uses a different Ethertype value than 8100. You must reboot the switch for this command to take effect.
config protocol <protocol_name> [add delete]<protocol_type> <hex_value> {<protocol_type><hex_value>} ...	Configures a protocol filter. Supported <protocol_type> values include: <ul style="list-style-type: none">■ etype■ llc■ snap The variable <hex_value> is a hexadecimal number between 0 and FFFF that represents either the Ethernet protocol type (for EtherType), the DSAP/SSAP combination (for LLC), or the SNAP-encoded Ethernet protocol type (for SNAP).
config vlan <name> add port <portlist> {tagged untagged} {nobroadcast}	Adds one or more ports to a VLAN. You can specify tagged port(s), untagged port(s). Specify nobroadcast to prevent the switch from forwarding broadcast, multicast, and unknown unicast traffic. By default, ports are untagged.
config vlan <name> delete port <portlist> {tagged untagged} {nobroadcast}	Deletes one or more ports from a VLAN.
config vlan <name> ipaddress <ipaddress> {<mask>}	Assigns an IP address and an optional mask to the VLAN.
config vlan <name> protocol [<protocol_name> any]	Configures a protocol-based VLAN. If the keyword any is specified, then it becomes the default VLAN. All packets that cannot be classified into other protocol-based VLANs are assigned to the default VLAN of that port.
config vlan <name> qosprofile <qosname>	Configures a VLAN to use a particular QoS profile. Dynamic FDB entries associated with the VLAN are flushed once the change is committed.

Table 6-1: VLAN Configuration Commands (continued)

Command	Description
config vlan <name> tag <vlanid>	Assigns a numerical VLANid. The valid range is from 1 to 4095.
config vlan <old_name> name <new_name>	Renames a previously configured VLAN.
create protocol <protocol_name>	Creates a user-defined protocol.
create vlan <name>	Creates a named VLAN.
delete protocol <protocol>	Removes a protocol.
delete vlan <name>	Removes a VLAN.
unconfig vlan <name> ipaddress	Resets the IP address of the VLAN.

VLAN CONFIGURATION EXAMPLES

The following BlackDiamond switch example creates a port-based VLAN named *accounting*, assigns the IP address 132.15.121.1, and assigns slot 2, ports 1, 2, 3, and 6, and slot 4, ports 1 and 2 to it:

```
create vlan accounting
config accounting ipaddress 132.15.121.1
config default delete port 2:1-2:3,2:6,4:1,4:2
config accounting add port 2:1-2:3,2:6,4:1,4:2
```



Because VLAN names are unique, you do not need to enter the keyword `vlan` after you have created the unique VLAN name. You can use the VLAN name alone.

The following Summit switch example creates a tag-based VLAN named *video*. It assigns the VLANid 1000. Ports 4 through 8 are added as tagged ports to the VLAN.

```
create vlan video
config video tag 1000
config video add port 4-8 tagged
```

The following Summit switch example creates a VLAN named *sales*, with the VLANid 120. The VLAN uses both tagged and untagged ports. Ports 1 through 3 are tagged, and ports 4 and 7 are untagged. Note that when not explicitly specified, ports are added as untagged.

```
create vlan sales
config sales tag 120
config sales add port 1-3 tagged
config sales add port 4,7
```

The following BlackDiamond switch example creates a protocol-based VLAN named *ipsales*. Slot 5, ports 6 through 8, and slot 6, ports 1, 3, and 4-6 are assigned to the VLAN.

```
create vlan ipsales
config ipsales protocol ip
config ipsales add port 5:6-5:8,6:1,6:3-6:6
```

The following BlackDiamond switch example defines a protocol filter, *myprotocol* and applies it to the VLAN named *myvlan*. This is an example only, and has no real-world application.

```
create protocol myprotocol
config protocol myprotocol add etype 0xf0f0
config protocol myprotocol add etype 0xffff
create vlan myvlan
config myvlan protocol myprotocol
```

DISPLAYING VLAN SETTINGS

To display VLAN settings, use the following command:

```
show vlan {<name>} {detail}
```

The `show` command displays summary information about each VLAN, and includes the following:

- Name
- VLANid
- How the VLAN was created (manually or by GVRP)
- IP address

- IPX address (if configured)
- STPD information
- Protocol information
- QoS profile information
- Ports assigned
- Tagged/untagged status for each port
- How the ports were added to the VLAN (manually or by GVRP)

Use the `detail` option to display the detailed format.

To display protocol information, use the following command:

```
show protocol {<protocol>}
```

This `show` command displays protocol information, including the following:

- Protocol name
- List of protocol fields
- VLANs that use the protocol

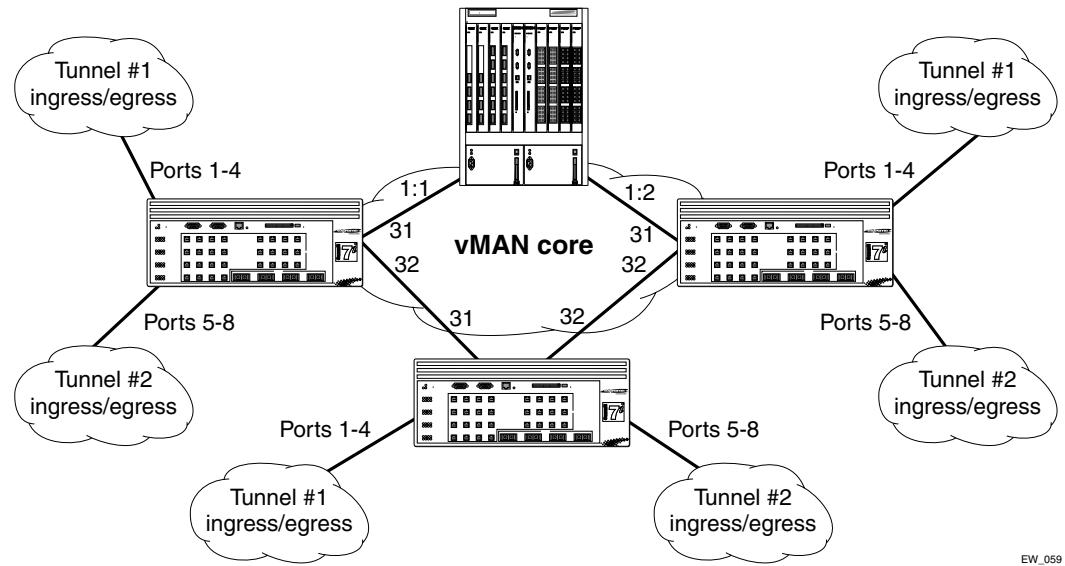
VLAN TUNNELING (vMANs)

You can "tunnel" any number of 802.1Q and/or Cisco ISL VLANs into a single VLAN that can be switched through an Extreme Ethernet infrastructure. A given tunnel is completely isolated from other tunnels or VLANs. This feature is useful in building transparent private networks (vMANs) that need point-to-point or point-to-multipoint connectivity across an Ethernet infrastructure. The VLAN tagging methods used within the vMAN tunnel are transparent to the tunnel. For the MAN provider, the tagging numbers and methods used by the customer are transparent to the provider.

The steps to configure a vMAN tunnel are:

- 1 Modify the 802.1Q Ethertype the switch uses to recognize tagged frames
- 2 Configure the switch to accept larger MTU size frames ("Jumbo" frames)
- 3 Create tunnels by creating VLANs and configuring member ports as tagged on switch-to-switch ports and untagged on the tunnel's ingress/egress ports.

Figure 6-7 illustrates a configuration with vMANs.



EW_059

Figure 6-7: vMAN example

Two tunnels are depicted that have ingress/egress ports on each Summit7i.

GENERIC VLAN REGISTRATION PROTOCOL

The Generic VLAN Registration Protocol (GVRP) allows a LAN device to signal other neighboring devices that it wishes to receive packets for one or more VLANs. The GVRP protocol is defined as part of the IEEE 802.1Q Virtual LANs draft standard. The main purpose of the protocol is to allow switches to automatically discover some of the VLAN information that would otherwise have to be manually configured in each switch. GVRP can also be run by network servers. These servers are usually configured to join several VLANs, and then signal the network switches of the VLANs they want to join.

Figure 6-8 illustrates a network using GVRP.

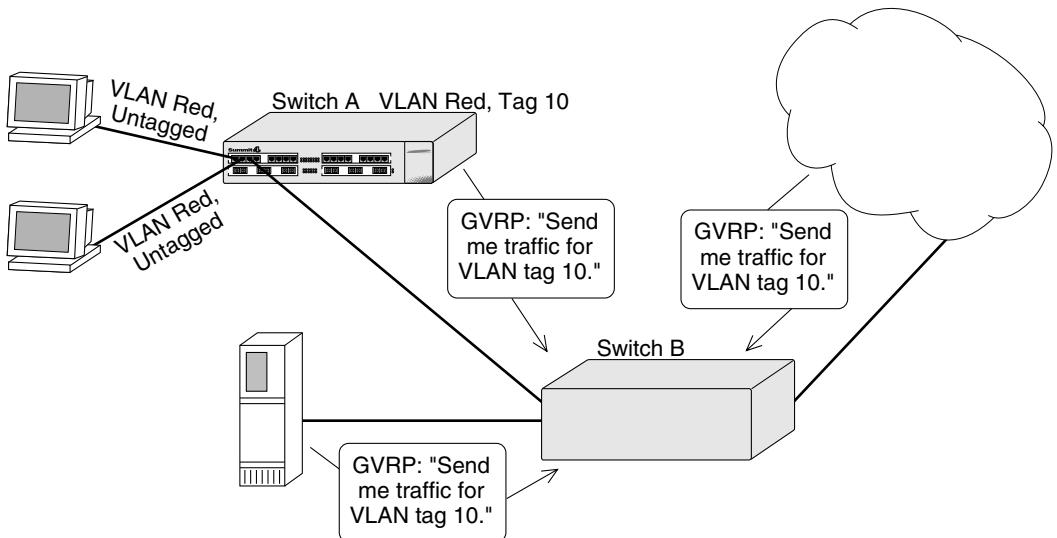


Figure 6-8: Network example using GVRP

In [Figure 6-8](#), Switch A is a member of VLAN Red. VLAN Red has the VLANid 10. Port 1 and port 2 on Switch A are added to the VLAN as untagged.

The configuration for Switch A is as follows:

```
create vlan red
config vlan red tag 10
config vlan red add port 1-2 untagged
enable gvrp
```

Switch B does not need to be configured with VLAN or tagging information. Instead, using GVRP, the server connected to Switch B, and the remainder of the network connected to Switch B provides Switch B with the information it needs to forward traffic. Switch A automatically adds port 3 to VLAN Red because Switch A now knows that there are other devices on port 3 that need access to VLAN Red.

VLANs that are automatically created using GVRP are given names in the format

```
gvrp vlan xxxx
```

where xxxx is the VLANid (in decimal) that is discovered by GVRP. These VLANs are not permanently stored in nonvolatile storage, and you cannot add or remove ports from these VLANs.

GVRP assumes that the VLANs for which it carries information operate using VLAN tags, unless explicitly configured otherwise. Typically, you must configure any untagged VLANs on the switches at the edges of the network, and the GVRP protocol is used across the core of the network to automatically configure other switches using tagged VLANs.

 *You cannot assign an IP address to a VLAN learned by way of GVRP.*

GVRP AND SPANNING TREE DOMAINS

Because GVRP-learned VLANs are dynamic, all VLANs created by GVRP use the system defaults and become members of the default Spanning Tree Domain (STPD), s0. Because two STPDs cannot exist on the same physical port, if two GVRP clients attempt to join two different VLANs that belong to two different STPDs, the second client is refused. You should configure all potential GVRP VLANs to be members of the same STPD. This configuration is done automatically, if you have not configured additional STPDs.

GVRP COMMANDS

[Table 6-2](#) describes GVRP commands.

Table 6-2: GVRP Commands

Command	Description
config gvrp [listen send both none] port <portlist>	Configures the sending and receiving GVRP information one or all a ports. Options include the following: <ul style="list-style-type: none">■ listen — Receive GVRP packets.■ send — Send GVRP packets.■ both — Send and receive GVRP packets.■ none — Disable the port from participating in GVRP operation. The default setting is both.
disable gvrp	Disables the Generic VLAN Registration Protocol (GVRP).
enable gvrp	Enables the Generic VLAN Registration Protocol (GVRP). The default setting is disabled.
show gvrp	Displays the current configuration and status of GVRP.

MAC-BASED VLANS

MAC-Based VLANs allow physical ports to be mapped to a VLAN based on the source MAC address learned in the FDB. This feature allows you to designate a set of ports that have their VLAN membership dynamically determined by the MAC address of the end station that plugs into the physical port. You may configure the source MAC address-to-VLAN mapping either offline, or dynamically on the switch. For example, you could use this application for a roaming user who wishes to connect to a network from a conference room. In each room, the user plugs into one of the designated ports on the switch and is mapped to the appropriate VLAN. Connectivity is maintained to the network with all of the benefits of the configured VLAN in terms of QoS, routing, and protocol support.

MAC-BASED VLAN GUIDELINES

When using the MAC-to-VLAN mapping, consider the following guidelines:

- A port can only accept connections from an endstation/host and should not be connected to a layer-2 repeater device. Connecting to a layer-2 repeater device can cause certain addresses to not be mapped to their respective VLAN if they are not correctly configured in the MAC-VLAN configuration database. If a repeater device is connected to a MAC-Based VLAN port, and the configured MAC-to-VLAN mapped station enters on the repeater, any endstation that is attached to the repeater can be mapped to that VLAN while the configured endstation is active in that VLAN. Upon removal of the configured MAC-to-VLAN endstation, all other endstations lose connectivity.
- Groups are used as a security measure to allow a MAC address to enter into a VLAN only when the group mapping matches the port mapping. As an example, the following configuration allows MAC 00:00:00:00:aa to enter into the VLAN only on ports 10 and 11 because of membership in group 100:

```
* Summit48:50 # show mac
Port      Vlan          Group        State
10       MacVlanDiscover 100         Discover
11       MacVlanDiscover 100         Discover
12       MacVlanDiscover any          Discover
13       MacVlanDiscover any          Discover
14       MacVlanDiscover any          Discover
Total Entries in Database:2
      Mac           Vlan     Group
00:00:00:00:00:aa   sales    100
00:00:00:00:00:01   sales    any
2 matching entries
```

- The group “any” is equivalent to the group “0”. Ports that are configured as “any” allow any MAC address to be assigned to a VLAN, regardless of group association.
- Partial configurations of the MAC to VLAN database can be downloaded to the switch using the timed download configuration feature.

MAC-BASED VLAN LIMITATIONS

The following list contains the limitations of MAC-based VLANs:

- Ports participating in MAC VLANs must first be removed from any static VLANs.
- The MAC- to-VLAN mapping can only be associated with VLANs that exist on the switch.

- A MAC address cannot be configured to associate with more than 1 VLAN. If this is attempted, the MAC address is associated with the most recent VLAN entry in the MAC-to-VLAN database.
- The feature is intended to support one client per physical port. Once a client MAC address has successfully registered, the VLAN association remains until the port connection is dropped or the FDB entry ages out.

MAC-BASED VLAN COMMANDS

[Table 6-3](#) describes MAC-based VLAN commands.

Table 6-3: MAC-Based VLAN Commands

Command	Description
config mac-vlan add mac-address [any <mac_address>] mac-group [any <group_number>] vlan <name>	Adds a MAC address to a MAC-based VLAN.
config mac-vlan delete [mac-address <mac_address> all]	Removes a MAC address from a MAC-based VLAN.
disable mac-vlan port <portlist>	Disables a port from using the MAC-based VLAN algorithm.
enable mac-vlan mac-group [any <group_number>] port <portlist>	enables a port to use the MAC-based VLAN algorithm.
show mac-vlan {configuration database}	Displays the MAC-based VLAN configuration and MAC address database content.

MAC-BASED VLAN EXAMPLE

In this following example, three VLANs are created: *engineering*, *marketing*, and *sales*. A single MAC address is associated with each VLAN. The MAC address 00:00:00:00:00:02 has a group number of “any” or “0” associated with it, allowing it to be plugged into any port that is in MacVlanDiscover mode (ports 10-15 in this case). The MAC address 00:00:00:00:00:01 has a group number of 10 associated with it, and can only be assigned to a VLAN if inserted into ports 16 or 17. The MAC address 00:00:00:00:00:03 has a group number of 200 associated with it and can only be inserted into ports 18 through 20.

```
enable mac-vlan mac-group any ports 10-15
enable mac-vlan mac-group 10 ports 16-17
enable mac-vlan mac-group 200 ports 18-20
config mac-vlan add mac-address 00:00:00:00:00:01 mac-group 10
engineering
config mac-vlan add mac-address 00:00:00:00:00:02 mac-group any
marketing
config mac-vlan add mac-address 00:00:00:00:00:03 mac-group 200 sales
```

TIMED CONFIGURATION DOWNLOAD FOR MAC-BASED VLANS

To allow centralized control of MAC-based VLANs over multiple switches, a timed TFTP configuration download allows you to download incremental configuration files from a primary or secondary server at specified time intervals. The timed downloads are configurable in 24 hour intervals. When a switch reboots, the configuration is automatically downloaded immediately after booting, per the configured primary and secondary servers.

To configure the primary and/or secondary server and file name, use the following command:

```
config download server [primary | secondary] <host_name> | <ip_address>
<filename>
```

To enable timed interval downloads, use the following command:

```
download configuration every <hour (0-23)>
```

To display timed download information, use the following command:

```
show switch
```

EXAMPLE

In relation to MAC-based VLANs, the downloaded file is an ASCII file that consists of CLI commands used to configure the most recent MAC-to-VLAN database. This feature is different from the normal download configuration command in that it allows incremental configuration without the automatic rebooting of the switch.

The following example shows an incremental configuration file for MAC-based VLAN information that updates the database and saves changes:

```
config mac-vlan add mac-address 00:00:00:00:00:01 mac-group any  
engineering  
config mac-vlan add mac-address 00:00:00:00:ab:02 mac-group any  
engineering  
config mac-vlan add mac-address 00:00:00:00:cd:04 mac-group any sales  
. .  
config mac-vlan add mac-address 00:00:00:00:ab:50 mac-group any sales  
config mac-vlan add mac-address 00:00:00:00:cd:60 mac-group any sales  
save
```




Forwarding Database (FDB)

This chapter describes the following topics:

- [Overview of the FDB on page 7-1](#)
- [Configuring FDB Entries on page 7-3](#)
- [Displaying FDB Entries on page 7-5](#)

OVERVIEW OF THE FDB

The switch maintains a database of all media access control (MAC) addresses received on all of its ports. It uses the information in this database to decide whether a frame should be forwarded or filtered.

FDB CONTENTS

Each FDB entry consists of the MAC address of the device, an identifier for the port on which it was received, and an identifier for the VLAN to which the device belongs. Frames destined for devices that are not in the FDB are flooded to all members of the VLAN.

FDB ENTRY TYPES

The following are four types of entries in the FDB:

- **Dynamic entries** — Initially, all entries in the database are dynamic. Entries in the database are removed (aged-out) if, after a period of time (aging time), the device has not transmitted. This prevents the database from becoming full with obsolete entries by ensuring that when a device is removed from the network, its entry is deleted from the database. Dynamic entries are deleted from the database if the switch is reset or a power off/on cycle occurs. For more information about setting the aging time, refer to the section “[Configuring FDB Entries](#),” later in this chapter.
- **Non-aging entries** — If the aging time is set to zero, all aging entries in the database are defined as static, non-aging entries. This means that they do not age, but they are still deleted if the switch is reset.
- **Permanent entries** — Permanent entries are retained in the database if the switch is reset or a power off/on cycle occurs. The system administrator must make entries permanent. A permanent entry can either be a unicast or multicast MAC address. All entries entered by way of the command-line interface are stored as permanent. The Summit switch can support a maximum of 64 permanent entries, the BlackDiamond switch supports a maximum of 254 permanent entries.

Once created, permanent entries stay the same as when they were created. For example, the permanent entry store is not updated when any of the following take place:

- A VLAN is deleted.
 - A VLANid is changed.
 - A port mode is changed (tagged/untagged).
 - A port is deleted from a VLAN.
 - A port is disabled.
 - A port enters blocking state.
 - A port QoS setting is changed.
 - A port goes down (link down).
- **Blackhole entries** — A blackhole entry configures the switch to discard packets with a specified MAC destination address. Blackhole entries are useful as a security measure or in special circumstances where a specific destination address must be discarded. Blackhole entries are treated like permanent entries in the event of a switch reset or power off/on cycle. Blackhole entries are never aged out of the database.

How FDB ENTRIES GET ADDED

Entries are added into the FDB in the following two ways:

- The switch can learn entries. The system updates its FDB with the source MAC address from a packet, the VLAN, and the port identifier on which the source packet is received.
- You can enter and update entries using a MIB browser, an SNMP Network Manager, or the command-line interface (CLI).

ASSOCIATING A QOS PROFILE WITH AN FDB ENTRY

You can associate a QoS profile with a MAC address (and VLAN) of a device that will be dynamically learned. The FDB treats the entry like a dynamic entry (it is learned, it can be aged out of the database, and so on). The switch applies the QoS profile as soon as the FDB entry is learned.



For more information on QoS, refer to [Chapter 9](#).

CONFIGURING FDB ENTRIES

To configure entries in the FDB, use the commands listed in [Table 7-1](#).

Table 7-1: FDB Configuration Commands

Command	Description
clear fdb {<mac_address> vlan <name> <portlist>}	Clears dynamic FDB entries that match the filter. When no options are specified, the command clears all FDB entries.
config fdb agingtime <number>	Configures the FDB aging time. The range is 15 through 1,000,000 seconds. The default value is 300 seconds. A value of 0 indicates that the entry should never be aged out.

Table 7-1: FDB Configuration Commands (continued)

Command	Description
create fdbentry <mac_address> vlan <name> {[blackhole <portlist> dynamic] {qosprofile <qosprofile>}}	<p>Creates an FDB entry. Specify the following:</p> <ul style="list-style-type: none"> ■ <code>mac_address</code> — Device MAC address, using colon separated bytes. ■ <code>name</code> — VLAN associated with MAC address. ■ <code>blackhole</code> — Configures the MAC address as a blackhole entry. ■ <code>portlist</code> — Port numbers associated with MAC address. ■ <code>dynamic</code> — Specifies that the entry will be learned dynamically. Used to associate a QoS profile with a dynamically learned entry. ■ <code>qosprofile</code> — QoS profile associated with MAC address. <p>If more than one port number is associated with a permanent MAC entry, packets are multicast to the multiple destinations.</p>
delete fdbentry <mac_address> vlan <name>	Deletes a permanent FDB entry.
disable learning port <portlist>	Disables MAC address learning on one or more ports for security purposes. If MAC address learning is disabled, only broadcast traffic, EDP traffic, and packets destined to a permanent MAC address matching that port number, are forwarded. The default setting is enabled.
enable learning port <portlist>	Enables MAC address learning on one or more ports.

FDB CONFIGURATION EXAMPLES

The following example adds a permanent entry to the FDB:

```
create fdbentry 00:E0:2B:12:34:56 vlan marketing port 3:4
```

The permanent entry has the following characteristics:

- MAC address is 00:E0:2B:12:34:56.
- VLAN name is *marketing*.

- Slot number for this device is 3.
- Port number for this device is 4.

This example associates the QoS profile *qp2* with a dynamic entry that will be learned by the FDB:

```
create fdbentry 00:A0:23:12:34:56 vlan net34 dynamic qosprofile qp2
```

This entry has the following characteristics:

- MAC address is 00A023123456.
- VLAN name is *net34*.
- The entry will be learned dynamically.
- QoS profile *qp2* will be applied when the entry is learned.

DISPLAYING FDB ENTRIES

To display FDB entries, use the command

```
show fdb {<mac_address> | vlan <name> | <portlist> | permanent}
```

where the following is true:

- *mac_address* — Displays the entry for a particular MAC address.
- *vlan <name>* — Displays the entries for a VLAN.
- *portlist* — Displays the entries for a slot and port combination.
- *permanent* — Displays all permanent entries.

With no options, the command displays all FDB entries.

Spanning Tree Protocol (STP)

This chapter covers the following topic:

- [Overview of the Spanning Tree Protocol on page 8-1](#)
- [Spanning Tree Domains on page 8-2](#)
- [STP Configurations on page 8-3](#)
- [Configuring STP on the Switch on page 8-6](#)
- [Displaying STP Settings on page 8-8](#)
- [Disabling and Resetting STP on page 8-9](#)

Using the Spanning Tree Protocol (STP) functionality of the switch makes your network more fault tolerant. The following sections explain more about STP and the STP features supported by ExtremeWare.



STP is a part of the 802.1D bridge specification defined by the IEEE Computer Society. To explain STP in terms used by the 802.1D specification, the Summit switch and the BlackDiamond switch will be referred to as a bridge.

OVERVIEW OF THE SPANNING TREE PROTOCOL

STP is a bridge-based mechanism for providing fault tolerance on networks. STP allows you to implement parallel paths for network traffic, and ensure that:

- Redundant paths are disabled when the main paths are operational.
- Redundant paths are enabled if the main path fails.

SPANNING TREE DOMAINS

The switch can be partitioned into multiple virtual bridges. Each virtual bridge can run an independent Spanning Tree instance. Each Spanning Tree instance is called a *Spanning Tree Domain* (STPD). Each STPD has its own Root Bridge and active path. Once the STPD is created, one or more VLANs can be assigned to it.

A port can belong to only one STPD. If a port is a member of multiple VLANs, then all those VLANs must belong to the same STPD.

The key points to remember when configuring VLANs and STP are the following:

- Each VLAN forms an independent broadcast domain.
- STP blocks paths to create a loop-free environment.
- When STP blocks a path, no data can be transmitted or received on the blocked port.
- Within any given STPD, all VLANs belonging to it use the same spanning tree.



Care must be taken to ensure that multiple STPD instances within a single switch do not see each other in the same broadcast domain. This could happen if, for example, another external bridge is used to connect VLANs belonging to separate STPDs.

If you delete a STPD, the VLANs that were members of that STPD are also deleted. You must remove all VLANs associated with the STP before deleting the STPD.



If no VLANs are configured to use the protocol filter any on a particular port, STP BPDUs are not flooded within a VLAN when STP is turned off. If you need STP to operate on this type of port, enable STP on the associated VLAN, so that it can participate.



On the BlackDiamond switch, in order to support more than 255 ports (as limited by the 802.1D specification), both the cost and port ID fields in the BPDU are used to uniquely identify each port.

STPD STATUS FOR GVRP-ADDED PORTS

If a port is added to a VLAN by GVRP, the newly added port reflects the SPTD membership and status of the VLAN to which it is added. For example, if VLAN Red is a member of STPD s0, and s0 is enabled, then all ports added to VLAN Red by GVRP have s0 enabled on those ports, as well. The command for disabling STP on a port basis has no permanent affect on ports controlled by GVRP.



For more information on GVRP, refer to [Chapter 6](#).

DEFAULTS

The default device configuration contains a single STPD called *s0*. The default VLAN is a member of STPD *s0*.

All STP parameters default to the IEEE 802.1D values, as appropriate.

STP CONFIGURATIONS

When you assign VLANs to an STPD, pay careful attention to the STP configuration and its effect on the forwarding of VLAN traffic.

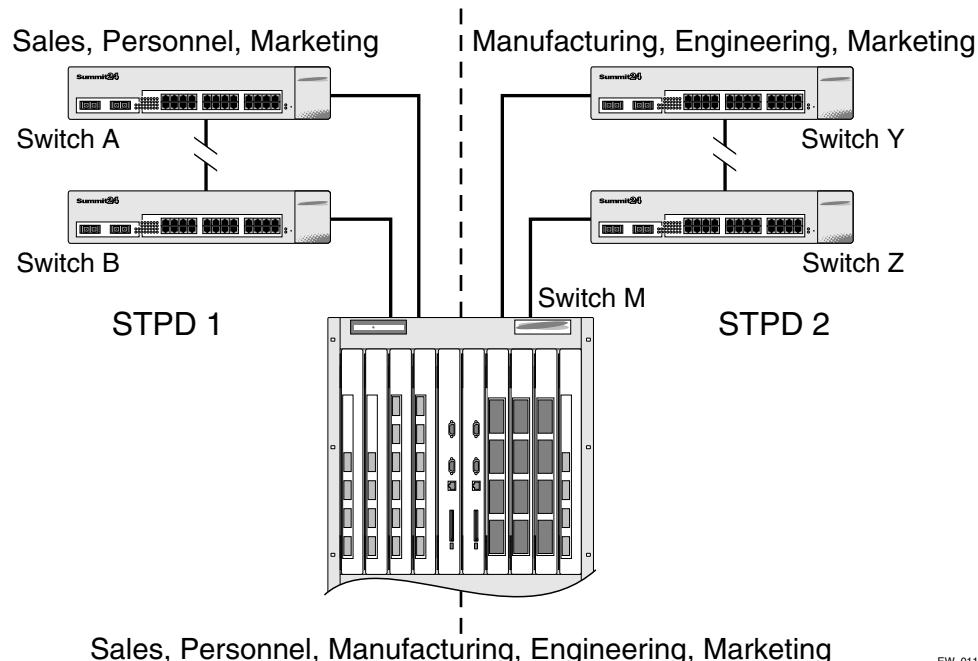
[Figure 8-1](#) illustrates a network that uses VLAN tagging for trunk connections. The following four VLANs have been defined:

- *Sales* is defined on Switch A, Switch B, and Switch M.
- *Personnel* is defined on Switch A, Switch B, and Switch M.
- *Manufacturing* is defined on Switch Y, Switch Z, and Switch M.
- *Engineering* is defined on Switch Y, Switch Z, and Switch M.
- *Marketing* is defined on all switches (Switch A, Switch B, Switch Y, Switch Z, and Switch M).

Two STPDs are defined:

- STPD1 contains VLANs *Sales* and *Personnel*.
- STPD2 contains VLANs *Manufacturing* and *Engineering*.

The VLAN *Marketing* is a member of the default STPD, but not assigned to either STPD1 or STPD2.



EW_011

Figure 8-1: Multiple Spanning Tree Domains

When the switches in this configuration start up, STP configures each STPD such that there are no active loops in the topology. STP could configure the topology in a number of ways to make it loop-free.

In [Figure 8-1](#), the connection between Switch A and Switch B is put into blocking state, and the connection between Switch Y and Switch Z is put into blocking state. After STP converges, all the VLANs can communicate, and all bridging loops are prevented.

The VLAN *Marketing*, which has not been assigned to either STPD1 or STPD2, communicates using all five switches. The topology has no loops, because STP has already blocked the port connection between Switch A and Switch B, and between Switch Y and Switch Z.

Within a single STPD, you must be extra careful when configuring your VLANs.

[Figure 8-2](#) illustrates a network that has been incorrectly set up using a single STPD so that the STP configuration disables the ability of the switches to forward VLAN traffic.

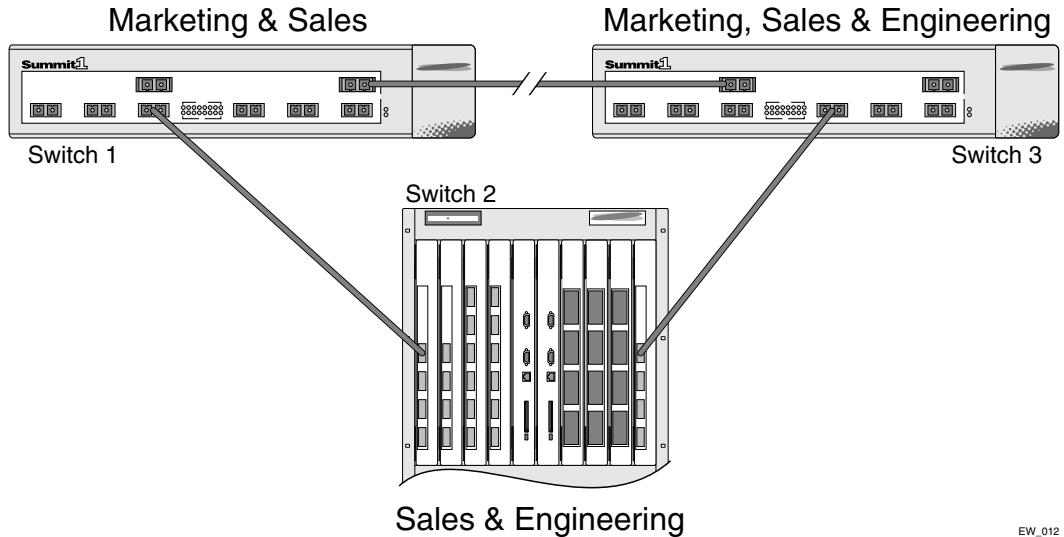


Figure 8-2: Tag-based STP configuration

The tag-based network in [Figure 8-2](#) has the following configuration:

- Switch 1 contains VLAN *Marketing* and VLAN *Sales*.
- Switch 2 contains VLAN *Engineering* and VLAN *Sales*.
- Switch 3 contains VLAN *Marketing*, VLAN *Engineering*, and VLAN *Sales*.
- The tagged trunk connections for three switches form a triangular loop that is not permitted in an STP topology.
- All VLANs in each switch are members of the same STPD.

STP may block traffic between Switch 1 and Switch 3 by disabling the trunk ports for that connection on each switch.

Switch 2 has no ports assigned to VLAN marketing. Therefore, if the trunk for VLAN marketing on Switches 1 and 3 is blocked, the traffic for VLAN marketing will not be able to traverse the switches.

CONFIGURING STP ON THE SWITCH

To configure STP you must perform the following actions:

- Create one or more STP domains using the following command:

```
create stpd <stpd_name>
```

 *STPD, VLAN, and QoS profile names must all be unique. For example, a name used to identify a VLAN cannot be used when you create an STPD or a QoS profile.*

- Add one or more VLANs to the STPD using the following command:

```
config stpd <stpd_name> add vlan <name>
```

- Enable STP for one or more STP domains using the following command:

```
enable stpd {<stpd_name>}
```

 *All VLANs belong to a STPD. If you do not want to run STP on a VLAN, you must add the VLAN to a STPD that is disabled.*

Once you have created the STPD, you can optionally configure STP parameters for the STPD.

 *You should not configure any STP parameters unless you have considerable knowledge and experience with STP. The default STP parameters are adequate for most networks.*

The following parameters can be configured on each STPD:

- Hello time
- Forward delay
- Max age
- Bridge priority

The following parameters can be configured on each port:

- Path cost
- Port priority

 *The device supports the RFC 1493 Bridge MIB. Parameters of only the s0 default STPD are accessible through this MIB.*

[Table 8-1](#) shows the commands used to configure STP.

Table 8-1: STP Configuration Commands

Command	Description
config stpd <stpd_name> add vlan <name>	Adds a VLAN to the STPD.
config stpd <stpd_name> forwarddelay <value>	Specifies the time (in seconds) that the ports in this STPD spend in the listening and learning states when the switch is the Root Bridge. The range is 4 through 30. The default setting is 15 seconds.
config stpd <stpd_name> hellotime <value>	Specifies the time delay (in seconds) between the transmission of BPDUs from this STPD when it is the Root Bridge. The range is 1 through 10. The default setting is 2 seconds.
config stpd <stpd_name> maxage <value>	Specifies the maximum age of a BPDU in this STPD. The range is 6 through 40. The default setting is 20 seconds.
config stpd <stpd_name> port cost <value> <portlist>	Note that the time must be greater than, or equal to $2 * (\text{Hello Time} + 1)$ and less than, or equal to $2 * (\text{Forward Delay} - 1)$. Specifies the path cost of the port in this STPD. The range is 1 through 65,535. The switch automatically assigns a default path cost based on the speed of the port, as follows: <ul style="list-style-type: none"> ■ For a 10Mbps port, the default cost is 100. ■ For a 100Mbps port, the default cost is 19. ■ For a 1000Mbps port, the default cost is 4.
config stpd <stpd_name> port priority <value> <portlist>	Specifies the priority of the port in this STPD. By changing the priority of the port, you can make it more or less likely to become the Root Port. The range is 0 through 31. The default setting is 16. A setting of 0 indicates the lowest priority.
config stpd <stpd_name> priority <value>	Specifies the priority of the STPD. By changing the priority of the STPD, you can make it more or less likely to become the Root Bridge. The range is 0 through 65,535. The default setting is 32,768. A setting of 0 indicates the highest priority.

Table 8-1: STP Configuration Commands (continued)

Command	Description
create stpd <stpd_name>	Creates an STPD. When created, an STPD has the following default parameters: <ul style="list-style-type: none"> ■ Bridge priority — 32,768 ■ Hello time — 2 seconds ■ Forward delay — 15 seconds
enable ignore-stp vlan <name>	Configures the switch to ignore the STP protocol, and not block traffic for the VLAN(s). This command is useful when multiple VLANs share the same physical ports, but only some of the VLANs require STP protection. The default setting is disabled.
enable stpd {<stpd_name>}	Enables the STP protocol for one or all STPDs. The default setting is disabled.
enable stpd port {<portlist>}	Enables the STP protocol on one or more ports. If STPD is enabled for a port, Bridge protocol Data Units (BPDUs) will be generated on that port if STP is enabled for the associated STPD. The default setting is enabled.

STP CONFIGURATION EXAMPLE

The following BlackDiamond switch example creates and enables an STPD named *Backbone_st*. It assigns the *Manufacturing* VLAN to the STPD. It disables STP on slot 2, ports 1 through 7, and slot 3 port 12.

```
create stpd backbone_st
config stpd backbone_st add vlan manufacturing
enable stpd backbone_st
disable stpd backbone_st port 2:1-2:7,3:12
```

DISPLAYING STP SETTINGS

To display STP settings, use the following command:

```
show stpd {<stpd_name>}
```

This command displays the following information:

- STPD name
- Bridge ID
- STPD configuration information

To display the STP state of a port, use the following command:

```
show stpd <stpd_name> port <portlist>
```

This command displays the following:

- STPD port configuration
- STPD state (Root Bridge, and so on)
- STPD port state (forwarding, blocking, and so on)

DISABLING AND RESETTING STP

To disable STP or return STP settings to their defaults, use the commands listed in [Table 8-2](#).

Table 8-2: STP Disable and Reset Commands

Command	Description
delete stpd <stpd_name>	Removes an STPD. An STPD can only be removed if all VLANs have been deleted from it. The default STPD, s0, cannot be deleted.
disable ignore-stp vlan <name>	Allows a VLAN to use STP port information.
disable stpd [<stpd_name> all]	Disables the STP mechanism on a particular STPD, or for all STPDs.
disable stpd port <portlist>	Disables STP on one or more ports. Disabling STP on one or more ports puts those ports in <i>forwarding</i> state; all BPDUs received on those ports will be disregarded.
unconfig stpd {<stpd_name>}	Restores default STP values to a particular STPD or to all STPDs.

Quality of Service (QoS)

This chapter covers the following topics:

- Overview of Policy-Based Quality of Service on page 9-2
- Applications and Types of QoS on page 9-3
- Assigning QoS Attributes on page 9-5
- QoS Profiles on page 9-6
- Traffic Groupings and Creating a QoS Policy on page 9-8
 - IP-Based Traffic Groupings on page 9-10
 - MAC-Based Traffic Groupings on page 9-10
 - Explicit Class of Service (802.1p and DiffServ) Traffic Groupings on page 9-12
 - Physical and Logical Groupings on page 9-20
- Verifying Configuration and Performance on page 9-21
- Modifying a QoS Policy on page 9-23
- Intra-Subnet QoS on page 9-24
- Bi-Directional Rate Shaping on page 9-25
- Dynamic Link Context System on page 9-27

Policy-Based Quality of Service (QoS) is a feature of ExtremeWare and the Extreme switch architecture that allows you to specify different service levels for traffic traversing the switch. Policy-Based QoS is an effective control mechanism for networks that have heterogeneous traffic patterns. Using Policy-Based QoS, you can specify the service level that a particular traffic type receives.

OVERVIEW OF POLICY-BASED QUALITY OF SERVICE

Policy-Based QoS allows you to protect bandwidth for important categories of applications or specifically limit the bandwidth associated with less critical traffic. For example, if voice-over-IP traffic requires a reserved amount of bandwidth to function properly, using Policy-Based QoS, you can reserve sufficient bandwidth critical to this type of application. Other applications deemed less critical can be limited so as to not consume excessive bandwidth. The switch contains separate hardware queues on every physical port. Each hardware queue is programmed by ExtremeWare with bandwidth management and prioritization parameters. The bandwidth management and prioritization parameters that modify the forwarding behavior of the switch affect how the switch transmits traffic for a given hardware queue on a physical port.

The switch tracks and enforces the minimum and maximum percentage of bandwidth utilization transmitted on every hardware queue for every port. When two or more hardware queues on the same physical port are contending for transmission, the switch prioritizes bandwidth use so long as their respective bandwidth management parameters are satisfied. Switch products with the “i” chipset can be configured with up to eight physical queues per port, while other Extreme switches can be configured with up to four physical queues per port.

 As with all Extreme Switch products, Policy-Based QoS has zero impact on switch performance. Using even the most complex traffic groupings is “costless” in terms of switch performance.

Policy-Based QoS can be configured to perform per-port Random Early Detection (RED) and drop-probability. Using this capability, the switch detects when traffic is filling up in any of the eight hardware queues, and performs a random discard on subsequent packets, based on the configured RED drop-probability.

Instead of dropping sessions during times when the queue depth is exceeded, RED causes the switch to lower session throughput. The destination node detects the dropped packet, and, using standard TCP windowing mechanisms, slows the transmission from the source node. RED drop-probability is configured on a system-wide basis, and has a valid range from 0% to 100%. Only switches and modules with the “i” chipset can use RED.

APPLICATIONS AND TYPES OF QoS

Different applications have different QoS requirements. The following applications are ones that you will most commonly encounter and need to prioritize:

- Voice applications
- Video applications
- Critical database applications
- Web browsing applications
- File server applications

General guidelines for each traffic type are given below and summarized in Table 8-1. Consider them as general guidelines and not strict recommendations. Once QoS parameters are set, you can monitor the performance of the application to determine if the actual behavior of the applications matches your expectations. It is very important to understand the needs and behavior of the particular applications you wish to protect or limit. Behavioral aspects to consider include bandwidth needs, sensitivity to latency and jitter, and sensitivity and impact of packet loss.

VOICE APPLICATIONS

Voice applications typically demand small amounts of bandwidth. However, the bandwidth must be constant and predictable because voice applications are typically sensitive to latency (inter-packet delay) and jitter (variation in inter-packet delay). The most important QoS parameter to establish for voice applications is minimum bandwidth, followed by priority.

VIDEO APPLICATIONS

Video applications are similar in needs to voice applications, with the exception that bandwidth requirements are somewhat larger, depending on the encoding. It is important to understand the behavior of the video application being used. For example, in the playback of stored video streams, some applications may transmit large amounts of data for multiple streams in one “spike,” with the expectation that the end-stations will buffer significant amounts of video-stream data. This can present a problem to the network infrastructure, because it must be capable of buffering the transmitted spikes where there are speed differences (for example, going from Gigabit Ethernet to Fast Ethernet). Key QoS parameters for video applications include minimum bandwidth, priority, and possibly buffering (depending upon the behavior of the application).

CRITICAL DATABASE APPLICATIONS

Database applications, such as those associated with ERP, typically do not demand significant bandwidth and are tolerant of delay. You can establish a minimum bandwidth using a priority less than that of delay-sensitive applications.

WEB BROWSING APPLICATIONS

QoS needs for Web browsing applications cannot be generalized into a single category. For example, ERP applications that use a browser front-end may be more important than retrieving daily news information. Traffic groupings can typically be distinguished from each other by their server source and destinations. Most browser-based applications are distinguished by the dataflow being asymmetric (small dataflows from the browser client, large dataflows from the server to the browser client).

An exception to this may be created by some JavaTM-based applications. In addition, Web-based applications are generally tolerant of latency, jitter, and some packet loss, however small packet-loss may have a large impact on perceived performance due to the nature of TCP. The relevant parameter for protecting browser applications is minimum bandwidth. The relevant parameter for preventing non-critical browser applications from overwhelming the network is maximum bandwidth. In addition, RED may be used to reduce session loss if the queue that floods web traffic becomes over-subscribed.

FILE SERVER APPLICATIONS

With some dependencies on the network operating system, file serving typically poses the greatest demand on bandwidth, although file server applications are very tolerant of latency, jitter, and some packet loss, depending on the network operating system and the use of TCP or UDP.



Full-duplex links should be used when deploying policy-based QoS. Half-duplex operation on links can make delivery of guaranteed minimum bandwidth impossible.

Table 9-1 summarizes QoS guidelines for the different types of network traffic.

Table 9-1: Traffic Type and QoS Guidelines

Traffic Type	Key QoS Parameters
Voice	Minimum bandwidth, priority
Video	Minimum bandwidth, priority, buffering (varies)
Database	Minimum bandwidth
Web browsing	Minimum bandwidth for critical applications, maximum bandwidth for non-critical applications, RED
File server	Minimum bandwidth

ASSIGNING QoS ATTRIBUTES

Assigning QoS attributes is a three-step process which consists of defining three interrelated QoS building blocks (defined below):

- Create a QoS profile.

QoS profile — A class of service that is defined through minimum and maximum bandwidth parameters, configuration of buffering and RED, and prioritization settings. The bandwidth and level of service that a particular type of traffic or “traffic grouping” receives is determined by assigning it to a QoS profile.

- Assign one or more traffic groupings to a QoS profile to create a QoS policy.

Traffic grouping — A classification or traffic type that has one or more attributes in common. These can range from a physical port to a VLAN to IP Layer 4 port information. Traffic groupings are assigned to QoS profiles to modify switch forwarding behavior. Traffic groupings transmitting out the same port that are assigned to a particular QoS profile share the assigned bandwidth and prioritization characteristics, and hence share the class of service.

QoS policy — The combination that results from assigning a traffic grouping to a QoS profile.

- Monitor the performance of the application with the QoS monitor to determine whether the policies are meeting the desired results.

The next sections describe each of these QoS components in detail.

QoS PROFILES

A QoS profile defines a class of service by specifying traffic behavior attributes, such as bandwidth. The parameters that make up a QoS profile include the following:

- **Minimum bandwidth** – The minimum percentage of total link bandwidth that is reserved for use by a hardware queue on a physical port. Bandwidth unused by the queue can be used by other queues. The minimum bandwidth for all queues should add up to less than 90%. The default value on all minimum bandwidth parameters is 0%.
- **Maximum bandwidth** – The maximum percentage of total link bandwidth that may be transmitted by a hardware queue on a physical port. The default value on all maximum bandwidth parameters is 100%.
- **Priority** – The level of priority assigned to a hardware queue on a physical port. Switch products that use the “*i*” chipset have eight different available priority settings. Other Extreme switches have four available priority settings. By default, each of the default QoS profiles is assigned a unique priority. You would use prioritization when two or more hardware queues on the same physical port are contending for transmission on the same physical port, only *after* their respective bandwidth management parameters have been satisfied. If two hardware queues on the same physical port have the same priority, a round-robin algorithm is used for transmission, depending on the available link bandwidth.
 - When configured to do so, the priority of a QoS profile may determine the 802.1p bits used in the priority field of a transmitted packet (described later).
 - On switch products using the “*i*” chipset, the priority of a QoS profile determines the DiffServ code point value used in an IP packet when the packet is transmitted (described later).
- **Buffer** – This parameter reserves buffer memory for use exclusively by a QoS profile across all affected ports. The default value for buffer settings is 0%. The sumvalue of all QoS profile buffer parameters should not exceed 100%. Reserving buffer memory for a QoS profile affects the dynamic buffer space available to other QoS profiles. You should not modify the buffer parameter unless specific situations and application behavior indicate.

A QoS profile does not alter the behavior of the switch until it is assigned to a traffic grouping. Recall that QoS profiles are linked to hardware queues. There are multiple hardware queues per physical port. By default, a QoS profile links to the identical hardware queue across all the physical ports of the switch.

Four or eight default QoS profiles are provided, depending on the chipset used in the switch. The default QoS profiles cannot be deleted. Also by default, a QoS profile maps directly to a specific hardware queue across all physical ports. The settings for the default QoS profiles for Summit chipset products are summarized in Table 9-2. The settings for the default QoS profiles for “*i*” chipset products are summarized in [Table 9-3](#).

Table 9-2: Default QoS Profiles for Summit Chipset Products

Profile Name	Hardware Queue	Priority	Buffer	Minimum Bandwidth	Maximum Bandwidth
Qp1	Q0	Low	0	0%	100%
Qp2	Q1	Normal	0	0%	100%
Qp3	Q2	Medium	0	0%	100%
Qp4	Q3	High	0	0%	100%

Table 9-3: Default QoS Profiles for “*i*” Chipset Products

Profile Name	Hardware Queue	Priority	Buffer	Minimum Bandwidth	Maximum Bandwidth
Qp1	Q0	Low	0	0%	100%
Qp2	Q1	Lowhi	0	0%	100%
Qp3	Q2	Normal	0	0%	100%
Qp4	Q3	Normalhi	0	0%	100%
Qp5	Q4	Medium	0	0%	100%
Qp6	Q5	Mediumhi	0	0%	100%
Qp7	Q6	High	0	0%	100%
Qp8	Q7	Highhi	0	0%	100%

CONFIGURING A QoS PROFILE

Table 9-4 lists the commands used to configure QoS.

Table 9-4: QoS Configuration Commands

Command	Description
config ports <portlist> qosprofile <qosprofile>	Configures one or more ports to use a particular QoS profile. Available only in ingress mode.
config qosprofile <qosprofile> {minbw <percent>} {maxbw <percent>} {priority <level>} {buffer <percent>} {<portlist>}	Configures a QoS profile. Specify: <ul style="list-style-type: none"> ■ minbw — The minimum bandwidth percentage guaranteed to be available to this queue for transmission. The default setting is 0. ■ maxbw — The maximum bandwidth percentage this queue is permitted to use for transmission. The default setting is 100. ■ priority — The service priority for this queue. Settings include low, normal, medium, and high. The default setting is low. Available only in egress mode.
config red drop-probability <percent>	Configures the Random Early Detect (RED) drop-probability. The percentage range is 0 - 100.
config vlan <name> qosprofile <qosprofile>	Allows you to configure a VLAN to use a particular QoS profile.
create qosprofile <qosprofile>	Creates a QoS profile.
delete qosprofile <qosprofile>	Deletes a QoS profile.
disable red ports	Disables RED on one or all ports.
enable red port <mgmt portlist>	Enables RED on a port.

TRAFFIC GROUPINGS AND CREATING A QOS POLICY

Once a QoS profile is modified for bandwidth and priority, you assign the profile to a particular traffic grouping. A QoS profile is assigned to a specific traffic grouping to create a QoS policy. A *traffic grouping* is a classification of traffic that has one or more attributes in common.

Traffic groupings are separated into the following categories for discussion:

- IP-based information, such as IP source/destination and TCP/UDP port information
- Destination MAC (MAC QoS groupings)
- Explicit packet class of service information, such as 802.1p or DiffServ (IP TOS)
- Physical/logical configuration (physical source port or VLAN association)

In the event that a given packet matches two or more grouping criteria, there is a predetermined precedence for which traffic grouping will apply. In general, the more specific traffic grouping takes precedence. By default, all traffic groupings are placed in the QoS profile Qp1. The supported traffic groupings are listed in [Table 9-5](#). The groupings are listed in order of precedence (highest to lowest). The four types of traffic groupings are described in detail on the following pages.

Table 9-5: Traffic Groupings by Precedence

IP Information (Access Lists) Groupings

- Access list precedence determined by user configuration
-

Destination Address MAC-based Groupings

- Permanent
 - Dynamic
 - Blackhole
 - Broadcast/unknown rate limiting
-

Explicit Packet Class of Service Groupings

- DiffServ (IP TOS)
 - 802.1P
-

Physical/Logical Groupings

- Source port
 - VLAN
-

IP-BASED TRAFFIC GROUPINGS

IP-based traffic groupings are based on any combination of:

- IP source or destination address
- TCP/UDP or other Layer 4 protocol
- TCP/UDP port information

IP-based traffic groupings are defined using access lists. Access lists are discussed in detail in [Chapter 16](#). By supplying a named QoS profile at the end of the access list command syntax, you can prescribe the bandwidth management and priority handling for that traffic grouping. This level of packet filtering has no impact on performance.

MAC-BASED TRAFFIC GROUPINGS

QoS profiles can be assigned to destination MAC addresses. MAC-based traffic groupings are configured using the following command:

```
create fdbentry <mac_address> vlan <name> [blackhole | port <portlist>
| dynamic] qosprofile <qosprofile>
```

The MAC address options, defined below, are as follows:

- Permanent
- Dynamic
- Blackhole
- Broadcast/unknown rate limiting

PERMANENT MAC ADDRESSES

Permanent MAC addresses can be assigned a QoS profile whenever traffic is destined to the MAC address. This can be done when you create a permanent (spell out) (FDB) entry. For example:

```
create fdbentry 00:11:22:33:44:55 vlan default port 4:1 qosprofile qp2
```

DYNAMIC MAC ADDRESSES

Dynamic MAC addresses can be assigned a QoS profile whenever traffic is destined to the MAC address. For any port on which the specified MAC address is learned in the specified VLAN, the port is assigned the specified QoS profile. For example:

```
create fdbentry 00:11:22:33:44:55 vlan default dynamic qosprofile qp3
```

The QoS profile is assigned when the MAC address is learned. If a client's location moves, the assigned QoS profile moves with the device. If the MAC address entry already exists in the FDB, you can clear the forwarding database so that the QoS profile can be applied when the entry is added again. The command to clear the FDB is as follows:

```
clear fdb
```

BLACKHOLE MAC ADDRESS

Using the `blackhole` option configures the switch to not forward any packets to the destination MAC address on any ports for the VLAN specified. The `blackhole` option is configured using the following command:

```
create fdbentry 00:11:22:33:44:55 vlan default blackhole
```

BROADCAST/UNKNOWN RATE LIMITING MAC ADDRESS

It is possible to assign broadcast and unknown destination packets to a QoS profile that has the desired priority and bandwidth parameters. Broadcast/unknown rate limiting is an extension of the QoS feature used for destination MAC addresses.

For example, if you want to limit broadcast and unknown traffic on the VLAN `default` to the bandwidth and priority defined in QoS profile `qp3`, the command is:

```
create fdbentry ff:ff:ff:ff:ff:ff vlan default dynamic qp3
```



IP multicast traffic is subject to broadcast and unknown rate limiting only when IGMP snooping is disabled.

VERIFYING MAC-BASED QoS SETTINGS

To verify any of the MAC-based QoS settings, use either the command

```
show fdb perm
```

or the command

```
show qosprofile <qosprofile>
```

EXPLICIT CLASS OF SERVICE (802.1P AND DIFFSERV) TRAFFIC GROUPINGS

This category of traffic groupings describes what is sometimes referred to as *explicit packet marking*, and refers to information contained within a packet intended to explicitly determine a class of service. That information includes:

- IP DiffServ code points, formerly known as IP TOS bits
- Prioritization bits used in IEEE 802.1p packets

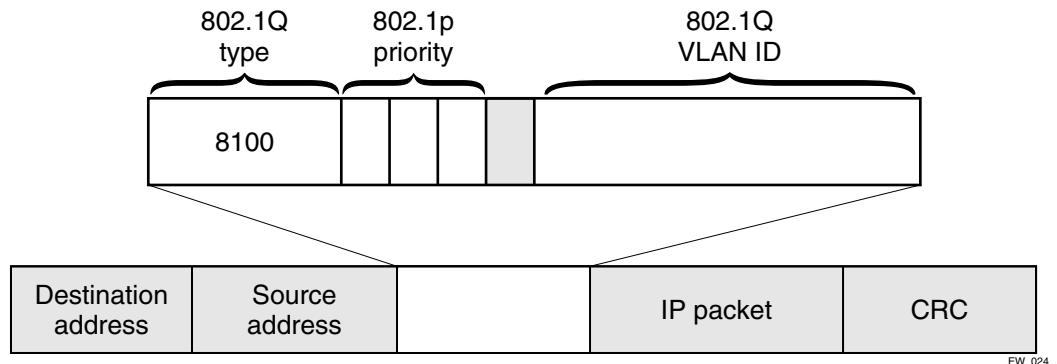
An advantage of explicit packet marking is that the class of service information can be carried throughout the network infrastructure, without repeating what may be complex traffic grouping policies at each switch location. Another advantage is that end stations can perform their own packet marking on an application-specific basis. Extreme switch products have the capability of observing and manipulating packet marking information with no performance penalty.

Extreme products that use “*i*” chipset support DiffServ capabilities. Products that do not use the “*i*” chipset s do not support DiffServ capabilities. The documented capabilities for 802.1p priority markings or DiffServ capabilities (if supported) are not impacted by the switching or routing configuration of the switch. For example, 802.1p information may be preserved across a routed switch boundary and DiffServ code points may be observed or overwritten across a layer 2 switch boundary.

CONFIGURING 802.1P PRIORITY

Extreme switches support the standard 802.1p priority bits that are part of a tagged Ethernet packet. The 802.1p bits can be used to prioritize the packet, and assign it to a particular QoS profile.

When a packet arrives at the switch, the switch examines the 802.1p priority field maps it to a specific hardware queue when subsequently transmitting the packet. The 802.1p priority field is located directly following the 802.1Q type field, and preceding the 802.1Q VLAN ID, as shown in [Figure 9-1](#).

**Figure 9-1:** Ethernet packet encapsulation**OBSERVING 802.1P INFORMATION**

When ingress traffic that contains 802.1p prioritization information is detected by the switch, the traffic is mapped to various hardware queues on the egress port of the switch. Switches that use the “*i*” chipset support eight hardware queues, all other products support four hardware queues. The transmitting hardware queue determines the bandwidth management and priority characteristics used when transmitting packets.

To control the mapping of 802.1p prioritization values to hardware queues, 802.1p prioritization values may be mapped to a QoS profile. The default mapping of each 802.1p priority value to QoS profile is shown in [Table 9-6](#).

Table 9-6: 802.1p Priority Value-to-QoS Profile Default Mapping

Priority Value	QoS Profile Summit Chipset	QoS Profile “ <i>i</i> ” Chipset
0	Qp1	Qp1
1	Qp1	Qp2
2	Qp2	Qp3
3	Qp2	Qp4
4	Qp3	Qp5
5	Qp3	Qp6
6	Qp4	Qp7
7	Qp4	Qp8

802.1P COMMANDS

[Table 9-7](#) shows the commands used to configure 802.1p priority. Two are explained in more detail in the following paragraphs.

Table 9-7: 802.1p Configuration Commands

Command	Description
config dot1p type <dot1p_priority> qosprofile <qosprofile>	Configures the default QoS profile to 802.1p priority mapping. The value for dot1p_priority is an integer between 0 and 7.
disable dot1p replacement ports [<portlist> all]	Disables the ability to overwrite 802.1p priority values for a given set of ports.
enable dot1p replacement ports [<portlist> all]	Enables the 802.1p priority field to be overwritten on egress according to the QoS profile to 802.1p priority mapping for a given set of ports.
show dot1p	Displays the 802.1p-to-QoS profile mappings.

CHANGING THE DEFAULT 802.1P MAPPING

By default, a QoS profile is mapped to a hardware queue, and each QoS profile has configurable bandwidth parameters and priority. In this way, an 802.1p priority value seen on ingress can be mapped to a particular QoS profile and with specific bandwidth management and priority behavior.

To change the default mappings of QoS profiles to 802.1p priority values, use the command:

```
config dot1p type <dot1p_priority> qosprofile <qosprofile>
```

REPLACING 802.1P PRIORITY INFORMATION

By default, 802.1p priority information is not replaced or manipulated, and the information observed on ingress is preserved when transmitting the packet. This behavior is not affected by the switching or routing configuration of the switch.

However, the switch is capable of inserting and/or overwriting 802.1p priority information when it transmits an 802.1Q tagged frame. If 802.1p replacement is enabled, the 802.1p priority information that is transmitted is determined by the hardware queue that is used when transmitting the packet. To replace 802.1p priority information, use the command:

```
enable dot1p replacement ports [<portlist> | all]
```

802.1p priority information is replaced according to the hardware queue that is used when transmitting from the switch. The mapping is described in [Table 9-8](#) for switches based on the “*i*” chipset and for other Extreme switches. This mapping cannot be changed.

Table 9-8: Queue to 802.1p Priority Replacement Value

Hardware Queue Summit Chipset	Hardware Queue “ <i>i</i> ” Chipset	802.1p Priority Replacement Value
Q0	Q0	0
	Q1	1
Q1	Q2	2
	Q3	3
Q2	Q4	4
	Q5	5
Q3	Q6	6
	Q7	7

CONFIGURING DIFFSERV

Contained in the header of every IP packet is a field for IP Type of Service (TOS), now also called the DiffServ field. The TOS field is used by the switch to determine the type of service provided to the packet.

Observing DiffServ code points as a traffic grouping mechanism for defining QoS policies and overwriting the Diffserv code point fields are supported in switches using the “*i*” chipset.

[Figure 9-2](#) shows the encapsulation of an IP packet header.

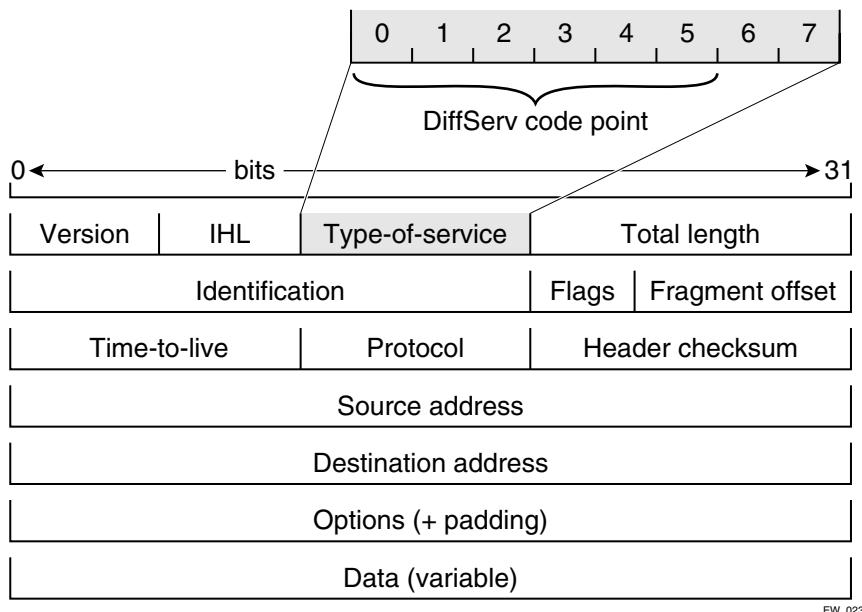
**Figure 9-2:** IP packet header encapsulation

Table 9-9 lists the commands used to configure DiffServ. Some of the commands are described in more detail in the following paragraphs.

Table 9-9: DiffServ Configuration Commands

Command	Description
config diffserv examination code-point <code_point> qosprofile <qosprofile> ports [<portlist> all]	Configures the default ingress diffserv code points to QoS profile mapping. The <code_point> is a 6-bit value in the IP-TOS byte in the IP header. You can specify up to 64 different code points for each port.
config diffserv replacement priority <vpri> code_point <code_point> ports [<portlist> all]	Configures the default egress diffserv replacement mapping.
disable diffserv examination ports [<portlist> all]	Disables the examination of the diffserv field in an IP packet.
disable diffserv replacement ports [<portlist> all]	Disables the replacement of diffserv code points in packets transmitted by the switch.

Table 9-9: DiffServ Configuration Commands (continued)

Command	Description
enable diffserv examination ports [<portlist> all]	Enables the diffserv field of an ingress IP packet to be examined by the switch in order to select a QoS profile. The default setting is disabled.
enable diffserv replacement ports [<portlist> all]	Enables the diffserv code point to be overwritten in packets transmitted by the switch. Eight user-defined code points can be configured on each port. The 802.1P priority bits (3-bits) are used to select one of the eight code points. The default setting is disabled.
unconfig diffserv examination ports [<portlist> all]	Removes the diffserv examination code point from a port.
unconfig diffserv replacement ports [<portlist> all]	Removes the diffserv replacement mapping from a port.

OBSERVING DIFFSERV INFORMATION

When a packet arrives at the switch on an ingress port, the switch examines the first six of eight TOS bits, called the *code point*. The switch can assign the QoS profile used to subsequently transmit the packet based on the code point. The QoS profile controls a hardware queue used when transmitting the packet out of the switch, and determines the forwarding characteristics of a particular code point. Viewing DiffServ information can be enabled or disabled; by default it is disabled. To view DiffServ information, use the command:

```
enable diffserv examination ports [<portlist> | all]
```

CHANGING DIFFSERV CODE POINT ASSIGNMENTS IN THE QoS PROFILE

Because the code point uses six bits, it has 64 possible values ($2^6 = 64$). By default, the values are grouped and assigned to the default QoS profiles listed in [Table 9-10](#).

Table 9-10: Default Code Point-to-QoS Profile Mapping

Code Point	QoS Profile
0-7	Qp1
8-15	Qp2
16-23	Qp3
24-31	Qp4

Table 9-10: Default Code Point-to-QoS Profile Mapping (continued)

Code Point	QoS Profile
32-39	Qp5
40-47	Qp6
48-55	Qp7
56-63	Qp8

You can change the QoS profile assignment for all 64 code points using the following command:

```
config diffserv examination code-point <code_point> qosprofile
<qosprofile> ports [<portlist> | all]
```

Once assigned, the rest of the switches in the network prioritize the packet using the characteristics specified by the QoS profile.

REPLACING DIFFSERV CODE POINTS

The switch can be configured to change the DiffServ code point in the packet prior to the packet being transmitted by the switch. This is done with no impact on switch performance.

The DiffServ code point value used in overwriting a packet is determined by the 802.1p priority value. The 802.1p priority value is, in turn, determined by the hardware queue used when transmitting a packet, as described in the section “[Replacing 802.1p Priority Information](#).”

It is not necessary to receive or transmit 802.1Q tagged frames, only to understand that the egress hardware queue, which also determines the 802.1p priority value, can also be configured to determine the DiffServ code point value if you want to replace the DiffServ code points.

To replace DiffServ code points you must enable both 802.1p replacement and DiffServ replacement using the following commands:

```
enable dot1p replacement ports [<portlist> | all]
enable diffserv replacement ports [<portlist> | all]
```

The default 802.1p priority value to code point mapping is described in [Table 9-11](#).

Table 9-11: Default 802.1p Priority Value-to-Code Point Mapping

Hardware Queue “i” Chipset	802.1p Priority value	Code Point
Q0	0	0
Q1	1	8
Q2	2	16
Q3	3	24
Q4	4	32
Q5	5	40
Q6	6	48
Q7	7	56

You then change the 802.1p priority to DiffServ code point mapping to any code point value using the following command:

```
config diffserv replacement priority <vpri> code_point <code_point>
ports [<portlist> | all]
```

By doing so, the hardware queue used to transmit a packet determines the DiffServ value replaced in the IP packet.

To verify the DiffServ configuration, use the command:

```
show ports <portlist> info {detail}
```

DIFFSERV EXAMPLE

In this example, we use DiffServ to signal a class of service throughput and assign any traffic coming from network 10.1.2.x with a specific DiffServ code point. This allows all other network switches to send and observe the Diffserv code point instead of repeating the same QoS policy on every network switch.

Configure the switch that handles incoming traffic from network 10.1.2.x as follows:

- 1 Configure parameters of the QoS profile QP3:

```
config qp3 min 10 max 100
```

- 2 Assign a traffic grouping for traffic from network 10.1.2.x to qp3:

```
create access-list TenOneTwo
config TenOneTwo 10.1.2.0/24 permit qp3
```

- 3 To enable the switch to overwrite the DiffServ code point:

```
enable dot1p replacement  
enable diffserv replacement
```

- 4 Configure the switch so that other switches may signal class of service that this switch should observe:

```
enable diffserv examination
```

Table 9-3 indicates that qp3 is tied to hardware queue Q2. We also know that when replacement is enabled all traffic sent out Q2 will contain code point value 16 (according to [Table 9-11](#)). If this is the desired code point to use, all traffic from 10.1.2.x will be sent out QP3 (at 10% minimum and 100% maximum) with a code point value of 16.

PHYSICAL AND LOGICAL GROUPINGS

Two traffic groupings exist in this category:

- Source port
- VLAN

SOURCE PORT

A source port traffic grouping implies that any traffic sourced from this physical port uses the indicated QoS profile when the traffic is transmitted out to any other port. To configure a source port traffic grouping, use the following command:

```
config ports <portlist> qosprofile <qosprofile>
```

In the following BlackDiamond switch example, all traffic sourced from slot 5 port 7 uses the QoS profile named *qp3* when being transmitted.

```
config ports 5:7 qosprofile qp3
```

VLAN

A VLAN traffic grouping indicates that all intra-VLAN switched traffic and all routed traffic sourced from the named VLAN uses the indicated QoS profile. To configure a VLAN traffic grouping, use the following command:

```
config vlan <name> qosprofile <qosprofile>
```

For example, all devices on VLAN *servnet* require use of the QoS profile *qp4*. The command to configure this example is as follows:

```
config vlan servnet qosprofile qp4
```

VERIFYING PHYSICAL AND LOGICAL GROUPINGS

To verify settings on ports or VLANs, use the command:

```
show qosprofile <qosprofile>
```

The same information is also available for ports or VLANs using the command:

```
show ports info
```

or

```
show vlan
```

VERIFYING CONFIGURATION AND PERFORMANCE

Once you have created QoS policies that manage the traffic through the switch, you can use the QoS monitor to determine whether the application performance meets your expectations.

QoS MONITOR

The QOS monitor is a utility that monitors the hardware queues associated with any port(s). The QOS monitor keeps track of the number of frames and the frames per second that a specific queue is responsible for transmitting on a physical port. Two options are available: a real-time display, and a separate option for retrieving information in the background and writing it to the log.

[Table 9-12](#) lists the QoS monitor commands. Some of the commands are described in more detail in the following paragraphs.

Table 9-12: QoS Monitor Commands

Command	Description
disable qosmonitor	Disables the QoS monitoring capability.
enable qosmonitor {port <port>}	Enables the QoS monitoring capability on the switch. When no port is specified, the QoS monitor automatically samples all the ports. Error messages are logged to the syslog if the traffic exceeds the parameters of the QoS profile(s). The default setting is disabled.
show ports {<portlist>} qosmonitor	Displays real-time QoS statistics for one or more ports.

REAL-TIME PERFORMANCE MONITORING

The real-time display scrolls through the given portlist to provide statistics. You can choose screens for packet count and packets per second. The specific port being monitored is indicated by an asterisk (*) appearing after the port number in the display.

To view real-time switch per-port performance, use the command:

```
show ports {<portlist>} qosmonitor
```

QoS monitor sampling is configured as follows:

- The port is monitored for 20 seconds before the switch moves on to the next port in the list.
- A port is sampled for five seconds before the packets per second (pps) value is displayed on the screen.

BACKGROUND PERFORMANCE MONITORING

Monitoring QoS in the background places transmit counter and any “overflow” information into the switch log. The log notification appears if one of the queues experiences an overflow condition since the last time it was sampled.

An overflow entry indicates that a queue was over-subscribed at least temporarily, and is useful for determining correct QoS settings and potential over-subscription issues.

DISPLAYING QoS PROFILE INFORMATION

The QoS monitor can also be used to verify the QoS configuration and monitor the use of the QoS policies that are in place. To display QoS information on the switch, use the following command:

```
show qosprofile <qosprofile>
```

Displayed information includes:

- QoS profile name
- Minimum bandwidth
- Maximum bandwidth
- Priority
- A list of all traffic groups to which the QoS profile is applied

Additionally, QoS information can be displayed from the traffic grouping perspective by using one or more of the following commands:

- `show fdb permanent` — Displays destination MAC entries and their QoS profiles.
- `show switch` — Displays information including PACE enable/disable information.
- `show vlan` — Displays the QoS profile assignments to the VLAN.
- `show ports info {detail}` — Displays information including QoS information for the port.

MODIFYING A QoS POLICY

If you make a change to the parameters of a QoS profile after a QoS policy has already been formed (by applying a QoS profile to a traffic grouping), the timing of the configuration change depends on the traffic grouping involved. To have a change in QoS profile effect a change in the QoS policy, the following rules apply:

- For destination MAC-based grouping (other than permanent), clear the MAC FDB using the command `clear fdb`. This command should also be issued after a policy is first formed, as the policy must be in place before an entry is made in the MAC FDB. For permanent destination MAC-based grouping, re-apply the QoS profile to the static FDB entry, as documented. You can also save and reboot the switch.
- For physical and logical groupings of a source port or VLAN, re-apply the QoS profile to the source port or VLAN, as documented. You can also save and reboot the switch.

INTRA-SUBNET QoS

Intra-Subnet QoS™ (ISQ) is used only on Extreme switches that do not use the “i” chipset. Using ISQ, it is possible to apply Layer 3 and Layer 4 access lists to traffic that is only being locally switched. Extreme products that use the “i” chipset are already capable of using Layer 3 and Layer 4 access lists without enabling ISQ, even though the switch is performing Layer 2 switching, only. The command syntax for all IP-related access list commands is described in [Chapter 16](#). ISQ is enabled on a per-VLAN basis.

Because ISQ instructs the switch to look at IP addresses within a VLAN, the normal MAC-based learning and refreshing for Layer 2 switching is altered for traffic that matches an IP-based traffic grouping. Instead, learning and refreshing is performed based on IP information in the packets. As a result, the FDB aging timer is automatically increased to a value comfortably above a normal ARP table refresh time; to 50 minutes (3,000 seconds). ISQ should not be used on VLANs with clients that have statically-defined ARP tables.

ISQ may also be used for the application of Layer 3 and Layer 4 QoS policies for traffic on a switch that is destined outside the subnet served by the Layer 2-only switch. This is a useful feature in conjunction with ExtremeWare Enterprise Manager (EEM) v2.0 or above and the Dynamic Link Context System (DLCS) feature documented below. To configure this capability, you need the MAC address of the next-hop router (or the MAC address of the WINS server, if the server is on the same subnet), and the list of the host IP addresses. The IP packets to this MAC address are then snooped.

ISQ commands are described in [Table 9-13](#).

Table 9-13: ISQ Configuration Commands

Command	Description
config isq-server <servers-listname> add ipaddress <remote destination ipaddress1>	Adds a remote destination IP address.
config isq-server <servers-listname> add mac <mac-address-of-next-hop> vlan <vlan name>	Adds the MAC address of the next hop router.
config isq-server <servers-listname> delete ipaddress <remote destination ipaddress1>	Deletes a remote destination IP address.
config isq-server <servers-listname> delete mac <mac-address-of-next-hop> vlan <vlan name>	Deletes the MAC address of the next hop router.

Table 9-13: ISQ Configuration Commands (continued)

Command	Description
create isq-server <servers-listname>	Creates a remote destination that should be snooped.
delete isq-server <servers-listname>	Deletes a remote destination.
disable isq <vlan name>	Disables ISQ.
enable isq <vlan name>	Enables ISQ.

Bi-DIRECTIONAL RATE SHAPING

Bi-directional rate shaping allows you to perform bandwidth management for traffic flowing both to and from devices attached to the switch. Bi-directional control is provided by defining minimum and maximum bandwidth parameters in order to build true “committed information rate” capabilities. These rates can be symmetric or asymmetric and use up to 8 queues and classes of service in both directions. All traffic grouping and bandwidth management capabilities associated with QoS can be used for both directions of traffic.

During configuration, an additional “loopback” port is defined on each VLAN for the purpose of queuing and bandwidth management. Ports added to the VLAN can then be configured to use the loopback port for rate-shaping and bandwidth management on ingress traffic. To configure QoS on ingress traffic, configure QoS normally on the loopback port. For example, a maximum bandwidth parameter and traffic grouping associated with a QoS profile on the loopback port will define a rate limit for matching ingress traffic on the VLAN’s ports configured to use the loopback port.

Bi-DIRECTIONAL RATE SHAPING PROPERTIES

Use the following guidelines for Bi-Directional Rate Shaping:

- A VLAN must also include a loopback port that is unused by an external device. This loopback port must be configured with a unique 802.1Q VLAN tag ID.
- Ports within the VLAN configured to use the loopback port will have ingress rate-shaping performed. Ports not configured to use the loopback port will not have ingress rate-shaping performed. Traffic from a rate-shaped port to a non-rate-shaped port within the VLAN is not subject to rate-shaping.
- The aggregate forwarding of all rate-shaped ports in a VLAN is determined by the setting of the queue parameters of the loopback port.

- The granularity of the maximum bandwidth settings are rounded up to one of the following percentage values: 5, 6, 8, 10, 12, 20, 30, 40, 55, 70, 85, 95, 100. In combination with port speed, this granularity translates to: 500, 600, 800 Kbps and 1, 2, 2.5, 3, 5, 6, 7.5, 8, 9.5, 10, 20, 25, 30, 50, 60, 75, 80, 95, 100, 200, 250, 300, 500, 600, 750, 950 Mbps.
- For 10/100 ports, the loopback port can be configured as a 10 Mbps port to achieve lower bandwidth values.

Bi-DIRECTIONAL RATE SHAPING LIMITATIONS

Consider the following limitations when configuring Bi-Directional Rate Shaping:

- All rate-shaped ports must be deleted before deleting the loopback port.
- If rate-shaped ports within a VLAN use differing bandwidth parameters, set the priority of all the QoS profiles on the loopback port and rate-shaped ports to low.
- Rate-shaping only takes effect with a single VLAN (L2) and cannot span multiple VLANs.
- On the BlackDiamond, the loopback port must be on the same I/O module as the rate-shaped ports.

Bi-DIRECTIONAL RATE SHAPING COMMANDS

To configure the loopback port in the VLAN, use the following command:

```
configure [vlan] <vlan_name> add port <port> loopback-vid <vlan_tag>
```

To enable the loopback port in the VLAN, use the following command:

```
Restart port <loopback_port>
```

To add rate-shaped ports to the VLAN, use the following command:

```
configure [vlan] vlan_name add [port] <portlist> soft-rate-limit
```

To delete rate-shaped ports from the VLAN, use the following command:

```
configure [vlan] vlan_name delete [port] portlist
```

To configure the rate-shaping parameters of the loopback port, use the normal QoS profile configuration command, as follows:

```
configure qosprofile <qosprofile_name> minbw <%> maxbw <%> priority
<low> <loopback port number>
```

To display Bi-Directional Rate Shaping configuration, use the following commands:

```
show <vlan> vlan_name
show vlan detail
```

Rate-shaped ports are displayed with an “R” and loopback ports are displayed with an “L” next to the port number in the show screen.

To set the port speed of a loopback port, use the normal port configuration command:

```
configure port <portlist> auto <off | on> speed <100 | 100> duplex <off
| on>
restart port <portlist>
```

DYNAMIC LINK CONTEXT SYSTEM

The Dynamic Link Context System (DLCS) is a feature that snoops WINS NetBIOS packets and creates a mapping between a user name, the IP address or MAC address, and the switch/port. Based on the information in the packet, DLCS can detect when an end station boots up or a user logs in or out, and dynamically maps the end station name to the current IP address and switch/port. This information is available for use by ExtremeWare Enterprise Manager (EEM) version 2.1, or later, in setting policies that may be applied to users and may dynamically follow a user's location. DLCS provides you with valuable information on a user's location and associated network attributes. For DLCS to operate within ExtremeWare, the user or end station must allow for automatic DLCS updates.

This feature should only be used in conjunction with the ExtremeWare Enterprise Manager Policy System. Refer to the ExtremeWare Enterprise Manager 2.0 documentation for more information.

DLCS GUIDELINES

Follow these guidelines when using DLCS:

- Only one user is allowed on one workstation at a given time.
- A user may be logged into many workstations simultaneously.
- An IP-address can be learned on only one port in the network at a given time.
- Multiple IP-addresses can be learned on the same port.
- DLCS mapping is flushed when a user logs in or logs out, or when an end-station is shutdown.

DLCS LIMITATIONS

Consider the following limitations concerning data received from WINS snooping:

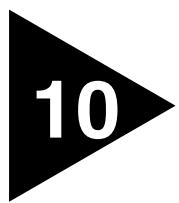
- DLCS does not work for the WINS server. This is because the WINS server does not send NETBIOS packets on the network (these packets are address to itself).
- When the IP address of a host is changed, and the host is not immediately rebooted, the old host-to-IP address mapping is never deleted. You must delete the mapping of the host-to-IP address through the EEM Policy Manager.
- When the host is moved from one port to another port on a switch, the old entry does not age out unless the host is rebooted or a user login operation is performed after the host is moved.
- DLCS information is dynamic, therefore, if the switch is rebooted, the information is lost. This information is still stored in the policy-server. To delete the information from the policy system, you must explicitly delete configuration parameters from the EEM Policy Applet user interface. As a workaround, you may delete the switch that was rebooted from the list of managed devices in the EEM Inventory Applet, and re-add the switch to the Inventory Manager.
- DLCS is not supported on hosts that have multiple NIC cards.
- IPQoS is not supported to a WINS server that is serving more than one VLAN. If you attempt to add a WINS server to serve more than one VLAN, and there are IPQoS rules defined for that server, the command to add the WINS server is rejected.

DLCS COMMANDS

The DLCS commands are described in [Table 9-14](#).

Table 9-14: DLCS Configuration Commands

Command	Description
clear dlcs	Clears learned DLCS data.
config isq-server <wins-servers-listname> add ipaddress <wins-server1>	Adds a WINS servername and IP address.
config isq-server <wins-servers-listname> add mac <mac-address-of-next-hop> vlan <vlan name>	Adds the MAC address of the next hop router.
config isq-server <wins-servers-listname> delete ipaddress <wins-server1>	Deletes a WINS servername and IP address.
config isq-server <wins-servers-listname> delete mac <mac-address-of-next-hop> vlan <vlan name>	Deletes the MAC address of the next hop router
create isq-server <wins-servers-listname>	Creates a WINS server to be snooped.
delete isq-server <wins-servers-listname>	Deletes a WINS server from being snooped.
disable dlcs	Disables snooping of DLCS packets.
disable dlcs ports <port-number>	Disables port on which DLCS packets are snooped.
enable dlcs	Enables snooping of DLCS packets.
enable dlcs ports <port-number>	Enables port on which DLCS packets are snooped.
show dlcs	Displays ports which are snooping WINS packets, along with the data that has been learned.



10 Extreme Standby Router Protocol

This chapter covers the following topics:

- [Overview on page 10-1](#)
- [ESRP Basics on page 10-2](#)
- [Determining the ESRP Master on page 10-3](#)
- [Grouping Blocks of 10/100 Ports on page 10-7](#)
- [ESRP Options on page 10-9](#)
- [ESRP and VLAN aggregation on page 10-13](#)
- [ESRP Commands on page 10-14](#)
- [Displaying ESRP Information on page 10-20](#)

OVERVIEW

ESRP is a feature of ExtremeWare that allows multiple switches to provide redundant routing services to users. From the workstation's perspective, there is only one default router (that has one IP address and one MAC address), so ARP cache entries in client workstations do not need to be refreshed or aged-out.

In addition to providing layer 3 routing redundancy for IP and IPX, ESRP also provides for layer 2 redundancy. These "layered" redundancy features can be used in combination or independently. You do not have to configure the switch for routing to make valuable use of ESRP. The layer 2 redundancy features of ESRP offer fast failure recovery and provide for dual-homed system design. In some instances, depending on

network system design, ESRP can provide better resiliency than using the Spanning Tree Protocol (STP).

It is highly recommended all switches participating in ESRP run the same version of ExtremeWare. Not all ESRP features are available in all ExtremeWare software releases.

ESRP-AWARE SWITCHES

Extreme switches that are not running ESRP, but are connected on a network that has other Extreme switches running ESRP are *ESRP-aware*. When ESRP-aware switches are attached to ESRP-enabled switches, the ESRP-aware switches reliably perform fail-over and fail-back scenarios in the prescribed recovery times. No configuration of this feature is necessary.

If Extreme switches running ESRP are connected to layer 2 switches that are not manufactured by Extreme Networks (or Extreme switches that are not running ExtremeWare 4.0 or above), the fail-over times seen for traffic local to the segment may appear longer, depending on the application involved and the FDB timer used by the other vendor's layer 2 switch. As such, ESRP can be used with layer 2 switches from other vendors, but the recovery times vary.

The VLANs associated with the ports connecting an ESRP-aware switch to an ESRP-enabled switch must be configured using an 802.1Q tag on the connecting port, or, if only a single VLAN is involved, as untagged using the protocol filter `any`. ESRP will not function correctly if the ESRP-aware switch interconnection port is configured for a protocol-sensitive VLAN using untagged traffic.

ESRP BASICS

ESRP is configured on a per-VLAN basis on each switch. A maximum of four switches can participate in providing redundant layer 3 or layer 2 services to a single VLAN. The switches exchange keep-alive packets for each VLAN independently. Only one switch can actively provide layer 3 routing and/or layer 2 switching for each VLAN. The switch performing the forwarding for a particular VLAN is considered the "master" for that VLAN. Other participating switches for the VLAN are in standby mode.

For a VLAN with ESRP enabled, each participating switch uses the same MAC address and must be configured with the same IP address or IPX NetID. It is possible for one switch to be master for one or more VLANs while being in standby for others, thus allowing the load to be split across participating switches.



If you configure OSPF and ESRP, you must manually configure an OSPF router identifier (ID). Be sure that you configure a unique OSPF router ID on each switch running ESRP. For more information on configuring OSPF, refer to Chapter 12.

To have two or more switches participate in ESRP, the following must be true:

- For each VLAN to be made redundant, the switches must have the ability to exchange packets on the same layer 2 broadcast domain for that VLAN. Multiple paths of exchange can be used, and typically exist in most network system designs that take advantage of ESRP.
- In order for a VLAN to be recognized as participating in ESRP, the assigned IP address or the IPX NETid for the separate switches must be *identical*. Other aspects of the VLAN, including its name, are ignored.
- ESRP must be enabled on the desired VLANs for each switch.



ESRP cannot be enabled on the VLAN default.

- Extreme Discovery Protocol (EDP) must be enabled on the ports that are members of the ESRP VLANs (The default setting is enabled.).

To verify EDP status, use the following command:

```
show ports <portlist> info {detail}
```

DETERMINING THE ESRP MASTER

The ESRP master switch (providing layer 3 routing and/or layer 2 switching services for a VLAN) is determined by the following factors:

- **Active ports**—The switch that has the greatest number of active ports takes highest precedence. A load-sharing port group is considered a single port.
- **Tracking information** — Various types of tracking are used to determine if the switch performing the master ESRP function has connectivity to the outside world. ExtremeWare supports the following types of tracking:
 - VLAN – Tracks any active port connectivity to one or more designated VLANs
 - IP route table entry – Tracks specific learned routes from the IP route table
 - Ping – Tracks ICMP ping connectivity to specified devices

If any of the configured tracking mechanisms fail, the master ESRP switch relinquishes status as master, and remains in standby mode for as long as the tracking mechanism continues to fail.

- **ESRP priority**—This is a user-defined field. The range of the priority value is 0 to 254; a higher number has higher priority. The default priority setting is 0. A priority setting of 255 loses the election and remains in standby mode.
- **System MAC address** —The switch with the higher MAC address has priority.

ESRP TRACKING

Tracking information is used to track various forms of connectivity from the ESRP switch to the outside world. This section describes ESRP tracking options.

ESRP VLAN TRACKING

You can configure ESRP to track connectivity to one or more specified VLANs as criteria for failover. If no active ports remain on the specified VLANs, the switch automatically relinquishes master status and remains in standby mode.

To add or delete a tracked VLAN, use the following command:

```
config vlan <name> [add | delete] track-vlan <vlan_tracked>
```

ESRP ROUTE TABLE TRACKING

You can configure ESRP to track specified routes in the route table as criteria for failover. If any of the configured routes are not available within the route table, the switch automatically relinquishes master status and remains in standby mode.

To participate in ESRP route table tracking, all ESRP switches must run ExtremeWare version 6.0 or above.

To add or delete a tracked route, use the following command:

```
config vlan <name> [add | delete] track-route <ipaddress/mask_length>
```

ESRP PING TRACKING

You can configure ESRP to track connectivity using a simple ping to any outside responder. The responder may represent the default route of the switch, or any device meaningful to network connectivity of the master ESRP switch. The switch

automatically relinquishes master status and remains in standby mode if a ping keepalive fails three consecutive times.

To participate in ESRP ping tracking, all ESRP switches must run ExtremeWare version 6.0 or above.

To view the status of tracked devices, use the following command:

```
show esrp
```

ESRP ELECTION ALGORITHMS

You configure the switch to use one of five different election algorithms to select the ESRP master. Each algorithm considers the election factors in a different order of precedence, as follows:

- `ports_track_priority_mac` — Active ports, tracking information, ESRP priority, MAC address (Default)
- `track_ports_priority_mac` — Tracking information, active ports, ESRP priority, MAC address
- `priority_ports_track_mac` — ESRP priority, active ports, tracking information, MAC address
- `priority_track_ports_mac` — ESRP priority, tracking information, active ports, MAC address
- `priority_mac_only` — ESRP priority, MAC address



All switches in the ESRP network must use the same election algorithm, otherwise loss of connectivity, broadcast storms, or other unpredictable behavior may occur.



Only the `ports_track_priority_mac` election algorithm is compatible with ExtremeWare releases prior to version 6.0.

MASTER SWITCH BEHAVIOR

If a switch is master, it actively provides layer 3 routing services to other VLANs, and layer 2 switching between all the ports of that VLAN. Additionally, the switch exchanges ESRP packets with other switches that are in standby mode.

STANDBY SWITCH BEHAVIOR

If a switch is in standby mode, it exchanges ESRP packets with other switches on that same VLAN. When a switch is in standby, it does not perform layer 3 routing or layer 2 switching services for the VLAN. From a layer 3 routing protocol perspective (for example, RIP or OSPF), when in standby for the VLAN, the switch marks the router interface associated with the VLAN as down. From a layer 2 switching perspective, no forwarding occurs between the member ports of the VLAN; this prevents loops and maintains redundancy.

ELECTING THE MASTER SWITCH

A new master can be elected in one of the following ways:

- A communicated parameter change
- Loss of communication between master and slave(s).

If a parameter that determines the master changes (for example, link loss or priority change), the election of the new master typically occurs within one timer cycle (2 seconds by default). If a switch in standby mode loses its connection with the master, a new election (using the same precedence order indicated previously) occurs. The new election typically takes place in three times the defined timer cycle (6 seconds by default).

FAILOVER TIME

Failover time is largely determined by the following factors:

- The ESRP timer setting.
- The routing protocol being used for inter-router connectivity if layer 3 redundancy is used. OSPF fail-over time is faster than RIP fail-over time.

The failover time associated with the ESRP protocol is dependent on the timer setting and the nature of the failure. The default timer setting is 2 seconds; the range is 1 to 255 seconds.

If routing is configured, the failover of the particular routing protocol (such as RIP V1, RIP V2, or OSPF) is added to the failover time associated with ESRP.

GROUPING BLOCKS OF 10/100 PORTS

Restrictions on port groupings apply only to switches that do not use the “*i*” chipset.

If you enable ESRP on a VLAN that contains 10/100 ports, a specific block of neighboring ports must also be participating in a VLAN running ESRP, or must not be used. The blocks of ports are physically adjacent, regardless of the switch module. For example, the blocks on a BlackDiamond F32T module consist of the following:

- Ports 1-4 and 17-20
- Ports 5-8 and 21-24
- Ports 9-12 and 25-28
- Ports 13-16 and 29-32

[Figure 10-1](#) through [Figure 10-5](#) illustrate the port blocks for each Extreme switch.

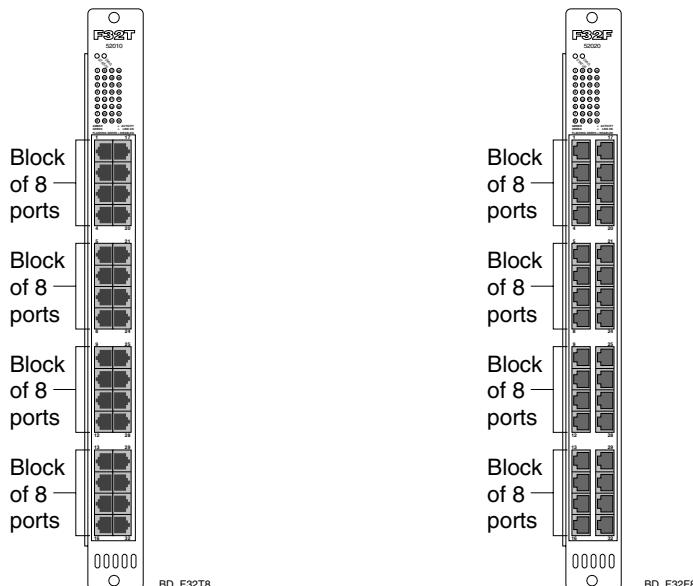
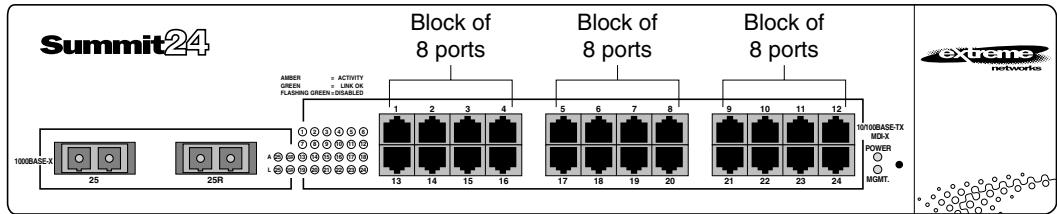
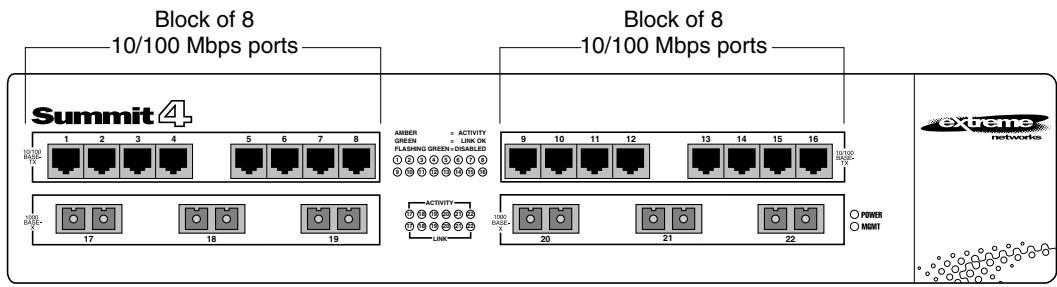


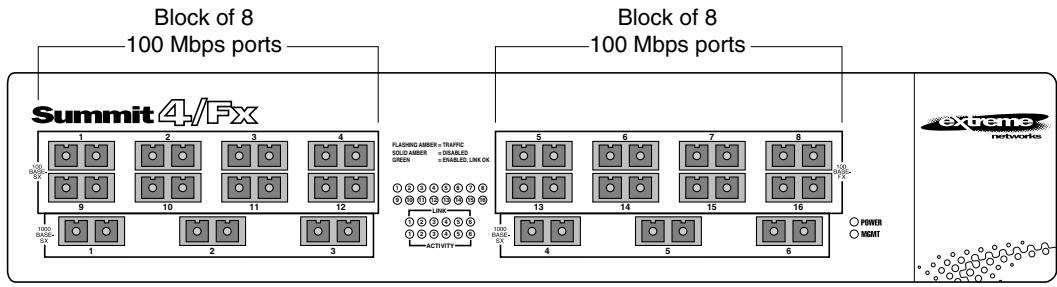
Figure 10-1: F32T and F32F ESRP port blocks



Sum24_8

Figure 10-2: Summit24 switch ESRP port blocks

SUG_4fr8

Figure 10-3: Summit4 switch ESRP port blocks

SUG_4FX8

Figure 10-4: Summit4/FX switch ESRP port blocks

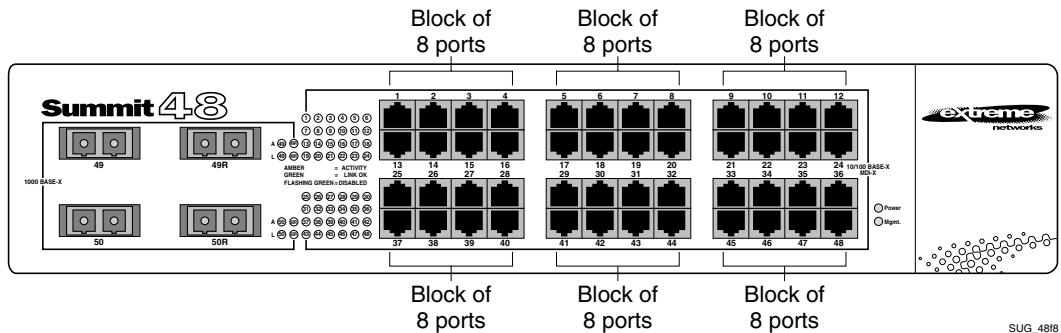


Figure 10-5: Summit48 switch ESRP port blocks

i For switches that do not use the “i” chipset, all VLANs using a port or port-block must enable ESRP. This requirement does not apply to switches that use the “i” chipset.

ESRP OPTIONS

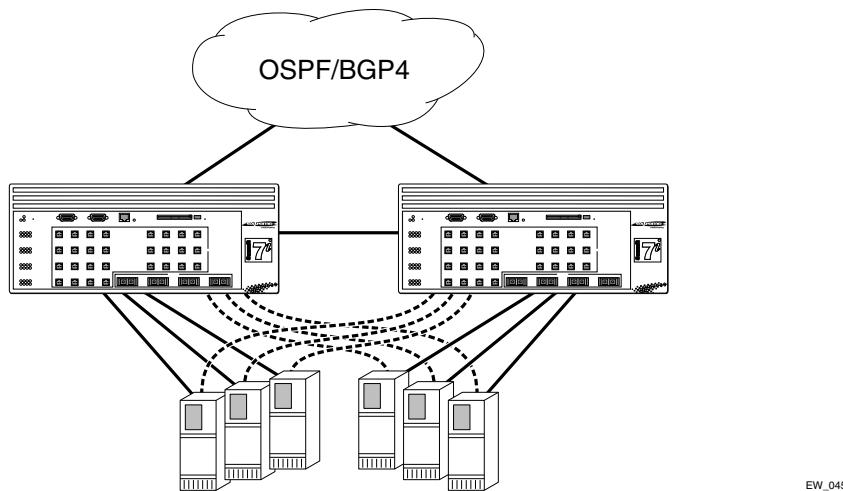
This section discusses the following ESRP options:

- ESRP Host Attach
- ESRP Domains
- ESRP Groups
- Linking ESRP Switches
- Configuring ESRP and Multinetting
- ESRP and Spanning Tree

ESRP HOST ATTACH

ESRP host attach (HA) is an optional ESRP configuration that allows you to connect active hosts directly to an ESRP master or standby switch. Normally, the Layer 2 redundancy and loop prevention capabilities of ESRP do not allow packet forwarding from the standby ESRP switch. ESRP HA allows configured ports that do not represent loops to the network to continue Layer 2 operation independent of their ESRP status.

The ESRP HA option is useful if you are using dual-homed network interface cards (NICs) for server farms, and in conjunction with high availability server load balancing (SLB) configurations, as shown in [Figure 10-6](#).



EW_045

Figure 10-6: ESRP host attach

Other applications allow lower cost redundant routing configurations, because hosts can be directly attached to the switch involved with ESRP. The ESRP HA feature is used only on switches and I/O modules that have the “*i*” series chipset. It also requires at least one link between the master and standby ESRP switch for carrying traffic and to exchange ESRP hello packets.

ESRP DOMAINS

ESRP Domains is an optional ESRP configuration that allows you to configure multiple VLANs under the control of a single instance of the ESRP protocol. By grouping multiple VLANs under one ESRP group, the ESRP protocol can scale to provide protection to large numbers of VLANs. All VLANs within an ESRP group simultaneously share the same active and standby router and failover. The ESRP group feature is used only on switches and I/O modules that have the “*i*” series chipset.

ESRP GROUPS

ExtremeWare supports running multiple instances of ESRP within the same VLAN or broadcast domain. This functionality is called an ESRP group. Though other uses exist, the most typical application for multiple ESRP groups is when two or more sets of ESRP switches are providing fast-failover protection within a subnet.

For example, two ESRP switches provide L2/L3 connectivity and redundancy for the subnet, while another two ESRP switches provide L2 connectivity and redundancy for a portion of the same subnet. [Figure 10-7](#) shows ESRP groups.

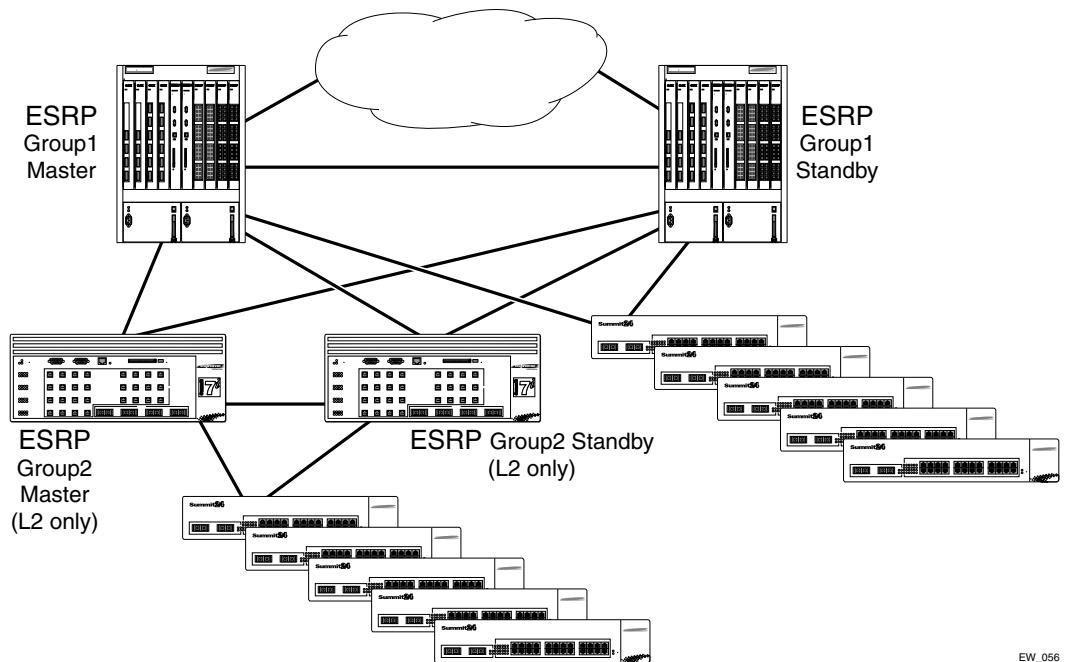


Figure 10-7: ESRP groups

 A switch cannot perform both master and slave functions on the same VLAN for separate instances of ESRP.

LINKING ESRP SWITCHES

When considering system design using ESRP, direct links between ESRP switches are useful under the following conditions:

- A direct link can provide a more direct routed path, if the ESRP switches are routing and supporting multiple VLANs where the master/standby configuration is split such that one switch is master for some VLANs and a second switch is master for other VLANs. The direct link can contain a unique router-to-router VLAN/subnet, so that the most direct routed path between two VLANs with different master switches uses a direct link, instead of forwarding through another set of connected routers.
- A direct link can be used as a highly reliable method to exchange ESRP hellos, so that the possibility of having multiple masters for the same VLAN is lessened, should all downstream Layer 2 switches fail.
- A direct link is necessary when the ESRP HA option. The direct link is used to provide Layer 2 forwarding services through an ESRP standby switch.

Direct links may contain a router-to-router VLAN, along with VLANs running ESRP. If multiple VLANs are used on the direct links, use 802.1Q tagging. The direct links may be aggregated into a load-shared group, if desired.

CONFIGURING ESRP AND MULTINETTING

When configuring ESRP and IP multinetting on the same switch, the parameters that affect the determination of the ESRP master must be configured identically for all the VLANs involved with IP multinetting. For example, the number of links in your configuration, the priority settings, and timer settings must be identical for all affected VLANs.

ESRP AND SPANNING TREE

A switch running ESRP should not simultaneously participate in the Spanning Tree Protocol (STP) for the same VLAN(s). Other switches in the VLAN being protected by ESRP may run STP and the switch running ESRP forwards, but does not filter, STP BPDUs. Therefore, you can combine ESRP and STP on a network and a VLAN, but you must do so on separate devices. You should be careful to maintain ESRP connectivity between ESRP master and standby switches when you design a network that uses ESRP and STP.

ESRP AND VLAN AGGREGATION

ESRP can be used to provide redundant default router protection to VLAN aggregation clients. ESRP is enabled on the super-VLAN *only* (not the sub-VLANs). The procedure is to add ports to the super-VLAN that is shared with the sub VLANs. To do so, the super-VLAN should be configured with an 802.1Q tag, and added as tagged with the sub-VLAN ports to avoid a protocol conflict. Lastly, enable ESRP on the super-VLAN.



For more information on VLAN aggregation, see Chapter 6.

The following example combines ESRP and VLAN aggregation for the super-VLAN *vsuper* and two sub-VLANs, *v1sub* and *v2sub*, that have ports 1 and 2 as members, respectively.

- 1 Create the VLANs and setup the super to sub-VLAN relationship

```
create vlan v1sub
create vlan v2sub
create vlan vsuper
config vsuper ipaddress 10.1.2.3/24
enable ipforwarding
enable ospf
config ospf add vsuper
config v1sub add port 1
config v2sub add port 2
config vsuper add subvlan v1sub
config vsuper add subvlan v2sub
```

- 2 Turn on ESRP for the VLAN *vsuper*.

```
config vsuper tag 1234
config vsuper add port 1,2 tagged
enable esrp vlan vsuper
```

Use these commands to verify the configuration:

- `show vlan {detail}`— Displays super- and sub-VLAN relationships, IP addresses, and port membership.
- `show esrp {detail}`—Verifies ESRP is enabled and operational.

ESRP COMMANDS

[Table 10-1](#) describes the commands used to configure ESRP.

Table 10-1: ESRP Commands

Command	Description
config esrp port-mode [host normal] ports <portlist>	Configures the ESRP port mode. A normal port does not accept or transmit traffic when the local ESRP device is a slave. The host port always switches user traffic, regardless of the ESRP state. The default setting is normal.
config vlan <name> add track-ping <ipaddress> frequency <seconds> miss <number>	Configures an ESRP-enabled VLAN to track an external gateway using ping. The switch will not be the ESRP master of the VLAN if the external gateway is not reachable.
config vlan <name> add track-route <ipaddress>/<masklength>	Configures an ESRP-enabled VLAN to track the condition of a route entry in the kernel route table. The switch cannot be the ESRP master if none of the specified routes are reachable.
config vlan <name> add track-vlan <vlan_tracked>	Configures an ESRP-enabled VLAN to track the condition of another VLAN.
config vlan <name> delete track-ping <ipaddress> frequency <seconds> miss <number>	Configures an ESRP-enabled VLAN to stop tracking an external gateway.
config vlan <name> delete track-route <ipaddress>/<masklength>	Disables route entry tracking for an ESRP-enabled VLAN.
config vlan <name> delete track-vlan <vlan_tracked>	Removes the tracking of a VLAN by an ESRP-enabled VLAN.

Table 10-1: ESRP Commands (continued)

Command	Description
config vlan <name> esrp election-algorithm [ports_track_priority_mac track_ports_priority_mac priority_ports_track_mac priority_track_ports_mac priority_mac_only]	<p>Configures the election algorithm on the switch. The algorithm must be the same on all switches for a particular VLAN. Specify one of the following:</p> <ul style="list-style-type: none"> ■ ports_track_priority_mac — Active ports, tracking information, ESRP priority, MAC address ■ track_ports_priority_mac — Tracking information, active ports, ESRP priority, MAC address ■ priority_ports_track_mac — ESRP priority, active ports, tracking information, MAC address ■ priority_track_ports_mac — ESRP priority, tracking information, active ports, MAC address ■ priority_mac — ESRP priority, MAC address <p>The default setting is ports_track_priority_mac. If no tracking information is configured for a particular field, the field is ignored.</p>
config vlan <name> esrp priority <value>	Configures the ESRP priority. The range is 0 to 255. The higher number has higher priority. The default setting is 0. A setting of 255 configures the switch to be in standby state.
config vlan <name> esrp timer <hello_timer>	Configures the time between ESRP updates. The range is 1 to 255 seconds. The default setting is 2 seconds. The timer setting must be configured identically for the VLAN across all participating switches.
config vlan <name> esrp-group <group number>	Configures the virtual MAC address to be used for the ESRP VLAN. The default group number is 0.
config vlan <super_ESRP_VLAN> add domain-member vlan <sub_ESRP_VLAN>	Adds a VLAN to an ESRP domain. ESRP is performed in the domain master VLAN, and not the other domain members.
config vlan <super_ESRP_VLAN> delete domain-member vlan <sub_ESRP_VLAN>	Deletes a VLAN from an ESRP domain.

Table 10-1: ESRP Commands (continued)

Command	Description
disable esrp vlan <name>	Disables ESRP on a VLAN.
enable esrp vlan <name>	Enables ESRP on a VLAN.
show esrp {detail}	Displays ESRP configuration information.
show esrp vlan <name>	Displays ESRP configuration information for a specific VLAN.

ESRP EXAMPLES

This section provides examples of ESRP configurations.

SINGLE VLAN USING LAYER 2 AND LAYER 3 REDUNDANCY

This example, shown in [Figure 10-8](#), uses a number of Summit switches that perform layer 2 switching for VLAN *Sales*. The Summit switches are dual-homed to the BlackDiamond switches. The BlackDiamond switches perform layer 2 switching between the Summit switches, and layer 3 routing to the outside world. Each Summit switch is dual-homed using active ports to two BlackDiamond switches (as many as four could be used). ESRP is enabled on each BlackDiamond switch only for the VLAN that interconnects to the Summit switches. Each BlackDiamond switch has the VLAN *Sales* configured using the identical IP address. The BlackDiamond switches then connect to the routed enterprise normally, using the desired routing protocol (for example RIP or OSPF).

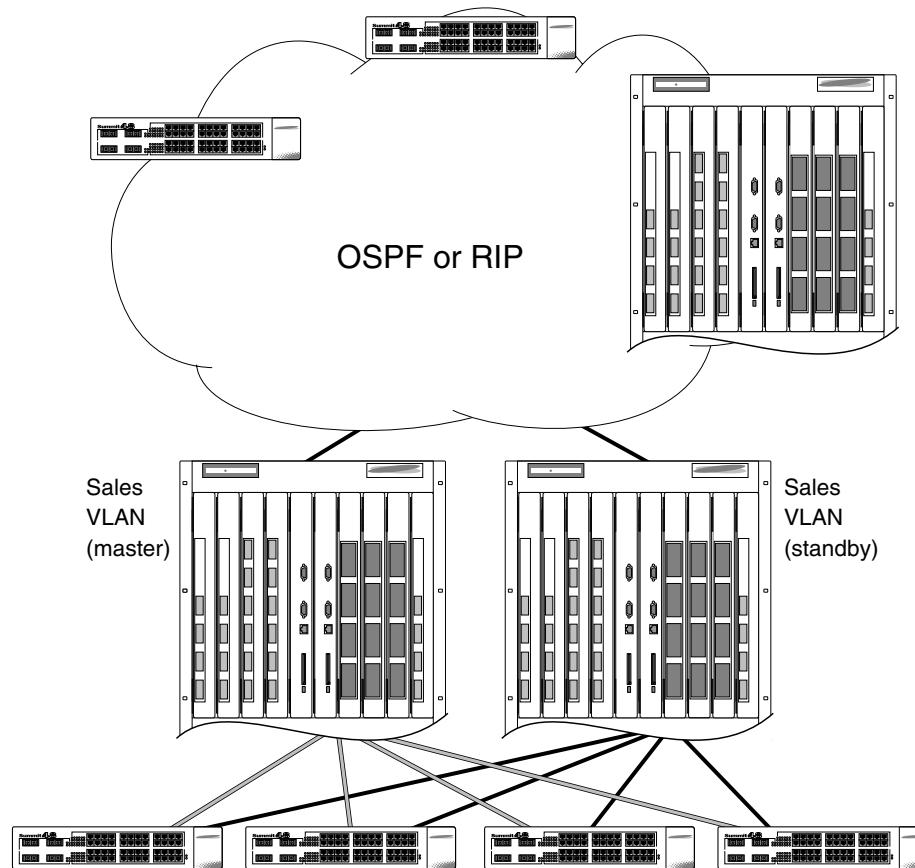


Figure 10-8: ESRP example using layer 2 and layer 3 redundancy

The BlackDiamond switch, acting as master for VLAN *Sales*, performs both layer 2 switching and layer 3 routing services for VLAN *Sales*. The BlackDiamond switch in standby mode for VLAN *Sales* performs neither, thus preventing bridging loops in the VLAN. The BlackDiamond switch in standby mode does, however, exchange ESRP packets with the master BlackDiamond switch.

There are four paths between the BlackDiamond switches on VLAN *Sales*. All the paths are used to send ESRP packets, allowing for four redundant paths for ESRP communication. The Summit switches, being ESRP-aware, allow traffic within the VLAN to fail-over quickly, as they will sense when a master/slave transition occurs and

flush FDB entries associated with the uplinks to the ESRP-enabled BlackDiamond switches.

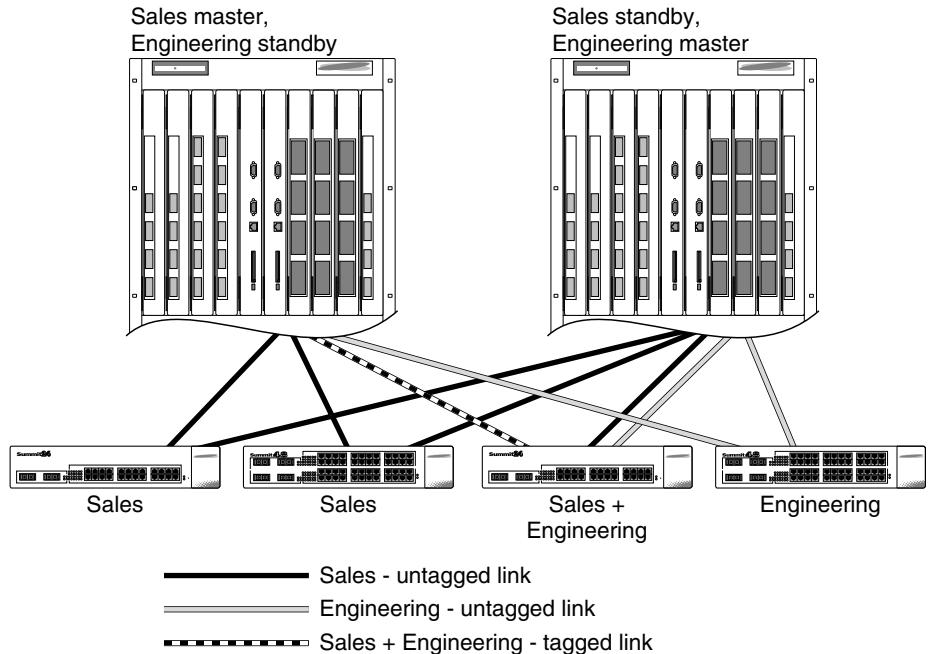
The following commands are used to configure both BlackDiamond switches. The assumption is that the inter-router backbone is running OSPF, with other routed VLANs already properly configured. Similar commands would be used to configure a switch on a network running RIP. The primary requirement is that the IP address for the VLAN(s) running ESRP must be identical. In this scenario, the master is determined by the programmed MAC address of the switch, because the number of active links for the VLAN and the priority are identical to both switches.

The commands used to configure the BlackDiamond switches are as follows:

```
create vlan sales
config sales add port 1:1-1:4
config sales ipaddr 10.1.2.3/24
enable ipforwarding
enable esrp sales
enable edp ports all
config ospf add vlan sales
enable ospf
```

MULTIPLE VLANs USING LAYER 2 REDUNDANCY

The example shown in [Figure 10-9](#) illustrates an ESRP configuration that has multiple VLANs using layer 2 redundancy.



EW_022

Figure 10-9: ESRP example using layer 2 redundancy

This example builds on the previous example, but eliminates the requirement of layer 3 redundancy. It has the following features:

- An additional VLAN, *Engineering*, is added that uses layer 2 redundancy.
- The VLAN *Sales* uses three active links to each BlackDiamond switch.
- The VLAN *Engineering* has two active links to each BlackDiamond switch.
- The third Summit switch carries traffic for both VLANs.
- The link between the third Summit switch and the first BlackDiamond switch uses 802.1Q tagging to carry traffic from both VLANs traffic on one link. The BlackDiamond switch counts the link active for each VLAN.
- The second BlackDiamond switch has a separate physical port for each VLAN connected to the third Summit switch.

In this example, the BlackDiamond switches are configured for ESRP such that the VLAN *Sales* normally uses the first BlackDiamond switch and the VLAN *Engineering* normally uses the second BlackDiamond switch. This is accomplished by manipulating the ESRP priority setting for each VLAN for the particular BlackDiamond switch.

Configuration commands for the first BlackDiamond switch are as follows:

```
create vlan sales
config sales tag 10
config sales add port 1:1-1:2
config sales add port 1:3 tagged
config sales ipaddr 10.1.2.3/24
create vlan eng
config eng tag 20
config eng add port 1:4
config eng add port 1:3 tagged
config eng ipaddr 10.4.5.6/24
enable esrp sales
enable esrp eng
enable edp ports all
config sales esrp priority 5
```

Configuration commands for the second BlackDiamond switch are as follows:

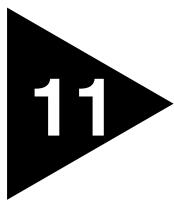
```
create vlan sales
config sales add port 1:1-1:3
config sales ipaddr 10.1.2.3/24
create vlan eng
config eng add port 1:4, 2:1
config eng ipaddr 10.4.5.6/24
enable esrp sales
enable esrp eng
config eng esrp priority 5
```

DISPLAYING ESRP INFORMATION

To verify the operational state of an ESRP VLAN and the state of its neighbor, use the following command:

```
show esrp
```

To view tracking information about a particular VLAN, including the VLANs tracked by it and a list of the VLANs tracking it, use the `show vlan` command.



IP Unicast Routing

This chapter describes the following topics:

- [Overview of IP Unicast Routing on page 11-2](#)
- [Proxy ARP on page 11-5](#)
- [Relative Route Priorities on page 11-6](#)
- [IP Multinetting on page 11-7](#)
- [Configuring IP Unicast Routing on page 11-10](#)
- [VLAN Aggregation on page 11-11](#)
- [Configuring DHCP/BOOTP Relay on page 11-16](#)
- [UDP-Forwarding on page 11-16](#)
- [IP Commands on page 11-19](#)
- [Routing Configuration Example on page 11-25](#)
- [Displaying Router Settings on page 11-27](#)
- [Resetting and Disabling Router Settings on page 11-28](#)

This chapter assumes that you are already familiar with IP unicast routing. If not, refer to the following publications for additional information:

- RFC 1256 — *ICMP Router Discovery Messages*
- RFC 1812 — *Requirements for IP Version 4 Routers*

 For more information on interior gateway protocols, refer to [Chapter 12](#). For information on exterior gateway protocols, refer to [Chapter 13](#).

OVERVIEW OF IP UNICAST ROUTING

The switch provides full layer 3, IP unicast routing. It exchanges routing information with other routers on the network using either the Routing Information Protocol (RIP) or the Open Shortest Path First (OSPF) protocol. The switch dynamically builds and maintains a routing table, and determines the best path for each of its routes.

Each host using the IP unicast routing functionality of the switch must have a unique IP address assigned. In addition, the default gateway assigned to the host must be the IP address of the router interface.

ROUTER INTERFACES

The routing software and hardware routes IP traffic between router interfaces. A router interface is simply a VLAN that has an IP address assigned to it.

As you create VLANs with IP addresses belonging to different IP subnets, you can also choose to route between the VLANs. Both the VLAN switching and IP routing function occur within the switch.

 *Each IP address and mask assigned to a VLAN must represent a unique IP subnet. You cannot configure the same IP address and subnet on different VLANs.*

In [Figure 11-1](#), a BlackDiamond switch is depicted with two VLANs defined; *Finance* and *Personnel*. All ports on slots 1 and 3 are assigned to *Finance*; all ports on slots 2 and 4 are assigned to *Personnel*. *Finance* belongs to the IP network 192.207.35.0; the router interface for *Finance* is assigned the IP address 192.206.35.1. *Personnel* belongs to the IP network 192.207.36.0; its router interface is assigned IP address 192.207.36.1. Traffic within each VLAN is switched using the Ethernet MAC addresses. Traffic between the two VLANs is routed using the IP addresses.

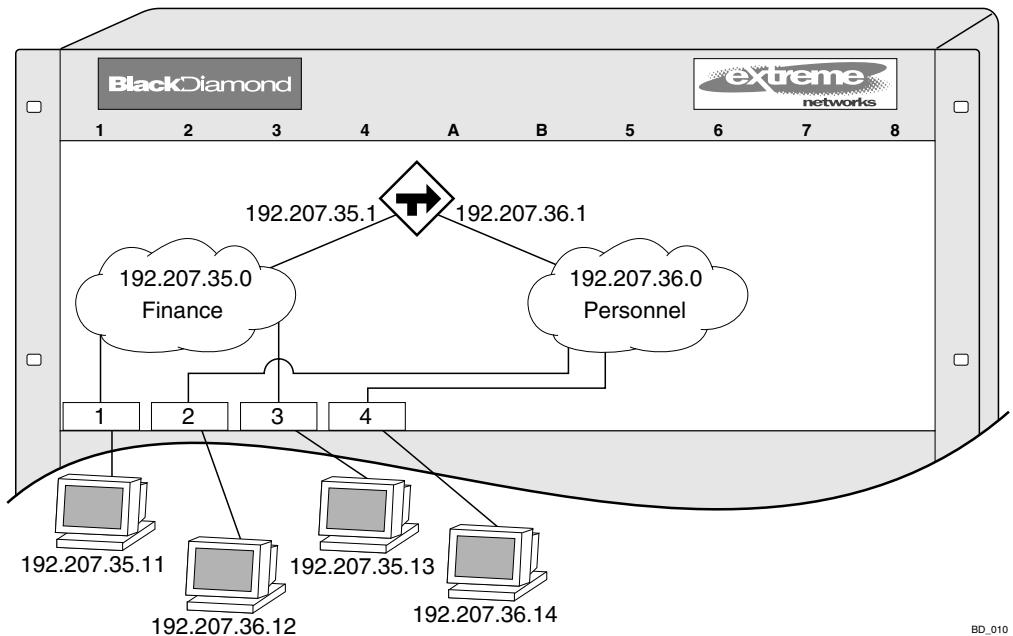


Figure 11-1: Routing between VLANs

POPULATING THE ROUTING TABLE

The switch maintains an IP routing table for both network routes and host routes. The table is populated from the following sources:

- Dynamically, by way of routing protocol packets or by ICMP redirects exchanged with other routers
- Statically, by way of routes entered by the administrator
 - Default routes, configured by the administrator
 - Locally, by way of interface addresses assigned to the system
 - By other static routes, as configured by the administrator

i *If you define a default route, and subsequently delete the VLAN on the subnet associated with the default route, the invalid default route entry remains. You must manually delete the configured default route.*

DYNAMIC ROUTES

Dynamic routes are typically learned by way of RIP or OSPF. Routers that use RIP or OSPF exchange information in their routing tables in the form of advertisements. Using dynamic routes, the routing table contains only networks that are reachable.

Dynamic routes are aged out of the table when an update for the network is not received for a period of time, as determined by the routing protocol.

STATIC ROUTES

Static routes are manually entered into the routing table. Static routes are used to reach networks not advertised by routers.

Static routes can also be used for security reasons, to control which routes you want advertised by the router. You can decide if you want all static routes to be advertised, using one of the following commands:

```
[enable | disable] rip export static  
[enable | disable] ospf export static
```

The default setting is enabled. Static routes are never aged out of the routing table.

A static route must be associated with a valid IP subnet. An IP subnet is associated with a single VLAN by its IP address and subnet mask. If the VLAN is subsequently deleted, the static route entries using that subnet must be deleted manually.

MULTIPLE ROUTES

When there are multiple, conflicting choices of a route to a particular destination, the router picks the route with the longest matching network mask. If these are still equal, the router picks the route using the following criteria (in the order specified):

- Directly attached network interfaces
- ICMP redirects (refer to [Table 11-6](#), later in this chapter)
- Static routes
- Directly attached network interfaces that are not active.

 *If you define multiple default routes, the route that has the lowest metric is used. If there are multiple default routes that have the same lowest metric, the system picks one of the routes.*

You can also configure *blackhole* routes — traffic to these destinations is silently dropped.

IP ROUTE SHARING

IP route sharing allows multiple equal-cost routes to be used concurrently. IP route sharing can be used with static routes or with OSPF routes. In OSPF, this capability is referred to as *equal cost multi-path* (ECMP) routing. To use IP route sharing, use the following command:

```
enable route sharing
```

Next, configure static routes and/or OSPF as you would normally. As many as five ECMP routes can be used for a given destination.

Route sharing is useful only in instances where you are constrained for bandwidth. This is typically not the case using Extreme switches. Using route sharing makes router troubleshooting more difficult because of the complexity in predicting the path over which the traffic will travel.

PROXY ARP

Proxy Address Resolution Protocol (ARP) was first invented so that ARP-capable devices could respond to ARP Request packets on behalf of ARP-incapable devices. Proxy ARP can also be used to achieve router redundancy and simplify IP client configuration. The switch supports proxy ARP for this type of network configuration. The section describes some example of how to use proxy ARP with the switch.

ARP-INCAPABLE DEVICES

To configure the switch to respond to ARP Requests on behalf of devices that are incapable of doing so, you must configure the IP address and MAC address of the ARP-incapable device using the use the following command:

```
config iparp add proxy <ipaddress> {<mask>} <mac_address> {always}
```

Once configured, the system responds to ARP Requests on behalf of the device as long as the following conditions are satisfied:

- The valid IP ARP Request is received on a router interface.
- The target IP address matches the IP address configured in the proxy ARP table.

- The proxy ARP table entry indicates that the system should always answer this ARP Request, regardless of the ingress VLAN (the `always` parameter must be applied).

Once all the proxy ARP conditions are met, the switch formulates an ARP Response using the configured MAC address in the packet.

PROXY ARP BETWEEN SUBNETS

In some networks, it is desirable to configure the IP host with a wider subnet than the actual subnet mask of the segment. Proxy ARP can be used so that the router answers ARP Requests for devices outside of the subnet. As a result, the host communicates as if all devices are local. In reality, communication with devices outside of the subnet are proxied by the router.

For example, an IP host is configured with a class B address of 100.101.102.103 and a mask of 255.255.0.0. The switch is configured with the IP address 100.101.102.1 and a mask of 255.255.255.0. The switch is also configured with a proxy ARP entry of IP address 100.101.0.0 and mask 255.255.0.0, *without* the `always` parameter.

When the IP host tries to communicate with the host at address 100.101.45.67, the IP hosts communicates as if the two hosts are on the same subnet, and sends out an IP ARP Request. The switch answers on behalf of the device at address 100.101.45.67, using its own MAC address. All subsequent data packets from 100.101.102.103 are sent to the switch, and the switch routes the packets to 100.101.45.67.

RELATIVE ROUTE PRIORITIES

Table 11-1 lists the relative priorities assigned to routes depending upon the learned source of the route.



Although these priorities can be changed, do not attempt any manipulation unless you are expertly familiar with the possible consequences.

Table 11-1: Relative Route Priorities

Route Origin	Priority
Direct	10
BlackHole	50
Static	1100

Table 11-1: Relative Route Priorities (continued)

Route Origin	Priority
ICMP	1200
OSPFIntra	2200
OSPFInter	2300
RIP	2400
OSPFExtern1	3200
OSPFExtern2	3300
BOOTP	5000

To change the relative route priority, use the following command:

```
config iproute priority [rip | bootp | icmp | static | ospf-intra |
ospf-inter | ospf-as-external | ospf-extern1 | ospf-extern2] <priority>
```

IP MULTINETTING

IP multinetting is used in many legacy IP networks when there is need to overlap multiple subnets onto the same physical segment. Though it can be a critical element in a transition strategy, due to the additional constraints introduced in troubleshooting and bandwidth, it is recommended that multinetting be used as a transitional tactic, and not as a long-term network design strategy.

On the switch, each subnet is represented by a different VLAN, and each of those VLANs has its own IP address. All of the VLANs share the same physical port(s). The BlackDiamond switch or Summit switch routes IP traffic from one subnet to another, all within the same physical port(s).

The following rules and comments apply when you are configuring IP multinetting:

- Multiple VLANs share the same physical ports; each of the VLANs is configured with an IP address.
- A maximum of four subnets (or VLANs) on multinetted ports is recommended.
- All VLANs used in the multinetting application must share the same port assignment.
- One VLAN is configured to use an IP protocol filter. This is considered the "primary" VLAN interface for the multinetted group.
- The "secondary" multinetted VLANs can be exported using the `export direct` command.

- The FDB aging timer is automatically set to 3,000 seconds (50 minutes).
- If you are using a UDP or DHCP relay function, only the "primary" VLAN that is configured with the IP protocol filter is capable of servicing these requests.
- The VLAN *default* should not be used for multinetting.

IP MULTINETTING OPERATION

To use IP multinetting, follow these steps:

- 1 Select a slot (BlackDiamond switch only) and port on which IP multinetting is to run.

For example, slot 1, port 2 on a BlackDiamond switch, or port 2 on a Summit switch.

- 2 Remove the port from the default VLAN, using the following command:

```
config default delete port 1:2 (BlackDiamond switch)
```

or

```
config default delete port 2 (Summit switch)
```

- 3 Create a dummy protocol, by using the following command:

```
create protocol mnet
```

- 4 Create the multinetted subnets, by using the following commands:

```
create vlan net21  
create vlan net22
```

- 5 Assign IP addresses to the net VLANs, by using the following commands:

```
config net21 ipaddress 123.45.21.1 255.255.255.0  
config net22 ipaddress 192.24.22.1 255.255.255.0
```

- 6 Assign one of the subnets to the IP protocol, by using the following command:

```
config net21 protocol ip
```

- 7 Assign the other subnets to the dummy protocol, by using the following command:

```
config net22 protocol mnet
```

- 8 Assign the subnets to a physical port, by using the following commands:

```
config net21 add port 1:2  
config net22 add port 1:2
```

- 9 Enable IP forwarding on the subnets, by using the following command:

```
enable ipforwarding
```

- 10** Enable IP multinetting, by using the following command:

```
enable multinetting
```

- 11** If you are using RIP, disable RIP on the dummy VLANs, by using the following command:

```
config rip delete net22
```



Multinetted VLAN groups must contain identical port assignments.

IP MULTINETTING EXAMPLES

The following example configures the BlackDiamond switch to have one multinetted segment (slot 5, port 5) that contains three subnets (192.67.34.0, 192.67.35.0, and 192.67.37.0).

```
config default delete port 5:5
create protocol mnet
create vlan net34
create vlan net35
create vlan net37
config net34 ipaddress 192.67.34.1
config net35 ipaddress 192.67.35.1
config net37 ipaddress 192.67.37.1
config net34 protocol ip
config net35 protocol mnet
config net37 protocol mnet
config net34 add port 5:5
config net35 add port 5:5
config net37 add port 5:5
enable ipforwarding
enable multinetting
```

The following example configures the BlackDiamond switch to have one multinetted segment (slot 5: port 5) that contains three subnets (192.67.34.0, 192.67.35.0, and 192.67.37.0). It also configures a second multinetted segment consisting of two subnets (192.67.36.0 and 192.99.45.0). The second multinetted segment spans three ports (slot1:port 8, slot2:port 9, and slot3:port 10). RIP is enabled on both multinetted segments.

```
config default delete port 5:5
create protocol mnet
create vlan net34
```

```
create vlan net35
create vlan net37
config net34 ipaddress 192.67.34.1
config net35 ipaddress 192.67.35.1
config net37 ipaddress 192.67.37.1
config net34 protocol ip
config net35 protocol mnet
config net37 protocol mnet
config net34 add port 5:5
config net35 add port 5:5
config net37 add port 5:5
config default delete port 1:8, 2:9, 3:10
create vlan net36
create vlan net45
config net36 ipaddress 192.67.36.1
config net45 ipaddress 192.99.45.1
config net36 protocol ip
config net45 protocol mnet
config net36 add port 1:8, 2:9, 3:10
config net45 add port 1:8, 2:9, 3:10
config rip add vlan net34
config rip add vlan net36
enable rip
enable ipforwarding
enable multinetting
```

CONFIGURING IP UNICAST ROUTING

This section describes the commands associated with configuring IP unicast routing on the switch. Configuring routing involves the following steps:

- 1 Create and configure two or more VLANs.

Although it is possible to enable IP forwarding and an IP routing protocol (such as RIP) with only one VLAN defined, the switch does not create or respond appropriately to ICMP messages unless at least two VLANs are created and configured.

- 2 Assign each VLAN that will be using routing an IP address, using the following command:

```
config vlan <name> ipaddress <ipaddress> {<mask>}
```

Ensure that each VLAN has a unique IP address.

- 3 Configure a default route, using the following command:

```
config iproute add default <gateway> {<metric>} {unicast-only |  
multicast-only}
```

Default routes are used when the router has no other dynamic or static route to the requested destination.

- 4 Turn on IP routing for one or all VLANs, using the following command:

```
enable ipforwarding {vlan <name>}
```

- 5 Turn on RIP or OSPF using one of the following commands:

```
enable rip  
enable ospf
```

VERIFYING THE IP UNICAST ROUTING CONFIGURATION

Use the `show iproute` command to display the current configuration of IP unicast routing for the switch, and for each VLAN. The `show iproute` command displays the currently configured routes, and includes how each route was learned.

Additional verification commands include the following:

- `show iparp` — Displays the IP ARP table of the system.
- `show ipfdb` — Displays the hosts that have been transmitting or receiving packets, and the port and VLAN for each host.
- `show ipconfig` — Displays configuration information for one or more VLANs.

VLAN AGGREGATION

VLAN aggregation is an ExtremeWare feature aimed primarily at Service Providers. The purpose of VLAN aggregation is to increase the efficiency of IP address space usage. It does this by allowing clients within the same IP subnet to use different broadcast domains while still using the same default router.

Using VLAN aggregation, a *super-VLAN* is defined with the desired IP address, but without any member ports (unless it is running ESRP). The sub-VLANs use the IP address of the super-VLAN as the default router address. Groups of clients are then assigned to sub-VLANs that have no IP address, but are members of the super-VLAN. In addition, clients can be informally allocated any valid IP addresses within the subnet. Optionally, you can prevent communication between sub-VLANs for isolation purposes.

As a result, sub-VLANs can be quite small, but allow for growth without re-defining subnet boundaries.

Without using VLAN aggregation, each VLAN has a default router address, and you need to use large subnet masks. The result of this is more unused IP address space.

Multiple secondary IP addresses can be assigned to the super-VLAN. These IP addresses are *only* used to respond to ICMP ping packets to verify connectivity.

Figure 11-2 illustrates VLAN aggregation.

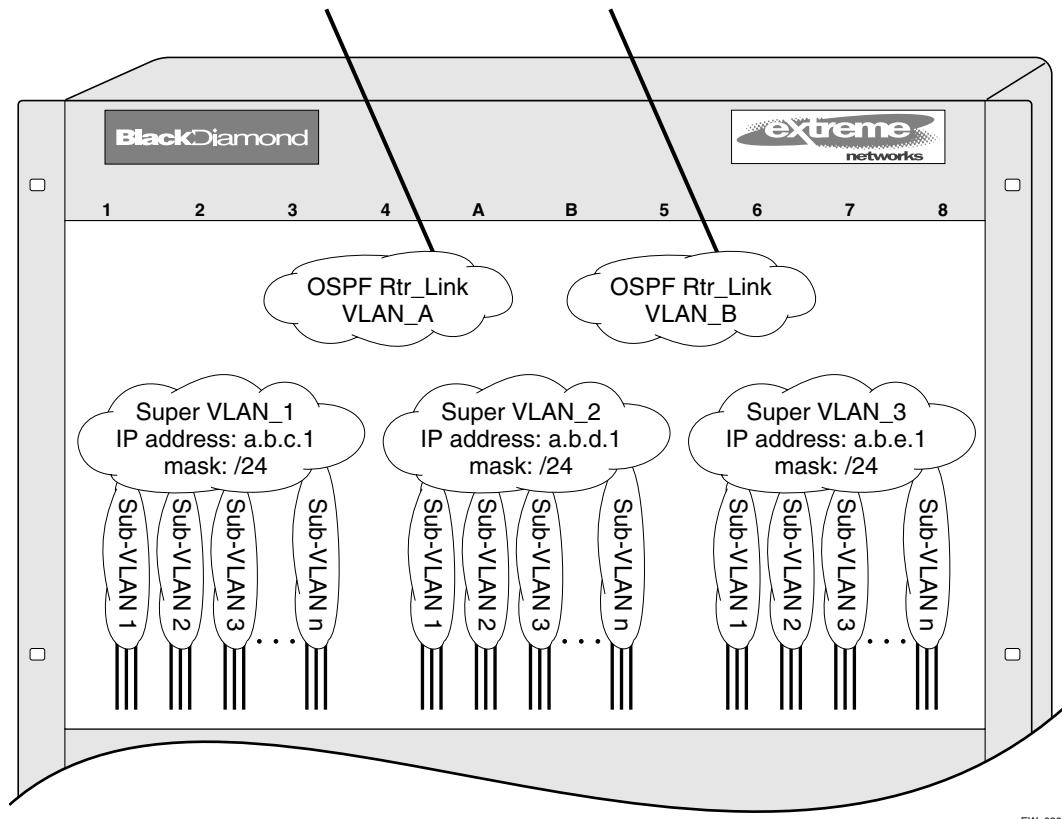


Figure 11-2: VLAN aggregation

In [Figure 11-2](#), all stations are configured to use the address 10.3.2.1 for the default router.

VLAN AGGREGATION PROPERTIES

VLAN aggregation is a very specific application, and the following properties apply to its operation:

- All broadcast and unknown traffic remains local to the sub-VLAN and does not cross the sub-VLAN boundary. All traffic within the sub-VLAN is switched by the sub-VLAN, allowing traffic separation between sub-VLANs (while using the same default router address among the sub-VLANs).
- Hosts are located on the sub-VLAN. Each host can assume any IP address within the address range of the super-VLAN router interface. Hosts on the sub-VLAN are expected to have the same network mask as the super-VLAN, and have their default router set to the IP address or the super-VLAN.
- All traffic (IP unicast and IP multicast) between sub-VLANs is routed through the super-VLAN. For example, no ICMP redirects are generated for traffic between sub-VLANs, because the super-VLAN is responsible for sub-VLAN routing. Unicast IP traffic across the sub-VLANs is facilitated by the automatic addition of an ARP entry (similar to a proxy ARP entry) when a sub-VLAN is added to a super-VLAN. This feature can be disabled for security purposes.
- IP multicast traffic between sub-VLANs is routed when an IP multicast routing protocol is enabled on the super-VLAN.

VLAN AGGREGATION LIMITATIONS

The following limitations apply to VLAN aggregation:

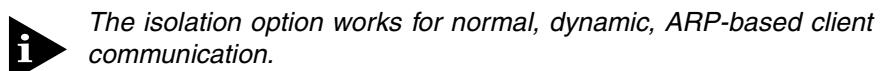
- No additional routers may be located in a sub-VLAN. This feature is only applicable for “leaves” of a network.
- A sub-VLAN cannot be a super-VLAN, and vice-versa.
- Sub-VLANs are not assigned an IP address.
- Typically, a super-VLAN has no ports associated with it, except in the case of running ESRP.
- If a client is moved from one sub-VLAN to another, you must clear the IP ARP cache at the client and the switch, in order to resume communication.

ISOLATION OPTION FOR COMMUNICATION BETWEEN SUB-VLANs

To facilitate communication between sub-VLANs, by default, an entry is made in the IP ARP table of the super-VLAN that performs a proxy ARP function. This allows clients on one sub-VLAN to communicate with clients on another sub-VLAN. In certain circumstances, intra-sub-VLAN communication may not be desired for isolation reasons.

To prevent normal communication between sub-VLANs, disable the automatic addition of the IP ARP entries on the super-VLAN, using the command:

```
disable subvlan-proxy-arp vlan <super-vlan name>
```



VLAN AGGREGATION COMMANDS

[Table 11-2](#) describes VLAN aggregation commands.

Table 11-2: VLAN Aggregation Commands

Command	Description
config vlan <super-vlan name> add secondary-ip <ipaddress> {<mask>}	Adds a secondary IP address to the super-VLAN for responding to ICMP ping requests.
config vlan <super-vlan name> add subvlan <sub-vlan name>	Adds a sub-VLAN to a super-VLAN.
config vlan <super-vlan name> delete secondary-ip <ipaddress> {<mask>}	Deletes a secondary IP address to the super-VLAN for responding to ICMP ping requests.
config vlan <super-vlan name> delete subvlan <sub-vlan name>	Deletes a sub-VLAN from a super-VLAN.
disable subvlan-proxy-arp vlan [<super-vlan name> all]	Disables sub-VLAN entries in the proxy ARP table.
enable subvlan-proxy-arp vlan [<super-vlan name> all]	Enables the automatic entry of sub-VLAN information in the proxy ARP table.

VLAN AGGREGATION EXAMPLE

The follow example illustrates how to configure VLAN aggregation. The VLAN *vsuper* is created as a super-VLAN, and sub-VLANs, *vsub1*, *vsub2*, and *vsub3* are added to it.

- 1 Create and assign an IP address to a VLAN designated as the super-VLAN. This VLAN should have no member ports. Be sure to enable IP forwarding, and any desired routing protocol, on the switch.

```
create vlan vsuper
config vsuper ipaddress 192.201.3.1/24
enable ipforwarding
enable ospf
config ospf add vsuper
```

- 2 Create and add ports to the sub-VLANs.

```
create vlan vsub1
con vsub1 add port 10-12
create vlan vsub2
config vsub2 add po 13-15
create vlan vsub3
config vsub3 add po 16-18
```

- 3 Configure the super-VLAN by adding the sub-VLANs.

```
config vsuper add subvlan vsub1
config vsuper add subvlan vsub2
config vsuper add subvlan vsub3
```

- 4 Optionally, disable communication among sub-VLANs.

```
disable subvlan-proxy-arp <super-VLAN name>
```

VERIFYING THE VLAN AGGREGATION CONFIGURATION

The following commands can be used to verify proper VLAN aggregation configuration.

- `show vlan` — Indicates the membership of a sub-VLANs in a super-VLAN.
- `show iparp` — Indicates an ARP entry that contains sub-VLAN information. Communication with a client on a sub-VLAN must have occurred in order for an entry to be made in the ARP table.

CONFIGURING DHCP/BOOTP RELAY

Once IP unicast routing is configured, you can configure the switch to forward Dynamic Host Configuration Protocol (DHCP) or BOOTP requests coming from clients on subnets being serviced by the switch and going to hosts on different subnets. This feature can be used in various applications, including DHCP services between Windows NT servers and clients running Windows 95. To configure the relay function, do the following:

- 1 Configure VLANs and IP unicast routing.
- 2 Enable the DHCP or BOOTP relay function, using the following command:

```
enable bootprelay
```

- 3 Configure the addresses to which DHCP or BOOTP requests should be directed, using the following command:

```
config bootprelay add <ipaddress>
```

To delete an entry, use the following command:

```
config bootprelay delete {<ipaddress> | all}
```

VERIFYING THE DHCP/BOOTP RELAY CONFIGURATION

To verify the DHCP/BOOTP relay configuration, use the following command:

```
show ipconfig
```

This command displays the configuration of the BOOTP relay service, and the addresses that are currently configured.

UDP-FORWARDING

UDP-forwarding is a flexible and generalized routing utility for handling the directed forwarding of broadcast UDP packets. UDP-forwarding allows applications, such as multiple DHCP relay services from differing sets of VLANs, to be directed to different DHCP servers. The following rules apply to UDP broadcast packets handled by this feature:

- If the UDP profile includes BOOTP or DHCP, it is handled according to guidelines in RFC 1542.

- If the UDP profile includes other types of traffic, these packets have the IP destination address modified as configured, and changes are made to the IP and UDP checksums and decrements to the TTL field, as appropriate.

If the UDP-forwarding is used for BOOTP or DHCP forwarding purposes, do not configure or use the existing `bootprelay` function. However, if the previous `bootprelay` functions are adequate, you may continue to use them.

CONFIGURING UDP-FORWARDING

To configure UDP-forwarding, the first thing you must do is create a UDP-forward destination profile. The profile describes the types of UDP packets (by port number) that are used, and where they are to be forwarded. You must give the profile a unique name, in the same manner as a VLAN, protocol filter, or Spanning Tree Domain.

Next, configure a VLAN to make use of the UDP-forwarding profile. As a result, all incoming traffic from the VLAN that matches the UDP profile is handled as specified in the UDP-forwarding profile.

A maximum of ten UDP-forwarding profiles can be defined. Each named profile may contain a maximum of eight “rules” defining the UDP port, and destination IP address or VLAN. A VLAN can make use of a single UDP-forwarding profile. UDP packets directed toward a VLAN use an all-ones broadcast on that VLAN.

UPD-FORWARDING EXAMPLE

In this example, the VLAN *Marketing* and the VLAN *Operations* are pointed toward a specific backbone DHCP server (with IP address 10.1.1.1) and a backup server (with IP address 10.1.1.2). Additionally, the VLAN *LabUser* is configured to use any responding DHCP server on a separate VLAN called *LabSvrs*.

The commands for this configuration are as follows:

```
create udp-profile backbonedhcp
create udp-profile labdhcp
config backbonedhcp add 67 ipaddress 10.1.1.1
config backbonedhcp add 67 ipaddress 10.1.1.2
config labdhcp add 67 vlan labsvrs
config marketing udp-profile backbonedhcp
config operations udp-profile backbonedhcp
config labuser udp-profile labdhcp
```

ICMP PACKET PROCESSING

As ICMP packets are routed or generated, you can take various actions to control distribution. For ICMP packets typically generated or observed as part of the routing function, you can assert control on a per-type, per-VLAN basis. You would alter the default settings for security reasons: to restrict the success of tools that can be used to find an important application, host, or topology information. The controls include the disabling of transmitting ICMP messages associated with unreachables, port-unreachables, time-exceeded, parameter-problems, redirects, time-stamp, and address-mask requests.

For ICMP packets that are typically routed, you can apply access lists to restrict forwarding behavior. Access lists are described in [Chapter 16](#).

UDP-FORWARDING COMMANDS

[Table 11-3](#) describes the commands used to configure UDP-forwarding.

Table 11-3: UDP-Forwarding Commands

Command	Description
config udp-profile <profile_name> add <udp_port> [<vlan <name>> ipaddress <dest_ipaddress>]	Adds a forwarding entry to the specified UDP-forwarding profile name. All broadcast packets sent to <udp_port> are forwarded to either the destination IP address (unicast or subnet directed broadcast) or to the specified VLAN as an all-ones broadcast.
config udp-profile <profile_name> delete <udp_port> [<vlan <name>> ipaddress <dest_ipaddress>]	Deletes a forwarding entry from the specified udp-profile name.
config vlan <name> udp-profile <profile_name>	Assigns a UDP-forwarding profile to the source VLAN. Once the UDP profile is associated with the VLAN, the switch picks up any broadcast UDP packets that matches with the user configured UDP port number, and forwards those packets to the user-defined destination. If the UDP port is the DHCP/BOOTP port number, appropriate DHCP/BOOTP proxy functions are invoked.

Table 11-3: UDP-Forwarding Commands (continued)

Command	Description
create udp-profile <profile_name>	Creates a UDP-forwarding profile. You must use a unique name for the UDP-forwarding profile.
delete udp-profile <profile_name>	Deletes a UDP-forwarding profile.
show udp-profile {<profile_name>}	Displays the profile names, input rules of UDP port, destination IP address, or VLAN and the source VLANs to which the profile is applied.
unconfig udp-profile vlan [<name> all]	Removes the UDP-forwarding profile configuration for one or all VLANs.

IP COMMANDS

[Table 11-4](#) describes the commands used to configure basic IP settings.

Table 11-4: Basic IP Commands

Command	Description
clear iparp {<ipaddress> <mask> vlan <name>}	Removes dynamic entries in the IP ARP table. Permanent IP ARP entries are not affected.
clear ipfdb {<ipaddress> vlan <name> }	Removes the dynamic entries in the IP forwarding database. If no options are specified, all dynamic IP FDB entries are removed.
config bootprelay add <ipaddress>	Adds the IP destination address to forward BOOTP packets.
config bootprelay delete [<ipaddress> all]	Removes one or all IP destination addresses for forwarding BOOTP packets.
config iparp add <ipaddress> <mac_address>	Adds a permanent entry to the ARP table. Specify the IP address and MAC address of the entry.

Table 11-4: Basic IP Commands (continued)

Command	Description
config iparp add proxy <ipaddress> {<mask>} {<mac_address>} {always}	Configures proxy ARP entries. When <code>mask</code> is not specified, an address with the mask 255.255.255.255 is assumed. When <code>mac_address</code> is not specified, the MAC address of the switch is used in the ARP Response. When <code>always</code> is specified, the switch answers ARP Requests without filtering requests that belong to the same subnet of the receiving router interface.
config iparp delete <ipaddress>	Deletes an entry from the ARP table. Specify the IP address of the entry.
config iparp delete proxy [<ipaddress> {<mask>} all]	Deletes one or all proxy ARP entries.
config iparp timeout <minutes>	Configures the IP ARP timeout period. The default setting is 20 minutes. A setting of 0 disables ARP aging. The maximum aging time is 32 minutes.
config tcp-sync-rate <number_sync_per_sec>	Configures a limit for the switch to process TCP connection requests. If the connection request rate is higher than the specified rate, or the total number of outstanding connection requests exceeds the system limit, the system ages out incomplete connection requests at a faster rate. The range is 5 to 200,000. The default setting is 25 connection requests per second.
disable bootp vlan [<name> all]	Disables the generation and processing of BOOTP packets.
disable bootprelay	Disables the forwarding of BOOTP requests.
disable ipforwarding {vlan <name>}	Disables routing for one or all VLANs.
disable ipforwarding broadcast {vlan <name>}	Disables routing of broadcasts to other networks.
disable loopback-mode vlan [<name> all]	Disables loopback-mode on an interface.
disable multinetting	Disables IP multinetting on the system.
enable bootp vlan [<name> all]	Enables the generation and processing of BOOTP packets on a VLAN to obtain an IP address for the VLAN from a BOOTP server. The default setting is enabled for all VLANs.

Table 11-4: Basic IP Commands (continued)

Command	Description
enable bootprelay	Enables the forwarding of BOOTP and Dynamic Host Configuration Protocol (DHCP) requests.
enable ipforwarding {vlan <name>}	Enables IP routing for one or all VLANs. If no argument is provided, enables routing for all VLANs that have been configured with an IP address. The default setting for ipforwarding is disabled.
enable ipforwarding broadcast {vlan <name>}	Enables forwarding IP broadcast traffic for one or all VLANs. If no argument is provided, enables broadcast forwarding for all VLANs. To enable, ipforwarding must be enabled on the VLAN. The default setting is disabled.
enable loopback-mode vlan [<name> all]	Enables a loopback mode on an interface. If loopback is enabled, the router interface remains in the UP state, even if no ports are defined in the VLAN. As a result, the subnet is always advertised as one of the available routes.
enable multinetting	Enables IP multinetting on the system.

Table 11-5 describes the commands used to configure the IP route table.

Table 11-5: Route Table Configuration Commands

Command	Description
config iproute add <ipaddress> <mask> <gateway> <metric> {unicast-only multicast-only}	Adds a static address to the routing table. Use a value of 255.255.255.255 for mask to indicate a host entry
config iproute add blackhole <ipaddress> <mask> {unicast-only multicast-only}	Adds a blackhole address to the routing table. All traffic destined for the configured IP address is dropped, and no Internet Control Message Protocol (ICMP) message is generated.

Table 11-5: Route Table Configuration Commands (continued)

Command	Description
config iproute add default <gateway> {<metric>} {unicast-only multicast-only}	Adds a default gateway to the routing table. A default gateway must be located on a configured IP interface. If no metric is specified, the default metric of 1 is used. Use the unicast-only or multicast-only options to specify a particular traffic type. If not specified, both unicast and multicast traffic uses the default route.
config iproute delete <ipaddress> <mask> <gateway>	Deletes a static address from the routing table.
config iproute delete blackhole <ipaddress> <mask>	Deletes a blackhole address from the routing table.
config iproute delete default <gateway>	Deletes a default gateway from the routing table.
config iproute priority [rip bootp icmp static ospf-intra ospf-inter ospf-as-external ospf-extern1 ospf-extern2] <priority>	Changes the priority for all routes from a particular route origin.
disable iproute sharing	Disables load sharing for multiple routes.
enable iproute sharing	Enables load sharing if multiple routes to the same destination are available. Only paths with the same lowest cost are shared. The default setting is disabled.
rtlookup [<ipaddress> <hostname>]	Performs a look-up in the route table to determine the best route to reach an IP address.

Table 11-6 describes the commands used to configure IP options and the ICMP protocol.

Table 11-6: ICMP Configuration Commands

Command	Description
config irdp [multicast broadcast]	Configures the destination address of the router advertisement messages. The default setting is <code>multicast</code> .

Table 11-6: ICMP Configuration Commands (continued)

Command	Description
config irdp <mininterval> <maxinterval> <lifetime> <preference>	<p>Configures the router advertisement message timers, using seconds. Specify:</p> <ul style="list-style-type: none"> ■ mininterval — The minimum amount of time between router advertisements. The default setting is 450 seconds. ■ maxinterval — The maximum time between router advertisements. The default setting is 600 seconds. ■ lifetime — The default setting is 1,800 seconds. ■ preference — The preference level of the router. An ICMP Router Discover Protocol (IRDP) client always uses the router with the highest preference level. Change this setting to encourage or discourage the use of this router. The default setting is 0.
disable icmp parameter-problem {vlan <name>}	Disables the generation of ICMP messages for the parameter problem packet type.
disable ip-option loose-source-route	Disables the loose source route IP option.
disable ip-option record-route	Disables the record route IP option.
disable ip-option record-timestamp	Disables the record timestamp IP option.
disable ip-option strict-source-route	Disables the strict source route IP option.
disable ip-option use-router-alert	Disables the generation of the router alert IP option.
enable icmp address-mask {vlan <name>}	Enables the generation of an ICMP address-mask reply (type 18, code 0) when an ICMP address mask request is received. The default setting is enabled. If a VLAN is not specified, the command applies to all IP interfaces.
enable icmp parameter-problem {vlan <name>}	Enables the generation of ICMP parameter problem packet (type 12) when the switch cannot properly process the IP header or IP option information.
enable icmp parameter-problem {vlan <name>}	Enables the generation of an ICMP parameter-problem message (type 12) when the switch cannot properly process the IP header or IP option information. The default setting is enabled. If a VLAN is not specified, the command applies to all IP interfaces.

Table 11-6: ICMP Configuration Commands (continued)

Command	Description
enable icmp port-unreachable {vlan <name>}	Enables the generation of ICMP port unreachable messages (type 3, code 3) when a TCP or UDP request is made to the switch, and no application is waiting for the request, or access policy denies the request. The default setting is enabled. If a VLAN is not specified, the command applies to all IP interfaces.
enable icmp redirects {vlan <name>}	Enables the generation of an ICMP redirect message (type 5) when a packet must be forwarded out on the ingress port. The default setting is enabled. If a VLAN is not specified, the command applies to all IP interfaces.
enable icmp time-exceeded {vlan <name>}	Enables the generation of an ICMP time exceeded message (type 11) when the TTL field expires during forwarding. IP multicast packets do not trigger ICMP time exceeded messages. The default setting is enabled. If a VLAN is not specified, the command applies to all IP interfaces.
enable icmp timestamp {vlan <name>}	Enables the generation of an ICMP timestamp response (type 14, code 0) when an ICMP timestamp request is received. The default setting is enabled. If a VLAN is not specified, the command applies to all IP interfaces.
enable icmp unreachable {vlan <name>}	Enables the generation of ICMP network unreachable messages (type 3, code 0), and host unreachable messages (type 3, code 1) when a packet cannot be forwarded to the destination because of unreachable route or host. ICMP packet processing on one or all VLANs. The default setting is enabled. If a VLAN is not specified, the command applies to all IP interfaces.
enable icmp useredirects	Enables the modification of route table information when an ICMP redirect message is received. This option applies to the switch when it is <i>not configured for routing</i> . The default setting is disabled.
enable ip-option loose-source-route	Enables the loose source route IP option.
enable ip-option record-route	Enables the record route IP option.
enable ip-option record-timestamp	Enables the record timestamp IP option.
enable ip-option strict-source-route	Enables the strict source route IP option.

Table 11-6: ICMP Configuration Commands (continued)

Command	Description
enable ip-option use-router-alert	Enables the switch to generate the router alert IP option with routing protocol packets.
enable irdp {vlan <name>}	Enables the generation of ICMP router advertisement messages on one or all VLANs. The default setting is enabled.
unconfig icmp	Resets all ICMP settings to the default values.
unconfig irdp	Resets all router advertisement settings to the default values.

ROUTING CONFIGURATION EXAMPLE

Figure 11-3 illustrates a BlackDiamond switch that has three VLANs defined as follows:

- *Finance*
 - Protocol-sensitive VLAN using the IP protocol
 - All ports on slots 1 and 3 have been assigned
 - IP address 192.207.35.1
- *Personnel*
 - Protocol-sensitive VLAN using the IP protocol
 - All ports on slots 2 and 4 have been assigned
 - IP address 192.207.36.1
- *MyCompany*
 - Port-based VLAN
 - All ports on slots 1 through 4 have been assigned

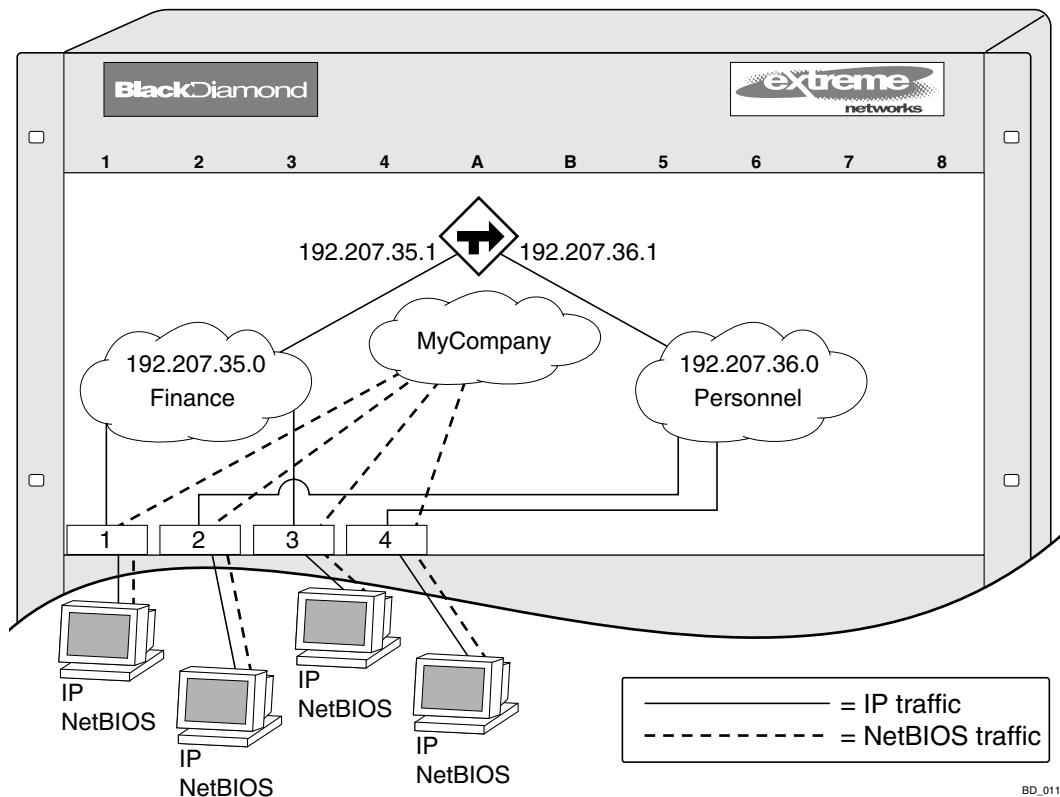


Figure 11-3: Unicast routing configuration example

The stations connected to the system generate a combination of IP traffic and NetBIOS traffic. The IP traffic is filtered by the protocol-sensitive VLANs. All other traffic is directed to the VLAN *MyCompany*.

In this configuration, all IP traffic from stations connected to slots 1 and 3 have access to the router by way of the VLAN *Finance*. Ports on slots 2 and 4 reach the router by way of the VLAN *Personnel*. All other traffic (NetBIOS) is part of the VLAN *MyCompany*.

The example in [Figure 11-3](#) is configured as follows:

```
create vlan Finance
create vlan Personnel
create vlan MyCompany
```

```

config Finance protocol ip
config Personnel protocol ip

config Finance add port 1:*,3:*
config Personnel add port 2:*,4:*
config MyCompany add port all

config Finance ipaddress 192.207.35.1
config Personnel ipaddress 192.207.36.1

config rip add vlan Finance
config rip add vlan Personnel

enable ipforwarding
enable rip

```

DISPLAYING ROUTER SETTINGS

To display settings for various IP routing components, use the commands listed in [Table 11-7](#).

Table 11-7: Router Show Commands

Command	Description
show iparp {<ipaddress vlan <name> permanent}	Displays the IP Address Resolution Protocol (ARP) table. You can filter the display by IP address, VLAN, or permanent entries.
show iparp proxy {<ipaddress> {<mask>}}	Displays the proxy ARP table.
show ipconfig {vlan <name>}	Displays configuration information for one or all VLANs.

Table 11-7: Router Show Commands (continued)

Command	Description
show ipfdb {<ipaddress> <netmask> vlan <name>} {sorted}	Displays the contents of the IP forwarding database (FDB) table. If no option is specified, all IP FDB entries are displayed.
show iproute {priority vlan <name> permanent <ipaddress> <mask> origin [direct static blackhole rip bootp icmp ospf-intra ospf-inter ospf-as-external ospf-extern1 ospf-extern2]} {sorted}	Displays the contents of the IP routing table or the route origin priority.
show ipstats {vlan <name>}	Displays IP statistics for the CPU of the system.

RESETTING AND DISABLING ROUTER SETTINGS

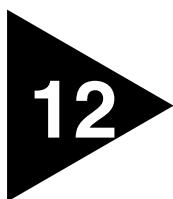
To return router settings to their defaults and disable routing functions, use the commands listed in [Table 11-8](#)

Table 11-8: Router Reset and Disable Commands

Command	Description
clear iparp {<ipaddress> vlan <name>}	Removes dynamic entries in the IP ARP table. Permanent IP ARP entries are not affected.
clear ipfdb {<ipaddress> <netmask> vlan <name>}	Removes the dynamic entries in the IP forwarding database. If no options are specified, all IP FDB entries are removed.
disable bootp vlan [<name> all]	Disables the generation and processing of BOOTP packets.
disable bootprelay	Disables the forwarding of BOOTP requests.
disable icmp address-mask {vlan <name>}	Disables the generation of an ICMP address-mask reply messages. If a VLAN is not specified, the command applies to all IP interfaces.
disable icmp parameter-problem {vlan <name>}	Disables the generation of ICMP parameter-problem messages. If a VLAN is not specified, the command applies to all IP interfaces.
disable icmp port-unreachables {vlan <name>}	Disables the generation of ICMP port unreachable messages. If a VLAN is not specified, the command applies to all IP interfaces.

Table 11-8: Router Reset and Disable Commands (continued)

Command	Description
disable icmp redirects {vlan <name>}	Disables the generation of ICMP redirect messages. If a VLAN is not specified, the command applies to all IP interfaces.
disable icmp time-exceeded {vlan <name>}	Disables the generation of ICMP time exceeded messages. If a VLAN is not specified, the command applies to all IP interfaces.
disable icmp timestamp {vlan <name>}	Disables the generation of ICMP timestamp response messages. If a VLAN is not specified, the command applies to all IP interfaces.
disable icmp unreachable {vlan <name>}	Disables the generation of ICMP network unreachable messages and host unreachable messages. If a VLAN is not specified, the command applies to all IP interfaces.
disable icmp useredirects	Disables the changing of routing table information when an ICMP redirect message is received.
disable ipforwarding {vlan <name>}	Disables routing for one or all VLANs.
disable ipforwarding broadcast {vlan <name>}	Disables routing of broadcasts to other networks.
disable irdp {vlan <name>}	Disables the generation of router advertisement messages on one or all VLANs.
unconfig icmp	Resets all ICMP settings to the default values.
unconfig irdp	Resets all router advertisement settings to the default values.



Interior Gateway Routing Protocols

This chapter describes the following topics:

- [Overview on page 12-2](#)
- [Overview of RIP on page 12-3](#)
- [Overview of OSPF on page 12-5](#)
- [Route Re-distribution on page 12-10](#)
- [Configuring RIP on page 12-14](#)
- [RIP Configuration Example on page 12-17](#)
- [Displaying RIP Settings on page 12-19](#)
- [Resetting and Disabling RIP on page 12-20](#)
- [Configuring OSPF on page 12-21](#)
- [OSPF Configuration Example on page 12-25](#)
- [Displaying OSPF Settings on page 12-28](#)
- [Resetting and Disabling OSPF Settings on page 12-28](#)

This chapter assumes that you are already familiar with IP unicast routing. If not, refer to the following publications for additional information:

- RFC 1058 — *Routing Information Protocol (RIP)*
- RFC 1723 — *RIP Version 2*
- RFC 2178 — *OSPF Version 2*

- *Interconnections: Bridges and Routers*
by Radia Perlman
ISBN 0-201-56332-0
Published by Addison-Wesley Publishing Company

OVERVIEW

The switch supports the use of two interior gateway protocols (IGPs); the Routing Information Protocol (RIP) and the Open Shortest Path First (OSPF) protocol for IP unicast routing.

RIP is a distance-vector protocol, based on the Bellman-Ford (or distance-vector) algorithm. The distance-vector algorithm has been in use for many years, and is widely deployed and understood.

OSPF is a link-state protocol, based on the Dijkstra link-state algorithm. OSPF is a newer Interior Gateway Protocol (IGP), and solves a number of problems associated with using RIP on today's complex networks.



Both RIP and OSPF can be enabled on a single VLAN.

RIP VERSUS OSPF

The distinction between RIP and OSPF lies in the fundamental differences between distance-vector protocols and link-state protocols. Using a distance-vector protocol, each router creates a unique routing table from summarized information obtained from neighboring routers. Using a link-state protocol, every router maintains an identical routing table created from information obtained from all routers in the autonomous system. Each router builds a shortest path tree, using itself as the root. The link-state protocol ensures that updates sent to neighboring routers are acknowledged by the neighbors, verifying that all routers have a consistent network map.

The biggest advantage of using RIP is that it is relatively simple to understand and implement, and it has been the *de facto* routing standard for many years.

RIP has a number of limitations that can cause problems in large networks, including the following:

- A limit of 15 hops between the source and destination networks
- A large amount of bandwidth taken up by periodic broadcasts of the entire routing table
- Slow convergence
- Routing decisions based on hop count; no concept of link costs or delay
- Flat networks; no concept of areas or boundaries

OSPF offers many advantages over RIP, including the following:

- No limitation on hop count
- Route updates multicast only when changes occur
- Faster convergence
- Support for load balancing to multiple routers based on the actual cost of the link
- Support for hierarchical topologies where the network is divided into areas

The details of RIP and OSPF are explained later in this chapter.

OVERVIEW OF RIP

RIP is an Interior Gateway Protocol (IGP) first used in computer routing in the Advanced Research Projects Agency Network (ARPAnet) as early as 1969. It is primarily intended for use in homogeneous networks of moderate size.

To determine the best path to a distant network, a router using RIP always selects the path that has the least number of hops. Each router that data must traverse is considered to be one hop.

ROUTING TABLE

The routing table in a router using RIP contains an entry for every known destination network. Each routing table entry contains the following information:

- IP address of the destination network
- Metric (hop count) to the destination network

- IP address of the next router
- Timer that tracks the amount of time since the entry was last updated

The router exchanges an update message with each neighbor every 30 seconds (default value), or if there is a change to the overall routed topology (also called *triggered updates*). If a router does not receive an update message from its neighbor within the route timeout period (180 seconds by default), the router assumes the connection between it and its neighbor is no longer available.

SPLIT HORIZON

Split horizon is a scheme for avoiding problems caused by including routes in updates sent to the router from which the route was learned. Split horizon omits routes learned from a neighbor in updates sent to that neighbor.

POISON REVERSE

Like split horizon, poison reverse is a scheme for eliminating the possibility of loops in the routed topology. In this case, a router advertises a route over the same interface that supplied the route, but the route uses a hop count of 16, defining it as unreachable.

TRIGGERED UPDATES

Triggered updates occur whenever a router changes the metric for a route, and it is required to send an update message immediately, even if it is not yet time for a regular update message to be sent. This will generally result in faster convergence, but may also result in more RIP-related traffic.

ROUTE ADVERTISEMENT OF VLANs

VLANs that are configured with an IP address, but are configured to not route IP or are not configured to run RIP, do not have their subnets advertised by RIP. Only those VLANs that are configured with an IP address and are configured to route IP and run RIP have their subnets advertised.

RIP VERSION 1 VERSUS RIP VERSION 2

A new version of RIP, called RIP version 2, expands the functionality of RIP version 1 to include the following:

- Variable-Length Subnet Masks (VLSMs)
- Support for next-hop addresses, which allows for optimization of routes in certain environments.
- Multicasting

RIP version 2 packets can be multicast instead of being broadcast, reducing the load on hosts that do not support routing protocols.



If you are using RIP with supernetting/Classless Inter-Domain Routing (CIDR), you must use RIPv2 only. In addition, RIP route aggregation must be turned off.

OVERVIEW OF OSPF

OSPF is a link-state protocol that distributes routing information between routers belonging to a single IP domain, also known as an *autonomous system* (AS). In a link-state routing protocol, each router maintains a database describing the topology of the autonomous system. Each participating router has an identical database maintained from the perspective of that router.

From the link-state database (LSDB), each router constructs a tree of shortest paths, using itself as the root. The shortest path tree provides the route to each destination in the autonomous system. When several equal-cost routes to a destination exist, traffic can be distributed among them. The cost of a route is described by a single metric.

LINK-STATE DATABASE

Upon initialization, each router transmits a link-state advertisement (LSA) on each of its interfaces. LSAs are collected by each router and entered into the LSDB of each router. OSPF uses flooding to distribute LSAs between routers. Any change in routing information is sent to all of the routers in the network. All routers within an area have the exact same LSDB. [Table 12-1](#) describes LSA type numbers.

Table 12-1: LSA Type Numbers

Type Number	Description
1	Router LSA
2	Network LSA
3	Summary LSA
4	AS summary LSA
5	AS external LSA
7	NSSA external LSA

AREAS

OSPF allows parts of a network to be grouped together into *areas*. The topology within an area is hidden from the rest of the autonomous system. Hiding this information enables a significant reduction in LSA traffic, and reduces the computations needed to maintain the LSDB. Routing within the area is determined only by the topology of the area.

The three types of routers defined by OSPF are as follows:

- Internal Router (IR)
An internal router has all of its interfaces within the same area.
- Area Border Router (ABR)
An ABR has interfaces in multiple areas. It is responsible for exchanging summary advertisements with other ABRs. You can create a maximum of 7 non-zero areas.
- Autonomous System Border Router (ASBR)
An ASBR acts as a gateway between OSPF and other routing protocols, or other autonomous systems.

AREA 0

Any OSPF network that contains more than one area is required to have an area configured as area 0, also called the *backbone*. All areas in an autonomous system must be connected to the backbone. When designing networks, you should start with area 0, and then expand into other areas.

The backbone allows summary information to be exchanged between ABRs. Every ABR hears the area summaries from all other ABRs. The ABR then forms a picture of the distance to all networks outside of its area by examining the collected advertisements, and adding in the backbone distance to each advertising router.

When a VLAN is configured to run OSPF, you must configure the area for the VLAN. If you want to configure the VLAN to be part of a different OSPF area, use the following command:

```
config ospf vlan <name> area <areaid>
```

If this is the first instance of the OSPF area being used, you must create the area first using the following command:

```
create ospf area <areaid>
```

STUB AREAS

OSPF allows certain areas to be configured as *stub areas*. A stub area is connected to only one other area. The area that connects to a stub area can be the backbone area. External route information is not distributed into stub areas. Stub areas are used to reduce memory consumption and computation requirements on OSPF routers.

NOT-SO-STUBBY-AREAS (NSSA)

NSSAs are similar to the existing OSPF stub area configuration option, but have the following two additional capabilities:

- External routes originating from an ASBR connected to the NSSA can be advertised within the NSSA.
- External routes originating from the NSSA can be propagated to other areas, including the backbone area.

The CLI command to control the NSSA function is similar to the command used for configuring a stub area, as follows:

```
config ospf area <area_id> nssa {summary | nosummary} stub-default-cost <cost> {translate}
```

The translate option determines whether type 7 LSAs are translated into type 5 LSAs. When configuring an OSPF area as an NSSA, the translate should only be used on NSSA border routers, where translation is to be enforced. If translate is not used on any NSSA border router in a NSSA, one of the ABRs for that NSSA is elected to

perform translation (as indicated in the NSSA specification). The option should not be used on NSSA internal routers. Doing so inhibits correct operation of the election algorithm.

NORMAL AREA

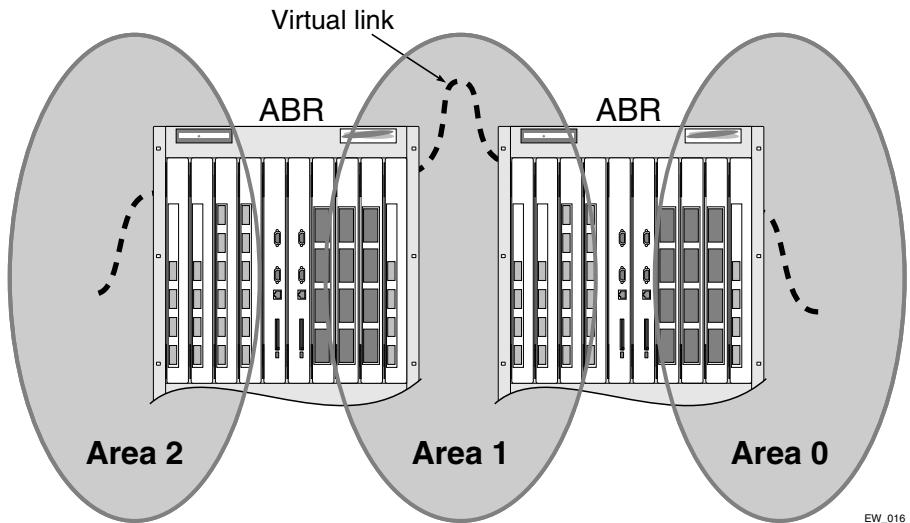
A normal area is an area that is not any of the following:

- Area 0
- Stub area
- NSSA

Virtual links can be configured through normal areas. External routes can be distributed into normal areas.

VIRTUAL LINKS

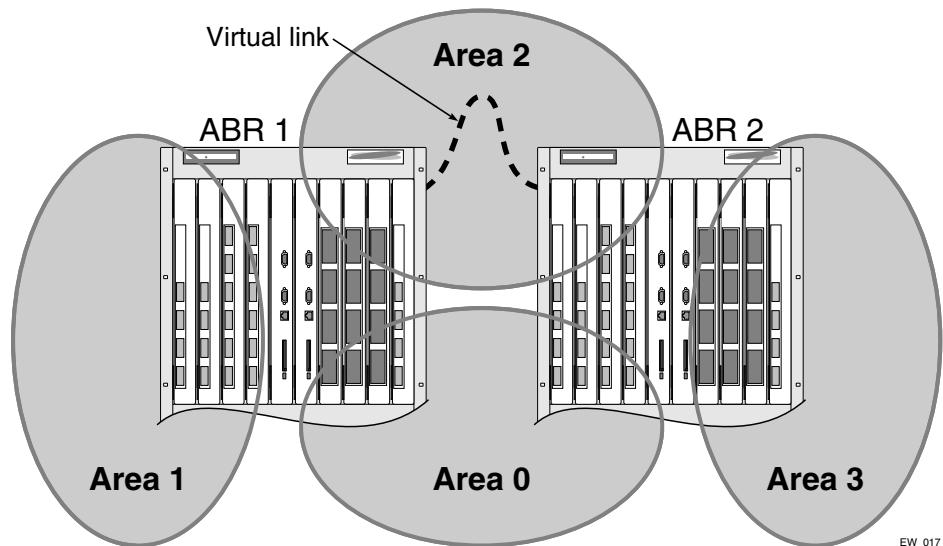
In the situation when a new area is introduced that does not have a direct physical attachment to the backbone, a *virtual link* is used. A virtual link provides a logical path between the ABR of the disconnected area and the ABR of the normal area that connects to the backbone. A virtual link must be established between two ABRs that have a common area, with one ABR connected to the backbone. [Figure 12-1](#) illustrates a virtual link.



EW_016

Figure 12-1: Virtual link using Area 1 as a transit area

Virtual links are also used to repair a discontiguous backbone area. For example, in [Figure 12-2](#), if the connection between ABR1 and the backbone fails, the connection using ABR2 provides redundancy so that the discontiguous area can continue to communicate with the backbone using the virtual link.



EW_017

Figure 12-2: Virtual link providing redundancy

ROUTE RE-DISTRIBUTION

Both RIP and OSPF can be enabled simultaneously on the switch. Route re-distribution allows the switch to exchange routes, including static routes, between the two routing protocols. [Figure 12-3](#) shows an example of route re-distribution between an OSPF autonomous system and a RIP autonomous system.

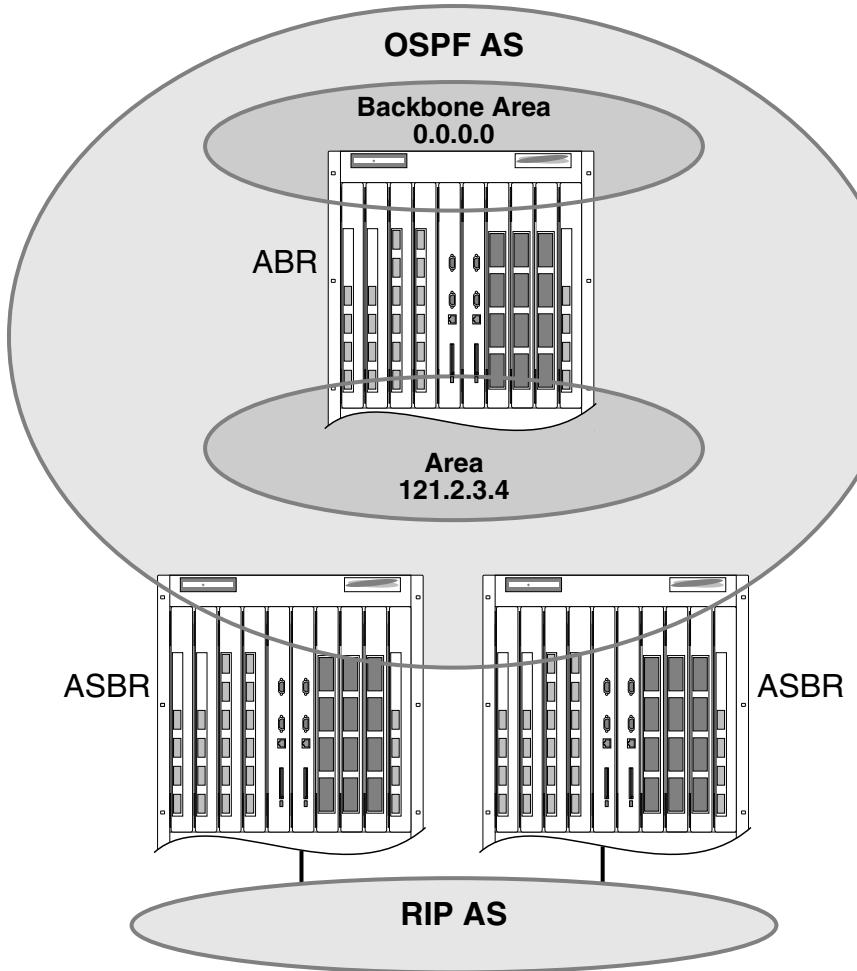


Figure 12-3: Route re-distribution

CONFIGURING ROUTE RE-DISTRIBUTION

Exporting routes from OSPF to RIP, and from RIP to OSPF, are discrete configuration functions. To run OSPF and RIP simultaneously, you must first configure both protocols and then verify the independent operation of each. Then you can configure the routes to export from OSPF to RIP and the routes to export from RIP to OSPF.

RE-DISTRIBUTING ROUTES INTO OSPF

Enable or disable the exporting of RIP, static, and direct (interface) routes to OSPF, using the following commands:

```
enable ospf export [static | rip | direct] cost <metric> [ase-type-1 |  
ase-type-2] {tag <number>}
```

```
disable ospf export [static | rip | direct]
```

These commands enable or disable the exporting of RIP, static, and direct routes by way of LSA to other OSPF routers as AS-external type 1 or type 2 routes. The default setting is disabled.

The cost metric is inserted for all RIP-learned, static, and direct routes injected into OSPF. If the cost metric is set to zero, the cost is inserted from the route. The tag value is used only by special routing applications. Use the number zero if you do not have specific requirements for using a tag. The tag value in this instance has no relationship with 802.1Q VLAN tagging.

Enable or disable the export of Virtual IP addresses to other OSPF routers, using the following commands:

```
enable ospf export vip cost <metric> [ase-type-1 | ase-type-2] {tag  
<number>}
```

```
disable ospf export vip
```

Verify the configuration using the command:

```
show ospf
```

PREVIOUS RELEASE ISSUES WITH OSPF RE-DISTRIBUTION

When re-distributing RIP routes you should turn off RIP aggregation unless you are expertly familiar with the possible consequences and impact. By default, new configurations of RIP using ExtremeWare 4.0 and above disable RIP aggregation. In previous ExtremeWare versions, RIP aggregation is enabled by default. This configuration is preserved when upgrading to ExtremeWare 4.0. Verify the configuration using the command `show rip`.



In versions of ExtremeWare prior to release 6.0, direct routes corresponding to the interfaces on which RIP was enabled were exported into OSPF as part of RIP routes, using the command enable ospf export rip. Using ExtremeWare 6.0 and above, you must configure ExtremeWare to export these direct routes to OSPF. You can use an access profile to filter unnecessary direct routes, using the command config ospf direct-filter [<access-profile> | none].

RE-DISTRIBUTING ROUTES INTO RIP

Enable or disable the exporting of static, direct, and OSPF-learned routes into the RIP domain, using the following commands:

```
enable rip export [static | direct | ospf | ospf-intra | ospf-inter | ospf-extern1 | ospf-extern2 | vip] cost <metric> tag <number>
```

```
disable rip export [static | direct | ospf | ospf-intra | ospf-inter | ospf-extern1 | ospf-extern2 | vip]
```

These commands enable or disable the exporting of static, direct, and OSPF-learned routes into the RIP domain. You can choose which types of OSPF routes are injected, or you can simply choose `ospf`, which will inject all learned OSPF routes regardless of type. The default setting is disabled.

OSPF TIMERS AND AUTHENTICATION

Configuring OSPF timers and authentication on a per-area basis is a shorthand for applying the timers and authentication to each VLAN in the area at the time of configuration. If you add more VLANs to the area, you must configure the timers and authentication for the new VLANs explicitly.

CONFIGURING RIP

[Table 12-2](#) describes the commands used to configure RIP.

Table 12-2: RIP Configuration Commands

Command	Description
config rip add vlan [<name> all]	Configures RIP on an IP interface. When an IP interface is created, per-interface RIP configuration is disabled by default.
config rip delete vlan [<name> all]	Disables RIP on an IP interface. When RIP is disabled on the interface, the parameters are not reset to their defaults.
config rip garbagetime {<seconds>}	Configures the RIP garbage time. The timer granularity is 10 seconds. The default setting is 120 seconds.
config rip routetimeout {<seconds>}	Configures the route timeout. The default setting is 180 seconds.
config rip rxmode [none v1only v2only any] {vlan <name>}	Changes the RIP receive mode for one or all VLANs. Specify: <ul style="list-style-type: none"> ■ none — Drop all received RIP packets. ■ v1only — Accept only RIP v1 format packets. ■ v2only — Accept only RIP v2 format packets. ■ any — Accept both RIP v1 and v2 packets. If no VLAN is specified, the setting is applied to all VLANs. The default setting is any.

Table 12-2: RIP Configuration Commands (continued)

Command	Description
config rip txmode [none v1only v1comp v2only] {vlan <name>}	Changes the RIP transmission mode for one or all VLANs. Specify: <ul style="list-style-type: none">■ none — Do not transmit any packets on this interface.■ v1only — Transmit RIP v1 format packets to the broadcast address.■ v1comp — Transmit RIP v2 format packets to the broadcast address.■ v2only — Transmit RIP v2 format packets to the RIP multicast address. If no VLAN is specified, the setting is applied to all VLANs. The default setting is v2only.
config rip updatetime {<seconds>}	Changes the periodic RIP update timer. The default setting is 30 seconds.
config rip vlan [<name> all] cost <number>	Configures the cost (metric) of the interface. The default setting is 1.
enable rip	Enables RIP. The default setting is disabled.
enable rip aggregation	Enables aggregation of subnet information on interfaces configured to send RIP v2 or RIP v2-compatible traffic. The switch summarizes subnet routes to the nearest class network route. The following rules apply when using RIP aggregation: <ul style="list-style-type: none">■ Subnet routes are aggregated to the nearest class network route when crossing a class boundary.■ Within a class boundary, no routes are aggregated.■ If aggregation is enabled, the behavior is the same as in RIP v1.■ If aggregation is disabled, subnet routes are never aggregated, even when crossing a class boundary. The default setting is disabled.

Table 12-2: RIP Configuration Commands (continued)

Command	Description
enable rip export [static direct ospf ospf-intra ospf-inter ospf-extern1 ospf-extern2 static vip] metric <metric> {tag <number>}	<p>Enables RIP to redistribute routes from other routing functions. Specify one of the following:</p> <ul style="list-style-type: none"> ■ static — Static routes ■ direct — Interface routes (only interfaces that have IP forwarding enabled are exported) ■ ospf — All OSPF routes ■ ospf-intra — OSPF intra-area routes ■ ospf-inter — OSPF inter-area routes ■ ospf-extern1 — OSPF AS-external route type 1 ■ ospf-extern2 — OSPF AS-external route type 2 ■ vip — Virtual IP <p>The <code>metric</code> range is 0-15. If set to 0, RIP uses the route metric obtained from the route origin.</p>
enable rip originate-default {always} cost <metric> {tag <number>}	Configures a default route to be advertised by RIP if no other default route is advertised. If <code>always</code> is specified, RIP always advertises the default route to its neighbors. If <code>always</code> is not specified, RIP adds a default route if there is a reachable default route in the route table.
enable rip poisonreverse	Enables the split horizon with poison-reverse algorithm for RIP. The default setting is enabled. If you enable poison reverse and split horizon, poison reverse takes precedence.
enable rip splithorizon	Enables the split horizon algorithm for RIP. Default setting is enabled.
enable rip triggerupdates	Enables triggered updates. <i>Triggered updates</i> are a mechanism for immediately notifying a router's neighbors when the router adds or deletes routes, or changes the metric of a route. The default setting is enabled.

RIP CONFIGURATION EXAMPLE

[Figure 12-4](#) illustrates a BlackDiamond switch that has three VLANs defined as follows:

- *Finance*
 - Protocol-sensitive VLAN using the IP protocol
 - All ports on slots 1 and 3 have been assigned
 - IP address 192.207.35.1
- *Personnel*
 - Protocol-sensitive VLAN using the IP protocol
 - All ports on slots 2 and 4 have been assigned
 - IP address 192.207.36.1
- *MyCompany*
 - Port-based VLAN
 - All ports on slots 1 through 4 have been assigned

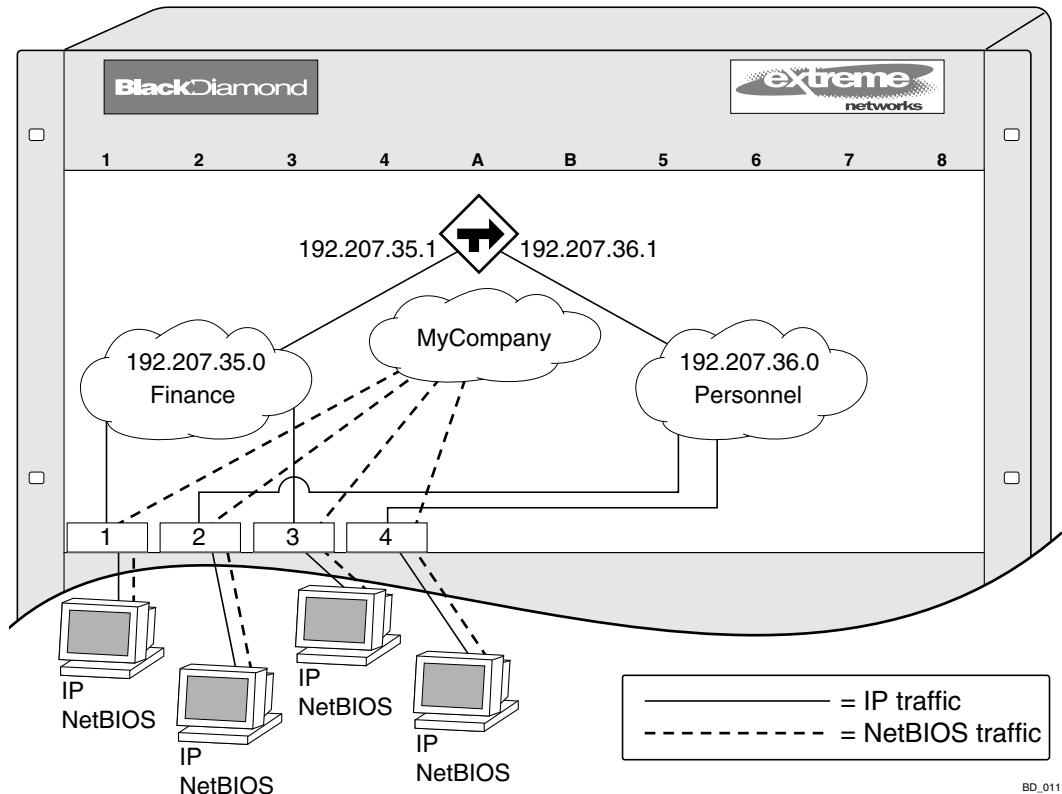


Figure 12-4: RIP configuration example

The stations connected to the system generate a combination of IP traffic and NetBIOS traffic. The IP traffic is filtered by the protocol-sensitive VLANs. All other traffic is directed to the VLAN *MyCompany*.

In this configuration, all IP traffic from stations connected to slots 1 and 3 have access to the router by way of the VLAN *Finance*. Ports on slots 2 and 4 reach the router by way of the VLAN *Personnel*. All other traffic (NetBIOS) is part of the VLAN *MyCompany*.

The example in [Figure 12-4](#) is configured as follows:

```

create vlan Finance
create vlan Personnel
create vlan MyCompany

config Finance protocol ip
config Personnel protocol ip

config Finance add port 1:*,3:*
config Personnel add port 2:*,4:*
config MyCompany add port all

config Finance ipaddress 192.207.35.1
config Personnel ipaddress 192.207.36.1

enable ipforwarding
config rip add vlan all
enable rip

```

DISPLAYING RIP SETTINGS

To display settings for RIP, use the commands listed in [Table 12-3](#).

Table 12-3: RIP Show Commands

Command	Description
show rip {detail}	Displays RIP configuration and statistics for all VLANs.
show rip stat {detail}	Displays RIP-specific statistics for all VLANs.
show rip stat vlan <name>	Displays RIP-specific statistics for a VLAN.
show rip vlan <name>	Displays RIP configuration and statistics for a VLAN.

RESETTING AND DISABLING RIP

To return RIP settings to their defaults, or to disable RIP, use the commands listed in [Table 12-4](#).

Table 12-4: RIP Reset and Disable Commands

Command	Description
config rip delete [vlan <name> all]	Disables RIP on an IP interface. When RIP is disabled on the interface, the parameters are not reset to their defaults.
disable rip	Disables RIP.
disable rip aggregation	Disables the RIP aggregation of subnet information on a RIP v2 interface.
disable rip export [static direct ospf ospf-intra ospf-inter ospf-extern1 ospf-extern2 static vip] metric <metric> {tag <number>}	Disables the distribution of non-RIP routes into the RIP domain.
disable rip originate-default	Disables the advertisement of a default route.
disable rip poisonreverse	Disables poison reverse.
disable rip splithorizon	Disables split horizon.
disable rip triggerupdates	Disables triggered updates.
unconfig rip {vlan <name>}	Resets all RIP parameters to match the default VLAN. Does not change the enable/disable state of the RIP settings. If no VLAN is specified, all VLANs are reset.

CONFIGURING OSPF

Each switch that is configured to run OSPF must have a unique router ID. It is recommended that you manually set the router ID of the switches participating in OSPF, instead of having the switch automatically choose its router ID based on the highest interface IP address. Not performing this configuration in larger, dynamic environments could result in an older link state database remaining in use.

[Table 12-5](#) describes the commands used to configure OSPF.

Table 12-5: OSPF Configuration Commands

Command	Description
config ospf [area <areaid> vlan [<name> all]] cost [automatic <number>]	Configures the cost metric of one or all VLAN(s). If an area is specified, the cost metric is applied to all VLANs currently within that area. When <code>automatic</code> is specified, the advertised cost is determined from the OSPF metric table and corresponds to the active highest bandwidth port in the VLAN.
config ospf [area <areaid> vlan [<name> all]] priority <number>	Configures the priority used in the designated router-election algorithm for one or all IP interface(s) (VLANs) for all VLANs currently within the area. The range is 0 through 255, and the default setting is 1. Setting the value to 0 ensures that the router is never selected as the designated router or backup designated router.
config ospf [vlan <name> area <areaid> virtual-link <routerid> <areaid>] authentication [simple-password <password> md5 <md5_key_id> <md5_key> none]	Specifies the authentication password (up to eight characters) or Message Digest 5 (MD5) key for one or all interfaces (VLANs) in an area. The <code>md5_key</code> is a numeric value with the range 0 to 65,536. When the OSPF area is specified, authentication information is applied to all OSPF interfaces within the area.

Table 12-5: OSPF Configuration Commands (continued)

Command	Description
config ospf [vlan <name> area <areaid> virtual-link <routerrid> <areaid>] timer <retransmission_interval> <transmission_delay> <hello_interval> <dead_interval>	<p>Configures the timers for one interface or all interfaces in the same OSPF area. The following default, minimum, and maximum values (in seconds) are used:</p> <ul style="list-style-type: none"> ■ <code>retransmission_interval</code> <ul style="list-style-type: none"> Default: 5 Minimum: 0 Maximum: 3,600 ■ <code>transmission_delay</code> <ul style="list-style-type: none"> Default: 1 Minimum: 0 Maximum: 3,600 ■ <code>hello_interval</code> <ul style="list-style-type: none"> Default: 10 Minimum: 1 Maximum: 65,535 ■ <code>dead_interval</code> <ul style="list-style-type: none"> Default: 40 Minimum: 1 Maximum: 2,147,483,647
config ospf add virtual-link <routerrid> <areaid>	<p>Adds a virtual link to another ABR. Specify the following:</p> <ul style="list-style-type: none"> ■ <code>routerrid</code> — Far-end router interface number. ■ <code>areaid</code> — Transit area used for connecting the two end-points.
config ospf add vlan [<name> all] area <areaid>	<p>Enables OSPF on one or all VLANs (router interfaces). The <code><areaid></code> specifies the area to which the VLAN is assigned.</p>
config ospf area <areaid> add range <ipaddress> <mask> [advertise noadvertise] [type 3 type 7]	<p>Configures a range of IP addresses in an OSPF area. If advertised, the range is exported as a single LSA by the ABR.</p>
config ospf area <areaid> delete range <ipaddress> <mask>	<p>Deletes a range of IP addresses in an OSPF area.</p>
config ospf area <areaid> normal	<p>Configures an OSPF area as a normal area. The default setting is <code>normal</code>.</p>

Table 12-5: OSPF Configuration Commands (continued)

Command	Description
config ospf area <areaid> nssa [summary nosummary] stub-default-cost <cost> {translate}	Configures an OSPF area as a NSSA.
config ospf area <areaid> stub [summary nosummary] stub-default-cost <cost>	Configures an OSPF area as a stub area.
config ospf asbr-filter [<access_profile> none]	Configures a route filter for non-OSPF routes exported into OSPF. If none is specified, no RIP and static routes are filtered.
config ospf ase-summary add <ipaddress> <mask> cost <cost> {<tag_number>}	Configures an aggregated OSPF external route using the IP addresses specified.
config ospf ase-summary delete <ipaddress> <mask>	Deletes an aggregated OSPF external route.
config ospf delete virtual-link <routerid> <areaid>	Removes a virtual link.
config ospf delete vlan [<name> all]	Disables OSPF on one or all VLANs (router interfaces).
config ospf direct-filter [<access_profile> none]	Configures a route filter for direct routes. If none is specified, all direct routes are exported if ospf export direct is enabled.
config ospf lsa-batching-timer <timer_value>	Configures the OSPF LSA batching timer value. The range is between 0 (disabled) and 600 seconds, using multiples of 5 seconds. The LSAs added to the LSDB during the interval are batched together for refresh or timeout. The default setting is 30 seconds.
config ospf metric-table <10M_cost> <100M_cost> <1G_cost>	Configures the automatic interface costs for 10 Mbps, 100 Mbps, and 1 Gbps interfaces. The default cost for 10 Mbps is 10, for 100 Mbps is 5, and for 1 Gbps is 1.
config ospf originate-default {always} cost <metric> [ase-type-1 ase-type-2] {tag <number>}	Configures a default external LSA to be generated by OSPF, if no other default route is originated by OSPF by way of RIP and static route re-distribution. If always is specified, OSPF always advertises the default route. If always is not specified, OSPF adds the default LSA if there is a reachable default route in the route table.
config ospf routerid [automatic <routerid>]	Configures the OSPF router ID. If automatic is specified, the switch uses the largest IP interface address as the OSPF router ID. The default setting is automatic.

Table 12-5: OSPF Configuration Commands (continued)

Command	Description
config ospf spf-hold-time {<seconds>}	Configures the minimum number of seconds between Shortest Path First (SPF) recalculations. The default setting is 3 seconds.
config ospf vlan <name> area <areaid>	Changes the area ID of an OSPF interface (VLAN).
create ospf area <areaid>	Creates an OSPF area. Area 0.0.0.0 does not need to be created. It exists by default.
disable ospf export [bgp i-bgp e-bgp]	Disables OSPF exporting of BGP-related routes.
enable ospf	Enables OSPF process for the router.
enable ospf export [bgp i-bgp e-bgp] cost <metric> [ase-type-1 ase-type-2] {tag <number>}	Enables OSPF to export BGP-related routes using LSAs to other OSPF routers. The default tag number is 0. The default setting is disabled.
enable ospf export direct cost <metric> [ase-type-1 ase-type-2] {tag <number>}	Enables the distribution of local interface (direct) routes into the OSPF domain. Once enabled, the OSPF router is considered to be an ASBR. The default tag number is 0. The default setting is disabled. Interface routes which correspond to the interface that has OSPF enabled are ignored.
enable ospf export rip cost <metric> [ase-type-1 ase-type-2] {tag <number>}	Enables the distribution of RIP routes into the OSPF domain. Once enabled, the OSPF router is considered to be an ASBR. The default tag number is 0. The default setting is disabled.
enable ospf export static cost <metric> [ase-type-1 ase-type-2] {tag <number>}	Enables the distribution of static routes into the OSPF domain. Once enabled, the OSPF router is considered to be an ASBR. The default tag number is 0. The default setting is disabled.
enable ospf export vip cost <metric> [ase-type-1 ase-type-2] {tag <number>}	Enables the distribution of virtual IP addresses into the OSPF domain. The default tag number is 0. The default setting is disabled.

OSPF CONFIGURATION EXAMPLE

Figure 12-5 shows an example of an autonomous system using OSPF routers. The details of this network follow.

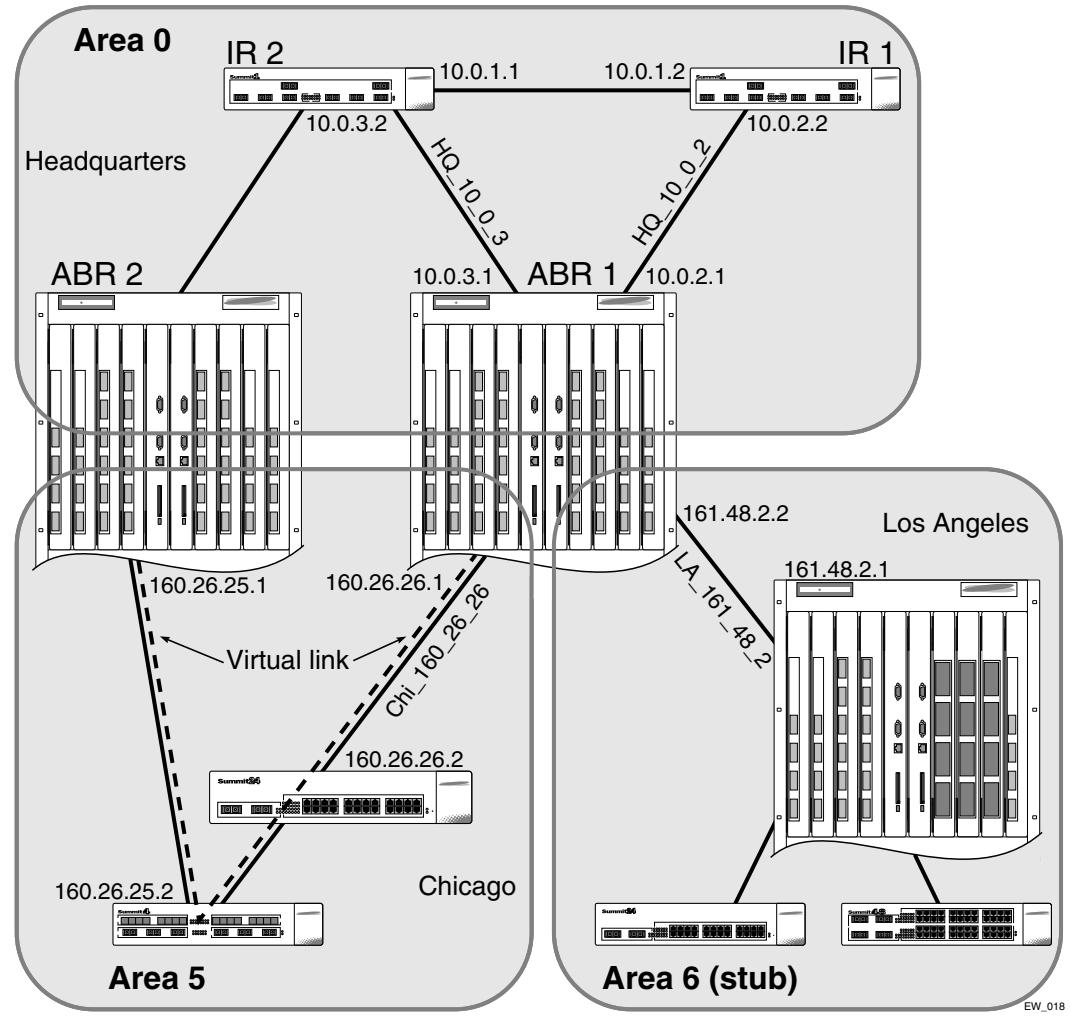


Figure 12-5: OSPF configuration example

Area 0 is the backbone area. It is located at the headquarters and has the following characteristics:

- 2 internal routers (IR1 and IR2)
- 2 area border routers (ABR1 and ABR2)
- Network number 10.0.x.x
- 2 identified VLANs (HQ_10_0_2 and HQ_10_0_3)

Area 5 is connected to the backbone area by way of ABR1 and ABR2. It is located in Chicago and has the following characteristics:

- Network number 160.26.x.x
- 1 identified VLAN (Chi_160_26_26)
- 2 internal routers
- A virtual link from ABR1 to ABR2 that traverses both internal routers.

In the event that the link between either ABR and the backbone fails, the virtual link provides a connection for all routers that become discontiguous from the backbone.

Area 6 is a stub area connected to the backbone by way of ABR1. It is located in Los Angeles and has the following characteristics:

- Network number 161.48.x.x
- 1 identified VLAN (LA_161_48_2)
- 3 internal routers
- Uses default routes for inter-area routing

Two router configurations for the example in [Figure 12-5](#) are provided in the following section.

CONFIGURATION FOR ABR1

The following is the configuration for the router labeled ABR1:

```
create vlan HQ_10_0_2
create vlan HQ_10_0_3
create vlan LA_161_48_2
create vlan Chi_160_26_26

config vlan HQ_10_0_2 ipaddress 10.0.2.1 255.255.255.0
config vlan HQ_10_0_3 ipaddress 10.0.3.1 255.255.255.0
config vlan LA_161_48_26 ipaddress 161.48.2.26 255.255.255.0
config vlan Chi_160_26_26 ipaddress 160.26.2.1 255.255.255.0

create ospf area 0.0.0.5
create ospf area 0.0.0.6

enable ipforwarding

config ospf area 0.0.0.6 stub nosummary stub-default-cost 10
config ospf vlan LA_161_48_2 area 0.0.0.6
config ospf vlan Chi_160_26_26 area 0.0.0.5
config ospf add virtual-link 160.26.25.1 0.0.0.5
config ospf add vlan all

enable ospf
```

CONFIGURATION FOR IR1

The following is the configuration for the router labeled IR1:

```
config vlan HQ_10_0_1 ipaddress 10.0.1.2 255.255.255.0
config vlan HQ_10_0_2 ipaddress 10.0.2.2 255.255.255.0
config ospf add vlan all
enable ipforwarding
enable ospf
```

DISPLAYING OSPF SETTINGS

To display settings for OSPF, use the commands listed in [Table 12-6](#).

Table 12-6: OSPF Show Commands

Command	Description
show ospf	Displays global OSPF information.
show ospf area {detail}	Displays information about all OSPF areas.
show ospf area <areaid>	Displays information about a particular OSPF area.
show ospf ase-summary	Displays the OSPF external route aggregation configuration.
show ospf interfaces {detail}	Displays information about all OSPF interfaces.
show ospf interfaces {vlan <name> area <areaid>}	Displays information about one or all OSPF interfaces.
show ospf lsdb {detail} area [<areaid> all] [router network summary-net summary-asb as-external external-type7 all]	Displays a table of the current LSDB. You can filter the display using the area ID and LSA type. The default setting is <code>all</code> with no detail. If <code>detail</code> is specified, each entry includes complete LSA information.
show ospf virtual-link {<areaid> <routerid> }	Displays virtual link information about a particular router or all routers.

RESETTING AND DISABLING OSPF SETTINGS

To return OSPF settings to their defaults, use the commands listed in [Table 12-7](#).

Table 12-7: OSPF Reset and Disable Commands

Command	Description
delete ospf area [<areaid> all]	Deletes an OSPF area. Once an OSPF area is removed, the associated OSPF area and OSPF interface information is removed. The backbone area cannot be deleted. A non-empty area cannot be deleted.
disable ospf	Disables OSPF process in the router.

Table 12-7: OSPF Reset and Disable Commands

Command	Description
disable ospf export direct	Disables exporting of local interface (direct) routes into the OSPF domain.
disable ospf export rip	Disables exporting of RIP routes in the OSPF domain.
disable ospf export static	Disables exporting of statically configured routes into the OSPF domain.
disable ospf export vip	Disables exporting of virtual IP addresses into the OSPF domain.
unconfig ospf {vlan <name> area <areaid>}	Resets one or all OSPF interfaces to the default settings.

13

Exterior Gateway Routing Protocols

This chapter covers the following topics:

- [Overview on page 13-2](#)
- [BGP Attributes on page 13-2](#)
- [BGP Communities on page 13-3](#)
- [BGP Features on page 13-3](#)
- [Configuring BGP on page 13-10](#)
- [Displaying BGP Settings on page 13-15](#)
- [Resetting and Disabling BGP on page 13-15](#)

This chapter describes how to configure the Border Gateway Protocol (BGP), an exterior routing protocol available on the switch.

For more information on BGP, refer to the following documents:

- RFC 1771 – *Border Gateway Protocol version 4 (BGP-4)*
- RFC 1965 – *Autonomous System Confederations for BGP*
- RFC 1966 – *BGP Route Reflection*
- RFC 1997 – *BGP Communities Attribute*
- RFC 1745 – *BGP/OSPF Interaction*



ExtremeWare supports BGP version 4 only.

OVERVIEW

BGP is an exterior routing protocol that was developed for use in TCP/IP networks. The primary function of BGP is to allow different autonomous systems (ASs) to exchange network reachability information.

An autonomous system is a set of routers that are under a single technical administration. This set of routers uses a different routing protocol (such as OSPF) for intra-AS routing. One or more routers in the AS are configured to be border routers, exchanging information with other border routers (in different autonomous systems) on behalf of all of the intra-AS routers.

BGP can be used as an exterior gateway protocol (E-BGP), or it can be used within an AS, as an interior gateway protocol (I-BGP).

BGP ATTRIBUTES

The following well-known BGP attributes are supported by the switch:

- Origin – Defines the origin of the route. Possible values are IGP, EGP, and incomplete.
- AS_Path – The list of ASs that are traversed for this route.
- Next_hop – The IP address of the next hop BGP router to reach the destination listed in the NLRI field.
- Multi_Exit_Discriminator – Used to select a particular border router in another AS when multiple border routers exist.
- Local_Preference – Used to advertise this router's degree of preference to other routers within the AS.
- Atomic_aggregate – Indicates that the sending border router is used a route aggregate prefix in the route update.
- Aggregator – Identifies the BGP router AS number and IP address that performed route aggregation.

- Community – Identifies a group of destinations that share one or more common attributes.
- Cluster_ID – Specifies a 4 byte field used by a route reflector to recognize updates from other route reflectors in the same cluster.

BGP COMMUNITIES

A BGP community is a group of BGP destinations that require common handling. ExtremeWare supports the following well-known BGP community attributes:

- no-export
- no-advertise
- internet

BGP FEATURES

This section describes the following BGP features supported by ExtremeWare:

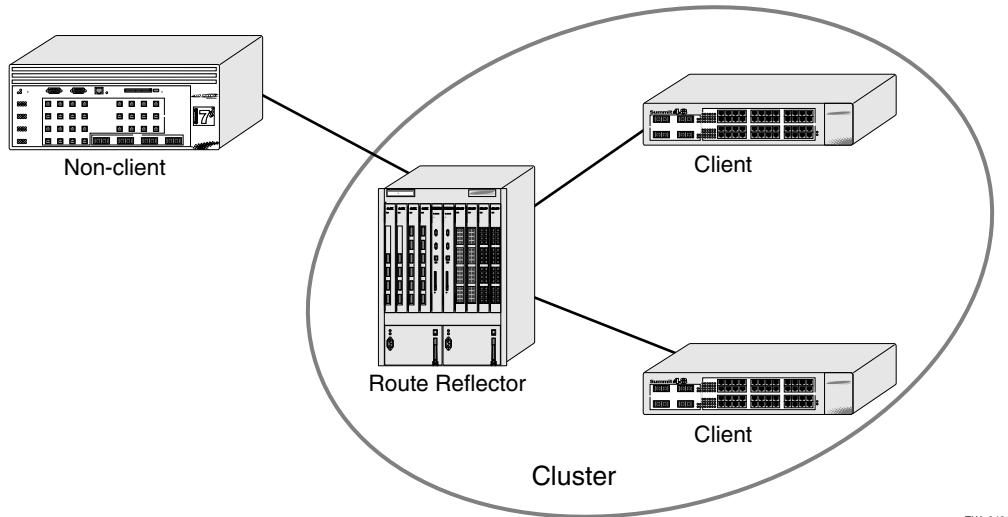
- Route Reflectors
- Route Confederations
- Route Aggregation
- IGP Synchronization
- Using the Loopback Interface
- OSPF to BGP Redistribution

ROUTE REFLECTORS

Another way to overcome the difficulties of creating a fully-meshed AS is to use *route reflectors*. Route reflectors allow a single router to serve as a central routing point for the AS or sub-AS.

A *cluster* is formed by the route reflector and its client routers. Peer routers that are not part of the cluster must be fully meshed according to the rules of BGP.

A BGP cluster, including the route reflector and its clients, is shown in [Figure 13-1](#).



EW_042

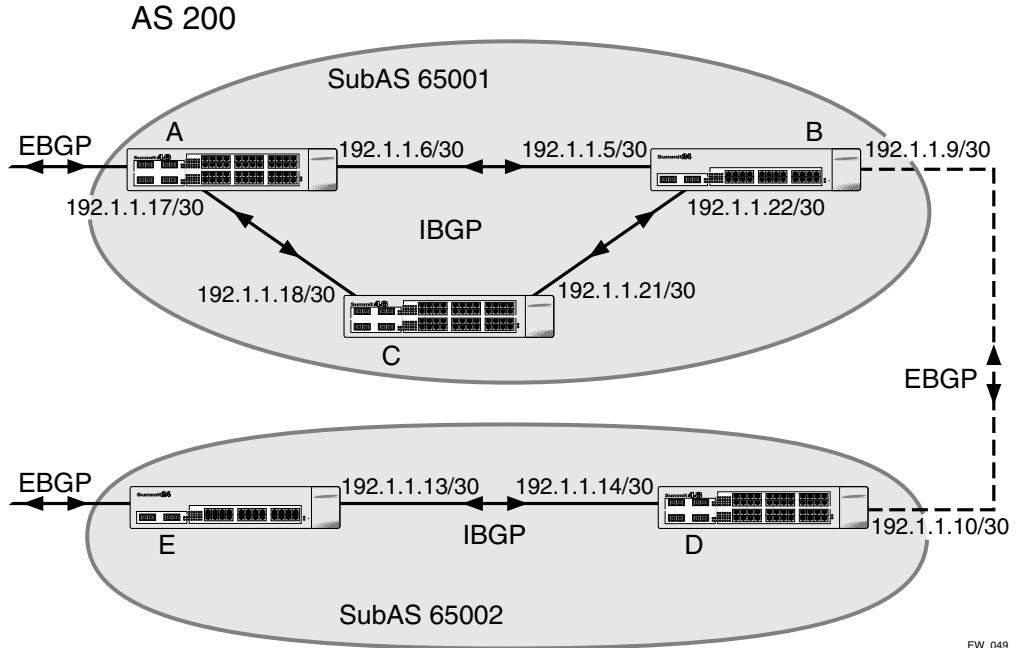
Figure 13-1: Route reflectors

ROUTE CONFEDERATIONS

BGP requires networks to use a fully-meshed router configuration. This requirement does not scale well, especially when BGP is used as an interior gateway protocol. One way to reduce the size of a fully-meshed AS is to divide the AS into multiple sub-autonomous systems and group them into a *routing confederation*. Within the confederation, each sub-AS must be fully-meshed. The confederation is advertised to other networks as a single AS.

ROUTE CONFEDERATION EXAMPLE

[Figure 13-2](#) shows an example of a confederation.



EW_049

Figure 13-2: Routing confederation

In this example, AS 200 has five BGP speakers. Without a confederation, BGP would require that the routes in AS 200 be fully meshed. Using the confederation, AS 200 is split into two sub-ASs: AS65001 and AS65002. Each sub-AS is fully meshed, and IBGP is running among its members. EBGP is used between sub-AS 65001 and sub-AS 65002. Router B and Router D are EBGP peers. EBGP is also used between the confederation and outside ASs.

To configure Router A, use the following commands:

```

create vlan ab
config vlan ab add port 1
config vlan ab ipaddress 192.1.1.6/30
enable ipforwarding vlan ab
config ospf add vlan ab area 0.0.0.0

create vlan ac
config vlan ac add port 2
config vlan ac ipaddress 192.1.1.17/30

```

```
enable ipforwarding vlan ac
config ospf add vlan ac area 0.0.0.0

disable bgp
config bgp as-number 65001
config bgp routerid 192.1.1.17
config bgp confederation-id 200
enable bgp

create bgp neighbor 192.1.1.5 as-number remote-AS-number 65001
create bgp neighbor 192.1.1.18 as-number remote-AS-number 65001
enable bgp neighbor all
```

To configure Router B, use the following commands:

```
create vlan ba
config vlan ba add port 1
config vlan ba ipaddress 192.1.1.5/30
enable ipforwarding vlan ba
config ospf add vlan ba area 0.0.0.0

create vlan bc
config vlan bc add port 2
config vlan bc ipaddress 192.1.1.22/30
enable ipforwarding vlan bc
config ospf add vlan bc area 0.0.0.0

create vlan bd
config vlan bd add port 3
config vlan bd ipaddress 192.1.1.9/30
enable ipforwarding vlan bd
config ospf add vlan bd area 0.0.0.0

disable bgp
config bgp as-number 65001
config bgp routerid 192.1.1.22
config bgp confederation-id 200
enable bgp

create bgp neighbor 192.1.1.6 as-number remote-AS-number 65001
create bgp neighbor 192.1.1.21 as-number remote-AS-number 65001
create bgp neighbor 192.1.1.10 as-number remote-AS-number 65002
enable bgp neighbor all
config bgp add confederation-peer sub-AS-number 65002
```

To configure Router C, use the following commands:

```
create vlan ca
config vlan ca add port 1
config vlan ca ipaddress 192.1.1.18/30
enable ipforwarding vlan ca
config ospf add vlan ca area 0.0.0.0

create vlan cb
config vlan cb add port 2
config vlan cb ipaddress 192.1.1.21/30
enable ipforwarding vlan cb
config ospf add vlan cb area 0.0.0.0

disable bgp
config bgp as-number 65001
config bgp routerid 192.1.1.21
config bgp confederation-id 200
enable bgp

create bgp neighbor 192.1.1.22 as-number remote-AS-number 65001
create bgp neighbor 192.1.1.17 as-number remote-AS-number 65001
enable bgp neighbor all
```

To configure Router D, use the following commands:

```
create vlan db
config vlan db add port 1
config vlan db ipaddress 192.1.1.10/30
enable ipforwarding vlan db
config ospf add vlan db area 0.0.0.0

create vlan de
config vlan de add port 2
config vlan de ipaddress 192.1.1.14/30
enable ipforwarding vlan de
config ospf add vlan de area 0.0.0.0

disable bgp
config bgp as-number 65002
config bgp routerid 192.1.1.14
config bgp confederation-id 200
enable bgp
```

```
create bgp neighbor 192.1.1.9 as-number remote-AS-number 65001
create bgp neighbor 192.1.1.13 as-number remote-AS-number 65002
enable bgp neighbor all
config bgp add confederation-peer sub-AS-number 65001
```

To configure Router E, use the following commands:

```
create vlan ed
config vlan ed add port 1
config vlan ed ipaddress 192.1.1.13/30
enable ipforwarding vlan ed
config ospf add vlan ed area 0.0.0.0

disable bgp
config bgp as-number 65002
config bgp routerid 192.1.1.13
config bgp confederation-id 200
enable bgp

create bgp neighbor 192.1.1.14 as-number remote-AS-number 65002
enable bgp neighbor 192.1.1.14
```

ROUTE AGGREGATION

Route aggregation is the process of combining the characteristics of several routes so that they are advertised as a single route. Aggregation reduces the amount of information that a BGP speaker must store and exchange with other BGP speakers. Reducing the information that is stored and exchanged also reduces the size of the routing table.

USING ROUTE AGGREGATION

To use BGP route aggregation, you must do the following:

- 1 Enable aggregation using the following command:

```
enable bgp aggregation
```

- 2 Create an aggregate route, using the following commands:

```
config bgp add aggregate-address <ipaddress>/<masklength> {as-set}
{summary-only} {advertise-route-map <route-map>} {attribute-route-map
<route-map>}
```

IGP SYNCHRONIZATION

You can configure an AS to be a transit AS, so that it can pass traffic through from one AS to a third AS. When you configure a transit AS, it is important that the routes advertised by BGP are consistent with the routes that are available within the AS using its interior gateway protocol. To ensure consistency, BGP should be synchronized with the IGP used within the AS. This will ensure that the routes advertised by BGP are, in fact, reachable within the AS. IGP synchronization is enabled by default.

USING THE LOOPBACK INTERFACE

If you are using BGP as your interior gateway protocol, you may decide to advertise the interface as available, regardless of the status of any particular interface. The loopback interface can also be used for EBGP multihop. Using the loopback interface eliminates multiple, unnecessary route changes.

OSPF TO BGP ROUTE RE-DISTRIBUTION

Both BGP and OSPF can be enabled simultaneously on the switch. Route re-distribution allows the switch to exchange routes, including static routes, between the two routing protocols.

Exporting routes from OSPF to BGP, and from BGP to OSPF, are discreet configuration functions. To run OSPF and BGP simultaneously, you must first configure both protocols and then verify the independent operation of each. Then you can configure the routes to export from OSPF to BGP and the routes to export from BGP to OSPF.

CONFIGURING BGP

[Table 13-1](#) describes the commands used to configure BGP.

Table 13-1: BGP Configuration Commands

Command	Description
<pre>config bgp add aggregate-address <ipaddress>/<masklength> {as-set} {summary-only} {advertise-route-map <route-map>} {attribute-route-map <route-map>}</pre>	<p>Configures an aggregate route. Options include the following:</p> <ul style="list-style-type: none"> ■ as-set – Aggregates only the path attributes of the aggregate routes. ■ summary-only – Sends both aggregated and non-aggregated routes to the neighbors. ■ advertise-route-map – Specifies the route map used to select routes for this aggregated route. ■ attribute-route-map – Specifies the route map used to set the attributes of the aggregated route.
<pre>config bgp add confederation-peer <sub_as_list></pre>	Specifies the list of sub-ASs that belong to a confederation. A maximum of 16 AS numbers can be specified.

Table 13-1: BGP Configuration Commands (continued)

Command	Description
config bgp add dampening <halflife> <reuse> <suppress> <max_suppression> bgp-policy	<p>Configures BGP route flap dampening used to suppress the advertisement of routes when the routes are changing rapidly. Specify the following:</p> <ul style="list-style-type: none"> ■ halflife — The time after which the penalty is decreased to half the original amount. The range is 1 to 45 minutes. The default setting is 15 minutes.
	<ul style="list-style-type: none"> ■ reuse — If the penalty of a flapping route is decreased below this number, the suppression on this route is removed. The range is 1 to 32,000. The default setting is 750.
	<ul style="list-style-type: none"> ■ suppress — If the penalty of a flapping route is above this number, this route is suppressed. The setting for suppress must be larger than the setting for reuse. The range is 1 to 32,000. The default setting is 2000.
	<ul style="list-style-type: none"> ■ max_suppression — A multiplier used to determine the maximum amount of time that a route can be suppressed. The maximum suppression time is <halflife> * <max_suppression>.
	<ul style="list-style-type: none"> ■ bgp_policy — Specifies the BGP route.
config bgp add network <ipaddress>/<mask_length> {<route_map>}	<p>Adds a network to be originated from this router. The network must be reachable by the router.</p>
config bgp as-number <as_number>	<p>Changes the local AS number used by BGP. BGP must be disabled before the AS number can be changed.</p>
config bgp cluster-id <cluster_id>	<p>Appends a BGP route reflector cluster ID to the cluster list of a route. Used when multiple router reflectors are used within the same cluster of clients.</p>
	<p>BGP must be disabled before configuring the cluster ID.</p>
config bgp confederation-id <confederation_id>	

Table 13-1: BGP Configuration Commands (continued)

Command	Description
config bgp dampening <halflife> <reuse> <suppress> <max_suppression>	<p>Configures BGP route flap dampening used to suppress the advertisement of routes when the routes are changing rapidly. Specify the following:</p> <ul style="list-style-type: none"> ■ halflife — The time after which the penalty is decreased to half the original amount. The range is 1 to 45 minutes. The default setting is 15 minutes.
	<ul style="list-style-type: none"> ■ reuse — If the penalty of a flapping route is decreased below this number, the suppression on this route is removed. The range is 1 to 32,000. The default setting is 750.
	<ul style="list-style-type: none"> ■ suppress — If the penalty of a flapping route is above this number, this route is suppressed. The setting for suppress must be larger than the setting for reuse. The range is 1 to 32,000. The default setting is 2000.
	<ul style="list-style-type: none"> ■ max_suppression — A multiplier used to determine the maximum amount of time that a route can be suppressed. The maximum suppression time is <halflife> * <max_suppression>.
	<ul style="list-style-type: none"> ■ bgp_policy — Specifies the BGP route.
config bgp delete aggregate-address [<ipaddress/masklength> all]	<p>Deletes one or all aggregate routes.</p>
config bgp delete dampening bgp-policy <bgp_policy>	<p>Deletes BGP route flap dampening from a route.</p>
config bgp export [ospf ospf-intro ospf-inter ospf-extern1 ospf-extern2] {<route_map>}	<p>Configures BGP to export OSPF-related routes to other BGP peers. BGP attributes associated with the OSPF routes can be applied using an optional route map.</p>
config bgp local-preference <local_preference>	<p>Changes the default local preference attribute. The range is 0 to 4294967295. The default value is 100.</p>
config bgp multi-exist-discriminator [<number> none]	

Table 13-1: BGP Configuration Commands (continued)

Command	Description
config bgp neighbor [<ipaddress> all] [route-reflector-client no-route-reflector-client]	Configures a BGP neighbor to be a route reflector client. Implicitly defines the router to be a route reflector. The neighbor must be in the same AS as the router.
config bgp neighbor [<ipaddress> all] [send-communities don't-send-communities]	
config bgp neighbor [<ipaddress> all] as-path-filter [in out] [none <access_profile>]	Configures an AS path filter for a neighbor. The filter is defined using the access-profile mechanism and can be installed on the input side and/or the output side. Use the none keyword to remove the filter.
config bgp neighbor [<ipaddress> all] nlri-filter [in out] [none <access_profile>]	Configures an NLRI filter for a neighbor. The filter is defined using the access-profile mechanism, and can be installed on the input side and/or the output side. Use the none keyword to remove the filter.
config bgp neighbor [<ipaddress> all] route-map-filter [in out] [none <route_map>]	Configures a route map for a neighbor. The route map can be installed on the input or output side, and while exchanging updates with the neighbor, it is used to modify or filter the NLRI information and the path attributes associated with it. The route map is removed using the none keyword.
config bgp neighbor [<ipaddress> all] soft-reset {input output}	Applies the current input and/or output routing policy to the routing information already exchanged with the neighbor. The input/output routing policy is determined by the nlri-filter, as-path-filter, and the route map configured for the neighbor in the input-output side. This command does not affect the switch configuration.
config bgp neighbor [<ipaddress> all] source-interface [any vlan <name>]	Changes the BGP source interface for TCP connections. The default setting is any.
config bgp neighbor [<ipaddress> all] timer <keepalive> <holdtime>	Configures the BGP neighbor timers. The range for <keepalive> is 0 to 65535 seconds. The default keep alive setting is 60. The range for <holdtime> is 0 to 21845 seconds. The default hold time is 90.

Table 13-1: BGP Configuration Commands (continued)

Command	Description
config bgp neighbor [<ipaddress> all] weight <weight>	Assigns a locally used weight to a neighbor connection for the route selection algorithm. All routes learned from this peer are assigned the same weight. The route with the highest weight is more preferable when multiple routes are available to the same network. The range is 0 to 4294967295. The default setting is 0.
config bgp router-id <router_id>	Changes the router ID. BGP must be disabled before changing the router ID.
config bgp soft-reconfiguration	Applies the route map associated with the network command, aggregation and redistribution, immediately. This command does not affect the switch configuration.
create bgp neighbor <ipaddress> remote-as <as_number> {multihop}	Creates a new BGP peer. Use the <code>multihop</code> keyword for EBGP peers that are not directly connected.
disable bgp aggregation	Disables BGP route aggregation filtering.
disable bgp always-compare-med	
disable bgp export [ospf ospf-intro ospf-inter ospf-extern1 ospf-extern2]	Disables BGP from export OSPF-related routes to other BGP peers.
enable bgp	Enables BGP.
enable bgp aggregation	Enables BGP route aggregation.
enable bgp aggregation	Enables BGP route aggregation filtering.
enable bgp always-compare-med	Enables BGP to use the Multi Exit Discriminator (MED) from neighbors in different autonomous systems in the route selection algorithm. MED is only used when comparing paths from the same AS. The default setting is enabled.
enable bgp compare-as-path	Enables using the As path as a factor in the route selection algorithm. The default setting is enabled.
enable bgp compare-med-within-as-only	Enables BGP to use the MED as a factor in the route selection algorithm from confederation peers. The default setting is disabled.
enable bgp neighbor [<ipaddress> all]	Enables the BGP session. The neighbor must be created before the BGP session can be enabled.

Table 13-1: BGP Configuration Commands (continued)

Command	Description
enable bgp synchronization	Enables synchronization between BGP and IGP. When enabled, BGP waits for IGP to provide the next-hop reachability before advertising the route to an external neighbor. The default setting is enabled.

DISPLAYING BGP SETTINGS

To display settings for BGP, use the commands listed in [Table 13-2](#).

Table 13-2: BGP Show Commands

Command	Description
show bgp	Displays BGP configuration information.
show bgp bgp-policy {<name>}	Displays the BGP policy parameters.
show bgp neighbor {detail}	Disables BGP neighbor information
show bgp neighbor <ipaddress>	Displays information about a specified neighbor.

RESETTING AND DISABLING BGP

To return BGP settings to their defaults, or to disable BGP, use the commands listed in [Table 13-3](#).

Table 13-3: BGP Reset and Disable Commands

Command	Description
delete bgp neighbor [<ipaddress> all]	Deletes one or all BGP neighbors.
disable bgp	Disables BGP.
disable bgp aggregation	Disables BGP route aggregation.
disable bgp always-compare-med	Disables MED from being used in the route selection algorithm.
disable bgp neighbor [<ipaddress> all]	Disables the BGP session. Once disabled, all the Adj-RIB-In for the neighbor will be flushed out.
disable bgp synchronization	Disables the synchronization between BGP and IGP. Default is enabled.

This chapter covers the following topics:

- [Overview on page 14-2](#)
- [Configuring IP Multicasting Routing on page 14-4](#)
- [Configuration Examples on page 14-9](#)
- [Displaying IP Multicast Routing Settings on page 14-13](#)
- [Deleting and Resetting IP Multicast Settings on page 14-14](#)

For more information on IP multicasting, refer to the following publications:

- RFC 1112 – *Host Extension for IP Multicasting*
- RFC 2236 – *Internet Group Management Protocol, Version 2*
- DVMRP Version 3 – *draft_ietf_dvmrp_v3_07*
- PIM-DM Version 2 – *draft_ietf_pim_v2_dm_03*
- RFC 2326 – *Protocol Independent Multicast-Sparse Mode*

The following URLs point to the Web sites for the IETF Working Groups:

- IETF DVMRP Working Group – <http://www.ietf.org/html.charters/idmr-charter.html>
- IETF PIM Working Group – <http://www.ietf.org/html.charters/pim-charter.html>

OVERVIEW

IP multicast routing is a function that allows a single IP host to send a packet to a group of IP hosts. This group of hosts can include devices that reside on the local network, within a private network, or outside of the local network.

IP multicast routing consists of the following functions:

- A router that can forward IP multicast packets.
- A router-to-router multicast routing protocol (for example, Distance Vector Multicast Routing Protocol (DVMRP) or Protocol Independent Multicast (PIM)).
- A method for the IP host to communicate its multicast group membership to a router (for example, Internet Group Management Protocol (IGMP)).

 *You should configured IP unicast routing before you configure IP multicast routing.*

DVMRP OVERVIEW

DVMRP is a distance vector protocol that is used to exchange routing and multicast information between routers. Like RIP, DVMRP periodically sends the entire routing table to its neighbors.

DVMRP has a mechanism that allows it to prune and graft multicast trees to reduce the bandwidth consumed by IP multicast traffic.

PIM OVERVIEW

The switch supports both dense mode and sparse mode operation. You can configure dense mode or sparse mode on a per-interface basis. Once enabled, some interfaces can run dense mode, while others run sparse mode.

 *You can run either DVMRP or PIM on the switch, but not both simultaneously.*

PIM DENSE MODE

Protocol Independent Multicast- Dense Mode (PIM-DM) is a multicast routing protocol that is similar to DVMRP. PIM-DM routers perform reverse path multicasting (RPM). However, instead of exchanging its own unicast route tables for the RPM algorithm, PIM-DM uses the existing unicast route table for the reverse path. As a result, PIM-DM requires less system memory.

PIM-DM is a broadcast and prune protocol. Using PIM-DM, multicast routes are pruned and grafted in the same way as DVMRP.

PIM SPARSE MODE (PIM-SM)

Unlike PIM-DM, PIM-SM is an explicit join and prune protocol, and supports shared trees as well as shortest path trees (SPTs). The routers must explicitly join the group(s) in which they are interested in becoming a member. This is beneficial for large networks that have group members who are sparsely distributed.

Using PIM-SM, the router sends a join message to the rendezvous point (RP). The RP is a central multicast router that is responsible for receiving and distributing multicast packets.

When a router that has a multicast packet to distribute, it encapsulates the packet in a unicast message and sends it to the RP. The RP decapsulates the multicast packet and distributes it among all member routers.

When a router determines that the multicast rate from a particular originating router (not the RP) has exceeded a configured threshold, that router can send an explicit join to the originating router. Once this occurs, the receiving router gets the multicast directly from the sending router, and bypasses the RP.

 You can run either PIM-DM or PIM-SM per VLAN.

IGMP OVERVIEW

IGMP is a protocol used by an IP host to register its IP multicast group membership with a router. Periodically, the router queries the multicast group to see if the group is still in use. If the group is still active, a single IP host responds to the query, and group registration is maintained.

IGMP is enabled by default on the switch. However, the switch can be configured to disable the generation of period IGMP query packets. IGMP query should be enabled when the switch is configured to perform IP unicast or IP multicast routing.

IGMP SNOOPING

IGMP snooping is a layer-2 function of the switch. It does not require multicast routing to be enabled. The feature reduces the flooding of IP multicast traffic. IGMP snooping optimizes the usage of network bandwidth, and prevents multicast traffic from being flooded to parts of the network that do not need it. The switch does not reduce any IP multicast traffic in the local multicast domain (224.0.0.x).

IGMP snooping is enabled by default on the switch. If you are using multicast routing, IGMP snooping must be enabled. If IGMP snooping is disabled, all IGMP and IP multicast traffic floods within a given VLAN. IGMP snooping expects at least one device in the network to periodically generate IGMP query messages. Without an IGMP querier, the switch stops forwarding IP multicast packets to any port. An optional optimization for IGMP snooping is the strict recognition of multicast routers only if the remote devices has joined the DVMRP (224.0.0.4) or PIM (244.0.0.13) multicast groups.

IGMP configuration commands can be found in [Table 14-2](#).

CONFIGURING IP MULTICASTING ROUTING

To configure IP multicast routing, you must do the following:

- 1 Configure the system for IP unicast routing.
- 2 Enable multicast routing on the interface, using the following command:

```
enable ipmcforwarding {vlan <name>}
```

- 3 Enable DVMRP or PIM on all IP multicast routing interfaces, using one of the following commands:

```
config dvmrp add vlan [<name> | all]
```

```
config pim add vlan [<name> | all] {dense | sparse}
```

- 4 Enable DVMRP or PIM on the router, using one of the following commands:

```
enable dvmrp
enable pim
```

[Table 14-1](#) describes the commands used to configure IP multicast routing.

Table 14-1: IP Multicast Routing Configuration Commands

Command	Description
config dvmrp add vlan [<name> all]	Enables DVMRP one or all IP interfaces. If no VLAN is specified, DVMRP is enabled on all IP interfaces. When an IP interface is created, DVMRP is disabled by default.
config dvmrp delete vlan [<name> all]	Disables DVMRP on one or all IP interfaces. If no VLAN is specified, DVMRP is disabled on all IP interfaces.
config dvmrp timer <route_report_interval> <route_replacement_time>	Configures the global DVMRP timers. Specify the following: <ul style="list-style-type: none"> ■ route_report_interval — The amount of time the system waits between transmitting periodic route report packets. The range is 1 to 2,147,483,647 seconds (68 years). The default setting is 60 seconds. Because triggered update is always enabled, the route report will always be transmitted prior to the expiration of the route report interval. ■ route_replacement_time — The hold-down time before a new route is learned, once the previous route has been deleted. The range is 1 to 2,147,483,647 seconds (68 years). The default setting is 140 seconds.
config dvmrp vlan [<name> all] cost <number>	Configures the cost (metric) of the interface. The default setting is 1.
config dvmrp vlan [<name> all] export-filter [<access_profile> none>]	Configures DVMRP to filter out routes specified in the export filter when sending out route advertisements.
config dvmrp vlan [<name> all] import-filter [<access_profile> none>]	Configures DVMRP to filter out certain routes (defined by the access profile) received from a neighbor.
config dvmrp vlan [<name> all] trusted-gateway [<access_profile> none>]	Configures the DVMRP trusted gateway, based on the access profile.

Table 14-1: IP Multicast Routing Configuration Commands (continued)

Command	Description
config dvmrp vlan <name> timer <probe_interval> <neighbor_timeout_interval>	<p>Configures DVMRP interface timers. Specify the following:</p> <ul style="list-style-type: none"> ■ <code>probe_interval</code> — The amount of time that the system waits between transmitting DVMRP probe messages. The range is 1 to 2,147,483,647 seconds (68 years). The default setting is 10 seconds. ■ <code>neighbor_timeout_interval</code> — The amount of time before a DVMRP neighbor route is declared to be down. The range is 1 to 2,147,483,647 seconds (68 years). The default setting is 35 seconds.
config pim add vlan [<vlan> all] {dense sparse}	<p>Enables PIM on an IP interface. When an IP interface is created, per-interface PIM configuration is disabled by default. The default PIM mode is dense.</p>
config pim cbsr [vlan <name> priority <priority> none]	<p>Configures a candidate bootstrap router for PIM sparse-mode operation. The range is 0 - 255. The default setting is 0 and indicates the lowest priority. To delete a CSBR, use the keyword <code>none</code> as the priority.</p>
config pim crp timer <crp_adv_interval>	<p>Configures the candidate rendezvous point advertising interval. The default is 60 seconds.</p>
config pim crp vlan <name> access_policy <access_policy> {<priority>}	<p>Configures the candidate rendezvous point for PIM sparse-mode operation. The access policy contains the list of multicast group accesses serviced by this RP. The range is 0 - 255. The default setting is 0 and indicates the highest priority. To delete a CRP, use the keyword <code>none</code> as the access policy.</p>
config pim delete vlan [<name> all]	<p>Disables PIM on an interface.</p>
config pim spt-threshold <leaf_threshold> <rp_threshold>	<p>Configures the threshold, in kbps, for switching to SPT. On leaf routers, this setting is based on data packets. On the RP, this setting is based on register packets. The default setting is 0.</p>

Table 14-1: IP Multicast Routing Configuration Commands (continued)

Command	Description
config pim timer <hello_interval> <jp_interval> vlan [<vlan> all]	Configures the global PIM timers. Specify the following: <ul style="list-style-type: none"> ■ <code>hello_interval</code> — The amount of time before a hello message is sent out by the PIM router. The range is 1 to 65,519 seconds. The default setting is 30 seconds. ■ <code>jp_interval</code> — The join/prune interval. The range is 1 to 65,519 seconds. The default setting is 60 seconds.
config pim vlan [<name> all] trusted-gateway [<access_profile> none]	Configures PIM to use the access-profile to determine the PIM trusted gateway.
enable dvmrp	Enables DVMRP on the system. The default setting is disabled.
enable dvmrp rxmode vlan [<name> all]	Enables receiving of DVMRP packets on a per-VLAN basis.
enable dvmrp txmode vlan [<name> all]	Enables transmitting of DVMRP packets on a per-VLAN basis.
enable ipmcforwarding {<vlan <name>}	Enables IP multicast forwarding on an IP interface. If no options are specified, all configured IP interfaces are affected. When new IP interfaces are added, <code>ipmcforwarding</code> is disabled by default.
enable pim	Enables PIM on the system. The default setting is disabled.

[Table 14-2](#) describes the commands used to configure the Internet Gateway Message Protocol (IGMP).

Table 14-2: IGMP Configuration Commands

Command	Description
config igmp <query_interval> <query_response_interval> <last_member_query_interval>	<p>Configures the IGMP timers. Timers are based on RFC2236. Specify the following:</p> <ul style="list-style-type: none"> ■ <code>query_interval</code> — The amount of time, in seconds, the system waits between sending out General Queries. The range is 1 to 2,147,483,647 seconds (68 years). The default setting is 125 seconds. ■ <code>query_response_interval</code> — The maximum response time inserted into the periodic General Queries. The range is 1 to 25 seconds. The default setting is 10 seconds. ■ <code>last_member_query_interval</code> — The maximum response time inserted into a Group-Specific Query sent in response to a Leave group message. The range is 1 to 25 seconds. The default setting is 1 second.
config igmp snooping <router_timeout> <host_timeout>	<p>Configures the IGMP snooping timers. Timers should be set to approximately 2.5 times the router query interval in use on the network. Specify the following:</p> <ul style="list-style-type: none"> ■ <code>router_timeout</code> — The interval, in seconds, between the last time the router was discovered and the current time. The range is 10 to 2,147,483,647 seconds (68 years). The default setting is 260 seconds. ■ <code>host_timeout</code> — The interval, in seconds, between the last IGMP group report message from the host and the current time. The range is 10 to 2,147,483,647 seconds (68 years). The default setting is 260 seconds.
enable igmp {vlan <name>}	<p>Enables IGMP on a router interface. If no VLAN is specified, IGMP is enabled on all router interfaces. The default setting is enabled.</p>
enable igmp snooping {forward-mcrouter-only}	<p>Enables IGMP snooping on the switch. If <code>forward-mcrouter-only</code> is specified, the switch forwards all multicast traffic to the multicast router, only. Otherwise, the switch forwards all multicast traffic to any IP router.</p>

CONFIGURATION EXAMPLES

Figure 14-1 and Figure 13-2 are used in Chapter 12 to describe the OSPF configuration on a switch. Refer to Chapter 12 for more information about configuring OSPF. In the first example, the system labeled IR1 is configured for IP multicast routing, using PIM-DM. In the second example, the system labeled ABR1 is configured for IP multicast routing using PIM-SM.

PIM-DM CONFIGURATION EXAMPLE

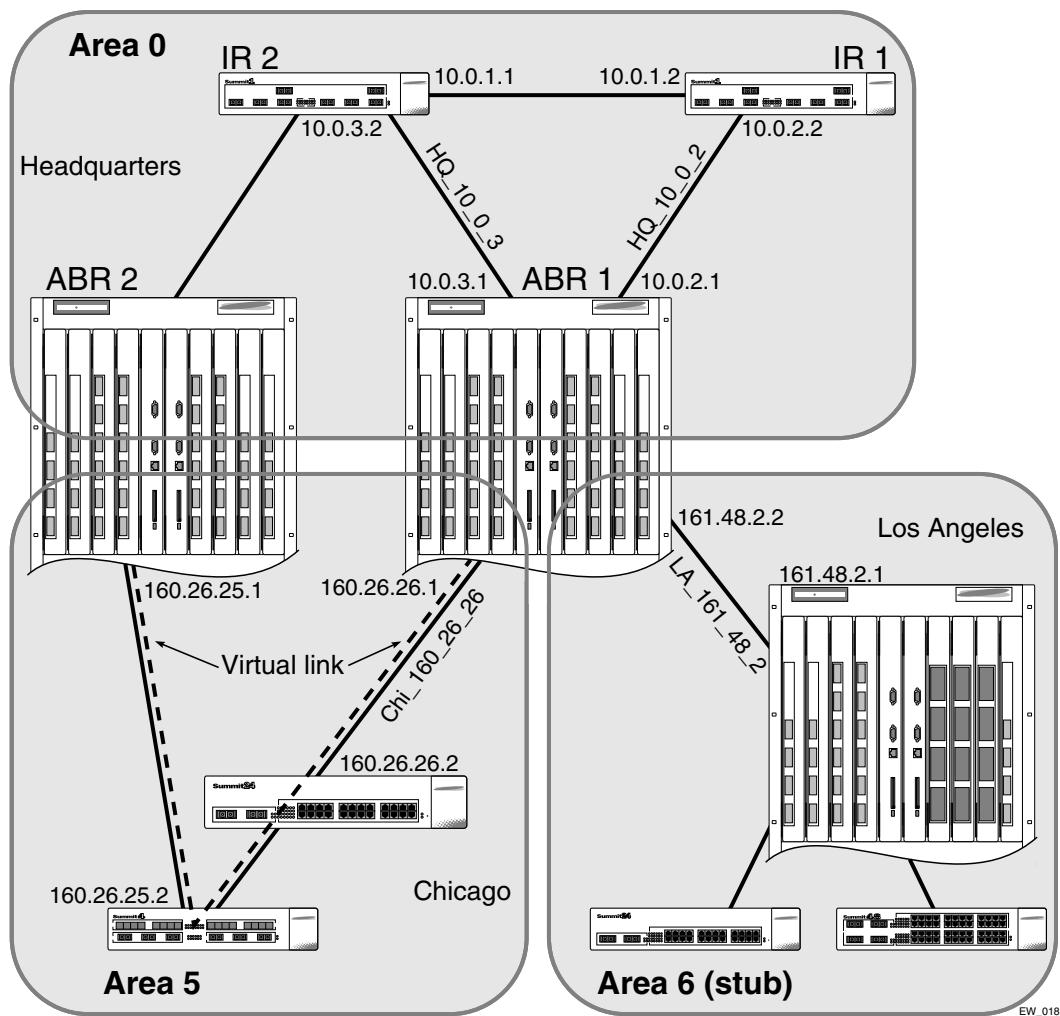


Figure 14-1: IP multicast routing using PIM-DM configuration example

CONFIGURATION FOR IR1

The following is the configuration for the router labeled IR1:

```
config vlan HQ_10_0_1 ipaddress 10.0.1.2 255.255.255.0
config vlan HQ_10_0_2 ipaddress 10.0.2.2 255.255.255.0
config ospf add vlan all
enable ipforwarding
enable ospf
enable ipmcforwarding
config pim add vlan all dense
config pim spt-threshold 16 8
enable pim
```

The following example configures PIM-SM.

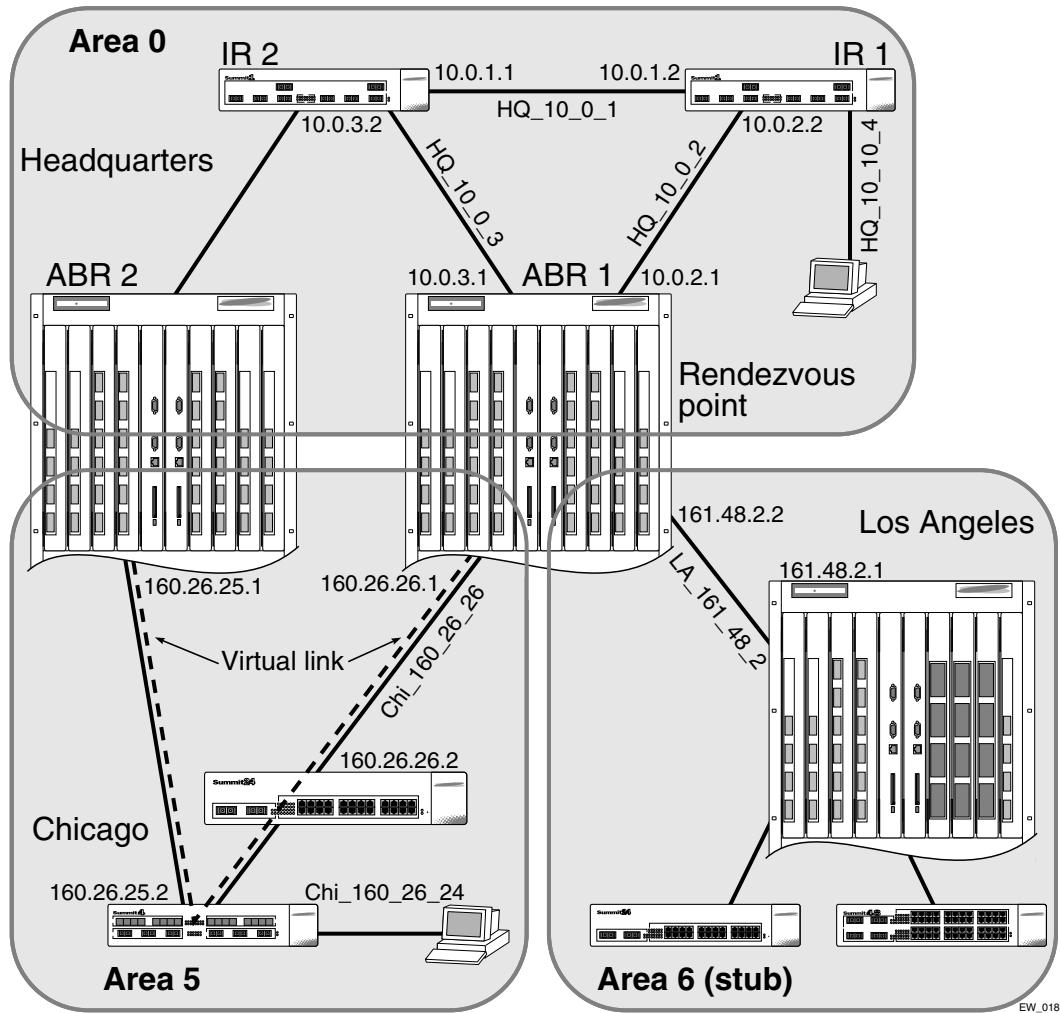


Figure 14-2: IP multicast routing using PIM-SM configuration example

CONFIGURATION FOR ABR1

The following is the configuration for the router labeled ABR1:

```
config vlan HQ_10_0_2 ipaddress 10.0.2.1 255.255.255.0
config vlan HQ_10_0_3 ipaddress 10.0.3.1 255.255.255.0
config vlan LA_161_48_2 ipaddress 161.48.2.2 255.255.255.0
config vlan CHI_160_26_26 ipaddress 160.26.26.1 255.255.255.0
config ospf add vlan all
enable ipforwarding
enable ipmcforwarding
config pim add vlan all sparse
create access-profile rp-list ipaddress
config rp-list add ipaddress 224.0.0.0 240.0.0.0
enable loopback HQ_10_0_3
config pim crp HQ_10_0_3 rp-list 30
config pim csbr HQ_10_0_3 30

config pim spt-threshold 16 8
```

DISPLAYING IP MULTICAST ROUTING SETTINGS

To display settings for IP multicast routing components, use the commands listed in [Table 14-3](#).

Table 14-3: IP Multicast Routing Show Commands

Command	Description
show dvmrp {vlan <name> route {detail}}	Displays the DVMRP configuration and statistics, or the unicast route table. The default setting is all.
show igmp snooping {<vlan <name>>}	Displays IGMP snooping registration information, and a summary of all IGMP timers and states.
show ipmc cache {detail} {<group>} {<src_ipaddress> <mask>}}	Displays the IP multicast forwarding cache.
show pim {vlan <name>}	Displays the PIM configuration and statistics. If no VLAN is specified, the configuration is displayed for all PIM interfaces.
show pim rp-set {group}	Displays the RP-set for one or all groups.

DELETING AND RESETTING IP MULTICAST SETTINGS

To return IP multicast routing settings to their defaults and disable IP multicast routing functions, use the commands listed in [Table 14-4](#).

Table 14-4: IP Multicast Routing Reset and Disable Commands

Command	Description
clear igmp snooping {vlan <name>}	Removes one or all IGMP snooping entries.
clear ipmc cache {<group> {<src_ipaddress> <mask>}}	Resets the IP multicast cache table. If no options are specified, all IP multicast cache entries are flushed.
config ipmc cache timeout <seconds>	Configures the aging time for multicast cache entries. The default setting is 300 seconds.
disable dvmrp	Disables DVMRP on the system.
disable dvmrp rxmode vlan [<name> all]	Disables receiving of DVMRP packets on a per-VLAN basis.
disable dvmrp txmode vlan [<name> all]	Disables transmitting of DVMRP packets on a per-VLAN basis.
disable igmp {vlan <name>}	Disables the router-side IGMP processing on a router interface. No IGMP query is generated, but the switch continues to respond to IGMP queries received from other devices. If no VLAN is specified, IGMP is disabled on all router interfaces.
disable igmp snooping	Disables IGMP snooping. IGMP snooping can be disabled only if IP multicast routing is not being used. Disabling IGMP snooping allows all IGMP and IP multicast traffic to flood within a given VLAN.
disable ipmcforwarding {vlan <name>}	Disables IP multicast forwarding.
disable pim	Disables PIM on the system.
unconfig dvmrp {vlan <name>}	Resets the DVMRP timers to their default settings. If no VLAN is specified, all interfaces are reset.
unconfig igmp	Resets all IGMP settings to their default values and clears the IGMP group table.
unconfig pim {vlan <name>}	Resets all PIM settings to their default values.



IPX Routing

This chapter describes the following topics:

- [Overview of IPX on page 15-1](#)
- [IPX/RIP Routing on page 15-4](#)
- [Configuring IPX on page 15-6](#)
- [IPX Commands on page 15-7](#)
- [IPX Configuration Example on page 15-11](#)
- [Displaying IPX Settings on page 15-13](#)
- [Resetting and Disabling IPX on page 15-14](#)

This chapter assumes that you are already familiar with IPX. If not, refer to your Novell™ documentation.

OVERVIEW OF IPX

The switch provides support for the IPX, IPX/RIP, and IPX/SAP protocols. The switch dynamically builds and maintains an IPX routing table and an IPX service table.

ROUTER INTERFACES

The routing software and hardware routes IPX traffic between IPX router interfaces. A router interface is simply a VLAN that has an IPX network identifier (NetID) and IPX encapsulation type assigned to it.

As you create VLANs with different IPX NetIDs the switch automatically routes between them. Both the VLAN switching and IPX routing function occur within the switch.

i A VLAN can be configured with either an IPX NetID or an IP address. A VLAN cannot be configured for both IPX and IP routing simultaneously.

Figure 15-1 shows the same BlackDiamond switch discussed in earlier chapters. In Figure 15-1, IPX routing has been added to the BlackDiamond switch, and two additional VLANs have been defined; *Exec*, and *Support*. Both VLANs have been configured as protocol-specific VLANs, using IPX.

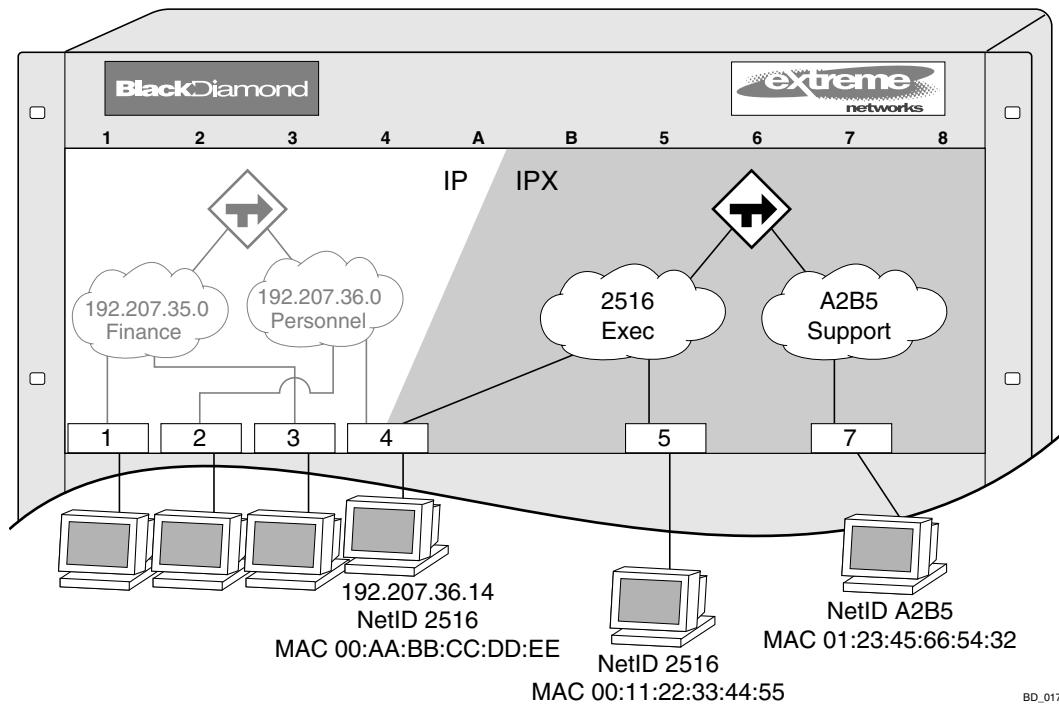


Figure 15-1: IPX VLAN configuration

Exec has been assigned the IPX NetID 2516. *Support* has been assigned the IPX NetID A2B5. All ports on slot 5 are assigned to *Exec*; all ports on slot 7 are assigned to *Support*. In addition, all ports on slot 4 have been assigned to *Exec*. Thus, the ports on slot 4 belong to both the *Personnel* VLAN (running IP) and the *Exec* VLAN (running IPX).

Traffic within each VLAN is switched using the Ethernet MAC address. Traffic between *Exec* and *Support* is routed using the IPX NetID. Traffic cannot be sent between the IP VLANs (*Finance* and *Personnel*) and the IPX VLANs (*Exec* and *Support*).

IPX ROUTING PERFORMANCE

To use IPX routing, you must have a switch that has the “*i*” chipset. Switches that have the “*i*” chipset are capable of performing IPX routing at wire-speed.

Switches that do not have the “*i*” chipset no longer support IPX routing capabilities. Previous versions of ExtremeWare supported CPU-based IPX routing on switches that did not have the “*i*” chipset. CPU-based IPX routing has been removed on switches that do not use the “*i*” chipset to support other features in ExtremeWare.

IPX ENCAPSULATION TYPES

Novell NetWare™ supports four types of frame encapsulation. The ExtremeWare term for each type is described in Table 15-1.

Table 15-1: IPX Encapsulation Types

Name	Description
ENET_II	The frame uses the standard Ethernet 2 header.
ENET_8023	The frame includes the IEEE 802.3 length field, but does not include the IEEE 802.2 Logical Link Control (LLC) header. This encapsulation is used by NetWare version 2.x and the original 3.x version.
ENET_8022	The frame uses the standard IEEE format and includes the IEEE 802.2 LLC header. This encapsulation is used by NetWare version 3.12 and 4.x.
ENET_SNAP	The frame adds a Subnetwork Access Protocol (SNAP) header to the IEEE 802.2 LLC header.

To configure a VLAN to use a particular encapsulation type, use the following command:

```
config vlan <name> xnetid <netid> [enet_ii | enet_8023 | enet_8022 | enet_snap]
```

POPULATING THE ROUTING TABLE

The switch builds and maintains an IPX routing table. As in the case of IP, the table is populated using dynamic and static entries.

DYNAMIC ROUTES

Dynamic routes are typically learned by way of IPX/RIP. Routers that use IPX/RIP exchange information in their routing tables in the form of advertisements. Using dynamic routes, the routing table contains only networks that are reachable.

Dynamic routes are aged out of the table when an update for the network is not received for a period of time, as determined by the routing protocol.

STATIC ROUTES

Static routes are manually entered into the routing table. Static routes are used to reach networks not advertised by routers. You can configure up to 64 static IPX routes on the switch. Static routes are never aged out of the routing table. Static routes are advertised to the network using IPX/RIP.

IPX/RIP ROUTING

The switch supports the use of IPX/RIP for unicast routing. IPX/RIP is different from IP/RIP. However, many of the concepts are the same. ExtremeWare supports the following IPX/RIP features:

- Split horizon
- Poison reverse
- Triggered Updates

Route information is entered into the IPX route table in one of the following two ways:

- Dynamically, by way of RIP
- Statically, using the command:

```
config ipxroute add [<dest_netid> | default] next_hop_netid  
next_hop_node_addr <hops> <ticks>
```

IPX/RIP is automatically enabled when a NetID is assigned to the VLAN. To remove the advertisement of an IPX VLAN, use the command:

```
config ipxrip delete {vlan <name> | all}
```

GNS SUPPORT

ExtremeWare support the Get Nearest Server (GNS) reply function. When a NetID is assigned to the switch, the GNS reply service is automatically enabled. When a station requests a particular service on the network (for example, locating a print server), the station sends a GNS request and the switch responds to the request. If GNS-reply is disabled, the switch drops the request.

To disable GNS-reply, use the following command:

```
disable ipxsap gns-reply {vlan <name>}
```

ROUTING SAP ADVERTISEMENTS

The switch contains an IPX Service Table, and propagates SAP advertisements to other IPX routers on the network. Each SAP advertisement contains the following:

- Service type
- Server name
- Server NetID
- Server node address

The service information is entered into the IPX Service Table in one of the following two ways:

- Dynamically, by way of SAP
- Statically, using the following command:

```
config ipxservice add <service_type> <service_name> <netid>
<mac_address> <socket> <hops>
```

CONFIGURING IPX

This section describes the commands associated with configuring IPX, IPX/RIP, and IPX/SAP on the switch. Configuring IPX routing involves the following steps:

- 1 Create at least two VLANs.
- 2 If you are combining an IPX VLAN with another VLAN on the same port(s), you must use a protocol filter on one of the VLANs, or use 802.1Q tagging.
- 3 Assign each VLAN a NetID and encapsulation type, using the following command:

```
config vlan <name> xnetid <netid> [enet_ii | enet_8023 | enet_8022 | enet_snap]
```

Ensure that each VLAN has a unique IPX NetID and that the encapsulation type matches the VLAN protocol.

Once you configure the IPX VLAN information, IPX forwarding automatically begins to function. Specifically, configuring the IPX VLAN automatically enables the IPX/RIP, IPX/SAP, and SAP GNS services.

VERIFYING IPX ROUTER CONFIGURATION

You can use the following commands to verify the IPX routing configuration:

- `show vlan` — In addition to other information, this command displays the IPX NetID setting and encapsulation type.
- `show ipxconfig` — This command is analogous to the `show ipconfig` command for the IP protocol. It displays summary global IPX configuration information followed by per-VLAN information. Information includes enable/disable status for IPX/RIP, IPX/SAP, IPX route sharing, IPX service sharing, and so on.
- `show ipxrouting` — This command is analogous to the `show iproute` command for the IP protocol. It displays static and learned routes, along with information about the VLAN that uses the route, hop count, age of the route, and so on.
- `show ipxsap` — This command displays the enable status of IPX/SAP for the VLAN, and its operational and administrative status (including the GNS reply service). It also lists any identified IPX/SAP neighbors, SAP packet statistics, and several other timer settings.
- `show ipxrip` — This command displays the enable status of IPX/RIP for the VLAN, including operational and administrative status. It also lists any identified IPX/RIP neighbors, RIP packet statistics, and several other timer settings.
- `show ipxservice` — This command displays the contents of the IPX Service Table.

PROTOCOL-BASED VLANs FOR IPX

When combining IPX VLANs with other VLANs on the same physical port, it may be necessary to assign a protocol filter to the VLAN. This is especially true if it is not possible to use 802.1Q VLAN tagging. For convenience, IPX-specific protocol filters have been defined and named in the default configuration of the switch. Each filter is associated with a protocol encapsulation type. The IPX-specific protocol filters and the associated encapsulation type of each are described in [Table 15-2](#).

Table 15-2: IPX Protocol Filters and Encapsulation Types

Protocol Name	Protocol Filter	Used for Filtering IPX Encapsulation Type
IPX	eypte 0x8137	enet_ii
IPX_8022	llc 0xe0e0	enet_802_2
IPX_snap	SNAP 0x8137	enet_snap

It is not possible to define a protocol-sensitive VLAN for filtering the IPX `enet_8023` encapsulation type. Instead, use a protocol-sensitive filter on the other VLANs that share the same ports, leaving the `enet_8023` encapsulation VLAN configured using the `any` protocol.

IPX COMMANDS

[Table 15-3](#) describes the commands used to configure basic IPX settings.

Table 15-3: Basic IPX Commands

Command	Description
<code>config ipxmaxhops <number></code>	Configures the IPX maximum hop count when forwarding IPX packets. The default setting is 16. Change this only if NetWare Link Services Protocol (NLSP) is running in the IPX network.

Table 15-3: Basic IPX Commands (continued)

Command	Description
config ipxroute add [<dest_netid> default]<next_hop_id> <next_hop_node_addr> <hops><tics>	<p>Adds a static IPX route entry in the IPX route table. Specify:</p> <ul style="list-style-type: none"> ■ <code>next_hop_id</code> — The NetID of the neighbor IPX network. ■ <code>next_hop_node_addr</code> — The node address of the next IPX router. ■ <code>hops</code> — The maximum hop count. ■ <code>tics</code> — The timer delay value. <p>Up to 64 static routes can be entered.</p>
config ipxroute delete [<dest_netid> default]<next_hop_netid> <next_hop_node_addr>	Removes a static IPX route entry from the route table.
config ipxservice add <service_type><service_name> <netid> <mac_address> <socket><hops>	<p>Adds a static entry to the IPX service table. Specify:</p> <ul style="list-style-type: none"> ■ <code>service_type</code> — The service type. ■ <code>service_name</code> — The service name. ■ <code>netid</code> — The IPX network identifier of the server. ■ <code>mac_address</code> — The MAC address of the server. ■ <code>socket</code> — The IPX port number on the server. ■ <code>hops</code> — The number of hops (for SAP routing purposes).
config ipxservice delete <service_type><service_name> <netid> <mac_address> <socket>	Deletes an IPX service from the service table.

Table 15-3: Basic IPX Commands (continued)

Command	Description
config vlan <name> xnetid <netid> [enet_ii enet_8023 enet_8022 enet_snap]	Configures a VLAN to run IPX routing. Specify: <ul style="list-style-type: none">■ enet_ii — Uses standard Ethernet 2 header.■ enet_8023 — Uses IEEE 802.3 length field, but does not include the IEEE 802.2 LLC header.■ enet_8022 — Uses standard IEEE format and uses IEEE 802.2 LLC header.■ enet_snap — Adds Subnetwork Access Protocol (SNAP) header to IEEE 802.2 LLC header.
enable type20 forwarding {vlan <name>}	Enables the forwarding of IPX type 20 (NetBIOS inside IPX) packets from one or more ingress VLANs. The default setting is disabled.
xping {continuous} {size <n>} <netid> <node_address>	Pings an IPX node. If continuous is not specified, 4 pings are sent. The default ping packet size is 256 data bytes. The size can be configured between 1 and 1,484 bytes.

Table 15-4 describes the commands used to configure the IPX route table.

Table 15-4: IPX/RIP Configuration Commands

Command	Description
config ipxrip add vlan [<name> all]	Configures one or all IPX VLANs to run IPX/RIP. IPX/RIP is enabled by default when you configure the IPX VLAN.
config ipxrip delete vlan [<name> all]	Disables IPX/RIP on one or all interfaces.
config ipxrip vlan [<name> all] delay <msec>	Configures the time between each IPX/RIP packet within an update interval. The default setting is 55 milliseconds.
config ipxrip vlan [<name> all] max-packet-size <size>	Configures the maximum transmission unit (MTU) size of the IPX/RIP packet. the default setting is 432 bytes.

Table 15-4: IPX/RIP Configuration Commands (continued)

Command	Description
config ipxrip vlan [<name> all] update-interval <time> {hold-multiplier <number>}	Configures the update interval and hold multiplier for IPX/RIP updates. This setting affects both the periodic update interval of IPX/RIP and the aging interval of learned routes. The default update interval is 60 seconds. The aging period is calculated using the formula (update-interval * multiplier). The default multiplier is 3.
enable ipxrip	Enables IPX/RIP on the router.

[Table 15-5](#) describes the commands used to configure IPX/SAP.

Table 15-5: IPX/SAP Configuration Commands

Command	Description
config ipxsap add vlan [<name> all]	Configures an IPX VLAN to run IPX/SAP routing. If no VLAN is specified, all VLANs are configured to run IPX/SAP routing. IPX/SAP routing is enabled by default when the IPX VLAN is configured.
config ipxsap delete vlan [<name> all]	Disables IPX/SAP on an interface.
config ipxsap vlan [<name> all] delay <msec>	Configures the time between each SAP packet within an update interval. The default setting is 55 milliseconds.
config ipxsap vlan [<name> all] max-packet-size <number>	Configures the MTU size of the IPX/SAP packets. The default setting is 432 bytes.
config ipxsap vlan [<name> all] update-interval <time> {hold-multiplier <number>}	Configures the update interval and hold multiplier for IPX/SAP updates. This setting affects both the periodic update interval of SAP and the aging interval of learned routes. The default update interval is 60 seconds. The aging period is calculated using the formula (update-interval * multiplier). The default multiplier is 3. Triggered update is always enabled; therefore, new information is processed and propagated immediately.
config ipxsap vlan <name> gns-delay <msec>	Configures the amount of time the switch waits before answering a GNS request. By default, the switch answers a GNS request as soon as possible (0 milliseconds).

Table 15-5: IPX/SAP Configuration Commands (continued)

Command	Description
enable ipxsap	Enables IPX/SAP on the router.
enable ipxsap gns-reply {vlan <name>}	Enables GNS reply on one or all IPX interfaces. If no VLAN is specified, GNS reply is enabled on all IPX interfaces. The default setting is enabled.

IPX CONFIGURATION EXAMPLE

Figure 15-2 builds on the example showing the IP/RIP configuration that was used in earlier chapters. Now, in addition to having IP VLANs configured, this example illustrates a switch that has the following IPX VLANs defined:

- *Exec*
 - Protocol-sensitive VLAN using the IPX protocol with the filter IPX_8022
 - All ports on slot 4 and slot 5 have been assigned to *Exec*
 - *Exec* is configured for IPX NetID 2516 and IPX encapsulation type 802.2
- *Support*
 - All ports on slot 7 have been assigned to *Support*
 - *Support* is configured for IPX NetID A2B5 and IPX encapsulation type 802.2

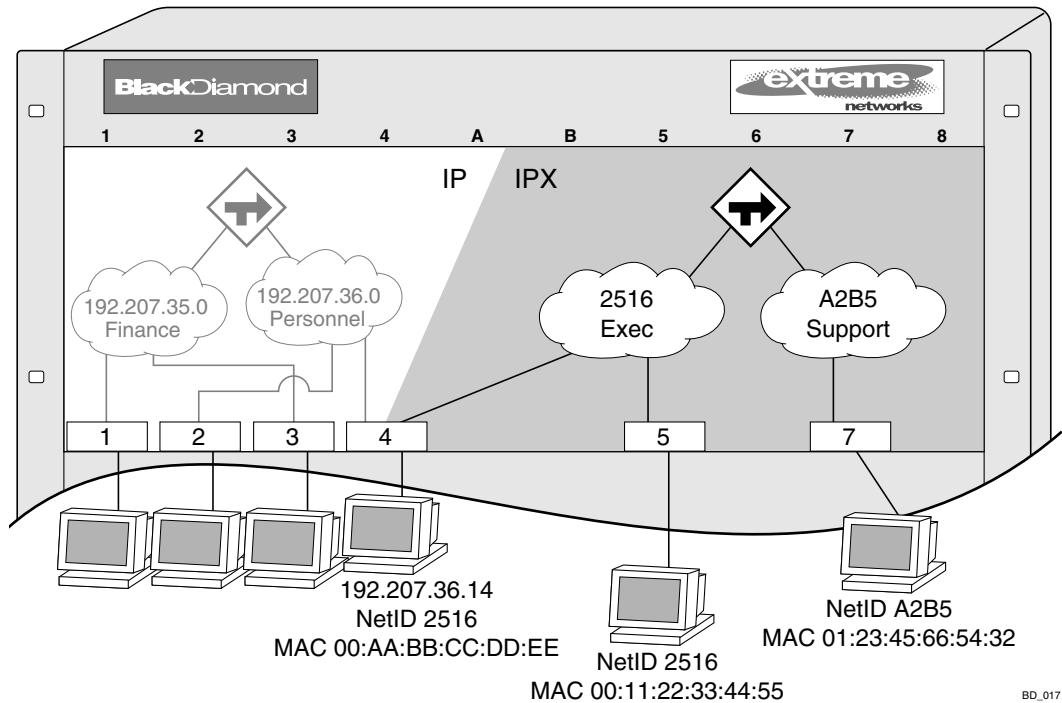


Figure 15-2: IPX routing configuration example

The stations connected to the system generate a combination of IP traffic and IPX traffic. The IP traffic is filtered by the IP VLANs. IPX traffic is filtered by the IPX VLANs.

In this configuration, all IP traffic from stations connected to slots 1 and 3 have access to the IP router by way of the VLAN *Finance*. IP traffic on ports on slots 2 and 4 reach the IP router by way of the VLAN *Personnel*.

Similarly, IPX traffic from stations connected to slots 4 and 5 have access to the IPX router by way of the VLAN *Exec*. IPX traffic on port 7 reach the IPX router by way of the VLAN *Support*. Both *Exec* and *Support* use enet_8022 as the encapsulation type.

The IPX configuration shown in example in [Figure 15-2](#) is as follows:

```
create vlan Exec
create vlan Support

config Exec protocol ipx_8022

config Exec add port 4:*,5:*
config Support add port 7:*

config Exec xnetid 2516 enet_8022
config Support xnetid A2B5 enet_8022
```

DISPLAYING IPX SETTINGS

To display settings for various IPX components, use the commands listed in [Table 15-6](#).

Table 15-6: IPX Show Commands

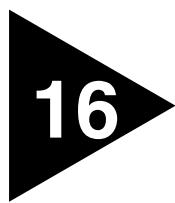
Command	Description
show ipxconfig {vlan <name>}	Displays IPX configuration information for one or all VLANs.
show ipxfdb {vlan <name> xnetid <netid>}	Displays the hardware IPX FDB information.
show ipxrip {vlan <name>}	Displays IPX/RIP configuration and statistics for one or all VLANs.
show ipxrout {vlan <name> xnetid <netid> origin [static rip local]}	Displays the IPX routes in the route table.
show ipxsap {vlan <name>}	Displays IPX/SAP configuration and status for one or all VLANs.
show ipxservice {vlan <name> xnetid <netid> origin [static sap local]}	Displays IPX services learned by way of SAP.
show ipxstats {vlan <name>}	Displays IPX packet statistics for the IPX router, and one or all VLANs.

RESETTING AND DISABLING IPX

To return IPX settings to their defaults and disable IPX functions, use the commands listed in [Table 15-7](#).

Table 15-7: IPX Reset and Disable Commands

Command	Description
disable ipxrip	Disables IPX/RIP on the router.
disable ipxsap	Disables IPX/SAP on the router.
disable ipxsap gns-reply {vlan <name>}	Disables GNS reply on one or all IPX interfaces.
disable type20 forwarding {vlan <name>}	Disables the forwarding of IPX type 20 packets.
unconfig ipxrip {vlan <name>}	Resets the IPX/RIP settings on one or all VLANs to the default. Removes import and export filters, and resets the MTU size, update interval, and inter-packet delay.
unconfig ipxsap {vlan <name>}	Resets the IPX/SAP settings on one or all VLANs to the default. Removes import and export filters, and resets the MTU size, update interval, and inter-packet delay.
unconfig vlan <name> xnetid	Removes the IPX NetID of a VLAN.



Access Policies

This chapter describes the following topics:

- [Overview of Access Policies on page 16-1](#)
- [Using IP Access Lists on page 16-2](#)
- [Using Routing Access Policies on page 16-15](#)
- [Making Changes to a Routing Access Policy on page 16-25](#)
- [Removing a Routing Access Policy on page 16-26](#)
- [Routing Access Policy Commands on page 16-26](#)
- [Using Route Maps on page 16-29](#)

OVERVIEW OF ACCESS POLICIES

Access policies are a generalized category of features that impact forwarding and route forwarding decisions. Access policies are used primarily for security and quality of service (QoS) purposes.

There are three categories of access policies:

- Access lists
- Routing access policies
- Route maps

IP ACCESS LISTS

IP access lists consist of IP access rules, and are used to perform packet filtering and forwarding decisions on incoming traffic. Each packet arriving on an ingress port is compared to the access list in sequential order, and is either forwarded to a specified QoS profile or dropped. Using access lists has no impact on switch performance.

Access lists are typically applied to traffic that crosses layer 3 router boundaries, but it is possible to use access lists within a layer 2 VLAN. Products that use the “*i*” chipset are capable of performing this function with no additional configuration. Products that do not use the “*i*” chipset require the enabling of Intra-subnet QoS (ISQ), to perform this function. For more information on ISQ, refer to [Chapter 9](#).

ROUTING ACCESS POLICIES

Routing access policies are used to control the advertisement or recognition of routing protocols, such as RIP, OSPF, or BGP. Routing access policies can be used to ‘hide’ entire networks, or to trust only specific sources for routes or ranges of routes. The capabilities of routing access policies are specific to the type of routing protocol involved, but are sometimes more efficient and easier to implement than access lists.

The following sections describe IP access lists first, and then describe the details of routing access policies.

ROUTE MAPS

Route maps are used to modify or filter routes redistributed into BGP. They are also used to modify or filter the routing information exchanged with BGP neighbors.

USING IP ACCESS LISTS

Each entry that makes up an IP access list contains a unique name. It can also contain an optional, unique precedence number. The rules of an IP access list consist of a combination of the following six components:

- IP source address and mask
- IP destination address and mask
- TCP or UDP source port range
- TCP or UDP destination port range

- Physical source port
- Precedence number (optional)

How IP Access Lists Work

When a packet arrives on an ingress port, the packet is compared with the access list rules to determine a match. When a match is found, the packet is processed. If the access list is of type deny, the packet is dropped. If the list is of type permit, the packet is forwarded. A permit access list can also apply a QoS profile to the packet.

Precedence Numbers

The precedence number is optional, and determines the order in which each rule is examined by the switch. Access list entries that contain a precedence number are evaluated from highest to lowest. Precedence numbers range from 1 to 25,600, with the *number 1 having the highest precedence*.

You can specify overlapping rules; however, if you are using precedence numbers, overlapping rules without precedence numbers are ignored. Therefore, the precedence numbers must be specified among all overlapping rules. If a new rule without a precedence number is entered, and this rule overlaps with already existing rules, the switch rejects the new rule and resolves the precedences among all remaining overlapping rules.

Specifying a Default Rule

To begin constructing an access list, you should specify a default rule. A *default rule* is a rule that contains wildcards for destination and source IP address, with no Layer 4 information. A default rule determines if the behavior of the access list is an “implicit deny” or “implicit accept”. If no access list entry is satisfied, the default rule is used to determine whether the packet is forwarded or dropped. If no default rule is specified, the default implicit behavior is to forward the packet.

The following example shows a default entry that is used to specify an explicit deny:

```
create access-list denyall ip dest 0.0.0.0/0 source 0.0.0.0/0 deny  
ports any
```

Once the default behavior of the access list is established, you may create additional entries using precedence numbers.

The access-list example, below, performs packet filtering in the following sequence, as determined by the precedence number:

- Deny UDP port 32 and TCP port 23 traffic to the 10.2.XX network.
- All other TCP port 23 traffic destined for other 10.X.X.X networks is permitted using QoS profile Qp4.
- All remaining traffic to 10.2.0.0 uses QoS profile Qp3.

With no default rule specified, all remaining traffic is allowed using the default QoS profile.

```
create access-list deny102_32 udp dest 10.2.0.0/16 ip-port 32 source any ip-port any deny ports any precedence 10
```

```
create access-list deny102_23 tcp dest 10.2.0.0/16 ip-port 23 source any ip-port any deny ports any precedence 20
```

```
create access-list allow10_23 tcp dest 10.0.0.0/8 ip-port 23 source any ip-port any permit qosprofile qp4 ports any precedence 30
```

```
create access-list allow102 ip dest 10.2.0.0/16 source 0.0.0.0/0 permit qosprofile qp3 ports any precedence 40
```

THE PERMIT-ESTABLISHED KEYWORD

The `permit-established` keyword is used to directionally control attempts to open a TCP session. Session initiation can be explicitly blocked using this keyword.



For an example of using the `permit-established` keyword, refer to [Using the Permit-Established Keyword](#), on page 16-11.

ADDING AND DELETING ACCESS LIST ENTRIES

Entries may be added and deleted to the access list. To add an entry, you must supply a unique name and, optionally, a unique precedence number. To modify an existing entry, you must delete the entry and retype it, or create a new entry with a new unique name.

To delete an access list entry, use the command:

```
delete access-list <name>
```

MAXIMUM ENTRIES

A maximum of 255 entries with an assigned precedence can be used. In addition to the 255 entries, entries that do not use precedence can also be created, with the following restrictions:

- A source IP address must use wildcards or be completely specified (32 bit mask).
- The layer 4 source and destination ports must use wildcards or be completely specified (no ranges).
- No physical source port can be specified.
- Access list rules that apply to all physical ports are implemented on all BlackDiamond I/O modules.

On a BlackDiamond 6808 switch the maximum number of access list entries is 255 entries per I/O module. One way to economize on the number of entries on a BlackDiamond switch is to provide a physical ingress port as a component of an access list rule. In this case, the rule is implemented only on the I/O modules that contain the specified ports. By restricting rules to specific I/O modules, you can extend the number of access list rules to 2,040 ($255 * 8$).

ACCESS LISTS FOR ICMP

Access lists for ICMP traffic processing are handled in a slightly different manner. An access list for ICMP is only effective for traffic routed by the switch. ICMP traffic may either be forwarded (routed) by the switch or discarded, but cannot contain options for assigning a QoS profile. Other included configuration options for filtering ICMP include:

- IP source and destination address and mask
- ICMP type code
- Physical source port (optional)
- Numbered precedence (optional)

VERIFYING ACCESS LIST CONFIGURATIONS

To verify access list settings you can view the access list configuration and see real-time statistics on which access list entries are being accessed when processing traffic.

To view the access list configuration and statistics screen use the following command:

```
show access-list {name | port <port>}
```

To refresh the access list statistics display, use the following command:

```
show access-list-monitor
```

ACCESS LIST COMMANDS

[Table 16-1](#) describes the commands used to configure IP access lists.

Table 16-1: Access List Configuration Commands

Command	Description
<pre>create access-list <name> ip destination [<dst_ipaddress>/<dst_mask> any] source [<src_ipaddress>/<src_mask> any] [permit <qosprofile> deny] ports [<portlist> any] {precedence <precedence_num>} {log}</pre>	<p>Creates a named IP access list. The access list is applied to all ingress packets. Options include:</p> <ul style="list-style-type: none"> ■ <name> — Specifies the access list name. The access list name can be between 1 and 16 characters. ■ ip — Specifies that the rule applies to IP traffic. ■ destination — Specifies an IP destination address and subnet mask. A mask length of 32 indicates a host entry. ■ source — Specifies an IP source address and subnet mask. ■ permit — Specifies the packets that match the access list description are permitted to be forward by this switch. An optional QoS profile can be assigned to the access list, so that the switch can prioritize packets accordingly. ■ deny — Specifies the packets that match the access list description are filtered (dropped) by the switch. ■ ports — Specifies the ingress port(s) on which this rule is applied. ■ precedence — Specifies the access list precedence number. The range is 1 to 25,600. ■ log — Logs a message to the Syslog facility for each packet that matches the access-list description. The message details the properties of the packet.

Table 16-1: Access List Configuration Commands (continued)

Command	Description
<pre>create access-list <name> tcp destination [<dst_ipaddress>/<dst_mask> any] ip-port [<dst_port> range <dst_port_min> <dst_port_max> any] source [<src_ipaddress>/<src_mask> any] ip-port [<src_port> range <src_port_min> <src_port_max> any] [permit <qosprofile> permit-established deny] ports [<portlist> any] {precedence <precedence_num>} {log}</pre>	<p>Creates a named IP access list. The access list is applied to all ingress packets. Options include:</p> <ul style="list-style-type: none"> ■ <name> — Specifies the access list name. The access list name can be between 1 and 16 characters. ■ tcp — Specifies that the rule applies to TCP traffic. ■ destination — Specifies an IP destination address and subnet mask. A mask length of 32 indicates a host entry. ■ source — Specifies an IP source address and subnet mask. ■ permit-established — Specifies a uni-directional session establishment is allowed. ■ permit — Specifies the packets that match the access list description are permitted to be forwarded by this switch. An optional QoS profile can be assigned to the access list, so that the switch can prioritize packets accordingly. ■ range — Specifies the TCP or UDP port range. ■ deny — Specifies the packets that match the access list description are filtered (dropped) by the switch. ■ ports — Specifies the ingress port(s) on which this rule is applied. ■ precedence — Specifies the access list precedence number. The range is 1 to 25,600. ■ log — Logs a message to the Syslog facility for each packet that matches the access-list description. The message details the properties of the packet.

Table 16-1: Access List Configuration Commands (continued)

Command	Description
<pre>create access-list <name> udp destination [<dst_ipaddress>/<dst_mask> any] ip-port [<dst_port> range <dst_port_min> <dst_port_max> any] source [<src_ipaddress>/<src_mask> any] ip-port [<src_port> range <src_port_min> <src_port_max> any] [permit <qosprofile> deny] ports [<portlist> any] {precedence <precedence_num>} {log}</pre>	<p>Creates a named IP access list. The access list is applied to all ingress packets. Options include:</p> <ul style="list-style-type: none"> ■ <name> — Specifies the access list name. The access list name can be between 1 and 16 characters. ■ udp — Specifies that the rule applies to UDP traffic. ■ destination — Specifies an IP destination address and subnet mask. A mask length of 32 indicates a host entry. ■ source — Specifies an IP source address and subnet mask. ■ permit — Specifies the packets that match the access list description are permitted to be forwarded by this switch. An optional QoS profile can be assigned to the access list, so that the switch can prioritize packets accordingly. ■ range — Specifies the TCP or UDP port range. ■ deny — Specifies the packets that match the access list description are filtered (dropped) by the switch. ■ ports — Specifies the ingress port(s) on which this rule is applied. ■ precedence — Specifies the access list precedence number. The range is 1 to 25,600. ■ log — Logs a message to the Syslog facility for each packet that matches the access-list description. The message details the properties of the packet.

Table 16-1: Access List Configuration Commands (continued)

Command	Description
create access-list icmp destination [<dest_ipaddress>/<mask> any] source [<src_ipaddress>/<source_mask> any] type <icmp_type> code <icmp_code> [permit deny] {<portlist>} {log}	<p>Creates a named IP access list. The access list is applied to all ingress packets. Options include:</p> <ul style="list-style-type: none"> ■ <name> — Specifies the access list name. The access list name can be between 1 and 16 characters. ■ icmp — Specifies an ICMP access list. ■ destination — Specifies an IP destination address and subnet mask. A mask length of 32 indicates a host entry. ■ source — Specifies an IP source address and subnet mask. ■ type — Specifies the ICMP_TYPE number. The ICMP type is a number from 0 to 255. ■ code — Specifies the ICMP_CODE number. The ICMP code is a number from 0 to 255. ■ permit — Specifies the packets that match the access list description are permitted to be forward by this switch. An optional QoS profile can be assigned to the access list, so that the switch can prioritize packets accordingly. ■ deny — Specifies the packets that match the access list description are filtered (dropped) by the switch. ■ log — Logs a message to the Syslog facility for each packet that matches the access-list description. The message details the properties of the packet.
delete access-list <name>	Deletes an access list.
disable access-list <name> counter	Disables the collection of access-list statistics.
enable access-list <name> counter	Enables the collection of access-list statistics. The default setting is enabled.
show access-list {<name> ports <portlist>}	Displays access-list information.
show access-list-fdb	Displays the hardware access control list mapping.
show access-list-monitor	Refreshes the access-list information display.

IP ACCESS LIST EXAMPLES

This section presents two IP access list examples:

- Using the permit-established keyword
- Filtering ICMP packets

USING THE PERMIT-ESTABLISHED KEYWORD

This example uses an access list that permits TCP sessions (Telnet, FTP, and HTTP) to be established in one direction.

The Summit7i, shown in [Figure 16-1](#), is configured as follows:

- Two vlans, NET10 VLAN and NET20 VLAN are defined.
- The IP addresses for NET10 VLAN is 10.10.10.1/24.
- The IP address for NET20 VLAN is 10.10.20.1/24.
- The workstations are configured using addresses 10.10.10.100 and 10.10.20.100.
- IPForwarding is enabled.

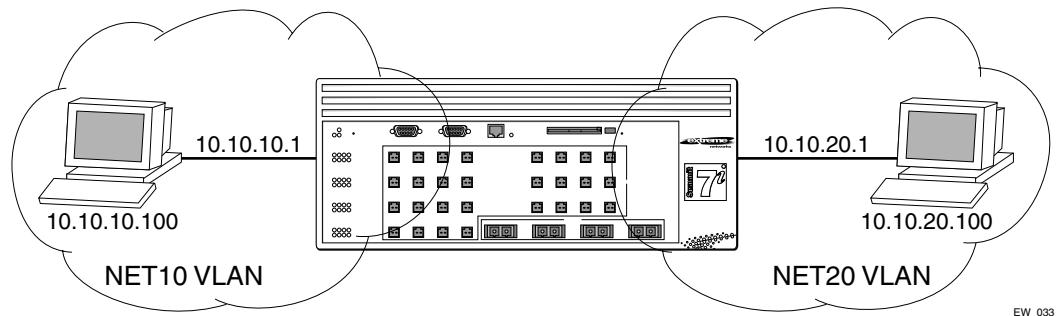


Figure 16-1: Permit-established access list example topology

The following sections detail the steps used to configure the example.

Step 1 – Deny IP Traffic.

First, create an access-list that blocks all IP-related traffic. This includes any TCP- and UDP-based traffic. Although ICMP is used in conjunction with IP, it is technically not an IP data packet. Thus, ICMP data traffic, such as ping traffic, is not affected.

The following command creates the access-list:

```
create access-list denyall ip destination any source any deny ports any
```

[Figure 16-2](#) illustrates the outcome of the access list.

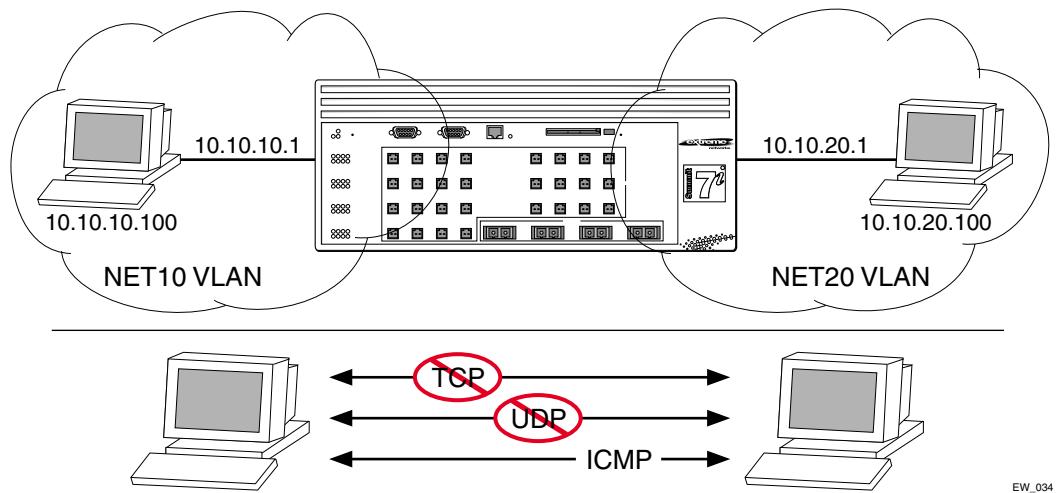


Figure 16-2: Access list denies all TCP and UDP traffic

Step 2 – Allow TCP traffic.

The next set of access-list commands permits TCP-based traffic to flow. Because each session is bi-directional, an access-list must be defined for each direction of the traffic flow. UDP traffic is still blocked.

The following commands create the access list:

```
create access-list tcp1 tcp destination 10.10.20.100/32 ip any source
10.10.10.100/32 ip any permit qpl ports any precedence 20
```

```
create access-list tcp2 tcp destination 10.10.10.100/32 ip any source
10.10.20.100/32 ip any permit qpl ports any precedence 21
```

[Figure 16-3](#) illustrates the outcome of this access list.

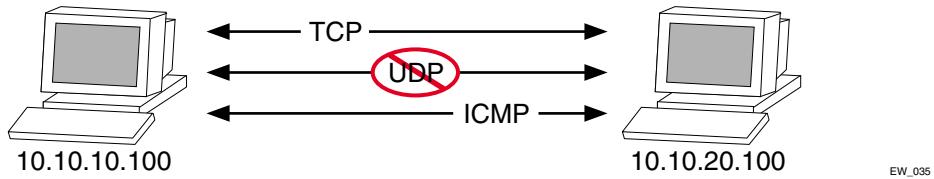


Figure 16-3: Access list allows TCP traffic

Step 3 - Permit-Established Access List.

When a TCP session begins, there is a 3-way handshake that includes a sequence of a SYN, SYN/ACK and ACK packets. [Figure 16-4](#) shows an illustration of the handshake that occurs when Host A initiates a TCP session to Host B. After this sequence, actual data can be passed.

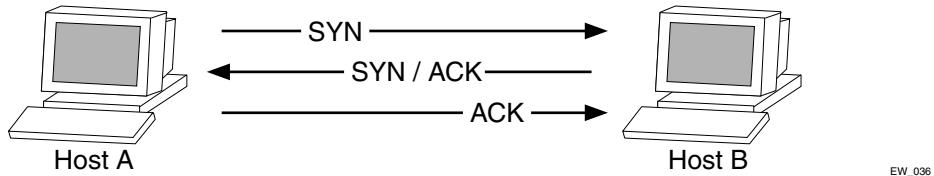


Figure 16-4: Host A initiates a TCP session to Host B

An access list that uses the permit-established keyword filters the SYN packet in one direction.

Use the permit-established keyword to allow only Host A to be able to establish a TCP session to Host B and to prevent any TCP sessions from being initiated by Host B, as illustrated in Figure 16-4. The syntax for this access-list is as follows:

```
create access-list <name> tcp destination HostA ip-port 23 source HOSTB
ip-port any permit-established ports any pre 8
```



This step may not be intuitive. Pay attention to the destination and source address, and the desired affect.

The exact command line entry for this example is as follows:

```
create access-list telnet-allow tcp destination 10.10.10.100/32 ip-port  
23 source any ip-port any permit-established ports any pre 8
```



This rule has a higher precedence than the rule "tcp2."

[Figure 16-5](#) shows the final outcome of this access list.

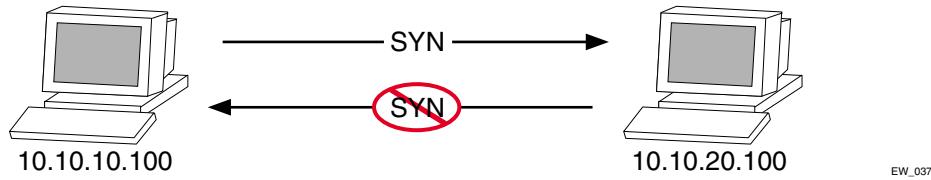


Figure 16-5: Permit-established access list filters out SYN packet to destination

EXAMPLE 2: FILTER ICMP PACKETS

This example creates an access list that filters out ping (ICMP echo) packets. ICMP echo packets are defined as type 8 code 0.

The command line syntax to create this access list is as follows:

```
create access-list denyping icmp destination any source any type 8 code  
0 deny ports any
```

The output for this access list is shown in [Figure 16-6](#).

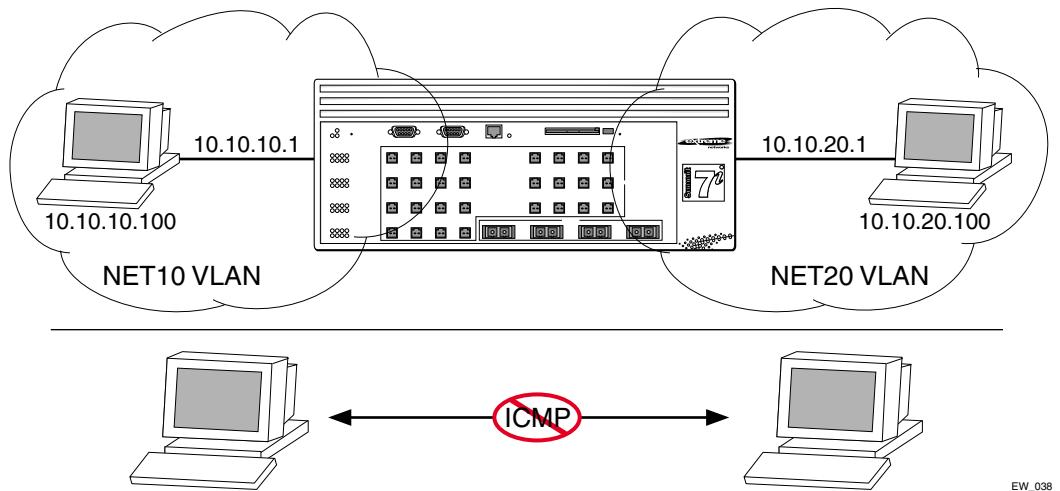


Figure 16-6: ICMP packets are filtered out

EW_038

USING ROUTING ACCESS POLICIES

To use routing access policies, you must perform the following steps:

- 1 Create an access profile.
- 2 Configure the access profile to be of type *permit*, *deny*, or *none*.
- 3 Add entries to the access profile. Entries can be one of the following types:
 - IP addresses and subnet masks
 - VLANs (BlackDiamond switch only)
 - Autonomous system path expressions (as-paths) (BGP only)
 - BGP communities (BGP only)
- 4 Apply the access profile.

CREATING AN ACCESS PROFILE

The first thing to do when using routing access policies is to create an *access profile*. An access profile has a unique name, and contains one of the following entry types:

- A list of IP addresses and associated subnet masks
- One or more autonomous system path expressions (BGP only)
- One or more BGP community numbers (BGP only)

You must give the access profile a unique name (in the same manner as naming a VLAN, protocol filter, or Spanning Tree Domain). To create an access profile, use the following command:

```
create access-profile <access_profile> type [ipaddress | as-path |  
bgp-community]
```

CONFIGURING AN ACCESS PROFILE MODE

After the access profile is created, you must configure the access profile mode. The access profile mode determines whether the items in the list are to be permitted access or denied access.

There are three available modes:

- Permit — The permit access profile mode permits the operation, as long as it matches any entry in the access profile. If the operation does not match any entries in the list, the operation is denied.
- Deny — The deny access profile mode denies the operation, as long as it matches any entry in the access profile. If it does not match all specified entries in the list, the operation is permitted.
- None — Using the none mode, the access profile can contain a combination of permit and deny entries. Each entry must have a permit or deny attribute. The operation is compared with each entry in the list. Once a match is found, the operation is either permitted or denied, depending on the configuration of the matched entry. If no match is found, the operation is implicitly denied.

To configure the access profile mode, use the following command:

```
config access-profile <access_profile> mode [permit | deny | none]
```

ADDING AN ACCESS PROFILE ENTRY

Next, configure the access profile by adding or deleting IP addresses, autonomous system path expressions, or BGP communities, using the following command:

```
config access-profile <access_profile> add {<seq_number>} {permit | deny} [ipaddress <ipaddress> <mask> {exact} | as-path <path-expression> | bgp-community [internet | no-export | no-advertise | no-export-subconfed | <as_no:number> | number <community>]]
```

The following sections describe the config access-profile add command.

SPECIFYING SUBNET MASKS

The subnet mask specified in the access profile command is interpreted as a *reverse mask*. A reverse mask indicates the bits that are significant in the IP address. In other words, a reverse mask specifies the part of the address that must match the IP address to which the profile is applied.

If you configure an IP address that is an exact match that is specifically denied or permitted, use a mask of /32 (for example, 141.251.24.28/32). If the IP address represents all addresses in a subnet address that you wish to deny or permit, then configure the mask to cover only the subnet portion (for example, 141.251.10.0/24). The keyword *exact* may be used when you wish to match only against the subnet address, and ignore all addresses within the subnet.

If you are using off-byte boundary subnet masking, the same logic applies, but the configuration is more tricky. For example, the address 141.251.24.128/27 represents any host from subnet 141.251.24.128.

SEQUENCE NUMBERING

You can specify the sequence number for each access profile entry. If you do not specify a sequence number, entries are sequenced in the order they are added. Each entry is assigned a value of 5 more than the sequence number of the last entry.

PERMIT AND DENY ENTRIES

If you have configured the access profile mode to be none, you must specify each entry type as either 'permit' or 'deny'. If you do not specify the entry type, it is added as a permit entry. If you have configured the access profile mode to be permit or deny, it is not necessary to specify a type for each entry.

AUTONOMOUS SYSTEM EXPRESSIONS

The AS-path keyword uses a regular expression string to match against the AS path. Regular expression notation can include any of the characters listed in [Table 16-2](#).

Table 16-2: Regular Expression Notation

Character	Definition
[.]	Specifies a range of numbers to be matched.
.	Matches any number
^	Matches the beginning of the AS path
\$	Matches the end of the AS path
-	Matches the beginning or end, or a space
-	Separates the beginning and end of a range of numbers
*	Matches 0 or more instances
+	Matches 1 or more instances
?	Matches 0 or 1 instance

DELETING AN ACCESS PROFILE ENTRY

To delete an access profile entry, use the following command:

```
config access-profile <access_profile> delete <seq_number>
```

APPLYING ACCESS PROFILES

Once the access profile is defined, apply it to one or more routing protocols or VLANs. When an access profile is applied to a protocol function (for example, the export of RIP routes) or a VLAN, this forms an access policy. A profile can be used by multiple routing protocol functions or VLANs, but a protocol function or VLAN can use only one access profile.

ROUTING ACCESS POLICIES FOR RIP

If you are using the RIP protocol, the switch can be configured to use an access profile to determine any of the following:

- **Trusted Neighbor** — Use an access profile to determine trusted RIP router neighbors for the VLAN on the switch running RIP. To configure a trusted neighbor policy, use the following command:

```
config rip vlan [<name> | all] trusted-gateway [<access_profile> | none]
```

- **Import Filter** — Use an access profile to determine which RIP routes are accepted as valid routes. This policy can be combined with the trusted neighbor policy to accept selected routes only from a set of trusted neighbors. To configure an import filter policy, use the following command:

```
config rip vlan [<name> | all] import-filter [<access_profile> | none]
```

- **Export Filter** — Use an access profile to determine which RIP routes are advertised into a particular VLAN, using the following command:

```
config rip vlan [<name> | all] export-filter [<access_profile> | none]
```

EXAMPLES

In the example shown in [Figure 16-7](#), a switch is configured with two VLANs, *Engsvrs* and *Backbone*. The RIP protocol is used to communicate with other routers on the network. The administrator wants to allow all internal access to the VLANs on the switch, but no access to the router that connects to the Internet. The remote router that connects to the Internet has a local interface connected to the corporate backbone. The IP address of the local interface connected to the corporate backbone is 10.0.0.10/24.

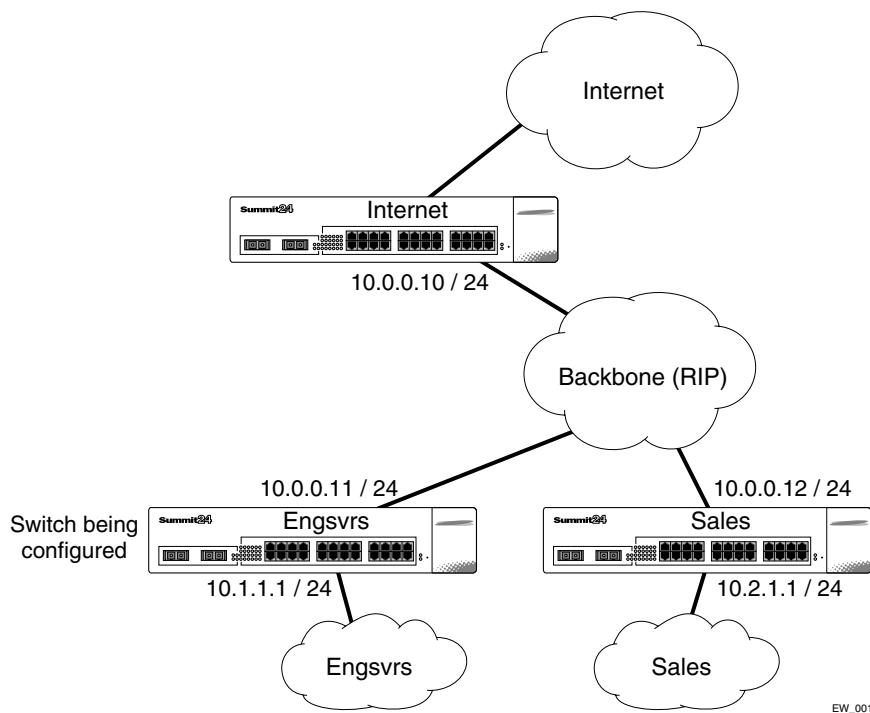


Figure 16-7: RIP access policy example

Assuming the backbone VLAN interconnects all the routers in the company (and, therefore, the Internet router does not have the best routes for other local subnets), the commands to build the access policy for the switch would be the following:

```
create access-profile nointernet ipaddress
config access-profile nointernet mode deny
config access-profile nointernet add 10.0.0.10/32
config rip vlan backbone trusted-gateway nointernet
```

In addition, if the administrator wants to restrict any user belonging to the VLAN *Engsrs* from reaching the VLAN *Sales* (IP address 10.2.1.0/24) , the additional access policy commands to build the access policy would be as follows:

```
create access-profile nosales ipaddress
config access-profile nosales mode deny
config access-profile nosales add 10.2.1.0/24
config rip vlan backbone import-filter nosales
```

This configuration results in the switch having no route back to the VLAN *Sales*.

ROUTING ACCESS POLICIES FOR OSPF

Because OSPF is a link-state protocol, the access policies associated with OSPF are different in nature than those associated with RIP. Access policies for OSPF are intended to extend the existing filtering and security capabilities of OSPF (for example, link authentication and the use of IP address ranges). If you are using the OSPF protocol, the switch can be configured to use an access profile to determine any of the following:

- **Inter-area Filter** — For switches configured to support multiple OSPF areas (an ABR function), an access profile can be applied to an OSPF area that filters a set of OSPF inter-area routes from being sourced from any other areas. To configure an inter-area filter policy, use the following command:

```
config ospf area <area_id> interarea-filter [<access_profile> | none]
```

- **External Filter** — For switches configured to support multiple OSPF areas (an ABR function), an access profile can be applied to an OSPF area that filters a set of OSPF external routes from being advertised into that area. To configure an external filter policy, use the following command:

```
config ospf area <area_id> external-filter [<access_profile> | none]
```



If any of the external routes specified in the filter have already been advertised, those routes will remain until the associated LSAs in that area time-out.

- **ASBR Filter** — For switches configured to support RIP and static route re-distribution into OSPF, an access profile can be used to limit the routes that are advertised into OSPF for the switch as a whole. To configure an ASBR filter policy, use the following command:

```
config ospf asbr-filter [<access_profile> | none]
```

- **Direct Filter** — For switches configured to support direct route re-distribution into OSPF, an access profile can be used to limit the routes that are advertised into OSPF

for the switch as a whole. To configure a direct filter policy, use the following command:

```
config ospf direct-filter [<access_profile> | none]
```

EXAMPLE

Figure 16-8 illustrates an OSPF network that is similar to the network used previously in the RIP example. In this example, access to the Internet is accomplished by using the ASBR function on the switch labeled Internet. As a result, all routes to the Internet will be done through external routes. Suppose the network administrator wishes to only allow access to certain internet addresses falling within the range 192.1.1.0/24 to the internal backbone.

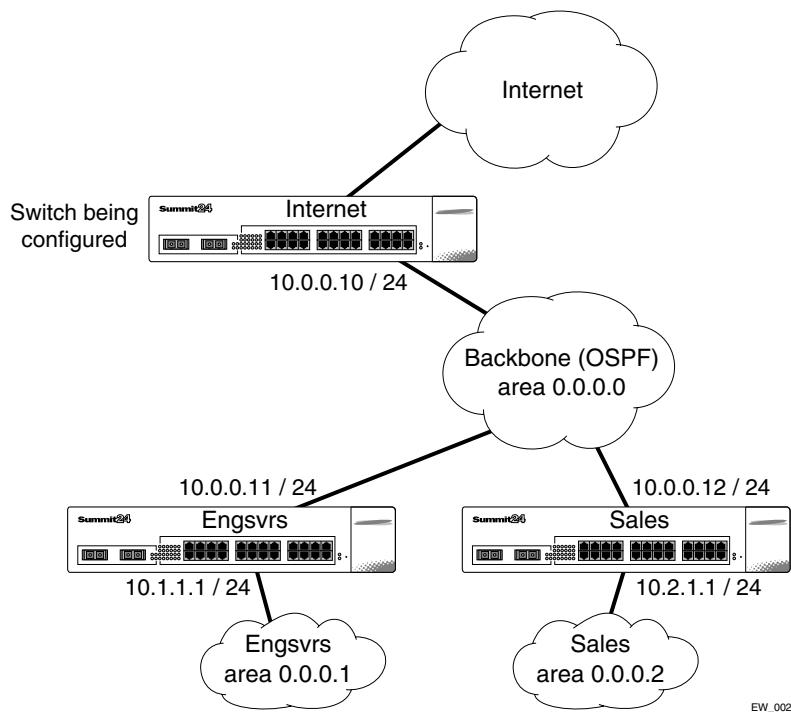


Figure 16-8: OSPF access policy example

To configure the switch labeled Internet, the commands would be as follows:

```
create access-profile okinternet ipaddress
config access-profile okinternet mode permit
config access-profile okinternet add 192.1.1.0/24
config ospf asbr-filter okinternet
```

ROUTING ACCESS POLICIES FOR DVMRP

The access policy capabilities for DVMRP are very similar to those for RIP. If you are using the DVMRP protocol is used for routing IP multicast traffic, you can configure the switch to use an access profile to determine any of the following:

- **Trusted Neighbor** — Use an access profile to determine trusted DVMRP router neighbors for the VLAN on the switch running DVMRP. To configure a trusted neighbor policy, use the following command:

```
config dvmrp vlan [<name> | all] trusted-gateway [<access_profile> | none]
```

- **Import Filter** — Use an access profile to determine which DVMRP routes are accepted as valid routes. To configure an import filter policy, use the following command:

```
config dvmrp vlan [<name> | all] import-filter [<access_profile> | none]
```

- **Export-Filter** — Use an access profile to determine which DVMRP routes are advertised into a particular VLAN, using the following command:

```
config dvmrp vlan [<name> | all] export-filter [<access_profile> | none]
```

EXAMPLE

In this example, the network used in the previous RIP example is configured to run DVMRP. The network administrator wants to disallow Internet access for multicast traffic to users on the VLAN *Engsvrs*. This is accomplished by preventing the learning of routes that originate from the switch labeled Internet by way of DVMRP on the switch labeled *Engsvrs*.

To configure the switch labeled Engsvrs, use the following commands:

```
create access-profile nointernet ipaddress
config access-profile nointernet mode deny
config access-profile nointernet add 10.0.0.10/32
config dvmrp vlan backbone trusted-gateway nointernet
```

In addition, suppose the administrator wants to preclude users on the VLAN *Engsvrs* from seeing any multicast streams that are generated by the VLAN *Sales* across the backbone. The additional configuration of the switch labeled Engsvrs is as follows:

```
create access-profile nosales ipaddress
config access-profile nosales mode deny
config access-profile nosales add 10.2.1.0/24
config dvmrp vlan backbone import-filter nosales
```

ROUTING ACCESS POLICIES FOR PIM

Because PIM leverages the unicast routing capability that is already present in the switch, the access policy capabilities are, by nature, different. If you are using the PIM protocol for routing IP multicast traffic, you can configure the switch to use an access profile to determine any of the following:

- **Trusted Neighbor** — Use an access profile to determine trusted PIM router neighbors for the VLAN on the switch running PIM. To configure a trusted neighbor policy, use the following command:

```
config pim vlan [<name> | all] trusted-gateway [<access_profile> | none]
```

EXAMPLE

Using PIM, the unicast access policies can be used to restrict multicast traffic. In this example, a network similar to the example used in the previous RIP example is also running PIM. The network administrator wants to disallow Internet access for multicast traffic to users on the VLAN *Engsvrs*. This is accomplished by preventing the learning of routes that originate from the switch labeled Internet by way of PIM on the switch labeled Engsvrs.

To configure the switch labeled Engsvrs, the commands would be as follows:

```
create access-profile nointernet ipaddress
config access-profile nointernet mode deny
config access-profile nointernet add 10.0.0.10/32
config pim vlan backbone trusted-gateway nointernet
```

ROUTING ACCESS POLICIES FOR BGP

If the BGP protocol is being used, the switch can be configured to use an access profile to determine any of the following:

- **NLRI filter** — Use an access profile to determine the NLRI information that must be exchanged with a neighbor. To configure an NLRI filter policy, use the following command:

```
config bgp neighbor [<ipaddress> | all] nlri-filter [in | out]
[<access_profile> | none]
```

The NLRI filter access policy can be applied to the ingress or egress updates, using the `in` and `out` keywords, respectively.

- **Autonomous system path filter** — Use an access profile to determine which NLRI information must be exchanged with a neighbor based on the AS path information present in the path attributes of the NLRI. To configure an autonomous system path filter policy, use the following command:

```
config bgp neighbor [<ipaddress> | all] as-path-filter [in | out]
[<access_profile> | none]
```

The autonomous system path filter can be applied to the ingress or egress updates, using the `in` and `out` keywords, respectively.

MAKING CHANGES TO A ROUTING ACCESS POLICY

You can change the routing access policy by changing the associated access profile. However, the propagation of the change depends on the protocol and policy involved. Propagation of changes applied to RIP, DVMRP, and PIM access policies depend on the respective protocol timers to age-out entries.

In BGP, the change to the policy is immediately effective on the routing information exchanged after the policy changes. The changes can be applied on the routing information that had been exchanged before the policy changes by issuing a soft reset

on the ingress or egress side, depending on the change. For soft resets to be applied on the ingress side, the changes must have been previously enabled on the neighbor.

-  *Changes to profiles applied to OSPF typically require rebooting the switch, or disabling and re-enabling OSPF on the switch.*

REMOVING A ROUTING ACCESS POLICY

To remove a routing access policy, you must remove the access profile from the routing protocol or VLAN. All the commands that apply an access profile to form an access policy also have the option of choosing none as the access profile. Using the none option removes any access profile of that particular type from the protocol or VLAN, and, therefore, removes the access policy.

ROUTING ACCESS POLICY COMMANDS

Table 16-3 describes the commands used to configure routing access policies.

Table 16-3: Routing Access Policy Configuration Commands

Command	Description
<pre>config access-profile <access_profile> add {<seq-number>} {permit deny} [<ipaddress> <mask> {exact} as-path <path_expression> bgp-community [internet no-advertise no-export no-export-subconfed <as_no:number> number <community>]]]</pre>	<p>Adds an entry to the access profile. The explicit sequence number, and permit or deny attribute should be specified if the access profile mode is none.</p> <p>Specify one of the following:</p> <ul style="list-style-type: none"> ■ <seq-number> — The order of the entry within the access profile. If no sequence number is specified, the new entry is added to the end of the access-profile and is automatically assigned a value of 5 more than the sequence number of the last entry. ■ permit deny — Per-entry permit or deny specification. The per-entry attribute only takes effect if the access-profile mode is none. Otherwise, the overall access profile type takes precedence. ■ <ipaddress> <mask> — An IP address and mask. If the attribute “exact” is specified for an entry, then a exact match with address and mask is performed, subnets within the address range do not match entry against entry. ■ as-path — A regular expression string to match against the autonomous system path. ■ bgp-community — The BGP community number in as_no:number format, or as an unsigned 32-bit integer in decimal format. The BGP community internet matches against all routes, because all routes belong to the internet community.
<pre>config access-profile <access_profile> delete <seq_number></pre>	<p>, Deletes an access profile entry using the sequence number.</p>

Table 16-3: Routing Access Policy Configuration Commands (continued)

Command	Description
config access-profile <access_profile> mode [permit deny none]	<p>Configures the access profile to be one of the following:</p> <ul style="list-style-type: none"> ■ permit — Allows the addresses that match the access profile description. ■ deny — Denies the addresses that match the access profile description. ■ none — Permits and denies access on a per-entry basis. Each entry must be added to the profile as either type permit or deny. <p>The default setting is permit.</p>
config bgp neighbor [<ipaddress> all] as-path-filter [in out] [<access_profile> none]	Configures BGP to use the AS path filter for the routing information exchanged with the neighbor.
config bgp neighbor [<ipaddress> all] nlri-filter [in out] [<access_profile> none]	Configures BGP to use the NLRI filter for the routing information exchanged with the neighbor.
config dvmrp vlan [<name> all] export-filter [<access_profile> none]	Configures DVMRP to filter out certain routes when performing the route advertisement.
config dvmrp vlan [<name> all] import-filter [<access_profile> none]	Configures DVMRP to filter certain routes received from its neighbor.
config dvmrp vlan [<name> all] trusted-gateway [<access_profile> none]	Configures DVMRP to use the access policy to determine which DVMRP neighbor is trusted and to receive routes from.
config ospf area <area_id> external-filter [<access_profile> none]	Configures the router to use the access policy to determine which external routes are allowed to be exported into the area. This router must be an ABR.
config ospf area <area_id> interarea-filter [<access_profile> none]	Configures the router to use the access policy to determine which inter-area routes are allowed to be exported into the area. This router must be an ABR.
config ospf asbr-filter [<access_profile> none]	Configures the router to use the access policy to limit the routes that are advertised into OSPF for the switch as a whole for switches configured to support RIP and static route re-distribution into OSPF.
config ospf direct-filter [<access_profile> none]	Configures the router to use the access policy to limit the routes that are advertised into OSPF for the switch as a whole for switches configured to support direct route re-distribution into OSPF.

Table 16-3: Routing Access Policy Configuration Commands (continued)

Command	Description
config pim vlan [<name> all] trusted-gateway [<access-profile> none]	Configures PIM to use the access profile to determine which PIM neighbor is to receive or reject the routes.
config rip vlan [<name> all] export-filter [<access-profile> none]	Configures RIP to suppress certain routes when performing route advertisements.
config rip vlan [<name> all] import-filter [<access_profile> none]	Configures RIP to ignore certain routes received from its neighbor.
config rip vlan [<name> all] trusted-gateway [<access_profile> none]	Configures RIP to use the access list to determine which RIP neighbor to receive (or reject) the routes.
create access-profile <access_profile> type [ipaddress as-path bgp-community]	Creates an access profile. Once the access profile is created, one or more addresses can be added to it, and the profile can be used to control a specific routing protocol.
	<p>Specify one of the following:</p> <ul style="list-style-type: none"> ■ ipaddress — A list of IP address and mask pairs. ■ as-path — A list of AS path expressions. ■ bgp-community — A list of BGP community numbers.
delete access-profile <access_profile>	Deletes an access profile.
show access-profile <access_profile>	Displays access-profile related information for the switch.

USING ROUTE MAPS

Route maps are a mechanism that can be used to conditionally control the redistribution of routes between two routing domains, and to modify the routing information that is redistributed.

Route maps are used in conjunction with the match and set operations. A match operation specifies a criteria that must be matched. A set operation specifies a change that is made to the route when the match operation is successful.

To create a route map, do the following:

- 1 Create a route map.
- 2 Add entries to the route map.
- 3 Add statements to the route map entries.

CREATING A ROUTE MAP

To create a route map, use the following command:

```
create route-map <route-map>
```

ADD ENTRIES TO THE ROUTE MAP

To add entries to the route map, use the following command:

```
config route-map <route-map> add <sequence number> [permit | deny]
{match-one | match-all}
```

where the following is true:

- The sequence number uniquely identifies the entry, and determines the position of the entry in the route map. Route maps are evaluated sequentially.
- The permit keyword permits the route; the deny keyword denies the route and is applied only if the entry is successful.
- The match-one keyword is a logical “or”. The route map is successful as long as at least one of the matching statements is true.
- The match-all keyword is a logical “and”. The route map is successful when all match statements are true. This is the default setting.

ADD STATEMENTS TO THE ROUTE MAP ENTRIES

To add statements to the route map entries, use one of the following three commands:

```
config route-map <route-map> <sequence number> add match [nlri-list
<access_profile> | as-path [<access_profile> | <as_no>] | community
[access-profile <access_profile> | <as_num:number> | number <community>]
| next-hop <ipaddress> | med <number> | origin [igp | egp |
incomplete]]
```

```

config route-map <route-map> <sequence number> add set [as-path
<as_num> | community [remove | {add | delete} [access-profile
<access_profile> | <as_num:number> | number <number>] |] next-hop
<ipaddress> | med <number> | local-preference <number> | origin [igp |
egp | incomplete]

config route-map <route-map> <sequence number> add goto <route-map>

```

where the following is true:

- The `route-map` is the name of the route map.
- The `sequence number` identifies the entry in the route map to which this statement is being added.
- The `match`, `set`, and `goto` keywords specify the operations to be performed. Within a entry, the statements are sequenced in the order of their operation. The `match` statements are first, followed by `set`, and then `goto`.
- The `nlri-list`, `as-path`, `community`, `next-hop`, `med`, `origin`, and `weight` keywords specify the type of values that must be applied using the specified operation against the corresponding attributes as described in [Table 16-4](#) and [Table 16-5](#).

Table 16-4: Match Operation Keywords

Keyword	Description
<code>nlri-list <access_profile></code>	Matches the NLRI against the specified access profile.
<code>as-path [<access_profile> <as-no>]</code>	Matches the AS path in the path attributes against the specified access profile or AS number.
<code>community [<access_profile> <community>]</code>	Matches the communities in the path attribute against the specified BGP community access profile or the community number.
<code>next-hop <ipaddress></code>	Matches the next hop in the path attribute against the specified IP address.
<code>med <number></code>	Matches the MED in the path attribute against the specified MED number.
<code>origin [igp egp incomplete]</code>	Matches the origin in the path attribute against the specified origin.

Table 16-5: Set Operation Keywords

Keyword	Definition
as-path <as no>	Prepends the specified AS number to the AS path in the path attribute.
community <community>	Adds the specified community to the existing community in the path attribute.
next-hop <ipaddress>	Sets the next hop in the path attribute to the specified IP address.
med <number>	Sets the MED in the path attribute to the specified MED number.
local-preference <number>	Sets the local preference in the path attribute to the specified local preference number.
weight <number>	Sets the weight associated with the NLRI to the specified number.
origin	Sets the origin in the path attributes to the specified origin.

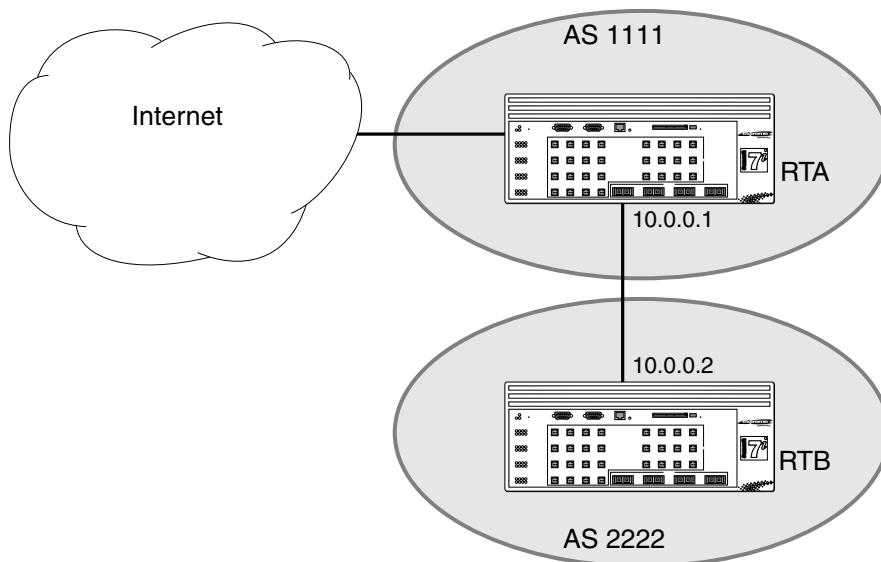
ROUTE MAP OPERATION

The entries in the route map are processed in the ascending order of the sequence number. Within the entry, the match statements are processed first. When the match operation is successful, the set and goto statements within the entry are processed, and the action associated with the entry is either applied, or else the next entry is processed. If the end of the route map is reached, it is implicitly denied.

When there are multiple match statement, the primitive match-one or match-all in the entry determines how many matches are required for success. When there are no match statements in an entry, the entry is considered a successful match.

ROUTE MAP EXAMPLE

[Figure 16-9](#) shows a topology in which route maps are used to filter or modify routing information that is exchanged between the neighbors RTA and RTB using BGP.



EW_048

Figure 16-9: Route maps

The following points apply to this example:

- RTA is a member of in AS 1111 and peers with a router in the Internet to receive the entire Internet routing table.
- RTB is a member of AS 2222, and has an EBGP connection with RTA through which it receives the Internet routing table.
- AS 1111 is acting as a transit AS for all traffic between AS 2222 and the Internet. If the administrator of AS 1111 wants to filter out route information about network 221.1.1.0 / 24 and its subnets from being passed on to AS 2222, s/he can configure a route-map on the egress side of RTA's EBGP connection with RTB and filter out the routes.

To configure RTA, use the following commands:

```
create access-profile iplist type ipaddress
config iplist add ipaddress 221.1.1.0 / 24

create route-map bgp-out
config bgp-out add 10 deny
config bgp-out 10 add match nlri-list iplist
```

```
config bgp-out add 20 permit  
  
config bgp neighbor 10.0.0.2 route-map-filter out bgp-out  
config bgp neighbor 10.0.0.2 soft-reset out
```

If you wish to modify the routing information originated from AS 300 to include a MED value of 200, the sequence of commands would be as follows:

```
create access-profile aslist type as-path  
config aslist add as-path "^300"  
  
config bgp-out add 15 permit  
config bgp-out 15 add match as-path access-profile aslist  
config bgp-out 15 add set med 200  
  
config bgp neighbor 10.0.0.2 soft-reset out
```

CHANGES TO ROUTE MAPS

Changes to the route maps used to modify or filter NLRI information exchanged with neighbors is immediately effective on the routing information exchanged after the policy changes. The changes can be applied on the NLRI information that had been exchanged before the policy changes by issuing a soft reset on the ingress or egress side, depending on the changes. For soft resets to be applied on the ingress side, the changes must be previously enabled on the neighbor.

Changes to the route maps associated with network aggregation or redistribution commands becomes effective after a maximum interval of 30 seconds. You can immediately apply them by using the soft reconfiguration command.

ROUTE MAPS IN BGP

Route maps are used in BGP to modify/filter NLRI information exchanged with neighbors. They are also used NLRI information that originates by way of network command, aggregation, or redistribution.

ROUTE MAP COMMANDS

[Table 16-6](#) describes route map commands.

Table 16-6: Route Map Commands

Command	Description
config route-map <route-map> <sequence number> add goto <route-map>	Configures a route map <code>goto</code> statement.
config route-map <route-map> <sequence number> add match [nlri-list <access_profile> as-path [<access_profile> <as_no>] community [access-profile <access_profile> <as_num:number number <community>] next-hop <ipaddress> med <number> origin [igp egp incomplete]]	Configures a route map <code>match</code> statement. Specify the following: <ul style="list-style-type: none">■ <code>route-map</code> – The name of the route map.■ <code>sequence number</code> – The statement in the route map to which this statement is being added.■ <code>nlri-list</code>, <code>as-path</code>, <code>community</code>, <code>next-hop</code>, <code>med</code>, and <code>origin</code> – Specify the type of values that must be applied using the specified operation against the corresponding attributes as described in Table 16-4.
config route-map <route-map> <sequence number> add set [as-path <as_num> community [remove {add delete} <access-profile <access_profile <as_num:number number <number> } next-hop <ipaddress> med <number> local-preference <number> origin [igp egp incomplete]	Configures a route map <code>set</code> statement. Specify the following: <ul style="list-style-type: none">■ <code>route-map</code> – The name of the route map.■ <code>sequence number</code> – The statement in the route map to which this statement is being added.■ <code>as-path</code>, <code>community</code>, <code>next-hop</code>, <code>med</code>, <code>local-preference</code>, and <code>origin</code> – Specify the type of values that must be applied using the specified operation against the corresponding attributes as described in Table 16-5.
config route-map <route-map> <sequence number> delete goto <route-map>	Deletes a route map <code>goto</code> statement.

Table 16-6: Route Map Commands (continued)

Command	Description
config route-map <route-map> <sequence number> delete match [nlri-list <access_profile> as-path [<access_profile> <as_no>] community [access-profile <access_profile> <as_num:number number <community>] next-hop <ipaddress> med <number> origin [igp egp incomplete]]]	Deletes a route-map match statement.
config route-map <route-map> <sequence number> delete set [as-path <as_num> community [remove {add delete} access-profile <access_profile <as_num:number> number <number>] next-hop <ipaddress> med <number> local-preference <number> origin [igp egp incomplete]]]	Deletes a route map set statement.
config route-map <route-map> add <sequence number> [permit deny] {match-one match-all}]	Adds a statement to the route map with the specified sequence number and action. The sequence number determines the order of the statement in the route map, and the action specifies the action to be taken on a successful match against the statements in the route map.
config route-map <route-map> delete <sequence number>	Deletes an statement from the route map.
create route-map <route-map>	Creates a route map statement.
delete route-map <route_map>	Deletes a route map statement from the route map.



Server Load Balancing (SLB)

This chapter describes the following topics:

- [Overview on page 17-2](#)
- [SLB Components on page 17-2](#)
- [Forwarding Modes on page 17-5](#)
- [VIP Network Advertisement on page 17-12](#)
- [Balancing Methods on page 17-13](#)
- [Basic SLB Commands on page 17-15](#)
- [Advanced SLB Application Example on page 17-18](#)
- [Health Checking on page 17-22](#)
- [Persistence on page 17-26](#)
- [Server Load Balancing with ESRP on page 17-27](#)
- [Using High Availability System Features on page 17-31](#)
- [3DNS Support on page 17-40](#)
- [Advanced SLB Commands on page 17-40](#)
- [Web Cache Redirection on page 17-46](#)

OVERVIEW

Server Load Balancing (SLB) is a feature of the switch that divides many client requests among several servers. This is done transparently to the client trying to use the resource. The main use for SLB is in the capacity of web hosting. Web hosting uses several redundant servers to increase the performance and reliability of busy websites.

Using SLB, the switch can manage and balance traffic for client equipment such as web servers, cache servers, routers, firewalls, and proxy servers. SLB has a variety of useful features that meet the special needs of e-commerce sites, Internet service providers, and managers of large intranets.

An introductory SLB application is shown in [Figure 17-1](#).

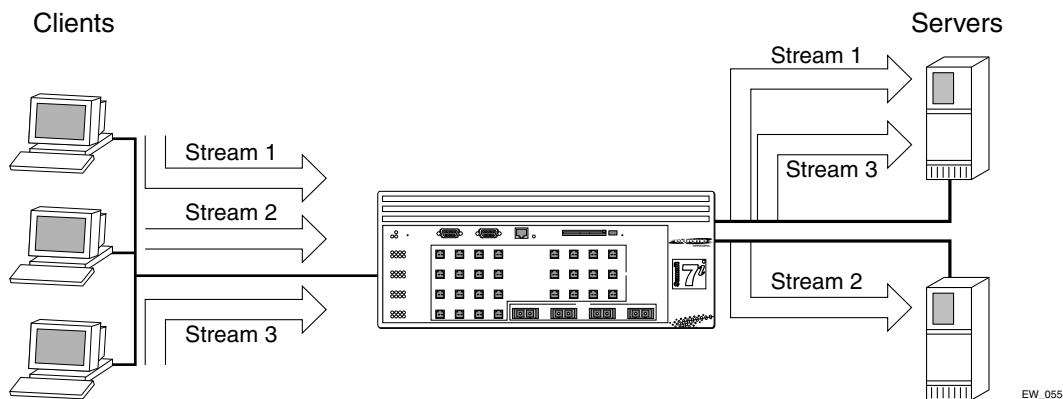


Figure 17-1: Introductory SLB application

SLB COMPONENTS

There are three components that comprise an SLB system:

- Nodes
- Pools
- Virtual Servers

All three components are required for every SLB configuration.

NODES

A node is an individual service on a physical server that consists of an IP addresses and a port number.

POOLS

A *pool* is a group of nodes that are mapped to a corresponding virtual server. Pools are used to more easily scale large networks that contain many nodes. Pools may be configured independently and associated with virtual servers in complex ways.

Each pool contains its own load balancing method. When associated with a virtual server, the pool can not be deleted from the SLB configuration. Pools must be added before, and deleted after, the virtual servers that reference them. If a pool is not associated with a virtual server, it is not used for load balancing.

To create a pool, use the following command:

```
create slb pool <poolname> {lb-method [round-robin | ratio | priority | least-connections]}
```

To add nodes to a pool, use the following command:

```
config slb pool <poolname> add <ipaddress>:<L4Port> {ratio <ratio> | priority <priority>}
```

To delete nodes from a pool, use the following command:

```
config slb pool <poolname> delete <ipaddress>:<L4Port>
```

VIRTUAL SERVERS

Virtual servers are the backbone of the SLB configuration. Virtual servers define the groups of servers or other network equipment that the switch load balances. Before you configure virtual servers, you need to know:

- The forwarding mode for your network design.
- The name of the pool.
- The virtual IP address.
- The virtual port number.

Once you know which virtual server options are useful in your network, you can:

- Define standard virtual servers.
- Define wildcard virtual servers.

USING STANDARD OR WILDCARD VIRTUAL SERVERS

Each virtual server maps to a single pool, which can be a group of content servers, firewalls, routers, or cache servers.

You can configure two different types of virtual servers:

- Standard virtual servers

A standard virtual server represents a site, such as a web site or an FTP site, and it provides load balancing for content servers. The virtual server IP address should be the same IP address that you register with the DNS for the site that the virtual server represents.

- Wildcard virtual servers

A wildcard virtual server load balances transparent network devices such as firewalls, routers, or cache servers. Wildcard virtual servers use a special wildcard IP address (0.0.0.0), and you can use them only if you have activated Transparent mode.

A virtual server is identified by a virtual IP address. To create a virtual server, use the following command:

```
create slb vip <vipname> pool <poolname> mode [transparent |  
translation | port-translation] <ipaddress>  
{-<upper_ipaddress>}:{<L4Port>}
```



For cache server applications, use Flow Redirection, described on [page 17-46](#).

FORWARDING MODES

The switch supports the following SLB forwarding modes:

- Transparent
- Translational
- Port Translation
- GoGo

[Table 17-1](#) summarizes the features supported by each forwarding mode.

Table 17-1: Forwarding Mode Feature Summary

	Transparent	Translational	Port Translation	GoGo
Performance	Hardware-based, server-to-client	CPU-based, bi-directional	CPU-based, bi-directional	Hardware-based, bi-directional
Load sharing algorithms	Round-robin, Ratio, Priority, Least Connections	Round-robin, Ratio, Priority, Least Connections	Round-robin, Ratio, Priority, Least Connections	Round-robin (hash)
Persistence	IPSA + Mask, IP list	IPSA + Mask, IP list	IPSA + Mask, IP list	IPSA
Health checking	L3, L4, L7, External	L3, L4, L7, External	L3, L4, L7, External	L1

TRANSPARENT MODE

Using transparent mode, the switch does not modify the IP addresses before sending the traffic on to the selected server. To accomplish this all servers must respond to the IP addresses associated with the virtual server. This virtual IP address is the address used by the clients to connect to the virtual server. The servers must have this address installed as a loopback address and have the address associated with the virtual server to be load balanced. As with any server load balancing application, the content must be duplicated on all physical servers.

In transparent mode, servers can be directly attached or have an L2 switch between the SLB switch and the server. It is not possible to have a router between the SLB switch and the servers being balanced.

To configure transparent mode, use the following command:

```
create slb vip <vipname> pool <poolname> mode transparent <ipaddress>
{- <upper_ipaddress>}:{<L4Port>}
```

Transparent mode is shown in [Figure 17-2](#).

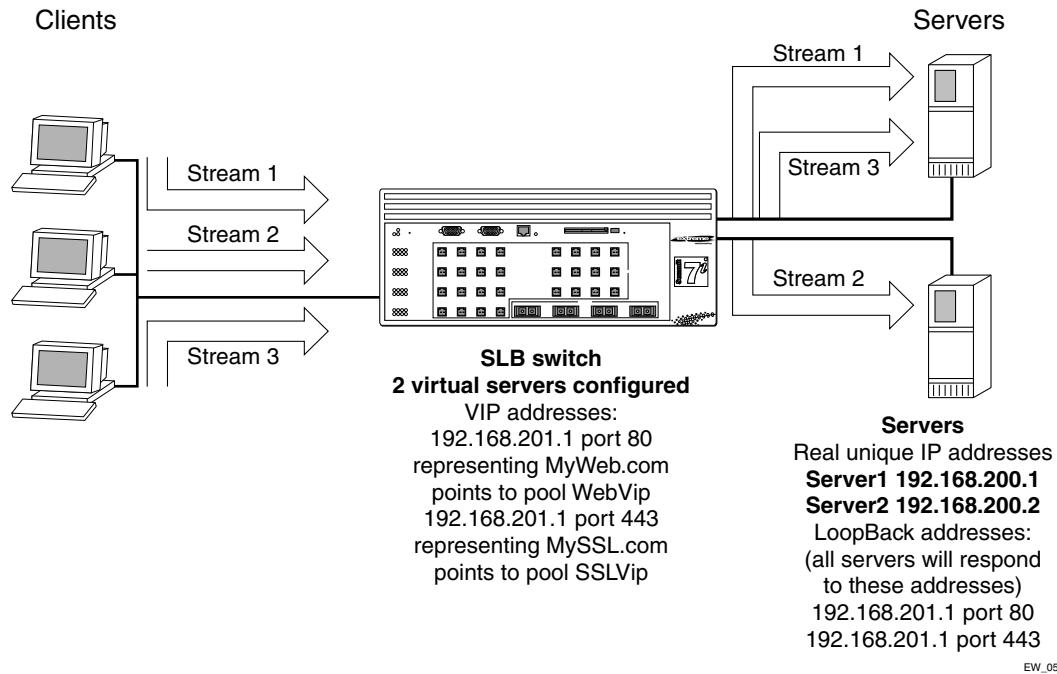


Figure 17-2: Transparent mode

In [Figure 17-2](#), the switch is configured to respond to requests for the VIP by forwarding them to the load balanced servers.

The servers are configured as follows:

- The interface for server 1 is 192.168.200.1.
- The interface for server 2 is 192.168.200.2.
- The loopback address on the servers is 192.168.201.1 (VIP).

- The service is configured to use the appropriate address and port, as specified in the switch configuration.

The commands used to configure the switch in [Figure 17-2](#) are described below.

The following commands configure the VLANs and the switch IP addresses and subnets:

```
create vlan srvr
create vlan clnt
create vlan vips
config srvr ipaddress 192.168.200.10 /24
config clnt ipaddress 10.1.1.1 /24
config vips ipaddress 192.168.201.1 /24
config server add port 29-32
config client add port 1-4
enable ipforwarding
```

The following commands create a round-robin pool, MyWeb, and add nodes to the new pool.

```
create slb pool MyWeb lb-method round
config slb pool MyWeb add 192.168.200.1:80
config slb pool MyWeb add 192.168.200.2:80
```

The following command creates a transparent mode VIP for the website and assigns the MyWeb pool to it:

```
create slb vip WebVip pool MyWeb mode transparent 192.168.201.1:80
```

The following commands create a round-robin pool, MySQL, and add nodes to the new pool.

```
create slb pool MySQL lb-method round-robin
config slb pool MySQL add 192.168.200.1:443
config slb pool MySQL add 192.168.200.2:443
```

The following command creates a transparent mode VIP for the website and assigns the MySQL pool to it.

```
create slb vip SSLVip pool MySQL mode transparent 192.168.201.1:443
```

The following commands enable SLB, configure the server VLAN to act as the server side, and configure the client VLAN to act as the client side.

```
enable slb
config vlan srvr slb-type server
config vlan clnt slb-type client
```

Individual servers require that a loopback address be configured for each IP address to which the server will respond.

TRANSLATIONAL MODE

In translational mode, requests coming in to the VIP are translated to the IP address of the server to be balanced. This mode does not require the configuration of a loopback address, each server only needs to use its own IP address. As with any server load balancing application, the content must be duplicated on all physical servers.

To configure translational mode, use the following command:

```
create slb vip <vipname> pool <poolname> mode translation <ipaddress>
{- <upper_ipaddress>}:{<L4Port>}
```

[Figure 17-3](#) shows translational mode.

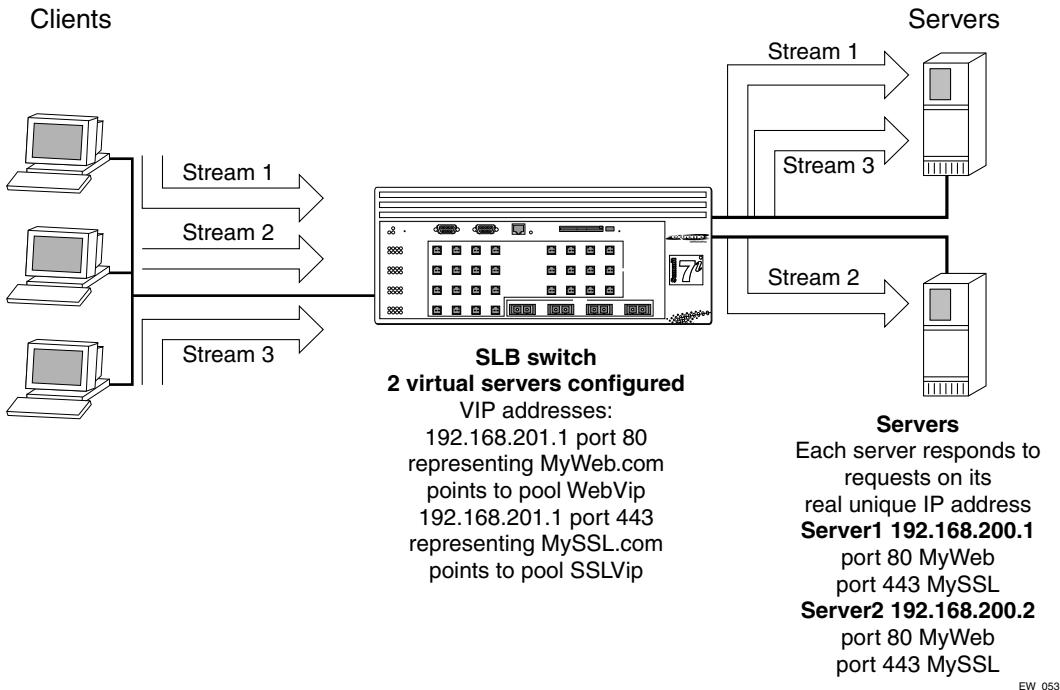


Figure 17-3: Translational mode

In [Figure 17-3](#), the switch is configured to respond to requests for the VIP by translating them and forwarding them to the load balanced servers. No additional server configuration is needed.

The commands used to configure the switch are listed below.

The following commands configure the VLANs and the switch IP addresses and subnets.

```
create vlan srvr
create vlan clnt
create vlan vips
config srvr ipaddress 192.168.200.10 /24
config clnt ipaddress 10.1.1.1 /24
config vips ipaddress 192.168.201.1 /24
config server add port 29-32
```

```
config client add port 1-4
enable ipforwarding
```

The following commands create a round-robin pool, MyWeb, and add nodes to the new pool:

```
create slb pool MyWeb lb-method round
config slb pool MyWeb add 192.168.200.1:80
config slb pool MyWeb add 192.168.200.2:80
```

The following command creates a translation mode VIP for the website and assigns the MyWeb pool to it:

```
create slb vip WebVip pool MyWeb mode translation 192.168.201.1:80
```

The following commands create a round-robin pool, MySSL, and add nodes to the new pool:

```
create slb pool MySSL lb-method round
config slb pool MySSL add 192.168.200.1:443
config slb pool MySSL add 192.168.200.2:443
```

The following command creates a translation mode VIP for the website and assigns the MySSL pool to it:

```
create slb vip SSLVip pool MySSL mode translation 192.168.201.1:443
```

The following commands enable SLB, configure the server VLAN to act as the server side, and configure the client VLAN to act as the client side.

```
enable slb
config vlan srvr slb-type server
config vlan clnt slb-type client
```

PORT TRANSLATION MODE

Port translation is essentially the same thing as translational mode, except that the Layer 4 port on the virtual server can be different from the Layer 4 port on the nodes being load balanced. The switch takes the traffic and changes the IP address and port address to that of the servers to be balanced. As with any server load balancing application, the content must be duplicated on all physical servers.

To configure port translation mode, use the following command:

```
create slb vip <vipname> pool <poolname> mode port-translation
<ipaddress> {- <upper_ipaddress>}:{<L4Port>}
```

GoGo MODE

GoGo mode is a very fast (line rate) method of server load balancing. GoGo mode forwards traffic without manipulating packet content. Session persistence is maintained using IP source address persistence information.

Traffic is optimally balanced across groups of 2, 4, or 8 directly attached servers. Because servers are always directly attached, there is no need to configure nodes, pools, or VIPs. In order to use GoGo mode, all servers are configured with the same MAC and IP addresses. As with any server load balancing application, the content must be duplicated on all physical servers.

In GoGo mode, the load balancing method is fixed, and is based on a hashing of the server IP address. All GoGo mode traffic exhibits persistence based on source IP information. That is, a given source address will be mapped to one and only one physical server.

[Figure 17-4](#) shows GoGo mode.

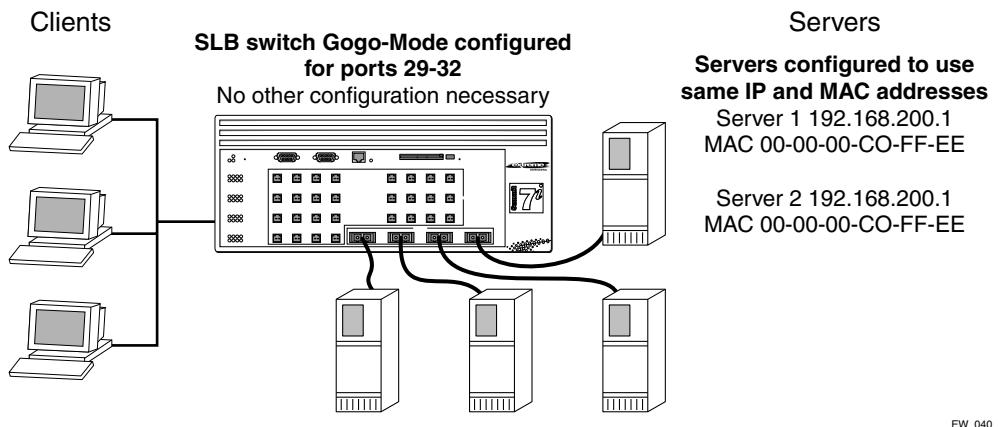


Figure 17-4: GoGo mode

In [Figure 17-4](#), the switch is configured to balance all traffic sent to the VIP based on the client IP address.

The servers are configured as follows:

- All servers have the same MAC address.
- All server have the same IP address.
- All servers must have the same content.

The commands used to configure the switch as indicated in the example are as follows:

```
create vlan server
create vlan client
config server ipaddress 10.1.1.1 /24
config client ipaddress 1.1.1.1 /24
config server add port 29-32
config client add port 1-4
enable slb gogo 29 grouping 29-32
enable ipforwarding
```

Separating clients and servers into separate VLANs is not a requirement in GoGo mode.

VIP NETWORK ADVERTISEMENT

There are three methods for controlling network connectivity to the VIPs. Depending on the subnet the VIP is a member of, the switch will adjust its behavior automatically.

- Proxy ARP

If the VIP is a member of an existing subnet the switch is directly attached to, the switch will respond to ARP requests on behalf of the VIP. This behavior is known as proxy ARP, and allows you to implement server load balancing on a layer 2 network. The VLAN containing the servers is a different subnet than the client VLAN's subnet. The VIP will appear to be a member of the client subnet.

- Host-Route

If the VIP created is not a member of an existing subnet the switch is directly attached to, a host-route entry will be added to the routing table for the switch. In this situation all clients will need to have a routed path to the VIP which points to the switch's IP address on the client VLAN.

- Subnet-Route

If your network configuration requires that the VIPs be propagated through a routing protocol by the switch, you will need to create a loopback vlan with the VIP(s) being valid members of the loopback VLAN's subnet. When a routing

protocol is enabled, the subnet containing the VIPs is propagated through the network.

BALANCING METHODS

A load balancing method defines, in part, the logic that the switch uses to determine which node should receive a connection hosted by a particular virtual server. Individual load balancing methods take into account one or more dynamic factors, such as current connection count. Because each application of SLB is unique, node performance depends on a number of different factors. We recommend that you experiment with different load balancing methods, and choose the one that offers the best performance in your particular environment.

The switch supports the following load balancing methods:

- Round-robin
- Ratio
- Least connections
- Priority

ROUND-ROBIN

The default load balancing method is round-robin, and it simply passes each new connection request to the next server in line, eventually distributing connections evenly across the array of machines being load balanced. Round-robin works well in most configurations, especially if the equipment that you are load balancing is roughly equal in processing speed and memory.

To configure round-robin, use this command:

```
config slb pool <poolname> lb-method round-robin
```

RATIO

If you are working with servers that differ significantly in processing speed and memory, you may want to switch to the ratio load balancing method. In ratio, the switch distributes connections among machines according to ratio weights that you set,

where the number of connections that each machine receives over time is proportionate to the ratio weight you defined for each machine.

The ratio method distributes new connections across server ports in proportion to a user-defined ratio. For example, if your array contained one new, high-speed server and two older servers, you could set the ratio so that the high-speed server receives twice as many connections as either of the two older servers.

To configure ratio, use this command:

```
config slb pool <poolname> lb-method ratio
```

RATIO WEIGHT

The ratio weight is the proportion of total connections that the node address should receive. The default ratio weight for a given node address is 1. If all node addresses use this default weight, the connections are distributed equally among the nodes. A ratio weight of 2 would result in twice as much traffic as a ratio weight of 1.

To configure a ratio weight, use this command:

```
config slb pool <poolname> add <ipaddress>:<L4Port> ratio {<ratio>}
```

LEAST CONNECTIONS

Least connections method is relatively simple in that the switch passes a new connection to the node having the least number of active sessions. The number of active sessions includes only those sessions occurring within the same VIP. Least connections works best in environments where the servers or other equipment you are load balancing have similar capabilities.

To configure least connections, use this command:

```
config slb pool <poolname> lb-method least-connections
```

PRIORITY

Priority mode is a variant of round-robin designed to provide redundant “standby” nodes within a pool. When you add a node to a pool, you can assign a priority level. Priority numbers range from 1 - 65535, with the highest number indicating the highest priority.

The switch will distribute traffic in round-robin fashion among the pools active nodes with the highest priority. If all nodes at that priority level go down or hit a session limit maximum, all new sessions will be directed to the nodes at the next lowest priority level. The switch continually monitors the status of the down nodes. As each node comes back up, the switch distributes traffic according to the priorities.

For example, with a pool that has six nodes divided evenly into two priority levels (2 and 1) all sessions will be evenly distributed via round-robin to the nodes at priority level 2. If one of the priority level 2 nodes goes down, all of the traffic will be assigned to the remaining level 2 nodes. If all of the priority level 2 nodes are down, all sessions will be directed to the priority level 1 nodes. If one of the level 2 nodes comes back up, all new sessions will be assigned to it.

BASIC SLB COMMANDS

[Table 17-2](#) describes basic SLB commands.

Table 17-2: Basic SLB Commands

Command	Description
clear slb connections [<vipname> <ipaddress>:{<L4Port>} all]	Clears the active connections.
config slb pool <poolname> add <ipaddress>:<L4Port> {ratio <ratio> priority <priority>}	Adds a physical server (node) to a server pool. When a new node is added, ping-check is automatically enabled.
config slb pool <poolname> delete <ipaddress>:<L4Port>	Deletes a physical server from a server pool.
config slb pool <poolname> lb-method [round-robin ratio priority least-connections]	Configures the SLB load balancing method.
config slb l4-port <L4Port> [treaper_timeout <seconds> udp-idle-timeout <seconds>]	Configures the inactive period for TCP or UDP before the connection is aged out.
config vlan <name> slb-type [server client both none]	Marks a VLAN to be either a server VLAN, or a client VLAN. If the server also originates connections to other servers, set the slb-type to both.

Table 17-2: Basic SLB Commands

Command	Description
create slb pool <poolname> {slb-method [round-robin ratio priority least-connections]}	Creates a server pool and optionally assigns a load-balancing method to the pool. The default load-balance method is round-robin. A pool represents a group of physical servers that is used to load-balance one or more VIPs.
create slb vip <vipname> pool <poolname> mode [transparent translation port-translation] <ipaddress> {- <upper_ipaddress>}:{<L4Port>}	Creates one or more new virtual IP addresses (VIPs) and attaches the VIP to a pool of physical servers. The server pool needs to be created before the VIP is created. If the <code>port</code> parameter is not specified, all requests to the VIP are forwarded to the server. If the <code>port</code> parameter is specified, only the specified TCP/UDP ports are allowed to reach the server. All other packets are dropped.
delete slb pool [<poolname> all]	Deletes a server pool.
delete slb vip [<vipname> all]	Deletes one or all VIPs.
disable slb	Disables SLB processing.
	Disabling SLB causes the following to occur:
	<ul style="list-style-type: none"> ■ Closes all connections. ■ Withdraws VIP routes or routes that do not respond with proxy ARP responses of VIP addresses. ■ Disconnects the switch from redundant SLB switches.
disable slb gogo-mode <master>	Disables GoGo mode processing.
disable slb node [<ipaddress>:<L4Port> all] {close-connections-now}	Disables one or more nodes from receiving new connection establishments. If <code>close-connections-now</code> is specified, all current open connections are torn down immediately.
disable slb l4-port [<L4Port> all]	Disables one or all L4 ports for SLB.
disable slb vip <ipaddress>:<L4Port>	Disables a single VIP port.

Table 17-2: Basic SLB Commands

Command	Description
disable slb vip <vipname> {close-connections-now}	Disables a VIP group. When disabled, no new connections are allowed to the real servers. If <code>close-connections-now</code> is specified, all existing connections are immediately closed. Otherwise, the existing connections are closed naturally, and are subject to connection reaping if idle for longer than the treaper-timeout configured on the SLB port.
enable slb	<p>Enables SLB processing on the switch, and activates the following functions for Transparent, Translational, and Port Translation modes:</p> <ul style="list-style-type: none"> ■ Exporting of VIP routes or proxy ARP for VIP addresses ■ Processing of VIP lookup and connection setup ■ Establishing communication with redundant SLB switches ■ Positively responding to MIB, 3DNS and SeelT requests <p>The default setting is disabled.</p>
enable slb gogo-mode <master> grouping <portlist>	Enables GoGo mode processing for a group of ports. There are no additional configuration commands for GoGo mode.
enable slb node [<ipaddress>:<L4Port> all]	Enables one or more nodes to receive data traffic. A node represents a physical server.
enable slb l4-port <L4Port>	Enables an L4 port to be used for SLB.
enable slb vip <ipaddress>:<L4Port>	Enables a single VIP port.
enable slb vip <vipname>	Enables a VIP group.
show slb	Displays the current SLB global configuration information, including:
	<ul style="list-style-type: none"> ■ Global enable/disable mode ■ Global modes ■ Default settings for health checker
show slb node [<ipaddress>:<L4Port> all] {detail}	Displays node-specific configuration information and status.

Table 17-2: Basic SLB Commands

Command	Description
show slb pool {detail}	Displays the current SLB pool configuration and statistics. If <code>detail</code> is not specified, the pool information is shown in a tabular format.
show slb pool <poolname> {detail}	Displays the configuration for the specified SLB pool.
show slb l4-port [<L4Port> all]	Displays the SLB configuration for one or all L4 ports.
show slb vip {detail}	Displays the current VIP configuration and statistics.
show slb vip <vipname> {detail}	Displays the configuration for the specified VIP.
unconfig slb all	Resets SLB global defaults and clears the SLB configuration.

ADVANCED SLB APPLICATION EXAMPLE

This example builds upon the Introductory SLB Application. The advanced concepts included in this example are the following:

- Multiple pools
- Multiple VIPs
- Multiple balancing algorithms
- Multiple types of health checking

[Figure 17-5](#) shows an example of an advanced SLB application.

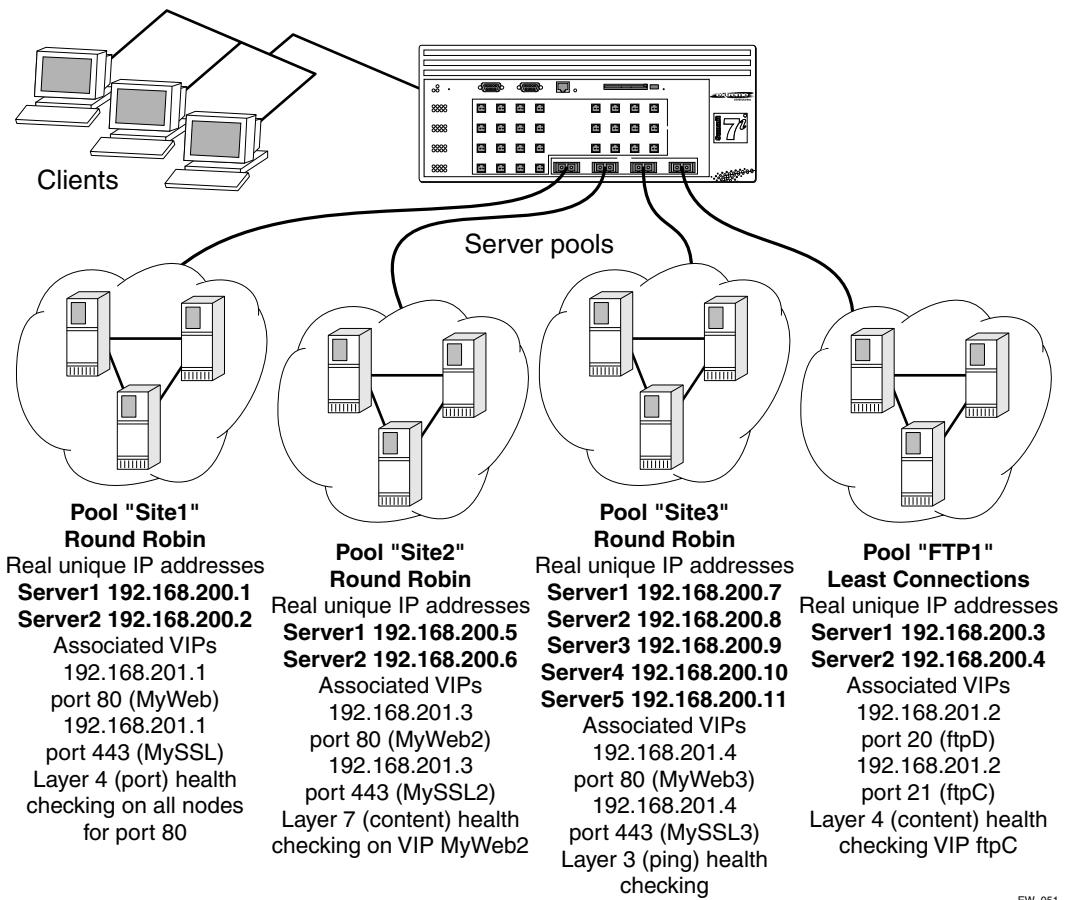


Figure 17-5: Advanced SLB configuration

The commands used to configure are described below.

The following commands create the VLAN from which outside connections will come.

```
create vlan outside
config vlan outside ipaddress 172.16.0.1 /16
config vlan outside add ports 1-8
```

To create is the virtual IP VLAN, use the following commands:

```
create vlan sites
config vlan sites ipaddress 192.168.201.254 /24
```

All VIPs will be configured to use this subnet. There are no ports associated with this VLAN.

The following commands create the VLAN *servers* and enable IP forwarding:

```
create vlan servers
config vlan servers ipaddress 192.168.200.254 /24
config vlan servers add ports 9-16
enable ipforwarding
```

The following series of commands creates a Web site. The site is defined as having 2 servers: 192.168.200.1 and 192.168.200.2, each with 2 services (HTTP and SSL). Two VIPs are then created to point at the appropriate pools. As a default, round-robin is used to load balance the services. Only one IP address is used for both VIPs; the difference is the port number. Finally, port checking is enabled to ensure fault tolerance on each of the servers.

```
create slb pool site1web
config slb site1 add 192.168.200.1:80
config slb site1 add 192.168.200.2:80
create slb pool site1ssl
config slb site1 add 192.168.200.1:443
config slb site1 add 192.168.200.2:443
create slb vip myweb pool site1web mode transparent 192.168.201.1:80
create slb vip myssl pool site1ssl mode transparent 192.168.201.1:443
enable slb node 192.168.200.1:80 tcp-port-check
enable slb node 192.168.200.2:80 tcp-port-check
enable slb node 192.168.200.1:443 tcp-port-check
enable slb node 192.168.200.2:443 tcp-port-check
```

The following series of commands creates a second Web site. This second site is similar to the first example; the difference is that content checking is enabled on this site. For this type of health checking, the server downloads a specified page (/testpage.htm) and looks for a specific string in the content ("test successful"). If it finds the string, the server considers the server up.

```

create slb pool site2web
config slb site2web add 192.168.200.5:80
config slb site2web add 192.168.200.6:80
create slb pool site2ssl
config slb site2ssl add 192.168.200.5:443
config slb site2ssl add 192.168.200.6:443
create slb vip myweb2 pool site2web mode transparent 192.168.201.3:80
create slb vip myssl2 pool site2ssl mode transparent 192.168.201.3:443
enable slb vip myweb2 service-check
config slb vip myweb2 service-check http url "/testpage.htm"
match-string "test successful"

```

The following series of commands creates a third Web site. This example creates 1 pool with a wildcard port specified. This means that the pool allows any port that is sent to it by the VIP. All five servers respond to requests on both port 80 and port 443.

```

create slb pool site3web
config slb site3web add 192.168.200.7:0
config slb site3web add 192.168.200.8:0
config slb site3web add 192.168.200.9:0
config slb site3web add 192.168.200.10:0
config slb site3web add 192.168.200.11:0
create slb vip myweb3 pool site3web mode transparent 192.168.201.4:80
create slb vip myssl3 pool site3web mode transparent 192.168.201.4:443

```

The following series of commands creates an FTP site. The site is defined as having two servers: 192.168.200.3 and 192.168.200.4. Only FTP is being serviced by the servers. The two different VIPs and port numbers refer to the control and data channels used by the FTP service. Two VIPs are then created to point at the appropriate pools.

As with the first site, the default load balancing method (round-robin) is used. Layer 7 health checking is used on the ftppc VIP. By using health checking, the switch logs in to the site as user *test* with the password *testpass*. If the login is successful, the server is labeled as “up” and is allowed to participate in load balancing. The account and password must be set up on all FTP servers.

```

create slb pool ftp1c
config slb ftp1c add 192.168.200.3:21
config slb ftp1c add 192.168.200.4:21
create slb pool ftp1d
config slb ftp1d add 192.168.200.3:20
config slb ftp1d add 192.168.200.4:20
create slb vip ftppc pool ftp1c mode transparent 192.168.201.2:21
create slb vip ftppd pool ftp1d mode transparent 192.168.201.2:20

```

```
enable slb vip ftpc service-check
config slb vip ftpc service-check ftp user test password testpass
```

Finally, enable SLB and configure the VLANs to be either client or server, using the following commands.

```
enable slb
config vlan outside slb-type client
config vlan servers slb-type server
```

HEALTH CHECKING

Three types of internal health checks are available:

- Ping-check
- Port-check
- Service-check

If any of the health checks enabled on a given node do not pass within the timeout specified, the node is considered down. When a node is down, no new connections will be established to that node until the node passes all configured health checks. If a health check fails and if the svcdown-reset parameter has been enabled on an associated VIP, existing connections for the VIP on this node will be closed by sending TCP Reset to the client and node.

In the command-line interface, the `show` commands for the pool and `vip show` individual node resources as up or down. New connections are only allowed if the VIP and node in question are both enabled and up. A node is assumed to be up unless health check is enabled and fails, in which case the node is marked down. A resource is also marked down if it has been disabled and the number of existing connections drops to zero. If a node is marked down for this reason, ping-checks and port-checks on this node are automatically stopped to conserve system resources, but resume if the node is enabled by the user.

The switch also supports external health checking. External health checking uses an external service configured by the user to perform health checks and uses SNMP as a mechanism to notify the switch of a server failure.

PING-CHECK

Ping-check is Layer 3 based pinging of the physical node. The default ping frequency is one ping generated to the node each 10 seconds. If the node does not respond to any ping within a timeout period of 30 seconds (3 ping intervals), then the node is considered down.

PING-CHECK COMMANDS

To enable ping-check, use this command:

```
enable slb node <ipaddress> ping-check
```

To disable ping-check, use this command:

```
disable slb node <ipaddress> ping-check
```

TCP-PORT-CHECK

TCP-port-check is Layer 4 based TCP port open/close testing of the physical node. The default frequency is 30 seconds and the default timeout is 90 seconds. Port-checking is useful when a node passes ping-checks, but a required TCP service (for example, httpd) has gone down. If the httpd daemon running on TCP port 80 crashed, that would cause a layer 4 port-check on port 80 to fail, because no TCP socket could be opened to that port. If this continues for the duration of the specified port-check timeout, the IP/port combination is considered down.

TCP-PORT-CHECK COMMANDS

To enable tcp-port-check, use this command:

```
enable slb node <ipaddress>:<L4Port> tcp-port-check
```

To disable tcp-port-check, use this command:

```
disable slb node <ipaddress>:{<L4Port> | all} tcp-port-check
```

SERVICE-CHECK

Service-check is Layer 7 based application-dependent checking defined on a VIP. Service-checking is performed on each node in the pool with which this VIP is associated. The default frequency is 60 seconds and the default timeout is 180 seconds. Each service check has associated parameters that you can set. These parameters are described in [Table 17-3](#).

Table 17-3: Service-Check Parameters

Service	Attribute	Global Default Value
HTTP	URL	“/”
	Match-string	Any-content
FTP	Userid	“anonymous”
	Password	“anonymous”
Telnet	Userid	“anonymous”
	Password	“anonymous”
SMTP	Dns-domain	Same as the switch DNS domain. If no DNS domain is configured for the switch, the value is “”.
NNTP	Newsgroup	“ebusiness”
POP3	Userid	“anonymous”
	Password	“anonymous”

If the service-check parameters are not specified on an individual node or VIP, the global default values for these parameters are used. The global service-check defaults themselves are configurable, so if you use the same value in many cases, change the global defaults accordingly.

In the case of HTTP service-checking, the URL of the Web page to be retrieved, such as “/index.html”, can be specified. A match-string that is expected to be in the retrieved Web page can be specified, such as “Welcome”. If the match-string is found in the first 1,000 bytes of the retrieved Web page, the service-check passes on the particular node. A match-string specified as keyword any-content will match any retrieved text. However, to distinguish valid data in the retrieved text from error text, specifying an actual string to match is suggested.

For FTP, Telnet, and POP3 service-check attempts to log on and off the application on the server using the specified userid and password.

For SMTP, service-check identifies the identity of the switch by providing the specified DNS domain. The SMTP server might not even use the specified DNS domain for authentication, only identification.

For NNTP, service-check queries the newsgroup specified.

Because service-checking is configured on a VIP basis, and multiple VIPs can use the same nodes, and you can run multiple service-checks against a particular node IP address and port number. It is possible for some of these service-checks to fail, while others pass. Therefore, when determining if a given node can accept a new connection for a VIP, the node must have passed the service-check configured for that VIP. When showing detailed VIP information, the status for individual nodes is shown with respect to that VIP.

SERVICE-CHECK COMMANDS

To enable service-check, use this command:

```
enable slb vip [<vipname> | all] service-check
```

To disable service-check, use this command:

```
disable slb vip [<vipname> | all] service-check
```

EXTERNAL HEALTH CHECKING

For server health checking that goes beyond the abilities of internal health checking, the switch also supports external health checking. The external health checking device sends the results of its check to the switch by way of SNMP MIB attributes. The specific MIB definitions for external health checking are available on the Extreme Networks Web site at: <http://www.extremenetworks.com/extreme/support/otherapps.htm>.

MAINTENANCE MODE

A node or VIP can be put into “maintenance mode” by simply disabling the node or VIP. In maintenance mode, existing connections remain active, but no new connections are permitted. The existing connections are either closed by the client and server, or are aged out if idle for more than 600 seconds.

PERSISTENCE

Using persistence, you can ensure that traffic flows do not span multiple servers. The switch supports two types of persistence:

- Client persistence
- Sticky persistence

CLIENT PERSISTENCE

Client persistence for a virtual server provides a persist mask feature. You can define a range of IP addresses that can be matched to a persistent connection. Any client whose source IP address falls within the range is considered a match for the given persistence entry.

To configure client persistence, use this command:

```
enable slb vip [<vipname> | all] client-persistence {timeout <seconds>}  
{mask <mask>}
```

STICKY PERSISTENCE

Sticky persistence provides a special type of persistence that is especially useful for cache servers. Similar to client persistence, sticky persistence keeps track of incoming clients' source and destination IP addresses. When a client is looking to make a repeat connection to a particular destination IP address, the switch directs the client to the same cache server or other transparent node that it previously used. Allowing clients to repeatedly use the same cache server can help you reduce the amount of content that might otherwise be duplicated on two or more cache servers in your network.

Sticky persistence provides the most benefits when you load balance caching proxy servers. A caching proxy server intercepts web requests and returns a cached web page if it is available. In order to improve the efficiency of the cache on these proxies, it is necessary to send similar requests to the same proxy server repeatedly. Sticky persistence can be used to cache a given web page on one proxy server instead of on every proxy server in an array. This saves the other proxies from having to duplicate the web page in their cache, wasting memory.

 *In order to prevent sticky entries from clumping on one server, use a static load balancing mode, such as round-robin.*

You can only activate sticky persistence on wildcard virtual servers.

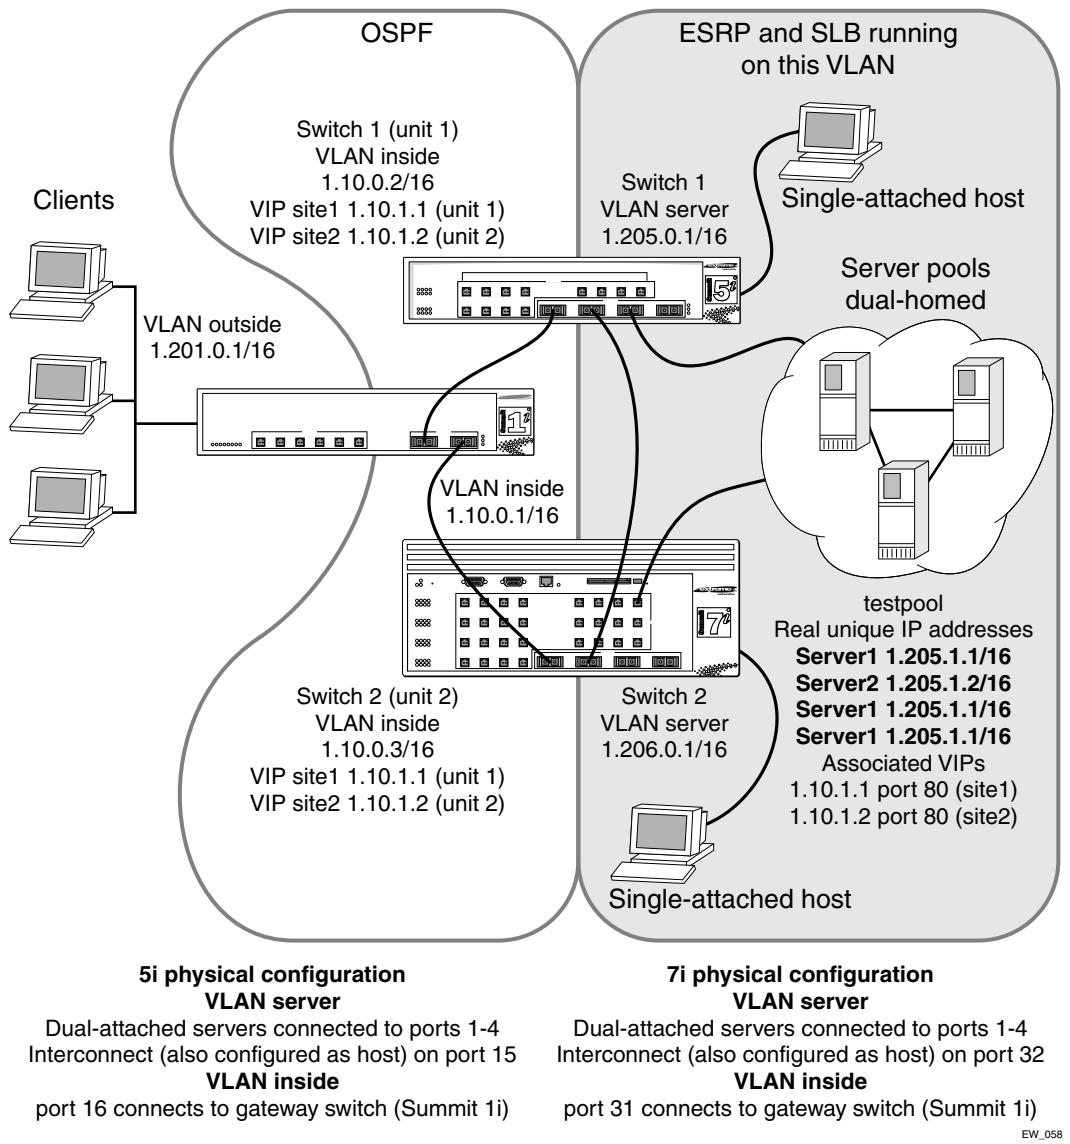
To configure sticky persistence, use this command:

```
enable slb vip [<vipname> | all] sticky-persistence {timeout <seconds>}
```

SERVER LOAD BALANCING WITH ESRP

By using ESRP, the SLB service is made redundant, along with the Layer 2 and Layer 3 services of the switch. This configuration allows single- or dual-attached servers to have redundant gateway services and very fast recovery from a fault. When ESRP is enabled, all servers can be online at the same time (as opposed to only the ones connected to the active switch in High Availability mode or having to introduce another interconnecting switch), and recovery from a switch failure occurs in less than 8 seconds.

[Figure 17-6](#) shows SLB enabled using ESRP and dual-homed servers.

**Figure 17-6:** SLB using ESRP and dual-homed servers

CONFIGURING THE SWITCHES FOR SLB AND ESRP

Two switches are used in the configuration shown in [Figure 17-6](#). The procedure used to configure the switches is described below.

Create the VLANs, using the following commands:

```
create vlan inside
create vlan server
```

Connect the gateway to the VLAN *inside*, using the following commands:

```
config inside ipaddress 1.10.0.2 /16
config inside add port 31
```

Configure the servers to connect to the VLAN *server* on ports 1-4, and configure port 32 to connect to the other ESRP switch, using the following commands:

```
configure server ipaddress 1.205.0.1 /16
configure server add port 1-4, 32
```

Enable ipforwarding, create a server pool called *testpool*, and add 4 servers to it using TCP port 80, using the following commands:

```
enable ipforwarding
create slb pool testpool
config slb pool testpool add 1.205.1.1:80
config slb pool testpool add 1.205.1.2:80
config slb pool testpool add 1.205.1.3:80
config slb pool testpool add 1.205.1.4:80
```

Create SLB VIP addresses for the two websites (*site1* and *site2*) and associate the server pool *testpool* with it, using the following commands:

```
create slb vip site1 pool testpool mode transparent 1.10.1.1:80
create slb vip site2 pool testpool mode transparent 1.10.1.2:80
```

Enable SLB and configure it for the appropriate VLANs (client connections enter from the VLAN *inside*), using the following commands:

```
enable slb
config inside slb client
config server slb server
```

Enable the routing protocol of choice (in this example, OSPF) and configure it appropriately, using the following commands:

```
enable ospf
```

Enable the ESRP protocol on the VLAN *server* and configure the ESRP direct-attached hosts mode to allow the proper failover of services, using the following commands:

```
enable esrp server
configure esrp port-mode host ports 1-4, 32
```

The interconnection between the switches is also configured as a host port.

Configure SLB to use the ESRP protocol, using the following command:

```
config slb esrp server add unit 1
```

NOTES REGARDING CONFIGURATION FOR SLB WITH ESRP

The following are things to note about the configurations for the switches running SLB and ESRP:

- All switch ports connected directly to the servers must be configured as ESRP host ports.
- The link between the two switches must be configured as an ESRP host port.
- The configuration uses transparent mode and HTTP services, but can be configured to support any of the currently supported load balancing protocols.
- Unlike the High Availability configuration, both switches are configured as unit 1.
- The SLB and ESRP configurations are identical on both switches, in relation to the ports being used.

WEB-SERVER CONFIGURATION

In [Figure 17-6](#), basic HTTP, configured at TCP port 80, is the only service being load balanced. The services must match those configured on the switch, for example HTTP services configured at TCP port 7080 on the switch require the servers to be able to allow connections at port 7080. You must ensure that the SLB connection is valid before trying to transfer the configuration to an ESRP/SLB configuration.

There are two main types of ESRP hosts that can be connected to the switches; single-attached hosts and dual-attached hosts. Single-attached hosts provide no server link redundancy, but allow hosts to be connected to the same VLAN as the web servers. Dual-attached hosts allow for redundant NICs in the servers, as well as connections to the switch. When configured as dual-attached hosts, the servers are supported fully by the ESRP redundant gateway services.



For information on specific NIC card configurations, please contact Extreme Networks Technical Support.

USING HIGH AVAILABILITY SYSTEM FEATURES

The switch supports several advanced redundant system features. Advanced redundant system features provide additional assurance that your content is available if a switch experiences a problem. The advanced redundant system options include:

- Redundant SLB
- Ping-check
- Active-active operation
- Manual fail-back
- SLB high availability

REDUNDANT SLB

The switch supports a failover process that uses a redundant configuration of two switches. If one switch fails, the second switch takes over the SLB duties of the first. By preparing a redundant switch for the possibility of failover, you effectively maintain your site's reliability and availability in advance.

The switches can be configured so that both perform SLB simultaneously. This type of operation is called *active-active*.

To configure failover, use the following commands:

```
config slb failover unit [1 | 2] remote-ip <ipaddress> local-ip
<ipaddress>:<L4Port> {alive-frequency <seconds> timeout <seconds>}
{dead-frequency <seconds>}

enable slb failover
```

The switches in a redundant SLB configuration should have identical SLB configurations except for the failover parameters. You can configure SLB on one switch, upload the configuration, edit it, and download it to the second switch to replicate the configuration.

USING PING-CHECK

Failover ping-check is used to determine if the currently active SLB server has the required network connectivity. If the specified IP address is unreachable for a specified duration, the ping-check triggers a failover to the redundant switch.

To configure ping-check, use the following commands:

```
config slb failover ping-check <ipaddress>
enable slb failover ping-check
```

 *The address being pinged should be that of a device other than the redundant SLB switch.*

CONFIGURING ACTIVE-ACTIVE OPERATION

Using active-active redundant SLB, you configure one switch as unit 1 and the other switch as unit 2. You then assign the VIPs either to unit 1 or to unit 2 (by default, a VIP is assigned to unit 1). When both switches are active, each switch performs SLB only for the VIPs assigned to it. If a switch fails, the other switch takes over the VIPs assigned to the failed switch.

The basic failover configure command assigns the switch's unit number:

```
config slb failover unit [1 | 2] remote-ip <ipaddress> local-ip
<ipaddress>:<L4Port> {alive-frequency <seconds> timeout <seconds>}
{dead-frequency <seconds>}
```

Where the following is true:

- **remote-ip** — Specifies the IP address of the redundant SLB switch.
- **local-ip** — Specifies the IP address of the switch you are configuring.

All VIPs with a given virtual IP address must be assigned to the same unit.

To assign a VIP to a unit, use the following command:

```
config slb vip <vipname> unit {1 | 2}
```

SAMPLE ACTIVE-ACTIVE CONFIGURATION

[Figure 17-7](#) shows an example of an active-active failover configuration.

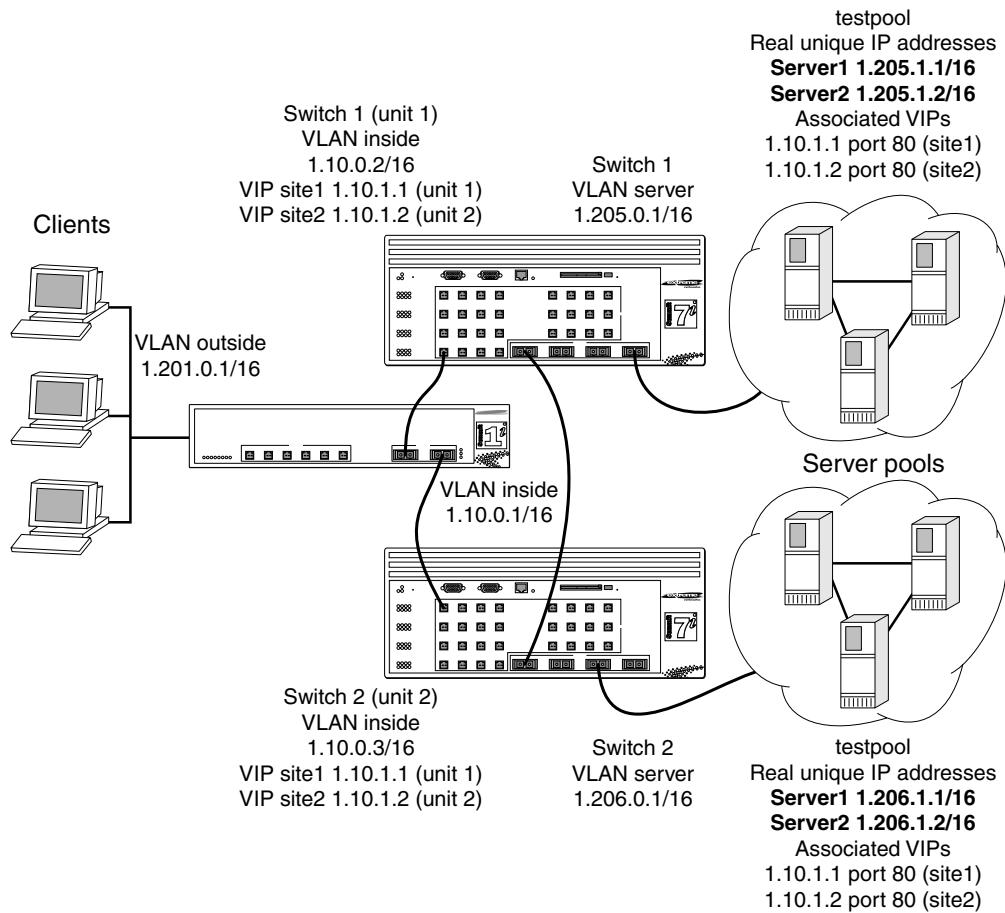


Figure 17-7: Active-active configuration

In this sample configuration, failover is enabled to ensure fault tolerance. To configure this example on the first switch, use the following commands:

EW_050

```
create vlan inside
create vlan server
config vlan inside ipaddress 1.10.0.2 /16
config vlan inside add port 31
config vlan server ipaddress 1.205.0.1 /16
config vlan server add port 29-30

enable ipforwarding

create slb pool testpool
config slb pool testpool add 1.205.1.1:80
config slb pool testpool add 1.205.1.2:80
create slb vip site1 pool testpool mode transparent 1.10.1.1:80
create slb vip site2 pool testpool mode transparent 1.10.1.2:80

enable slb
config vlan inside slb-type client
config vlan server slb-type server

config slb failover unit 1 remote 1.10.0.3 local 1.10.0.2:1028

enable slb failover

enable slb failover ping

config slb vip site1 unit 1
config slb vip site2 unit 2

config slb fail ping-check 1.10.0.1 freq 1
```

To configure this example on the second switch, use the following commands:

```
create vlan inside
create vlan server
config vlan inside ipaddress 1.10.0.3 /16
config vlan inside add port 31
config vlan server ipaddress 1.206.0.1 /16
config vlan server add port 29-30

enable ipforwarding
```

```
create slb pool testpool
config slb pool testpool add 1.206.1.1:80
config slb pool testpool add 1.206.1.2:80
create slb vip site1 pool testpool mode transparent 1.10.1.1:80
create slb vip site2 pool testpool mode transparent 1.10.1.2:80

enable slb
config vlan inside slb-type client
config vlan server slb-type server

config slb failover unit 2 remote 1.10.0.2 local 1.10.0.3:1028
enable slb failover
enable slb fail ping

config slb vip site1 unit 1
config slb vip site2 unit 2

config slb fail ping-check 1.10.0.1 freq 1
```

The differences between the configurations of these two switches are the IP addresses, and the designation of the first switch as the master of the active-active configuration.

USING MANUAL FAIL-BACK

In an active-active configuration, fail-back is the action of releasing the virtual servers that are assigned to a failed switch when that switch becomes operational again. By default, fail-back occurs automatically. If the minor disruption of fail-back makes automatic fail-back undesirable, you can enable manual fail-back. With manual fail-back, fail-back occurs only when the operator enters the fail-back command.

To enable manual fail-back, use the following command:

```
enable slb failover manual-failback
```

To execute a manual fail-back, use the following command:

```
configure slb failover failback-now
```

USING SLB HIGH AVAILABILITY

Using SLB High Availability (SLB H/A) provides redundancy in the case of an SLB service failure. Using SLB H/A, a site is configured with multiple servers spanning two switches. All servers are capable of responding to requests for content, but only those servers connected to the active switch receive requests. The other servers are idle or are used to serve another site.

Figure 17-8 shows an SLB failover configuration using SLB H/A.

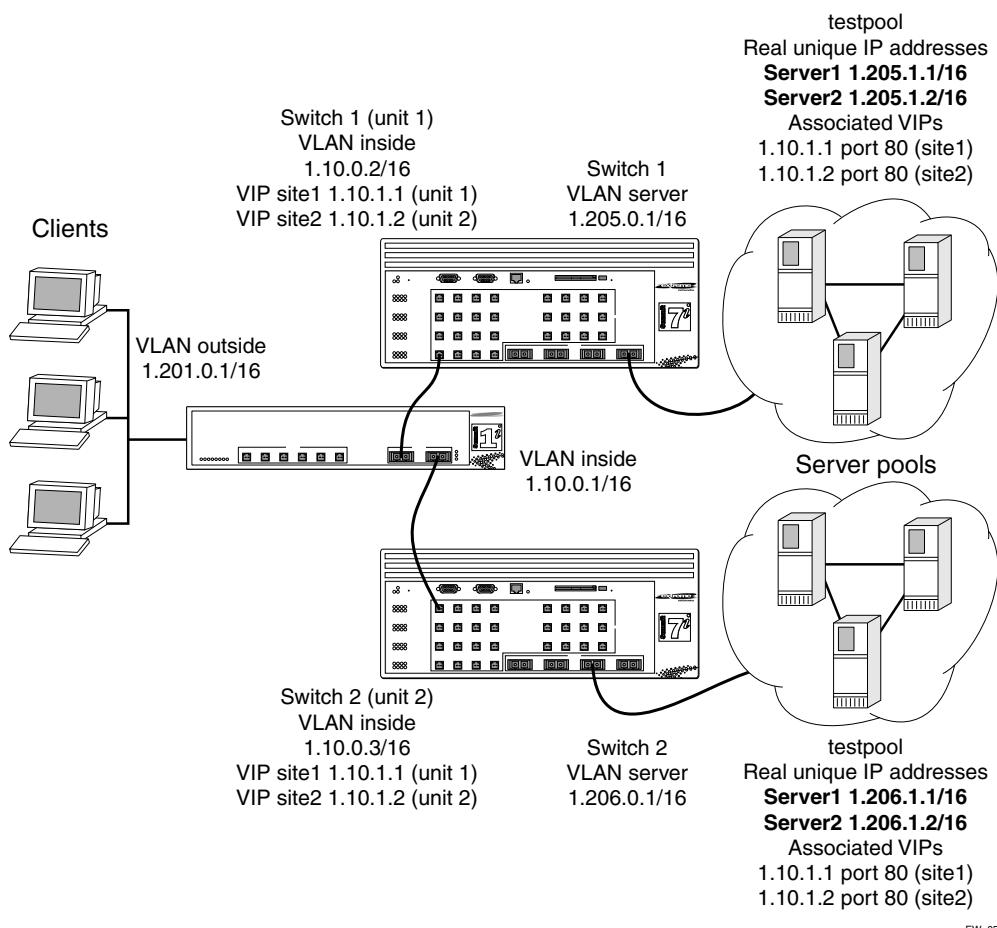


Figure 17-8: SLB failover configuration using SLB H/A

CONFIGURING CLIENTS

The configuration used to connect clients to SLB virtual sites with High Availability enabled is transparent to the accessing clients. As with normal SLB, the clients connect to the VIP believing that it is the physical address on a host server.

CONFIGURING SWITCHES FOR SLB H/A

Two switches are used in the configuration shown in [Figure 17-8](#). The procedure used to configure the switches is described below.

Create the VLANs, using the following commands:

```
create vlan inside
create vlan server
```

The VLAN *inside* connects to the gateway and the VLAN *server* contains all of the load balanced servers.

The gateway is connected to the VLAN *inside*, using the following commands:

```
config inside ipaddress 1.10.0.2 /16
config inside add port 31
```

Connect the servers the VLAN *server* on ports 29-30, using the following commands:

```
config server ipaddress 1.205.0.1 /16
config server add port 29-30
```

Enable IP forwarding, create a server pool called *testpool*, and add 2 servers to *testpool* using TCP port 80, using the following commands:

```
enable ipforwarding
create slb pool testpool
config slb pool testpool add 1.205.1.1:80
config slb pool testpool add 1.205.1.2:80
```



Two servers are connected to each High Availability switch.

Create SLB VIP addresses for the two websites (*site1* and *site2*) and associate the server pool *testpool* with them, using the following commands:

```
create slb vip site1 pool testpool mode transparent 1.10.1.1:80
```

```
create slb vip site2 pool testpool mode transparent 1.10.1.2:80
```

Enable SLB and configure it for the appropriate VLANs (client connections enter from the VLAN *inside*), using the following commands:

```
enable slb
config inside slb client
config server slb server
```

Configure SLB H/A for the switch, using the following command:

```
config slb failover unit 1 remote 1.10.0.3 local 1.10.0.2 14-port 1028
```

One switch in a High Availability pair is designated as unit 1 and the other is designated as unit 2. VIPs associated with the unit numbers are primarily serviced by the appropriate switch. The IP address of the remote switch in the failover pair is 1.10.0.3. The IP address of the local interface used by the High Availability protocol to communicate with the remote switch is 1.10.0.2. The Layer 4 port used by the High Availability protocol to exchange information is 1028.

In addition to performing normal status checking on the remote switch, the High Availability protocol pings the gateway to ensure that a connection to the client exists. If the connection to the gateway at IP address 1.10.0.1 fails, the remote switch services all of the connections. Configure status checking and enable failover using the following commands:

```
enable slb failover
config slb failover ping-check 1.10.0.1
enable slb failover ping
```

Configure the unit numbers on the two sites to determine which of the High Availability switches will actively serve the VIPs, using the following commands:

```
config slb vip site1 unit 1
config slb vip site2 unit 2
```

In this example, *site1* is serviced by the current switch and the remote switch (configured as unit 2) services *site2*. A switch configured as unit 1 services unit 2 VIPs only when the remote switch (configured as unit 2) fails.

NOTES REGARDING SLB/HA

The following are things to note about the configurations for SLB/HA:

- In the design shown in [Figure 17-8](#), only the servers directly connected to the switch that is actively servicing the VIP are used in the load balancing scheme. Without ESRP (discussed later), another switch interconnecting all the servers is necessary.
- One switch is designated as unit 1 and the other is unit 2. This designation determines which VIPs are active on each switch in the failover pair.
- In this configuration, *site1* is serviced by Switch 1 and has two servers that respond to client requests. *Site2* is serviced by the remote switch (Switch 2) and has two other servers that respond to client requests.
- If ping-check is enabled, it must not be directed at the remote switch. The remote switch is checked by the High Availability protocol. The ping-check works best when directed at a gateway to ensure that a path out of the network is available to the switch.
- The configuration uses transparent mode and HTTP services, but can be configured to support any of the currently supported load balancing protocols.
- The configurations for the High Availability switches are identical, with the exception of the failover command:

```
config slb failover unit 1 remote 1.10.0.3 local 1.10.0.2 14-port
1028
```

- The remote switch is set to unit 2, and the remote/local IP addresses are reversed to accurately describe the network, as shown in the following command:

```
config slb failover unit 2 remote 1.10.0.2 local 1.10.0.3 14-port
1028
```

WEB-SERVER CONFIGURATION

In the configuration shown in [Figure 17-8](#), basic HTTP, configured at TCP port 80, is the only service being load balanced. It is important the services match those configured on the switch, for example HTTP services configured at TCP port 7080 on the switch would require the servers to be able to allow connections at port 7080. You must also ensure that the SLB configuration is valid before enabling High Availability.

All 4 servers (2 local and 2 connected to the remote switch) should be identical in content, with the content for both *site1* and *site2* configured to be served.

This configuration uses transparent mode. Therefore, the VIPs need to be added to the servers as loopback addresses. This is done by configuring the network interfaces on the servers. A detailed description for doing this is provided later in this section.

3DNS SUPPORT

When you enable SLB, the switch reports health status to 3DNS using the iQuery™ protocol from F5 Networks®. The health status of the nodes within the server farm is based on L3, L4, L7, or external health checker mechanisms.

ADVANCED SLB COMMANDS

[Table 17-4](#) describes advanced SLB commands.

Table 17-4: Advanced SLB Commands

Command	Description
clear slb vip [<vipname> all] persistence	Resets all connection information in the persistence table. New connections opened are directed to a new server.
config slb 3dns-encryption-key {encrypted} <key>	Configures the encryption key for the iQuery protocol. The key can contain up to 100 characters.
config slb failover fallback-now	Configures the local SLB to release the remote SLB resources if the remote SLB is alive.
config slb failover ping-check ipaddress <ipaddress> {frequency <seconds>} timeout <seconds>	Configures the SLB device to actively determine if an external gateway is reachable by performing a ping. If the external gateway is not reachable, the VIPs failover to the remote SLB device. Specify the following: <ul style="list-style-type: none"> ■ ipaddress — The IP address of the external gateway. ■ frequency — The interval, in seconds, between pings sent to the remote gateway. The default setting is 1 second. ■ timeout — The amount of time before the local device declares the remote gateway is not reachable. The default setting is 3 seconds.

Table 17-4: Advanced SLB Commands

Command	Description
config slb failover unit [1 2] remote-ip <ipaddress> local-ip <ipaddress>:<L4Port> {alive-frequency <seconds> timeout <seconds>} {dead-frequency <seconds>}	<p>Configures the slb failover. Specify the following:</p> <ul style="list-style-type: none"> ■ remote-ip-address — The remote peer IP address. ■ local-ip-address — The address of a local IP interface used for the failover connection. ■ <L4Port> — The TCP port used for keep-alives between the failover peers. The default port is 1028. ■ unit — The unit number for this SLB device. The default unit number is 1. ■ alive-frequency — The frequency of how often the local SLB device polls the remote SLB device. The default setting is 1 second. ■ dead-frequency — The interval, in seconds, that the local switch attempts to re-establish communication with the remote (dead) switch. The default setting is 2 seconds. ■ timeout — The amount of time within which the switch must receive a response message from the remote switch. If no response is received in this period of time, the other switch is considered dead, and this switch takes over its VIPs. The default setting is 3 seconds.

Table 17-4: Advanced SLB Commands

Command	Description
config slb global [ping-check tcp-port-check service-check] frequency <seconds> timeout <seconds>	Configures default health checking frequency and timeout period. If the health check frequency and timeout are not specified for a specific node or VIP, the global values are used. Specify one of the following service checkers:
	<ul style="list-style-type: none"> ■ ping-check – L3-based pinging of the physical node. Default ping frequency is one ping generated to the node each 10 seconds. If the node does not respond any ping within a timeout period of 30 seconds (3 ping intervals), the node is considered inoperable. ■ tcp-port-check – L4-based TCP port open/close testing. Default values are 30 seconds for frequency and 90 seconds for timeout. ■ service-check – L7-based application-dependent checking. Default values are 60 seconds for frequency and 180 seconds for timeout.
config slb global ftp userid <userid> password {encrypted} {<password>}	Configures the default parameters for L7 service checking. If the password is not provided, you are prompted for the password twice.
config slb global http url <url_string> match-string [<match_string> any-content]	Configures the default parameters for L7 service checking.
config slb global nntp newsgroup <newsgroup>	Configures the default parameter for L7 service checking.
config slb global persistence-level [same-vip-same-port same-vip-any-port any-vip]	
config slb global persistence-method [per-packet per-session]	
config slb global pop3 userid <userid> password {encrypted} {<password>}	Configures the default parameter for L7 service checking.
config slb global smtp <dns_domain>	Configures the default parameter for L7 service checking.
config slb global synguard max-unacknowledge-SYNs <num_syns>	Configures the num_syns value that is used to trigger the SYN-guard feature.

Table 17-4: Advanced SLB Commands

Command	Description
config slb global telnet userid <userid> password {encrypted} {<password>}	Configures the default parameters for L7 service checking. If the password is not provided, you are prompted for the password twice.
config slb node <ipaddress>:{<L4Port>} max-connections <connections>	Configures the maximum number of simultaneous connections that can be established to a node. Use 0 to specify no limit. The default setting is 0.
config slb node <ipaddress>:{<L4Port>} tcp-port-check frequency <seconds> timeout <seconds>	Overrides the global default frequency and timeout values for this node. Use a value of 0 to restore the settings to the global default values.
config slb node <ipaddress> ping-check frequency <seconds> timeout <seconds>]	Overrides the global default frequency and timeout values for this node. Use a value of 0 to restore the settings to the global default values.
config slb vip <vipname> max-connections <connections>	Configures the maximum connections allowed to a particular VIP. A value of 0 indicates that no maximum is enforced. The default value is 0.
config slb vip <vipname> service-check frequency <seconds> timeout <seconds>	Configures the L7 service check frequency and timeout parameters for a particular VIP. To return to the global values, specify 0 for frequency and timeout.
config slb vip <vipname> service-check http {url <url> match-string [<match_string> any-content]}	Configures the service checker parameters on a per-VIP basis. Automatically enables the service checker. When the match-string option is specified, the string must be in the first 500 bytes of the returned Web page.
config slb vip <vipname> service-check ftp {userid <userid> password {encrypted} <password>}	
config slb vip <vipname> service-check telnet {userid <userid> password {encrypted} <password>}	
config slb vip <vipname> service-check smtp {<dns_domain>}	
config slb vip <vipname> service-check nntp <newsgroup>	
config slb vip <vipname> service-check pop3 userid <userid> password {encrypted} {<password>}	

Table 17-4: Advanced SLB Commands

Command	Description
config slb vip <vipname> unit {1 2}	Configures a unit number of a VIP name for active-active failover. The default unit number is 1.
disable slb 3dns iquery-client	Disables 3DNS support.
disable slb failover	Disables SLB failover.
disable slb failover manual-failback	Disables manual failback.
disable slb failover ping-check	Disables ping-check to an external gateway.
disable slb global synguard	Disables the TCP SYN-guard feature.
disable slb node <ipaddress>:{<L4Port> all} tcp-port-check	Disables L4 port checking.
disable slb node <ipaddress> ping-check	Disable L3 pinging.
disable slb vip [<vipname> all] client-persistence	Disables client-persistence.
disable slb vip [<vipname> all] service-check	Disables L7 service checking.
disable slb vip [<vipname> all] sticky-persistence	Disables sticky persistence.
disable slb vip [<vipname> all] svcdown-reset	Disables sysdown-reset.
enable slb 3dns iquery-client	Enables 3DNS support. The following 3DNS global balance modes are supported: completion, rate, global_availability, leastconn, null, packet_rate, random, ration, rr, and return_to_dns. The default setting is disabled.
enable slb failover	Enables the SLB failover mechanism. The default setting is disabled.
enable slb failover manual-failback	Enables manual failback.
enable slb failover ping-check	Enables ping-checking to an external gateway. The default setting is disabled.
enable slb global synguard	Enables the TCP SYN-guard feature. The SYN-guard feature minimizes the effect of the TCP-open type of denial-of-service attack by keeping track of all the half-open connections. When the number of half-open connections exceeds the num_syns value, the half-open connections are fast-aged out.

Table 17-4: Advanced SLB Commands

Command	Description
enable slb node <ipaddress> ping-check	Enables L3 pinging to the node address. Ping-check is automatically enabled when a node is added to a pool.
enable slb node <ipaddress>:<L4Port> tcp-port-check	Enables L4 port-check to the node address.
enable slb vip [<vipname> all] client-persistence {timeout <seconds>} {mask <mask>}	Enables client persistence and specifies the timeout and client address mask. In some circumstance, the client sets up multiple sessions to the virtual server and all the sessions needs to be connected to the same physical node.
	Enabling client persistence instructs the switch to forward new session requests from the same client (or clients from the same network using the <code>mask</code> argument) to the same node. The association between the client and physical node is removed after the specified timeout. The default setting is disabled.
enable slb vip [<vipname> all] service-check	Enables L7 service checking. The service checks performed are based on the following information:
	<ul style="list-style-type: none"> ■ If a service check is already configured, then it will use the user configured service-checking information. ■ If a service-check is not explicitly configured and a well known port is used when creating a VIP, then ExtremeWare will guess the application based on the well known port number and start the L7 service checker with the global default parameters.
enable slb vip [<vipname> all] sticky-persistence {timeout <seconds>}	Enables sticky persistence and specifies the timeout. Sticky persistence is usually used to load balance firewall and Web caches. When enabled, the switch forwards all traffic and new sessions toward a destination address (or address within certain subnet boundary specified by the <code>mask</code> argument) to the same physical node. The default setting is disabled.

Table 17-4: Advanced SLB Commands

Command	Description
enable slb vip [<vipname> all] svcdown-reset	Enables the svcdown-reset configuration. If enabled, the switch sends TCP RST to both the clients and the node, if the node associated with this VIP completely fails a ping-check, port-check, or service-check. Otherwise, the connections to the node are left as is, and are subject to connection reaping if idle for longer than the treaper-timeout configured on the SLB port. The default setting is disabled.
show slb 3dns members	Disables the current connection information between the switch and all the 3DNS queriers.
show slb failover	Disables the SLB failover configuration and status.
unconfig slb vip [<vipname> all] service-check	Disables and removes the service check configuration.

WEB CACHE REDIRECTION

Web cache redirection uses the TCP or UDP port number to redirect client requests to a target device (or group of devices). Web cache redirection transparently redirects traffic to web cache devices or to proxy servers and firewalls located in a demilitarized zone.

There are two ways to configure web cache redirection:

- Transparent mode SLB (described earlier in this chapter)
- Flow redirection

FLOW REDIRECTION

Flow redirection examines traffic and redirects it based on the following criteria:

- IP source address and mask
- IP destination address and mask
- Layer 4 port

FLOW REDIRECTION COMMANDS

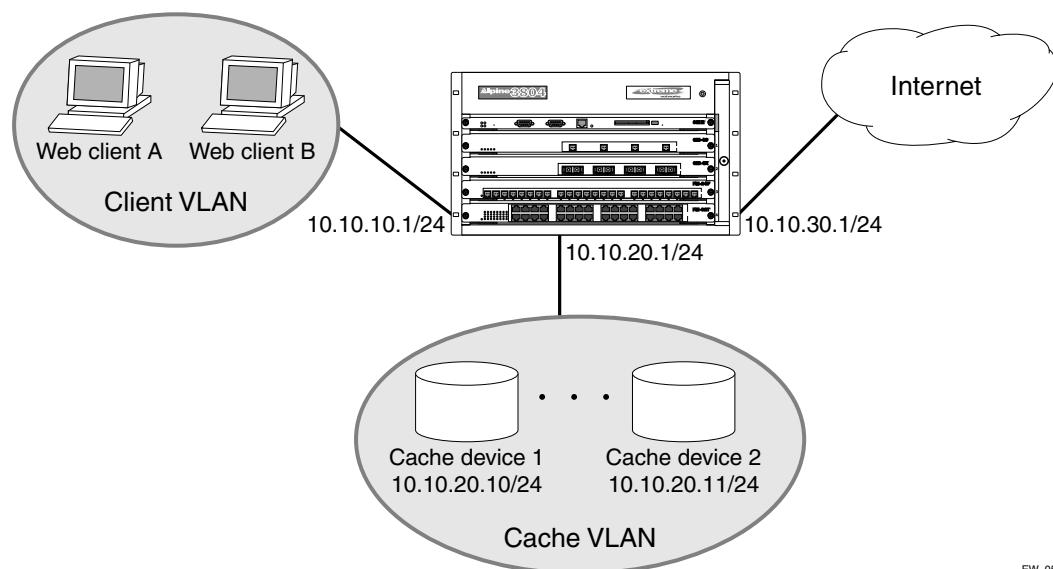
To configure flow redirection, use the commands listed in [Table 17-5](#).

Table 17-5: Flow Redirection Commands

Command	Description
config flow-redirection <flow_policy> add next-hop <ipaddress>	Adds the next hop host (gateway) that is to receive the packets that match the flow policy. By default, ping-based health checking is enabled.
config flow-redirection <flow_policy> delete next-hop <ipaddress>	Deletes the next hop host (gateway).
create flow-redirection <flow_policy> [tcp udp] destination {<ipaddress/mask> any} ip-port [<L4Port> any] source [<ipaddress/mask> any]	Creates a flow redirection policy.
delete flow-redirection <flow_policy>	Deletes a flow redirection policy.
show flow-redirection	Displays the current flow redirection configuration and statistics.

FLOW REDIRECTION EXAMPLE

[Figure 17-9](#) uses flow redirection to redirect Web traffic to Web cache servers. In this example, the clients and the cache devices are located on different networks. This is done by creating a different VLAN for the clients and cache devices.

**Figure 17-9: Flow-redirection example**

The following commands are used to configure the switch in this example:

```

create vlan client
config vlan client add port 1
config vlan client ipaddress 10.10.10.1/24

create vlan cache
config vlan cache add port 2
config vlan cache ipaddress 10.10.20.1/24

create vlan internet
config vlan internet add port 3
config vlan internet ipaddress 10.10.30.1/24

enable ipforwarding

create flow-redirection wcr tcp destination any ip-port 80 source any
config flow-redirection wcr add next-hop 10.10.20.10
config flow-redirection wcr add next-hop 10.10.20.11

```



18 Status Monitoring and Statistics

This chapter describes the following topics:

- [Status Monitoring on page 18-1](#)
- [Slot Diagnostics on page 18-3](#)
- [Port Statistics on page 18-4](#)
- [Port Errors on page 18-5](#)
- [Port Monitoring Display Keys on page 18-6](#)
- [Setting the System Recovery Level on page 18-7](#)
- [Logging on page 18-7](#)
- [RMON on page 18-12](#)

Viewing statistics on a regular basis allows you to see how well your network is performing. If you keep simple daily records, you will see trends emerging and notice problems arising before they cause major network faults. This way, statistics can help you get the best out of your network.

STATUS MONITORING

The status monitoring facility provides information about the switch. This information may be useful for your technical support representative if you have a problem. ExtremeWare includes many show commands that display information about different switch functions and facilities.

 For more information about show commands for a specific ExtremeWare feature, refer to the appropriate chapter in this guide.

[Table 18-1](#) describes show commands that are used to monitor the status of the switch.

Table 18-1: Status Monitoring Commands

Command	Description
show diag {<slot> msm-a msm-b}	Displays software diagnostics. For BlackDiamond switches, optionally specify a slot number of the MSM64i.
show log {<priority>}	Displays the current snapshot of the log. Options include: <ul style="list-style-type: none"> ■ priority — Filters the log to display message with the selected priority or higher (more critical). Priorities include critical, emergency, alert, error, warning, notice, info, and debug. If not specified, all messages are displayed.
show log config	Displays the log configuration, including the syslog host IP address, the priority level of messages being logged locally, and the priority level of messages being sent to the syslog host.
show memory {detail}	Displays the current system memory information. Specify the detail option to view task-specific memory usage.
show switch	Displays the current switch information, including: <ul style="list-style-type: none"> ■ sysName, sysLocation, sysContact ■ MAC address ■ Current time and time, system uptime, and time zone ■ Operating environment (temperature indication, fans, and power supply status) ■ NVRAM configuration information ■ MSM64i information (BlackDiamond switch only) ■ Scheduled reboot information ■ Software licensing information (Summit24, Summit48, Summit7i switches only)

Table 18-1: Status Monitoring Commands (continued)

Command	Description
show version	Displays the hardware and software versions currently running on the switch. Displays the switch serial number and version numbers of MSM64i and I/O modules (BlackDiamond switch).

SLOT DIAGNOSTICS

The BlackDiamond switch provides a facility for running normal or extended diagnostics on an I/O module or a Management Switch Fabric Module (MSM) without affecting the operation of the rest of the system.

If you select to run the diagnostic routine on an I/O module, that module is taken off-line while the diagnostic test is performed. Traffic to and from the ports on the module are temporarily unavailable. Once the diagnostic test is completed, the I/O module is reset and becomes operational again.

You can run normal or extended diagnostics on the slave MSM. The normal diagnostic routing is a short series of tests that do not test all the internal Application-Specific Integrated Circuit (ASIC) functions. The extended diagnostic routine tests coverage of all MSM components including the internal ASIC functions. The slave MSM is taken off-line while the diagnostic test is performed. It is reset and operational once the test is completed.

If you want to run the diagnostic routine on the master MSM, you must set it in diagnostic mode and then reboot the switch. When you reboot, the master MSM becomes the slave MSM (and vice versa) so that the diagnostic routing can be performed.

To run the diagnostic routine, use the command

```
run diag [normal | extended] [<slot> | msm-a | msm-b]
```

where the following is true:

- **normal** — Takes the switch fabric and ports offline, and performs a simple ASIC and packet loopback test on all ports. The test is completed in 30 seconds. CPU and out-of-band management ports are not tested in this mode. As a result, console and telnet access from the management port is available during this routine.

- **extended** — Takes the switch fabric and ports offline, and performs extensive ASIC, ASIC-memory, and packet loopback tests. Extended diagnostic tests take a maximum of 15 minutes. The CPU is not tested. Console access is available during extended diagnostics.
- **<slot>** — Specifies the slot number of an I/O module. Once the diagnostics test is complete, the system attempts to bring the I/O module back online. This parameter is applicable to the BlackDiamond switch, only.
- **msm-a | msm-b** — Specifies the slot letter of an MSM64i. If the master MSM is specified, the diagnostic routine is performed when the system reboots. Both switch fabric and management ports are taken offline during diagnostics. This parameter is applicable to the BlackDiamond switch, only.

PORt STATISTICS

ExtremeWare provides a facility for viewing port statistic information. The summary information lists values for the current counter against each port on each operational module in the system, and it is refreshed approximately every 2 seconds. Values are displayed to nine digits of accuracy.

To view port statistics, use the following command:

```
show ports <portlist> stats
```

The following port statistic information is collected by the switch:

- **Link Status** — The current status of the link. Options are
 - Ready (the port is ready to accept a link)
 - Active (the link is present at this port)
 - Chassis (the link is connected to a Summit Virtual Chassis)
- **Transmitted Packet Count (Tx Pkt Count)** — The number of packets that have been successfully transmitted by the port.
- **Transmitted Byte Count (Tx Byte Count)** — The total number of data bytes successfully transmitted by the port.
- **Received Packet Count (Rx Pkt Count)** — The total number of good packets that have been received by the port.
- **Received Byte Count (RX Byte Count)** — The total number of bytes that were received by the port, including bad or lost frames. This number includes bytes contained in the Frame Check Sequence (FCS), but excludes bytes in the preamble.

- **Received Broadcast (RX Bcast)** — The total number of frames received by the port that are addressed to a broadcast address.
- **Received Multicast (RX Mcast)** — The total number of frames received by the port that are addressed to a multicast address.

PORT ERRORS

The switch keeps track of errors for each port.

To view port transmit errors, use the following command:

```
show ports <portlist> txerrors
```

The following port transmit error information is collected by the system:

- **Port Number**
- **Link Status** — The current status of the link. Options are
 - Ready (the port is ready to accept a link)
 - Active (the link is present at this port)
- **Transmit Collisions (TX Coll)** — The total number of collisions seen by the port, regardless of whether a device connected to the port participated in any of the collisions.
- **Transmit Late Collisions (TX Late Coll)** — The total number of collisions that have occurred after the port's transmit window has expired.
- **Transmit Deferred Frames (TX Deferred)** — The total number of frames that were transmitted by the port after the first transmission attempt was deferred by other network traffic.
- **Transmit Errorred Frames (TX Error)** — The total number of frames that were not completely transmitted by the port because of network errors (such as late collisions or excessive collisions).
- **Transmit Parity Frames (TX Parity)** — The bit summation has a parity mismatch.

To view port receive errors, use the following command:

```
show ports <portlist> rxerrors
```

The following port receive error information is collected by the switch:

- **Receive Bad CRC Frames (RX CRC)** — The total number of frames received by the port that were of the correct length, but contained a bad FCS value.
- **Receive Oversize Frames (RX Over)** — The total number of good frames received by the port greater than the supported maximum length of 1,522 bytes. For products that use the “*i*” chipset, ports with jumbo frames enabled do no increment this counter.
- **Receive Undersize Frames (RX Under)** — The total number of frames received by the port that were less than 64 bytes long.
- **Receive Fragmented Frames (RX Frag)** — The total number of frames received by the port were of incorrect length and contained a bad FCS value.
- **Receive Jabber Frames (RX Jab)** — The total number of frames received by the port that was of greater than the support maximum length and had a Cyclic Redundancy Check (CRC) error.
- **Receive Alignment Errors (RX Align)** — The total number of frames received by the port that occurs if a frame has a CRC error and does not contain an integral number of octets.
- **Receive Frames Lost (RX Lost)** — The total number of frames received by the port that were lost because of buffer overflow in the switch.

PORt MONITORING DISPLAY KEYS

[Table 18-2](#) describes the keys used to control the displays that appear when you issue any of the `show port` commands.

Table 18-2: Port Monitoring Display Keys

Key(s)	Description
U	Displays the previous page of ports.
D	Displays the next page of ports.
[Esc] or [Return]	Exits from the screen.
0	Clears all counters.

Table 18-2: Port Monitoring Display Keys (continued)

Key(s)	Description
[Space]	Cycles through the following screens: <ul style="list-style-type: none"> ■ Packets per second ■ Bytes per second ■ Percentage of bandwidth Available using the <code>show port utilization</code> command only.

SETTING THE SYSTEM RECOVERY LEVEL

You can configure the system to automatically reboot after a software task exception, using the following command:

```
config sys-recovery-level [none | critical | all]
```

Where the following is true:

- `none` — Configures the level to no recovery.
- `critical` — Configures ExtremeWare to log an error into the syslog and automatically reboot the system after a critical task exception.
- `all` — Configures ExtremeWare to log an error into the syslog and automatically reboot the system after any task exception.

The default setting is `none`.

LOGGING

The switch log tracks all configuration and fault information pertaining to the device. Each entry in the log contains the following information:

- **Timestamp** — The timestamp records the month and day of the event, along with the time (hours, minutes, and seconds) in the form HH:MM:SS. If the event was caused by a user, the user name is also provided.
- **Fault level** — [Table 18-3](#) describes the three levels of importance that the system can assign to a fault.

Table 18-3: Fault Levels Assigned by the Switch

Level	Description
Critical	A desired switch function is inoperable. The switch may need to be reset.
Warning	A noncritical error that may lead to a function failure.
Informational	Actions and events that are consistent with expected behavior.
Debug	Information that is useful when performing detailed troubleshooting procedures.

By default, log entries that are assigned a critical or warning level remain in the log after a switch reboot. Issuing a clear log command does not remove these static entries. To remove log entries of all levels (including warning or critical), use the following command:

```
clear log static
```

- **Subsystem** — The subsystem refers to the specific functional area to which the error refers. [Table 18-4](#) describes the subsystems.

Table 18-4: Fault Log Subsystems

Subsystem	Description
Syst	General system-related information. Examples include memory, power supply, security violations, fan failure, overheat condition, and configuration mode.
STP	STP information. Examples include an STP state change.
Brdg	Bridge-related functionality. Examples include low table space and queue overflow.
SNMP	SNMP information. Examples include community string violations.
Telnet	Information related to Telnet login and configuration performed by way of a Telnet session.
VLAN	VLAN-related configuration information.
Port	Port management-related configuration. Examples include port statistics and errors.

- **Message** — The message contains the log information with text that is specific to the problem.

LOCAL LOGGING

The switch maintains 1,000 messages in its internal log. You can display a snapshot of the log at any time by using the command

```
show log {<priority>}
```

where the following is true:

- **priority** — Filters the log to display message with the selected priority or higher (more critical). Priorities include (in order) critical, emergency, alert, error, warning, notice, info, and debug. If not specified, all messages are displayed.

REAL-TIME DISPLAY

In addition to viewing a snapshot of the log, you can configure the system to maintain a running real-time display of log messages on the console. To turn on the log display, enter the following command:

```
enable log display
```

To configure the log display, use the following command:

```
config log display {<priority>}
```

If **priority** is not specified, only messages of critical priority are displayed.

If you enable the log display on a terminal connected to the console port, your settings will remain in effect even after your console session is ended (unless you explicitly disable the log display).

When using a Telnet connection, if your Telnet session is disconnected (because of the inactivity timer, or for other reasons), the log display is automatically halted. You must restart the log display by using the `enable log display` command.

REMOTE LOGGING

In addition to maintaining an internal log, the switch supports remote logging by way of the UNIX syslog host facility. To enable remote logging, do the following:

- Configure the syslog host to accept and log messages.
- Enable remote logging by using the following command:

```
enable syslog
```

- Configure remote logging by using the following command:

```
config syslog {add} <ipaddress> <facility> {<priority>}
```

Specify the following:

- ipaddress — The IP address of the syslog host.
- facility — The syslog facility level for local use. Options include local0 through local7.
- priority — Filters the log to display message with the selected priority or higher (more critical). Priorities include (in order) critical, emergency, alert, error, warning, notice, info, and debug. If not specified, only critical priority messages are sent to the syslog host.



Refer to your UNIX documentation for more information about the syslog host facility.

LOGGING CONFIGURATION CHANGES

ExtremeWare allows you to record all configuration changes and their sources that are made using the CLI by way of Telnet or the local console. The changes are logged to the system log. Each log entry includes the user account name that performed the change and the source IP address of the client (if Telnet was used). Configuration logging applies only to commands that result in a configuration change. To enable configuration logging, use the following command:

```
enable cli-config-logging
```

To disable configuration logging, use the following command:

```
disable cli-config-logging
```

CLI configuration logging is enabled by default.

LOGGING COMMANDS

The commands described in [Table 18-5](#) allow you to configure logging options, reset logging options, display the log, and clear the log.

Table 18-5: Logging Commands

Command	Description
clear counters	Clears all switch statistics and port counters.
clear log {static}	Clears the log. If <code>static</code> is specified, the critical log messages are also cleared.
config log display {<priority>}	Configures the real-time log display. Options include: <ul style="list-style-type: none"> ■ <code>priority</code> — Filters the log to display messages with the selected priority or higher (more critical). Priorities include critical, emergency, error, alert, warning, notice, info, and debug. If not specified, informational priority messages and higher are displayed.
config syslog {add} <ip_address> <facility> {<priority>}	Configures the syslog host address and filters messages sent to the syslog host. Up to 4 syslog servers can be configured. Options include: <ul style="list-style-type: none"> ■ <code>ipaddress</code> — The IP address of the syslog host. ■ <code>facility</code> — The syslog facility level for local use (local0 - local7). ■ <code>priority</code> — Filters the log to display messages with the selected priority or higher (more critical). Priorities include critical, emergency, alert, error, warning, notice, info, and debug. If not specified, only critical priority messages and are sent to the syslog host.
config syslog delete <ip_address>	Deletes a syslog host address.
disable cli-config-logging	Disables configuration logging.
disable log display	Disables the log display.
disable syslog	Disables logging to a remote syslog host.
enable cli-config-logging	Enables the logging of CLI configuration commands to the Syslog for auditing purposes. The default setting is enabled.

Table 18-5: Logging Commands (continued)

Command	Description
enable log display	Enables the log display.
enable syslog	Enables logging to a remote syslog host.
show log {<priority>}	Displays the current snapshot of the log. Options include: <ul style="list-style-type: none"> ■ priority — Filters the log to display message with the selected priority or higher (more critical). Priorities include critical, emergency, alert, error, warning, notice, info, and debug. If not specified, all messages are displayed.
show log config	Displays the log configuration, including the syslog host IP address, the priority level of messages being logged locally, and the priority level of messages being sent to the syslog host.

RMON

Using the Remote Monitoring (RMON) capabilities of the switch allows network administrators to improve system efficiency and reduce the load on the network.

The following sections explain more about the RMON concept and the RMON features supported by the switch.



You can only use the RMON features of the system if you have an RMON management application, and have enabled RMON on the switch.

ABOUT RMON

RMON is the common abbreviation for the Remote Monitoring Management Information Base (MIB) system defined by the Internet Engineering Task Force (IETF) documents RFC 1271 and RFC 1757, which allows you to monitor LANs remotely.

A typical RMON setup consists of the following two components:

- **RMON probe** — An intelligent, remotely controlled device or software agent that continually collects statistics about a LAN segment or VLAN. The probe transfers the information to a management workstation on request, or when a predefined threshold is crossed.

- **Management workstation** — Communicates with the RMON probe and collects the statistics from it. The workstation does not have to be on the same network as the probe, and can manage the probe by in-band or out-of-band connections.

RMON FEATURES OF THE SWITCH

The IETF defines nine groups of Ethernet RMON statistics. The switch supports the following four of these groups:

- Statistics
- History
- Alarms
- Events

This section describes these groups, and discusses how they can be used.

STATISTICS

The RMON Ethernet Statistics group provides traffic and error statistics showing packets, bytes, broadcasts, multicasts, and errors on a LAN segment or VLAN.

Information from the Statistics group is used to detect changes in traffic and error patterns in critical areas of the network.

HISTORY

The History group provides historical views of network performance by taking periodic samples of the counters supplied by the Statistics group. The group features user-defined sample intervals and bucket counters for complete customization of trend analysis.

The group is useful for analysis of traffic patterns and trends on a LAN segment or VLAN, and to establish baseline information indicating normal operating parameters.

ALARMS

The Alarms group provides a versatile, general mechanism for setting threshold and sampling intervals to generate events on any RMON variable. Both rising and falling thresholds are supported, and thresholds can be on the absolute value of a variable or its delta value. In addition, alarm thresholds may be autocalibrated or set manually.

Alarms inform you of a network performance problem and can trigger automated action responses through the Events group.

EVENTS

The Events group creates entries in an event log and/or sends SNMP traps to the management workstation. An event is triggered by an RMON alarm. The action taken can be configured to ignore it, to log the event, to send an SNMP trap to the receivers listed in the trap receiver table, or to both log and send a trap. The RMON traps are defined in RFC 1757 for rising and falling thresholds.

Effective use of the Events group saves you time. Rather than having to watch real-time graphs for important occurrences, you can depend on the Event group for notification. Through the SNMP traps, events can trigger other actions, providing a mechanism for an automated response to certain occurrences.

CONFIGURING RMON

RMON requires one probe per LAN segment, and standalone RMON probes have traditionally been expensive. Therefore, Extreme's approach has been to build an inexpensive RMON probe into the agent of each system. This allows RMON to be widely deployed around the network without costing more than traditional network management. The switch accurately maintains RMON statistics at the maximum line rate of all of its ports.

For example, statistics can be related to individual ports. Also, because a probe must be able to see all traffic, a stand-alone probe must be attached to a nonsecure port. Implementing RMON in the switch means that all ports can have security features enabled.

To enable or disable the collection of RMON statistics on the switch, use the following command:

```
[enable | disable] rmon
```

By default, RMON is disabled. However, even in the disabled state, the switch responds to RMON queries and sets for alarms and events. By enabling RMON, the switch begins the processes necessary for collecting switch statistics.

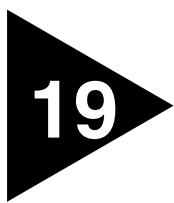
EVENT ACTIONS

The actions that you can define for each alarm are shown in [Table 18-6](#).

Table 18-6: Event Actions

Action	High Threshold
No action	
Notify only	Send trap to all trap receivers.
Notify and log	Send trap; place entry in RMON log.

To be notified of events using SNMP traps, you must configure one or more trap receivers, as described in [Chapter 3](#).



Using ExtremeWare Vista

This chapter covers the following topics:

- [Enabling and Disabling Web Access on page 19-2](#)
- [Setting Up Your Browser on page 19-2](#)
- [Accessing ExtremeWare Vista on page 19-3](#)
- [Navigating ExtremeWare Vista on page 19-4](#)
- [Saving Changes on page 19-6](#)
- [Filtering Information on page 19-6](#)
- [Do a GET When Configuring a VLAN on page 19-7](#)
- [Sending Screen Output to Extreme Networks on page 19-7](#)

ExtremeWare Vista is device-management software running in the switch that allows you to access the switch over a TCP/IP network, using a standard Web browser. Any properly configured standard Web browser that supports frames and JavaScript (such as Netscape Navigator 3.0 or higher, or Microsoft Internet Explorer 3.0 or higher) can be used to manage the system.

ExtremeWare Vista provides a subset of the command-line interface (CLI) commands available for configuring and monitoring the switch. If a particular command is not available using ExtremeWare Vista, you must use the CLI to access the desired functionality.

ENABLING AND DISABLING WEB ACCESS

By default, Web access is enabled on the switch. Use of ExtremeWare Vista Web access can be restricted through the use of an access profile. An access profile permits or denies a named list of IP addresses and subnet masks. To configure Vista Web access to use an access profile, use the following command:

```
enable web access-profile [<access-profile> | none] {port  
<tcp_port_number>}
```

Use the none option to remove a previously configured access profile.

To display the status of Web access, use the following command:

```
show management
```

To disable ExtremeWare Vista, use the following command:

```
disable web
```

To re-enable Web access, use the following command:

```
enable web {access-profile [<access-profile> | none]} {port  
<tcp_port_number>}
```

You will need to reboot the system in order for these changes to take effect.



For more information on rebooting, refer to [Chapter 20](#).

To use ExtremeWare Vista, at least one VLAN must be assigned an IP address.

SETTING UP YOUR BROWSER

In general, the default settings that come configured on your browser work well with ExtremeWare Vista. The following are recommended settings that you can use to improve the display features and functionality of ExtremeWare Vista:

- After downloading a newer version of the switch image, clear the browser disk and memory cache to see the updated menu screens. You must clear the cache while at the main ExtremeWare Vista Logon screen, so that all underlying .GIF files are updated.

- Check for newer versions of stored pages. Every visit to the page should be selected as a cache setting.

If you are using Netscape Navigator, configure the cache option to check for changes “Every Time” you request a page.

If you are using Microsoft Internet Explorer, configure the Temporary Internet Files setting to check for newer versions of stored pages by selecting “Every visit to the page.”

- Images must be auto-loaded.
- Use a high-resolution monitor to maximize the amount of information displayed in the content frame. The recommended resolution is 1024 x 768 pixels. You can also use 800 x 600 pixels.
- Turn off one or more of the browser toolbars to maximize the viewing space of the ExtremeWare Vista content screen.
- If you will be using ExtremeWare Vista to send an e-mail to the Extreme Networks Technical Support department, configure the e-mail settings in your browser.
- Configure the browser to use the following recommended fonts:
 - Proportional font—Times New Roman
 - Fixed-width font—Courier New

ACCESSING EXTREMEWARE VISTA

To access the default home page of the switch, enter the following URL in your browser:

`http://<ip_address>`

When you access the home page of the system, you are presented with the Login screen. Enter your user name and password in the appropriate fields, and click OK.

If you have entered the name and password of an administrator-level account, you have access to all ExtremeWare Vista pages. If you have used a user-level account name and password, you only have access to the Statistics and Support information.

If multiple people access the same switch using ExtremeWare Vista, you might see the following error message:

Web:server busy

To correct this situation, log out of the switch and log in again.

NAVIGATING EXTREMEWARE VISTA

After logging in to the switch, the ExtremeWare Vista home page is displayed.

ExtremeWare Vista divides the browser screen into the following sections:

- Task frame
- Content frame
- Standalone buttons

TASK FRAME

The task frame has two sections: menu buttons and submenu links. There are four task menu buttons, as follows:

- Configuration
- Statistics
- Support
- Logout

Below the task buttons are options. Options are specific to the task button that you select. When you select an option, the information displayed in the content frame changes. However, when you select a new task button, the content frame does not change until you select a new option.



Submitting a configuration page with no change will result in an asterisk () appearing at the CLI prompt, even though actual configuration values have not changed.*

CONTENT FRAME

The content frame contains the main body of information in ExtremeWare Vista. For example, if you select an option from the Configuration task button, enter configuration parameters in the content frame. If you select the Statistics task button, statistics are displayed in the content frame.

BROWSER CONTROLS

Browser controls include drop-down list boxes, check boxes, and multi-select list boxes. A multi-select list box has a scrollbar on the right side of the box. Using a multi-select list box, you can select a single item, all items, a set of contiguous items, or multiple non-contiguous items. [Table 19-1](#) describes how to make selections from a multi-select list box.

Table 19-1: Multi-Select List Box Key Definitions

Selection Type	Key Sequence
Single item	Click the item using the mouse.
All items	Click the first item, and drag to the last item.
Contiguous items	Click the first desired item, and drag to the last desired item.
Selected non-contiguous items	Hold down [Ctrl], click the first desired item, click the next desired item, and so on.

STATUS MESSAGES

Status messages are displayed at the top of the content frame. There are four types of status messages, as follows:

- **Information** — Displays information that is useful to know prior to, or as a result of, changing configuration options.
- **Warning** — Displays warnings about the switch configuration.
- **Error** — Displays errors caused by incorrectly configured settings.
- **Success** — Displays informational messages after you click Submit. The message displayed reads, “Request was submitted successfully.”

STANDALONE BUTTONS

At the bottom of some of the content frames is a section that contains standalone buttons. Standalone buttons are used to perform tasks that are not associated with a particular configuration option. An example of this is the Reboot Switch button.

SAVING CHANGES

There are two ways to save your changes to non-volatile storage using ExtremeWare Vista:

- Select Save Configuration from the Configuration task button, Switch option.

This field contains a drop-down list box that allows you to select either the primary or secondary configuration area. After you select the configuration area, click Submit to save the changes.

- Click the Logout button.

If you attempt to log out without saving your changes, ExtremeWare Vista prompts you to save your changes.

If you select Yes, the changes are saved to the selected configuration area. To change the selected configuration area, you must go to the Configuration task button, Switch option.

FILTERING INFORMATION

Some pages have a Filter button. The Filter button is used to display a subset of information on a given page. For example, on the OSPF configuration page, you can configure authentication based on the VLAN, area identifier, or virtual link. Once you select a filtering option and click the Filter button, the form that provides the configuration options displays the available interfaces in the drop-down menu, based on your filtering selection.

Similarly, in certain Configuration and Statistics pages, information is shown based on a particular slot.

Because the BlackDiamond switch allows you to preconfigure modules without having them physically available in the chassis, the configuration pages offer a drop-down menu to select any module card that has been configured on the system, whether or not the module is physically available. By default, information for the first configured module that is found in the chassis is displayed on the page. You can configure available slots and ports by filtering on a selected module from the Sort by Slot drop-down menu.

On the Statistics pages, you can only view information for cards that are configured and physically inserted into the BlackDiamond chassis. On these pages, the Sort by Slot drop-down menu displays only these modules.

DO A GET WHEN CONFIGURING A VLAN

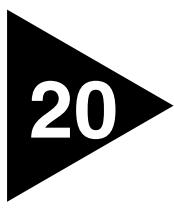
When configuring a VLAN using ExtremeWare Vista, prior to editing the VLAN configuration, you must first click the get button to ensure that subsequent edits are applied to the correct VLAN. If you do not click the get button and you submit the changes, the changes will be made to the VLAN that was previously displayed.

If you configure a VLAN and then delete it, the *default* VLAN is shown in the VLAN name window, but the VLAN information contained in the lower portion of the page is not updated. Click the get button to update the display.

SENDING SCREEN OUTPUT TO EXTREME NETWORKS

If Extreme Networks requests that you e-mail the output of a particular ExtremeWare Vista screen, do the following:

- 1 Click on the content frame of the screen that you must send.
- 2 From Netscape Navigator, select Save Frame As from the File menu, and enter a name for the file.
- 3 From Microsoft Internet Explorer 3.0, select Save As File from the File menu, and enter a name for the file.
- 4 From Microsoft Internet Explorer 4.0, right-click in the content frame, select View Source, and save the HTML text by copying it and pasting it into a text editor.
- 5 Attach the file to the e-mail message that you are sending to Extreme Networks.



20 Software Upgrade and Boot Options

This chapter describes the following topics:

- [Downloading a New Image on page 20-1](#)
- [Saving Configuration Changes on page 20-3](#)
- [Using TFTP to Upload the Configuration on page 20-4](#)
- [Using TFTP to Download the Configuration on page 20-5](#)
- [Synchronizing MSMs on page 20-7](#)
- [Upgrading and Accessing BootROM on page 20-7](#)
- [Boot Option Commands on page 20-8](#)

DOWNLOADING A NEW IMAGE

The image file contains the executable code that runs on the switch. It comes preinstalled from the factory. As new versions of the image are released, you should upgrade the software running on your system.

The image is upgraded by using a download procedure from either a Trivial File Transfer Protocol (TFTP) server on the network or from a PC connected to the serial port using the XMODEM protocol. Downloading a new image involves the following steps:

- Load the new image onto a TFTP server on your network (if you will be using TFTP).
- Load the new image onto a PC (if you will be using XMODEM).

- Download the new image to the switch using the command

```
download image [<ipaddress> | <hostname>] <filename> {primary | secondary}
```

where the following is true:

`ipaddress` — Is the IP address of the TFTP server.

`hostname` — Is the hostname of the TFTP server. (You must enable DNS to use this option.)

`filename` — Is the filename of the new image.

`primary` — Indicates the primary image.

`secondary` — Indicates the secondary image.

The switch can store up to two images; a primary and a secondary. When you download a new image, you must select into which image space (primary or secondary) the new image should be placed. If not indicated, the primary image space is used.

If two MSMs are installed in the BlackDiamond switch, the downloaded image is saved to the same location on each one.

You can select which image the switch will load on the next reboot by using the following command:

```
use image [primary | secondary]
```

REBOOTING THE SWITCH

To reboot the switch, use the following command:

```
reboot { time <date> <time> | cancel}
```

where `date` is the date and `time` is the time (using a 24-hour clock format) when the switch will be rebooted. The values use the following format:

`mm/dd/yyyy hh:mm:ss`

If you do not specify a reboot time, the reboot occurs immediately following the command, and any previously scheduled reboots are cancelled. To cancel a previously scheduled reboot, use the `cancel` option.

SAVING CONFIGURATION CHANGES

The configuration is the customized set of parameters that you have selected to run on the switch. As you make configuration changes, the new settings are stored in run-time memory. Settings that are stored in run-time memory are not retained by the switch when the switch is rebooted. To retain the settings, and have them load when you reboot the switch, you must save the configuration to nonvolatile storage.

The switch can store two different configurations: a primary and a secondary. When you save configuration changes, you can select to which configuration you want the changes saved. If you do not specify, the changes are saved to the configuration area currently in use.

If you have made a mistake, or you must revert to the configuration as it was before you started making changes, you can tell the switch to use the secondary configuration on the next reboot.

To save the configuration, use the following command:

```
save {configuration} {primary | secondary}
```

To use the configuration, use the following command:

```
use configuration [primary | secondary]
```

The configuration takes effect on the next reboot.



If the switch is rebooted while in the middle of a configuration save, the switch boots to factory default settings. The configuration that is not in the process of being saved is unaffected.

RETURNING TO FACTORY DEFAULTS

To return the switch configuration to factory defaults, use the following command:

```
unconfig switch
```

This command resets the entire configuration, with the exception of user accounts and passwords that have been configured, and the date and time.

To erase the currently selected configuration image and reset all switch parameters, use the following command:

```
unconfig switch all
```

USING TFTP TO UPLOAD THE CONFIGURATION

You can upload the current configuration to a TFTP server on your network. The uploaded ASCII file retains the command-line interface (CLI) format. This allows you to do the following:

- Modify the configuration using a text editor, and later download a copy of the file to the same switch, or to one or more different switches.
- Send a copy of the configuration file to the Extreme Networks Technical Support department for problem-solving purposes.
- Automatically upload the configuration file every day, so that the TFTP server can archive the configuration on a daily basis. Because the filename is not changed, the configured file stored in the TFTP server is overwritten every day.

To upload the configuration, use the command

```
upload configuration [<ipaddress> | <hostname>] <filename> {every  
<time>}
```

where the following is true:

- **ipaddress** — Is the IP address of the TFTP server.
- **hostname** — Is the hostname of the TFTP server. (You must enable DNS to use this option.)
- **filename** — Is the name of the ASCII file. The filename can be up to 255 characters long, and can not include any spaces, commas, quotation marks, or special characters.
- **every <time>** — Specifies the time of day you want the configuration automatically uploaded on a daily basis. If not specified, the current configuration is immediately uploaded to the TFTP server.

To cancel a previously scheduled configuration upload, use the command

```
upload configuration cancel
```

USING TFTP TO DOWNLOAD THE CONFIGURATION

You can download ASCII files that contain CLI commands to the switch to modify the switch configuration. There are three types of configuration scenarios that can be downloaded:

- Complete configuration
- Incremental configuration
- Scheduled incremental configuration

DOWNLOADING A COMPLETE CONFIGURATION

Downloading a complete configuration replicates or restores the entire configuration to the switch. You typically use this type of download in conjunction with the `upload config` command, which generates a complete switch configuration in an ASCII format. As part of the complete configuration download, the switch is automatically rebooted.

To download a complete configuration, use the following command:

```
download configuration [<hostname | ip_address>] <filename>
```

After the ASCII configuration is downloaded by way of TFTP, you are prompted to reboot the switch. The downloaded configuration file is stored in current switch memory during the rebooting process, and is not retained if the switch has a power failure.

When the switch completes booting, it treats the downloaded configuration file as a script of CLI commands, and automatically executes the commands. If your CLI connection is through a Telnet connection (and not the console port), your connection is terminated when the switch reboots, but the command executes normally.

DOWNLOADING AN INCREMENTAL CONFIGURATION

A partial or incremental change to the switch configuration may be accomplished by downloaded ASCII files that contain CLI commands. These commands are interpreted as a script of CLI commands, and take effect at the time of the download, without requiring a reboot of the switch.

To download an incremental configuration, use the following command:

```
download configuration <hostname | ip_address> <filename> {incremental}
```

SCHEDULED INCREMENTAL CONFIGURATION DOWNLOAD

You can schedule the switch to download a partial or incremental configuration on a regular basis. You could use this feature to update the configuration of the switch regularly from a centrally administered TFTP server. As part of the scheduled incremental download, you can optionally configuration a backup TFTP server.

To configure the primary and/or secondary TFTP server and filename, use the following command:

```
config download server [primary | secondary] <hostname | ip_address>  
<filename>
```

To enable scheduled incremental downloads, use the following command:

```
download configuration every <hour (0-23)>
```

To display scheduled download information, use the following command:

```
show switch
```

To cancel scheduled incremental downloads, use the following command:

```
download configuration cancel
```

REMEMBER TO SAVE

Regardless of which download option is used, configurations are downloaded into switch runtime memory, only. The configuration is saved only when the `save` command is issued, or if the configuration file, itself, contains the `save` command.

If the configuration currently running in the switch does not match the configuration that the switch used when it originally booted, an asterisk (*) appears before the command line prompt when using the CLI.

SYNCHRONIZING MSMS

On the BlackDiamond switch, you can take the master MSM configurations and images and replicate them on the slave MSM using the following command:

```
synchronize
```

In addition to replicating the configuration settings and images, this command also replicates which configuration or image the MSM should use on subsequent reboots. This command does not replicate the run-time configuration. You must use the save configuration command to store the run-time configuration first. It also does not replicate the BootROM imaged stored on the MSM.

UPGRADING AND ACCESSING BOOTROM

The BootROM of the switch initializes certain important switch variables during the boot process. If necessary, BootROM can be upgraded, after the switch has booted, using TFTP. In the event the switch does not boot properly, some boot option functions can be accessed through a special BootROM menu.

UPGRADING BOOTROM

Upgrading BootROM is done using TFTP (from the CLI), after the switch has booted. Upgrade the BootROM only when asked to do so by an Extreme Networks technical representative. To upgrade the BootROM, use the following command:

```
download bootrom [<host_name> | <ip_addr>]
```

ACCESSING THE BOOTROM MENU

Interaction with the BootROM menu is only required under special circumstances, and should be done only under the direction of Extreme Networks Customer Support. The necessity of using these functions implies a non-standard problem which requires the assistance of Extreme Networks Customer Support.

To access the BootROM menu, follow these steps:

- 1 Attach a serial cable to the console port of the switch.
- 2 Attach the other end of the serial cable to a properly configured terminal or terminal emulator, power cycle the switch while depressing the spacebar on the keyboard of the terminal.

As soon as you see the `BootROM->` prompt, release the spacebar. You can see a simple help menu by pressing `h`. Options in the menu include

- Selecting the image to boot from
- Booting to factory default configuration
- Performing a serial download of an image

For example, to change the image that the switch boots from in flash memory, press `1` for the image stored in primary or `2` for the image stored in secondary. Then, press the `f` key to boot from newly selected on-board flash memory.

To boot to factory default configuration, press the `d` key for default and the `f` key to boot from the configured on-board flash.

To perform a serial download, you can optionally change the baud rate to 38.4K using the `b` command, and then pressing the `s` key to prepare the switch for an image to be sent from your terminal using the XMODEM protocol. After this has completed, select the `g` command, to boot the image that is currently in RAM. The switch restores the console port to 9600 bps, and begins the boot process.



Doing a serial download does not store an image into flash, it only allows the switch to boot an operational image so that a normal TFTP upgrade from CLI can then be performed.

BOOT OPTION COMMANDS

[Table 20-1](#) lists the CLI commands associated with switch boot options.

Table 20-1: Boot Option Commands

Command	Description
<code>config download server [primary secondary] <hostname ipaddress> <filename></code>	Configures the TFTP server(s) used by a scheduled incremental configuration download.
<code>download bootrom [<ipaddress> <hostname>] <filename></code>	Downloads a BOOT ROM image from a TFTP server. The downloaded image replaces the BOOT ROM in the onboard FLASH memory.



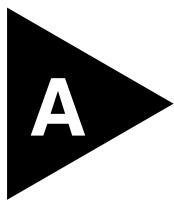
If this command does not complete successfully it could prevent the switch from booting.

Table 20-1: Boot Option Commands (continued)

Command	Description
download configuration <hostname ipaddress> <filename> {incremental}	Downloads a complete configuration. Use the incremental keyword to specify an incremental configuration download.
download configuration cancel	Cancels a previously scheduled configuration download.
download configuration every <hour>	Schedules a configuration download. Specify the hour using a 24-hour clock, where the range is 0 to 23.
download image [<ipaddress> <hostname>] <filename> {primary secondary}	Downloads a new image from a TFTP server over the network. If no parameters are specified, the image is saved to the current image.
reboot {time <date> <time> cancel}	Reboots the switch at the date and time specified. If you do not specify a reboot time, the reboot happens immediately following the command, and any previously scheduled reboots are cancelled. To cancel a previously scheduled reboot, use the cancel option.
save {configuration} {primary secondary}	Saves the current configuration to nonvolatile storage. You can specify the primary or secondary configuration area. If not specified, the configuration is saved to the primary configuration area.
show configuration	Displays the current configuration to the terminal. You can then capture the output and store it as a file.
synchronize	Replicates all saved images and configurations from the master MSM to the slave MSM. The run-time configuration is not copied, because it has not been saved to FLASH memory. Use the save configuration command to save the run-time information. (BlackDiamond switch only)
upload configuration [<ipaddress> <hostname>] <filename> {every <time>}	Uploads the current run-time configuration to the specified TFTP server. If every <time> is specified, the switch automatically saves the configuration to the server once per day, at the specified time. If the time option is not specified, the current configuration is immediately uploaded.
upload configuration cancel	Cancels a previously schedule configuration upload.

Table 20-1: Boot Option Commands (continued)

Command	Description
use configuration [primary secondary]	Configures the switch to use a particular configuration on the next reboot. Options include the primary configuration area or the secondary configuration area.
use image [primary secondary]	Configures the switch to use a particular image on the next reboot.



Supported Standards

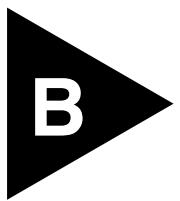
The following is a list of software standards supported by ExtremeWare.

Standards and Protocols

RFC 1058 RIP	RFC 783 TFTP
RFC 1723 RIP v2	RFC 1542 BootP
RFC 1112 IGMP	RFC 854 Telnet
RFC 2236 IGMP v2	RFC 768 UDP
DVMRP v3 - Draft IETF DVMRP v3-07	RFC 791 IP
PIM-DM v2 - Draft IETF PIM-DM v2-dm-01	RFC792 ICMP
RFC 1587-NSSA option	RFC 793 TCP
RFC 2178 OSPF	RFC 826 ARP
RFC 1122 Host requirements	RFC 2068 HTTP
IEEE 802.1D-1998 (802.1p) Packet priority	RFC 2131 BootP/DHCP relay
IEEE 802.1Q VLAN tagging	RFC 2030 - Simple Network Time Protocol
RFC 1256 Router discovery protocol	IPX RIP/SAP Router specification
RFC 1812 IP router requirement	Extreme Standby Router Protocol (ESRP)

Management and Security

RFC 1157 SNMP v1/v2c	RFC 1757 Four groups of RMON
RFC 1213 MIB II	RFC 2021 RMON probe configuration
RFC 1354 IP forwarding table MIB	RFC 2239 802.3 MAU MIB
RFC 1493 Bridge MIB	RFC 1724 RIP v2 MIB
RFC 2037 Entity MIB	ExtremeWare Enterprise MIB
RFC 1573 Evolution of Interface	HTML and Telnet management
RFC 1643 Ethernet MIB	RFC 2138 RADIUS



Troubleshooting

If you encounter problems when using the switch, this appendix may be helpful. If you have a problem not listed here or in the “Release Notes,” contact your local technical support representative.

LEDs

Power LED does not light:

Check that the power cable is firmly connected to the device and to the supply outlet.

On powering-up, the MGMT LED lights yellow:

The device has failed its Power On Self Test (POST) and you should contact your supplier for advice.

A link is connected, but the Status LED does not light:

Check that

- All connections are secure.
- Cables are free from damage.
- The devices at both ends of the link are powered-up.

- Both ends of the Gigabit link are set to the same autonegotiation state.

Both sides of the Gigabit link must be enabled or disabled. If the two are different, typically the side with autonegotiation disabled will have the link LED lit, and the side with autonegotiation enabled will not. The default configuration for a Gigabit port is autonegotiation enabled. This can be verified by entering the following command:

```
show port config
```

On power-on, some I/O modules do not boot:

Check if you are using 110V power input. The BlackDiamond powers only up to four modules if it is connected to a 110V outlet.

Error LED on the MSM64i turns amber:

Check the syslog message for a "critical" software errors.

Status LED on the I/O module turns amber:

Check the syslog message for a related I/O module error. If the error is an inserted an I/O module that conflicts with the software configuration, use one of the following commands to reset the slot configuration:

```
clear slot
```

```
config slot <slot> module [f32t | f32f | f48t | g4x | g6x | g8x | g12x]
```

Otherwise, contact Extreme Networks for further assistance.

ENV LED on the MSM64i turns amber:

Check each of the power supplies and all of the fans. Additionally, the status of these should be indicated in the display by entering "show switch" at the CLI. Look for the "Temperature" and "Power Supply" entries in the displayed information.

Switch does not power up:

All products manufactured by Extreme Networks use digital power supplies with surge protection. In the event of a power surge, the protection circuits shut down the power supply. To reset, unplug the switch for 1 minute, plug it back in, and attempt to power up the switch.

If this does not work, try using a different power source (different power strip/outlet) and power cord.

USING THE COMMAND-LINE INTERFACE

The initial welcome prompt does not display:

Check that your terminal or terminal emulator is correctly configured.

For console port access, you may need to press [Return] several times before the welcome prompt appears.

Check the settings on your terminal or terminal emulator. The settings are 9600 baud, 8 data bits, 1 stop bit, no parity, XON/OFF flow control enabled.

The SNMP Network Manager cannot access the device:

Check that the device IP address, subnet mask, and default router are correctly configured, and that the device has been reset.

Check that the device IP address is correctly recorded by the SNMP Network Manager (refer to the user documentation for the Network Manager).

Check that the community strings configured for the system and Network Manager are the same.

Check that SNMP access was not disabled for the system.

The Telnet workstation cannot access the device:

Check that the device IP address, subnet mask and default router are correctly configured, and that the device has been reset. Ensure that you enter the IP address of the switch correctly when invoking the Telnet facility. Check that Telnet access was not disabled for the switch. If you attempt to log in and the maximum number of Telnet sessions are being used, you should receive an error message indicating so.

Traps are not received by the SNMP Network Manager:

Check that the SNMP Network Manager's IP address and community string are correctly configured, and that the IP address of the Trap Receiver is configured properly on the system.

The SNMP Network Manager or Telnet workstation can no longer access the device:

Check that Telnet access or SNMP access is enabled.

Check that the port through which you are trying to access the device has not been disabled. If it is enabled, check the connections and network cabling at the port.

Check that the port through which you are trying to access the device is in a correctly configured VLAN.

Try accessing the device through a different port. If you can now access the device, a problem with the original port is indicated. Re-examine the connections and cabling.

A network problem may be preventing you accessing the device over the network. Try accessing the device through the console port.

Check that the community strings configured for the device and the Network Manager are the same.

Check that SNMP access was not disabled for the system.

Permanent entries remain in the FDB:

If you have made a permanent entry in the FDB (which requires you to specify the VLAN to which it belongs and then delete the VLAN), the FDB entry will remain. Though causing no harm, you must manually delete the entry from the FDB if you want to remove it.

Default and Static Routes:

If you have defined static or default routes, those routes will remain in the configuration independent of whether the VLAN and VLAN IP address that used them remains. You should manually delete the routes if no VLAN IP address is capable of using them.

You forget your password and cannot log in:

If you are not an administrator, another user having administrator access level can log in, delete your user name, and create a new user name for you, with a new password.

Alternatively, another user having administrator access level can log in and initialize the device. This will return all configuration information (including passwords) to the initial values.

In the case where no one knows a password for an administrator level user, contact your supplier.

PORT CONFIGURATION

No link light on 10/100 Base port:

If patching from a hub or switch to another hub or switch, ensure that you are using a CAT5 cross-over cable. This is a CAT5 cable that has pins 1&2 on one end connected to pins 3&6 on the other end.

Excessive RX CRC errors:

When a device that has auto-negotiation disabled is connected to a Extreme switch that has auto-negotiation enabled, the Extreme switch links at the correct speed, but in half duplex mode. The Extreme switch 10/100 physical interface uses a method called *parallel detection* to bring up the link. Because the other network device is not participating in auto-negotiation (and does not advertise its capabilities), parallel detection on the Extreme switch is only able to sense 10Mbps versus 100Mbps speed, and not the duplex mode. Therefore, the switch establishes the link in half duplex mode using the correct speed.

The only way to establish a full duplex link is to either force it at both sides, or run auto-negotiation on both sides (using full duplex as an advertised capability, which is the default setting on the Extreme switch).



A mismatch of duplex mode between the Extreme switch and another network device will cause poor network performance. Viewing statistics using the show port rx command on the Extreme switch may display a constant increment of CRC errors. This is characteristic of a duplex mismatch between devices. This is NOT a problem with the Extreme switch.

Always verify that the Extreme switch and the network device match in configuration for speed and duplex.

No link light on Gigabit fiber port:

Check to ensure that the transmit fiber goes to the receive fiber side of the other device, and vice-versa. All gigabit fiber cables are of the cross-over type.

The Extreme switch has auto-negotiation set to on by default for gigabit ports. These ports need to be set to auto off (using the command config port <port #> auto off) if you are connecting it to devices that do not support auto-negotiation.

Ensure that you are using multi-mode fiber (MMF) when using a 1000BASE-SX GBIC, and single mode fiber (SMF) when using a 1000BASE-LX GBIC. 1000BASE-SX does not work with SMF. 1000BASE-LX works with MMF, but requires the use of a mode conditioning patchcord (MCP).

VLANs

You cannot add a port to a VLAN:

If you attempt to add a port to a VLAN and get an error message similar to

```
localhost:7 # config vlan marketing add port 1:1,1:2  
ERROR: Protocol conflict on port 1:5
```

you already have a VLAN using untagged traffic on a port. Only one VLAN using untagged traffic can be configured on a single physical port.

VLAN configuration can be verified by using the following command:

```
show vlan <name>
```

The solution for this error is to remove ports 1 and 2 from the VLAN currently using untagged traffic on those ports. If this were the “default” VLAN, the command would be

```
localhost:23 # config vlan default del port 1:1,1:2
```

which should now allow you to re-enter the previous command without error as follows:

```
localhost:26 # config vlan red add port 1:1,1:2
```

VLAN names:

There are restrictions on VLAN names. They cannot contain whitespaces and cannot start with a numeric value unless you use quotation marks around the name. If a name contains whitespaces, starts with a number, or contains non-alphabetical characters, you must use quotation marks whenever referring to the VLAN name.

802.1Q links do not work correctly:

Remember that VLAN names are only locally significant through the command-line interface. For two switches to communicate across a 802.1Q link, the VLAN ID for the VLAN on one switch should have a corresponding VLAN ID for the VLAN on the other switch.

If you are connecting to a third-party device and have checked that the VLAN IDs are the same, the Ethertype field used to identify packets as 802.1Q packets may differ between the devices. The default value used by the switch is **8100**. If the third-party device differs from this and cannot be changed, you may change the 802.1Q Ethertype used by the switch with the following command:

```
config dot1p ethertype <ethertype>
```

Changing this parameter changes how the system recognizes all tagged frames received, as well as the value it inserts in all tagged frames it transmits.

VLANs, IP Addresses and default routes:

The system can have an IP address for each configured VLAN. It is necessary to have an IP address associated with a VLAN if you intend to manage (Telnet, SNMP, ping) through that VLAN or route IP traffic. You can also configure multiple default routes for the system. The system first tries the default route with the lowest cost metric.

STP

You have connected an endstation directly to the switch and the endstation fails to boot correctly:

The switch has STP enabled, and the endstation is booting before the STP initialization process is complete. Specify that STP has been disabled for that VLAN, or turn off STP for the switch ports of the endstation and devices to which it is attempting to connect, and then reboot the endstation.

The switch keeps aging out endstation entries in the switch Forwarding Database (FDB):

Reduce the number of topology changes by disabling STP on those systems that do not use redundant paths.

Specify that the endstation entries are static or permanent.

DEBUG TRACING

ExtremeWare includes a debug-tracing facility for the switch. The show debug-tracing command can be applied to one or all VLANs, as follows:

```
show debug-tracing {vlan <name>}
```

The debug commands should only be used under the guidance of Extreme Networks technical personnel.

TOP COMMAND

The `top` command is a utility that indicates CPU utilization by process.

CONTACTING EXTREME TECHNICAL SUPPORT

If you have a network issue that you are unable to resolve, contact Extreme Networks technical support. Extreme Networks maintains several Technical Assistance Centers (TACs) around the world to answer networking questions and resolve network problems. You can contact technical support by phone at:

- (800) 998-2408
- (408) 579-2826

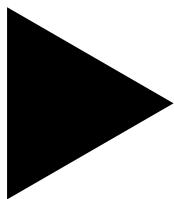
or by email at:

- support@extremenetworks.com

You can also visit the support website at:

- <http://www.extremenetworks.com/extreme/support/techsupport.asp>

to download software updates (requires a service contract) and documentation (including a .pdf version of this manual).



Index

Numerics

3DNS 17-40
802.1p configuration commands (table) 9-14

A

access levels 2-9
access lists
 BlackDiamond switch maximum entries 16-5
 configuration commands (table) 16-7
 deleting 16-4
 description 16-2
 examples 16-11
 ICMP filter example 16-14
 ICMP traffic 16-5
 maximum entries 16-5
 permit-established example 16-11
 permit-established keyword 16-4
 restrictions 16-5
 verifying settings 16-6
access policies, description 16-1
access profiles
 ExtremeWare Vista 3-9
 reverse mask 16-17
 SNMP 3-10
 Telnet 3-6
accounts, creating 2-12
admin account 2-11
aging entries, FDB 7-2
aging timer, FDB and ISQ 9-24
alarm actions 18-15
Alarms, RMON 18-13
area 0, OSPF 12-6
areas, OSPF 12-6

autonomous system, description 13-2

B

backbone area, OSPF 12-6
BGP
 attributes 13-2
 autonomous system 13-2
 autonomous system path 13-2
 cluster 13-3
 community 13-3
 configuration commands (table) 13-10
 description 13-2
 features 13-3
 IGP synchronization 13-9
 loopback interface 13-9
 redistributing to OSPF 13-9
 reset and disable commands (table) 13-15
 route aggregation 13-8
 route maps 16-34
 route reflectors 13-3
 routing access policies 16-25
 settings, displaying 13-15
 show commands (table) 13-15
Bi-directional rate shaping 9-25
BlackDiamond switch
 access list maximum entries 16-5
 autonegotiation 4-3
 configuring load sharing 4-9
 enabling and disabling ports 4-3
 jumbo frames 4-7
 load sharing example 4-11
 load sharing group combinations 4-9
 load sharing master port 4-9
 MSMs, synchronizing 20-7

port configuration 4-2
port-mirroring, virtual port 4-11
slot configuration 4-1
verifying load sharing 4-11
blackhole entries, FDB 7-2
boot option commands (table) 20-8
BOOTP and UDP-Forwarding 11-17
BOOTP relay, configuring 11-16
BOOTP, using 3-4
BootROM
 menu, accessing 20-7
 prompt 20-8
 upgrading 20-7
Border Gateway Protocol. *See* BGP
browser
 controls 19-5
 fonts 19-3
 setting up 19-2

C

CLI
 command history 2-6
 command shortcuts 2-2
 line-editing keys 2-5
 named components 2-4
 numerical ranges, BlackDiamond switch 2-3
 numerical ranges, Summit switch 2-4
 symbols 2-4
 syntax helper 2-2
 using
command
 history 2-6
 shortcuts 2-2
 syntax, understanding 2-1
Command-Line Interface. *See* CLI
common commands (table) 2-6
complete configuration download 20-5
config dvmrp vlan cost 14-5
configuration
 downloading 20-5
 downloading complete 20-5
 downloading incremental 20-5
 logging 18-10
 primary and secondary 20-3
 saving changes 20-3
 schedule download 20-6
 uploading to file 20-4
controlling Telnet access 3-6
conventions
 notice icons, About This Guide xx
 text, About This Guide xxi

D
database applications, and QoS 9-4
default
 passwords 2-11
 settings 1-8
 users 2-11
default STP domain 8-3
default VLAN 6-12
deleting a session 3-6
DHCP and UDP-Forwarding 11-16
DHCP relay, configuring 11-16
DiffServ, configuring 9-15
disabling a BlackDiamond switch port 4-3
disabling a Summit switch port 5-1
disabling route advertising (RIP) 12-4
disconnecting a Telnet session 3-6
Distance Vector Multicast Routing Protocol. *See* DVMRP
distance-vector protocol, description 12-2
DLCS
 configuration commands (table) 9-29
 description 9-27
 guidelines 9-28
 limitations 9-28
DNS
 configuration commands (table) 2-13
 description 2-13
Domain Name Service. *See* DNS
domains, Spanning Tree Protocol 8-2
downloading incremental configuration 20-5
DVMRP
 configuring 14-5
 description 14-2
 routing access policies 16-23
dynamic entries, FDB 7-2
Dynamic Link Context System. *See* DLCS 9-24
dynamic routes 11-4, 15-4

E

ECMP. *See* IP route sharing
EDP
 commands (table) 4-13, 5-12
 description 4-13, 5-12
enabling a BlackDiamond switch port 4-3
enabling a Summit switch port 5-1
Equal Cost Multi-Path (ECMP) routing. *See* IP
 route sharing
errors, port 18-5
ESRP
 and IP multinetting 10-12
 and STP 10-12
 and VLAN aggregation 10-13
 configuration commands (table) 10-14
 description 10-1
 direct link 10-12
 domains 10-10
 example 10-16
 failover time 10-6
 groups 10-11

host attach 10-9
linking switches 10-12
master
 behavior 10-5
 definition 10-2
 determining 10-3
 electing 10-6
 election algorithms 10-5
port blocks 10-7
standby mode
 behavior 10-6
 definition 10-2
super-VLAN 10-13
 using 10/100 ports 10-7

Events, RMON 18-14
external health checking, SLB 17-25
Extreme Discovery Protocol *See EDP*
Extreme Standby Router Protocol. *See ESRP*
ExtremeWare
 factory defaults 1-8
 features 1-1
ExtremeWare Vista
 accessing 19-3
 browser controls 19-5
 browser setup 19-2
 capturing screen output 19-7
 controlling access 3-9
 description 19-1
 fonts 19-3
 home page 3-8, 19-3
 navigating 19-4
 saving changes 19-6
 screen layout 19-4
 screen resolution 19-3
 status messages 19-5
 VLAN configuration 19-2

F

FDB
 adding an entry 7-3
 aging entries 7-2
 aging timer and ISQ 9-24
 blackhole entries 7-2
 configuration commands (table) 7-3
 configuring 7-3
 contents 7-1
 creating a permanent entry example 7-4
 displaying 7-5
 dynamic entries 7-2
 entries 7-1
 non-aging entries 7-2
 permanent entries 7-2
 QoS profile association 7-3
file server applications, and QoS 9-4
flow control 4-4, 5-2
flow redirection 17-46

flow redirection commands (table) 17-47
fonts, browser 19-3
Forwarding Database. *See FDB*
forwarding modes, SLB 17-5

G

GARP VLAN Registration Protocol. *See GVRP*
GoGo mode, SLB 17-11
Greenwich Mean Time Offsets (table) 3-23
GVRP
 configuration commands (table) 6-21, 11-14
 description 6-18
 example 6-18

H

history command 2-6
History, RMON 18-13
home page 3-8, 19-3

I

ICMP configuration commands (table) 11-22
ICMP Router Discovery Protocol. *See IRDP*
ICMP, access lists 16-5
IEEE 802.1Q 6-6
IGMP
 configuration commands (table) 14-8
 description 14-3
 snooping 14-4
image
 downloading 20-1
 primary and secondary 20-2
 upgrading 20-1
interfaces, router 11-2, 15-1
Internet Group Management Protocol. *See IGMP*
Intra-Subnet QoS. *See ISQ* 9-24
IP address, entering 3-4
IP multicast routing
 configuration commands (table) 14-5
 configuring 14-4
 description 1-4, 14-2
 disabling 14-14
 DVMRP
 configuring 14-5
 description 14-2
 example 14-9
 IGMP
 configuration commands (table) 14-8
 description 14-3
 snooping 14-4
 PIM 14-5
 PIM-DM 14-3
 PIM-SM 14-3
 reset and disable commands (table) 14-14

resetting 14-14
settings, displaying 14-13
show commands (table) 14-13

IP multinetting
description 11-7
example 11-9
primary VLAN interface 11-7
secondary VLAN interface 11-7
using 11-8

IP route sharing 11-5

IP TOS configuration commands (table) 9-16

IP unicast routing
basic IP commands (table) 11-19
BOOTP relay 11-16
configuration examples 11-25
configuring 11-10
default gateway 11-2
description 1-4
DHCP relay 11-16
disabling 11-28
ECMP
enabling 11-11
IP route sharing 11-5
multinetting, description 11-7
multinetting, example 11-9
proxy ARP 11-5
reset and disable commands (table) 11-28
resetting 11-28
router interfaces 11-2
router show commands (table) 11-27
routing table
configuration commands (table) 11-21
dynamic routes 11-4
multiple routes 11-4
populating 11-3
static routes 11-4

settings, displaying 11-27

verifying the configuration 11-11

IPX

configuration commands (table) 15-7
configuration example 15-11
configuring 15-6
disabling 15-14
protocol filters 15-7
protocol-based VLANs 15-7
reset and disable commands (table) 15-14
resetting 15-14
router interfaces 15-1
routing table
configuration commands (table) 15-9
dynamic routes 15-4
populating 15-4
static routes 15-4

service table
configuration commands (table) 15-10
settings, displaying 15-13
show commands (table) 15-13

verifying router configuration 15-6

IPX/RIP 15-14
configuring 15-6
disabling 15-14
reset and disable commands (table) 15-14
routing table configuration commands
(table) 15-9
routing table, populating 15-4
settings, displaying 15-13
show commands (table) 15-13

IPX/SAP 15-14
configuration commands (table) 15-10
configuring 15-6
disabling 15-14
reset and disable commands (table) 15-14
settings, displaying 15-13
show commands (table) 15-13

IRDP 11-25

ISQ
configuration commands (table) 9-24
FDB aging timer 9-24

J
jumbo frames 4-7, 5-5

K
keys
line-editing 2-5
port monitoring 18-6

L
line-editing keys 2-5
link-state database 12-5
link-state protocol, description 12-2
load balancing methods, SLB 17-13
load sharing
configuring on BlackDiamond switch 4-9
configuring on Summit switch 5-7
description on BlackDiamond switch 4-7
description on Summit switch 5-6
group combinations on BlackDiamond switch
(table) 4-9
group combinations on Summit switch
(table) 5-8
load-sharing group on BlackDiamond switch,
description 4-8
load-sharing group on Summit switch,
description 5-6
master port on BlackDiamond switch 4-9
master port on Summit switch 5-7
verifying the configuration on BlackDiamond
switch 4-11

verifying the configuration on Summit switch 5-10
local logging 18-9
log display 18-9
logging
and Telnet 18-9
commands (table) 18-11
configuration changes 18-10
description 18-7
fault level 18-7
local 18-9
message 18-8
QoS monitor 9-22
real-time display 18-9
remote 18-9
subsystem 18-8
timestamp 18-7
logging in 2-11

M

MAC-based VLAN
configuration commands (table) 6-23
example 6-23
timed configuration download 6-24
maintenance mode, SLB 17-25
management access 2-9
Management Switch Fabric Module. *See* MSM
master port
load sharing on BlackDiamond switch 4-9
load sharing on Summit switch 5-7
MIBs 3-10
mirroring. *See* port-mirroring
monitoring the switch 18-1
MSM 3-2
multinetting. *See* IP multinetting
multiple routes 11-4

N

names, VLANs 6-12
non-aging entries, FDB 7-2
Not-So-Stubby_Area. *See* NSSA
NSSA. *See* OSPF

O

Open Shortest Path First. *See* OSPF
OSPF
advantages 12-3
area 0 12-6
areas 12-6
backbone area 12-6
configuration commands (table) 12-21
configuration example 12-25
description 12-2, 12-5

disabling 12-28
enabling 11-11
hello interval 12-22
link-state database 12-5
normal area 12-8
NSSA 12-7
redistributing to BGP 13-9
reset and disable commands (table) 12-28
resetting 12-28
router types 12-6
routing access policies 16-21
settings, displaying 12-28
show commands (table) 12-28
stub area 12-7
virtual link 12-8

P

passwords
default 2-11
forgetting 2-12
permanent entries, FDB 7-2
permit-established keyword 16-4
persistence, SLB 17-26
PIM
configuration 14-5
PIM-DM
description 14-3
PIM-SM
description 14-3
rendezvous point 14-3
ping command 2-14
ping-check 17-23
poison reverse 12-4
port
autonegotiation on BlackDiamond switch 4-3
autonegotiation on Summit switch 5-2
BlackDiamond switch 4-2
BlackDiamond switch commands (table) 4-4
configuring on BlackDiamond switch 4-2
enabling and disabling on BlackDiamond switch 4-3
enabling and disabling on Summit switch 5-1
errors, viewing 18-5
load-sharing groups 5-8
master port on BlackDiamond switch 4-9
master port on Summit switch 5-7
monitoring display keys 18-6
priority, STP 8-6
receive errors 18-5
statistics, viewing 18-4
STP state, displaying 8-9
STPD membership 8-2
Summit switch commands (table) 5-3
transmit errors 18-5
port translation mode, SLB 17-10
port-based VLANs 6-2

port-mirroring
BlackDiamond switch configuration commands (table) 4-12
BlackDiamond switch example 4-12
description on BlackDiamond switch 4-11
description on Summit switch 5-10
example on Summit switch 5-12
Summit switch configuration commands (table) 5-11
virtual port on BlackDiamond switch 4-11
virtual port on Summit switch 5-11
primary image 20-2
profiles, QoS 9-6
protocol filters 6-10
protocol filters, IPX 15-7
Protocol Independent Multicast- Dense Mode. *See* PIM-DM
Protocol Independent Multicast- Sparse Mode. *See* PIM-SM
protocol-based VLANs 6-9
proxy ARP, and subnets 11-6
proxy ARP, description 11-5

Q

QoS
802.1p configuration commands (table) 9-14
802.1p priority 9-12
applications 9-3
Assigning QoS Service Levels 9-5
blackhole 9-11
buffer 9-6
configuration commands (table) 9-8
database applications 9-4
default QoS profiles 9-7
default QoS profiles (table) 9-7
description 1-3, 9-1
DiffServ, configuring 9-15
examples
 MAC address 9-10
 source port 9-20
 VLAN 9-21
FDB entry association 7-3
file server applications 9-4
IP TOS configuration commands (table) 9-16
maximum bandwidth 9-6
minimum bandwidth 9-6
policy, description 9-5
priority 9-6
profiles
 default 9-7
 default (table) 9-7
 description 9-5
 parameters 9-6
Random Early Detection (RED) 9-2
traffic groupings 9-8
 blackhole 9-11

broadcast/unknown rate limiting 9-11
description 9-5
explicit packet marking 9-12
IP address 9-10
MAC address 9-10
source port 9-20
VLAN 9-20
traffic groupings (table) 9-9
verifying 9-23
video applications 9-3
voice applications 9-3
web browsing applications 9-4
QoS monitor
 configuration commands (table) 9-22
 description 9-21
 logging 9-22
 real-time display 9-22
Quality of Service. *See* QoS 9-2

R

RADIUS commands (table) 3-15
Random Early Detection (RED) 9-2
receive errors 18-5
remote logging 18-9
Remote Monitoring. *See* RMON
reset to factory defaults 20-3
resetting 15-14
reverse mask 16-17
RIP
 advantages 12-2
 configuration commands (table) 12-14
 configuration example 12-17
 description 12-2, 12-3
 disabling route advertising 12-4
 enabling 11-11
 limitations 12-3
 poison reverse 12-4
 reset and disable commands (table) 12-20
 routing access policies 16-19
 routing table entries 12-3
 settings, displaying 12-19
 show commands (table) 12-19
 split horizon 12-4
 triggered updates 12-4
 version 2 12-5
RMON
 alarm actions 18-15
 Alarms group 18-13
 Events group 18-14
 features supported 18-13
 History group 18-13
 probe 18-12
 Statistics group 18-13
route maps
 BGP 16-34
 changing 16-34

configuration commands (table) 16-35
creating 16-30
description 16-2, 16-29
example 16-32
goto entries 16-31
match entries 16-31
match operation keywords (table) 16-31
processing 16-32
set entries 16-31
 set operation keywords (table) 16-32
route sharing. *See* IP route sharing
router interfaces 11-2, 15-1
router types, OSPF 12-6
routing access policies
 access profile
 applying 16-18
 changing 16-25
 configuring 16-16
 creating 16-16
 types 16-16
BGP 16-25
configuration commands (table) 16-27
deny 16-16
DVMRP 16-23
examples
 DVMRP 16-23
 OSPF 16-22
 PIM 16-24
 RIP 16-19
none 16-16
OSPF 16-21
permit 16-16
PIM 16-24
removing 16-26
RIP 16-19
 using 16-15
Routing Information Protocol. *See* RIP
routing table, populating 11-3
routing table, populating IPX 15-4
routing. *See* IP unicast routing

S

saving changes using ExtremeWare Vista 19-6
saving configuration changes 20-3
scheduling configuration download 20-6
screen resolution, ExtremeWare Vista 19-3
secondary image 20-2
Server Load Balancing *See* SLB
service-check 17-24
sessions, deleting 3-6
shortcuts, command 2-2
show ipxfdb 15-13
Simple Network Management Protocol. *See* SNMP
SLB
 3DNS support 17-40
 active-active 17-31

advanced configuration commands
 (table) 17-40
basic configuration commands (table) 17-15
client persistence 17-26
components 17-2
description 17-2
external health checking 17-25
failover 17-32
forwarding mode 17-5
GoGo mode 17-11
health checking 17-22
high availability 17-31
host-route 17-12
least connections 17-14
load balancing methods 17-13
maintenance mode 17-25
manual fail-back 17-35
nodes 17-3
persistence 17-26
ping-check 17-23
ping-checking 17-32
pools 17-3
port translation mode 17-10
priority mode 17-14
proxy ARP 17-12
ratio 17-13
ratio weight 17-14
redundant configuration 17-31
round-robin 17-13
service-check 17-24
standard virtual servers 17-4
sticky persistence 17-26
subnet-route 17-12
tcp-port-check 17-23
translational mode 17-8
transparent mode 17-5
VIPs 17-3
VIPs, creating 17-4
virtual servers 17-3
wildcard virtual servers 17-4

slot
 automatic configuration 4-1
 clearing 4-2
 manually configuring 4-2
 mismatch 4-2
smart redundancy 5-13
SNAP protocol 6-11
SNMP
 community strings 3-11
 configuration commands (table) 3-11
 configuring 3-10
 controlling access 3-10
 read access 3-10
 read/write access 3-10
 settings, displaying 3-13
 supported MIBs 3-10
 trap receivers 3-10

using 3-9
SNTP
 configuration commands (table) 3-25
 configuring 3-22
 Daylight Savings Time 3-22
 description 3-21
 example 3-25
 Greenwich Mean Time offset 3-22
 Greenwich Mean Time Offsets (table) 3-23
Spanning Tree Protocol. *See* STP
speed, ports on BlackDiamond switch 4-4
speed, ports on Summit switch 5-2
split horizon 12-4
static routes 11-4, 15-4
statistics, port 18-4
Statistics, RMON 18-13
status monitoring 18-1
status monitoring commands (table) 18-2
STP
 and ESRP 10-12
 and VLANs 8-2
 bridge priority 8-6
 configurable parameters 8-6
 configuration commands (table) 8-7
 configuration example 8-8
 configuring 8-6
 default domain 8-3
 description 1-3
 disable and reset commands (table) 8-9
 displaying settings 8-8
 domains 8-2
 examples 8-3
 forward delay 8-6
 hello time 8-6
 max age 8-6
 overview 8-1
 path cost 8-6
 port priority 8-6
 port state, displaying 8-9
stub area, OSPF 12-7
sub-VLAN 11-11
Summit switch
 autonegotiation 5-2
 disabling a port 5-1
 enabling a port 5-1
 jumbo frames 5-5
 load sharing example 5-10
 load sharing group combinations 5-8
 load sharing master port 5-7
 port-mirroring, virtual port 5-11
 verifying load sharing 5-10
super-VLAN 10-13, 11-11
switch
 logging 18-7
 monitoring 18-1
 RMON features 18-13
 synchronizing MSMs 20-7

syntax, understanding 2-1
syslog host 18-9

T

TACACS+
 commands (table) 3-20
 description 3-20
 servers, specifying 3-20
tagging, VLAN 6-6
tcp-port-check 17-23
technical support B-8
Telnet
 controlling access 3-6
 disconnecting a session 3-6
 logging 18-9
 using 3-3
Terminal Access Controller Access Control System Plus. *See* TACACS+
TFTP
 server 20-1
 using 20-4
timed configuration download, MAC-based VLAN 6-24
traceroute command 2-15
traffic groupings 9-8
translational mode, SLB 17-8
transmit errors 18-5
transparent mode, SLB 17-5
triggered updates 12-4
trunks 6-6

U

UDP-Forwarding
 and BOOTP 11-17
 and DHCP 11-16
 configuration commands (table) 11-18
 configuring 11-17
 description 11-16
 example 11-17
 profiles 11-17
 VLANs 11-17
upgrading the image 20-1
uploading the configuration 20-4
users
 access levels 2-9
 creating 2-12
 default 2-11
 viewing 2-12

V

video applications, and QoS 9-3
viewing accounts 2-12
VIPs, SLB 17-3

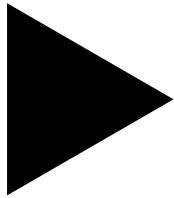
Virtual LANs. *See* VLANs
virtual link, OSPF 12-8
virtual port on Summit switch 5-11
VLAN aggregation
 commands (table) 11-14
 description 11-11
 limitations 11-13
 properties 11-13
 proxy ARP 11-14
 secondary IP address 11-12
 sub-VLAN 11-11
 super-VLAN 11-11
VLAN tagging 6-6
VLANs
 and ExtremeWare Vista 19-2
 and STP 8-2
 assigning a tag 6-6
 benefits 6-2
 configuration commands (table) 6-14
 configuration examples 6-15
 configuring 6-13
 default 6-12
 description 1-3
 disabling route advertising 12-4
 displaying settings 6-16
 ISQ 9-24
 mixing port-based and tagged 6-9
 names 6-12
 port-based 6-2
 protocol filters 6-10
 protocol-based 6-9
 protocol-based, IPX 15-7
 routing 11-10, 15-6
 tagged 6-6
 trunks 6-6
 types 6-2
 UDP-Forwarding 11-17
vMANs 6-17
voice applications, QoS 9-3

W

Web access, controlling 3-9
web browsing applications, and QoS 9-4
web cache redirection 17-46

X

xmodem 20-2



Index of Commands

C

clear counters 18-11
clear dlc5 9-29
clear fdb 7-3, 9-11
clear igmp snooping 14-14
clear iparp 11-19, 11-28
clear ipfdb 11-19, 11-28
clear ipmc cache 14-14
clear log 18-11
clear session 2-6, 3-6
clear slb connections 17-15
clear slb vip persistence 17-40
clear slot 4-2, 4-4
config access-profile 16-28
config access-profile add 16-17, 16-27
config access-profile delete 16-18, 16-27
config access-profile mode 16-17
config access-profile type 16-16
config account 2-6
config banner 2-6
config bgp add aggregate-address 13-8, 13-10
config bgp add confederation-peer 13-10
config bgp add dampening 13-11
config bgp add network 13-11
config bgp as-number 13-11
config bgp cluster-id 13-11
config bgp confederation-id 13-11
config bgp dampening 13-12
config bgp delete aggregate-address 13-12
config bgp delete dampening 13-12
config bgp export 13-12

config bgp local-preference 13-12
config bgp multi-exist-discriminator 13-12
config bgp neighbor as-path-filter 13-13, 16-25, 16-28
config bgp neighbor nlri-filter 13-13, 16-25, 16-28
config bgp neighbor route-map-filter 13-13
config bgp neighbor route-reflector-client 13-13
config bgp neighbor send-communities 13-13
config bgp neighbor soft-reset 13-13
config bgp neighbor source-interface 13-13
config bgp neighbor timer 13-13
config bgp neighbor weight 13-14
config bgp router-id 13-14
config bgp soft-reconfiguration 13-14
config bootprelay add 11-16, 11-19
config bootprelay delete 11-16, 11-19
config diffserv examination code-point 9-16, 9-18
config diffserv replacement 9-16
config diffserv replacement priority 9-19
config dns-client add 2-13
config dns-client default-domain 2-13
config dns-client delete 2-13
config dot1p type 9-14
config dot1q ethertype 6-14
config download server 6-24, 20-6, 20-8
config dvmrp add vlan 14-4, 14-5
config dvmrp delete vlan 14-5
config dvmrp timer 14-5
config dvmrp vlan 14-6
config dvmrp vlan export-filter 14-5, 16-23, 16-28
config dvmrp vlan import-filter 14-5, 16-23, 16-28
config dvmrp vlan trusted-gateway 14-5, 16-23,

16-28
config esrp port-mode 10-14
config fdb agingtime 7-3
config flow redirection add next-hop 17-47
config flow-redirection delete next-hop 17-47
config gvrp 6-21
config igmp query_interval 14-8
config igmp snooping 14-8
config iparp add 11-19
config iparp add proxy 11-5, 11-20
config iparp delete 11-20
config iparp delete proxy 11-20
config iparp timeout 11-20
config ipmc cache timeout 14-14
config iproute add 11-21
config iproute add blackhole 11-21
config iproute add default 11-11, 11-22
config iproute delete 11-22
config iproute delete blackhole 11-22
config iproute delete default 11-22
config iproute priority 11-7, 11-22
config ipxmaxhops 15-7
config ipxrip add vlan 15-9
config ipxrip delete 15-5
config ipxrip delete vlan 15-9
config ipxrip vlan delay 15-9
config ipxrip vlan max-packet-size 15-9
config ipxrip vlan update-interval 15-10
config ipxroute add 15-4, 15-8
config ipxroute delete 15-8
config ipxsap add vlan 15-10
config ipxsap delete vlan 15-10
config ipxsap vlan delay 15-10
config ipxsap vlan gns-delay 15-10
config ipxsap vlan max-packet-size 15-10
config ipxsap vlan update-interval 15-10
config ipxservice add 15-5, 15-8
config ipxservice delete 15-8
config irdp 11-22, 11-23
config isq-server add ipaddress 9-24, 9-29
config isq-server add mac 9-24, 9-29
config isq-server delete ipaddress 9-24, 9-29
config isq-server delete mac 9-24, 9-29
config jumbo-frame size 4-4, 4-7, 5-3, 5-5
config log display 18-9, 18-11
config mac-vlan add 6-23
config mac-vlan delete 6-23
config mirroring add 4-12, 5-11
config mirroring delete 4-12, 5-11
config ospf add virtual-link 12-22
config ospf add vlan 12-22
config ospf area add range 12-22
config ospf area delete range 12-22
config ospf area external-filter 16-21, 16-28
config ospf area interarea-filter 16-21, 16-28
config ospf area normal 12-22
config ospf area nssa 12-23
config ospf area stub 12-23
config ospf asbr-filter 12-23, 16-21, 16-28
config ospf ase-summary add 12-23
config ospf ase-summary delete 12-23
config ospf authentication 12-21
config ospf cost 12-21
config ospf delete virtual-link 12-23
config ospf delete vlan 12-23
config ospf direct-filter 12-23, 16-22, 16-28
config ospf lsa-batching-timer 12-23
config ospf metric-table 12-23
config ospf originate-default 12-23
config ospf priority 12-21
config ospf routerid 12-23
config ospf spf-hold-time 12-24
config ospf timer 12-22
config ospf vlan 12-24
config ospf vlan area 12-7
config pim add vlan 14-4, 14-6
config pim cbsr 14-6
config pim crp timer 14-6
config pim crp vlan 14-6
config pim delete vlan 14-6
config pim spt-threshold 14-6
config pim timer 14-7
config pim vlan trusted-gateway 14-7, 16-24, 16-29
config ports auto off 2-6, 4-4, 4-5, 5-2, 5-3
config ports auto on 4-4, 4-5, 5-2, 5-3
config ports display-string 4-5, 5-3
config ports qosprofile 4-5, 5-3, 9-8, 9-20
config protocol 6-14
config protocol add 6-11
config qosprofile 9-8
config radius server 3-15
config radius shared-secret 3-15
config radius-accounting 3-15
config radius-accounting shared-secret 3-16
config red drop-probability 9-8
config rip add 12-14
config rip delete 12-14, 12-20
config rip garbagetime 12-14
config rip routetimeout 12-14
config rip rxmode 12-14
config rip txmode 12-15
config rip updatetime 12-15
config rip vlan cost 12-15
config rip vlan export-filter 16-19, 16-29

config rip vlan import-filter 16-19, 16-29
config rip vlan trusted-gateway 16-19, 16-29
config route-map add 16-30, 16-36
config route-map add goto 16-31, 16-35
config route-map add match 16-30, 16-35
config route-map add set 16-31, 16-35
config route-map delete 16-36
config route-map delete goto 16-35
config route-map delete match 16-36
config route-map delete set 16-36
config slb 3dns-encryption-key 17-40
config slb failover 17-31, 17-32, 17-41
config slb failover fallback-now 17-35, 17-40
config slb failover ping-check 17-32
config slb failover ping-check ipaddress 17-40
config slb global frequency 17-42
config slb global ftp userid 17-42
config slb global http url 17-42
config slb global nntp 17-42
config slb global persistence-level 17-42
config slb global persistence-method 17-42
config slb global pop3 userid 17-42
config slb global smtp 17-42
config slb global synguard 17-42
config slb global telnet userid 17-43
config slb l4-port 17-15
config slb node max-connections 17-43
config slb node ping-check 17-43
config slb node tcp-port-check frequency 17-43
config slb pool 17-14
config slb pool add 17-3, 17-15
config slb pool delete 17-15
config slb pool lb-method 17-13, 17-14, 17-15
config slb vip max-connections 17-43
config slb vip service-check frequency 17-43
config slb vip service-check ftp 17-43
config slb vip service-check http 17-43
config slb vip service-check nntp 17-43
config slb vip service-check pop3 userid 17-43
config slb vip service-check smtp 17-43
config slb vip service-check telnet 17-43
config slb vip unit 17-33, 17-44
config slot 2-7, 4-2, 4-5
config snmp access-profile readonly 3-11
config snmp access-profile readwrite 3-11
config snmp add trapreceiver 3-11
config snmp community 3-11
config snmp delete trapreceiver 3-12
config snmp readonly access-profile 3-10
config snmp readwrite access-profile 3-10
config snmp syscontact 3-12
config snmp syslocation 3-12
config snmp sysname 3-12
config sntp-client 3-22
config sntp-client server 3-25
config sntp-client update-interval 3-22, 3-25
config ssh2 key 2-7, 3-8
config stpd add vlan 8-6, 8-7
config stpd forwarddelay 8-7
config stpd hello-time 8-7
config stpd max-age 8-7
config stpd port cost 8-7
config stpd port priority 8-7
config stpd priority 8-7
config syslog 18-10, 18-11
config sys-recovery-level 2-7, 18-7
config tacacs 3-20
config tacacs shared-secret 3-20
config tacacs-accounting 3-20
config tacacs-accounting shared-secret 3-20
config tcp-sync-rate 11-20
config time 2-7
config timezone 2-7, 3-22
config udp-profile add 11-18
config udp-profile delete 11-18
config vlan add domain-member 10-15
config vlan add port 6-14
config vlan add secondary-ip 11-14
config vlan add subvlan 11-14
config vlan add track-ping 10-14
config vlan add track-route 10-4, 10-14
config vlan add track-vlan 10-4, 10-14
config vlan delete domain-member 10-15
config vlan delete port 6-14
config vlan delete secondary-ip 11-14
config vlan delete subvlan 11-14
config vlan delete track-ping 10-14
config vlan delete track-route 10-4, 10-14
config vlan delete track-vlan 10-4, 10-14
config vlan esrp election-algorithm 10-15
config vlan esrp priority 10-15
config vlan esrp timer 10-15
config vlan esrp-group 10-15
config vlan ipaddress 2-7, 6-14, 11-10
config vlan name 6-13, 6-15
config vlan protocol 6-14
config vlan qosprofile 6-14, 9-8, 9-20
config vlan slb-type 17-15
config vlan tag 6-15
config vlan udp-profile 11-18
config vlan xnetid 15-3, 15-6, 15-9
create access-list icmp 16-10
create access-list ip destination 16-7
create access-list tcp destination 16-8

create access-list udp destination 16-9
create access-profile 16-29
create account 2-7, 2-12
create bgp neighbor 13-14
create fdbentry 7-4, 9-10
create flow-redirection 17-47
create isq-server 9-25, 9-29
create ospf area 12-7, 12-24
create protocol 6-15
create qosprofile 9-8
create route-map 16-30, 16-36
create slb pool 17-3, 17-16
create slb vip 17-4
create slb vip pool mode 17-6, 17-8, 17-11
create stpd 8-6, 8-8
create udp-profile 11-19
create vlan 2-7, 6-15

D
delete access-list 16-4, 16-10
delete access-profile 16-29
delete account 2-7
delete bgp neighbor 13-15
delete fdbentry 7-4
delete flow-redirection 17-47
delete isq server 9-25
delete isq-server 9-29
delete ospf area 12-28
delete protocol 6-15
delete qosprofile 9-8
delete route-map 16-36
delete slb pool 17-16
delete stpd 8-9
delete udp-profile 11-19
delete vlan 2-8, 6-15
disable access-list counter 16-10
disable bgp 13-15
disable bgp aggregation 13-14, 13-15
disable bgp always-compare-med 13-14, 13-15
disable bgp export 13-14
disable bgp neighbor 13-15
disable bgp synchronization 13-15
disable bootp 2-8, 11-20, 11-28
disable bootprelay 11-20, 11-28
disable cli-config-logging 2-8, 18-10, 18-11
disable clipaging 2-8
disable diffserv examination ports 9-16
disable diffserv replacement ports 9-16
disable dlcs 9-29
disable dlcs ports 9-29
disable dot1p replacement ports 9-14
disable dvmrp 14-14
disable dvmrp rxmode 14-14
disable dvmrp txmode 14-14
disable edp ports 4-13, 5-12
disable esrp vlan 10-16
disable gvrp 6-21
disable icmp 11-28
disable icmp address-mask 11-28
disable icmp parameter-problem 11-23
disable icmp port-unreachables 11-28
disable icmp redirects 11-29
disable icmp time-exceeded 11-29
disable icmp timestamp 11-29
disable icmp unreachables 11-29
disable icmp useredirects 11-29
disable idletimeout 2-8
disable igmp 14-14
disable igmp snooping 14-14
disable ignore-stp vlan 8-9
disable ipforwarding 11-20, 11-29
disable ipforwarding broadcast 11-20, 11-29
disable ipmcforwarding 14-14
disable ip-option loose-source-route 11-23
disable ip-option record-route 11-23
disable ip-option record-timestamp 11-23
disable ip-option strict-source-route 11-23
disable ip-option use-router-alert 11-23
disable iproute sharing 11-22
disable ipxrip 15-14
disable ipxsap 15-14
disable ipxsap gns-reply 15-5, 15-14
disable irdp 11-29
disable isq 9-25
disable jumbo-frame 4-5, 5-3
disable learning port 7-4
disable learning ports 4-5, 5-4
disable log display 18-11
disable loopback-mode vlan 11-20
disable mac-vlan 6-23
disable mirroring 4-12, 5-11
disable multinetting 11-20
disable ospf 12-28
disable ospf export 11-4, 12-24
disable ospf export direct 12-29
disable ospf export rip 12-12, 12-29
disable ospf export static 12-12, 12-29
disable ospf export vip 12-12, 12-29
disable pim 14-14
disable port 2-8
disable ports 4-3, 4-5, 5-1, 5-4
disable qosmonitor 9-22
disable radius 3-16

disable radius-accounting 3-16
disable red ports 9-8
disable rip 12-20
disable rip aggregation 12-20
disable rip export 11-4, 12-13, 12-20
disable rip originate-default 12-20
disable rip poisonreverse 12-20
disable rip splithorizon 12-20
disable rip triggerupdates 12-20
disable rmon 18-14
disable sharing 4-5, 4-10, 5-4, 5-10
disable slb 17-16
disable slb 3dns 17-44
disable slb failover 17-44
disable slb failover manual-fallback 17-44
disable slb failover ping-check 17-44
disable slb global synguard 17-44
disable slb gogo-mode 17-16
disable slb l4-port 17-16
disable slb node 17-16
disable slb node ping-check 17-23, 17-44
disable slb node tcp-port-check 17-23, 17-44
disable slb vip 17-17
disable slb vip client-persistence 17-44
disable slb vip service-check 17-25, 17-44
disable slb vip sticky-persistence 17-44
disable slb vip svcdown-reset 17-44
disable smartredundancy 5-4
disable snmp access 3-12
disable snmp traps 3-12
disable sntp-client 3-25
disable ssh2 2-8
disable stpd 8-9
disable stpd port 8-9
disable subvlan-proxy-arp 11-14
disable syslog 18-11
disable tacacs 3-20
disable tacacs-accounting 3-21
disable telnet 2-8, 3-7
disable type20 forwarding 15-14
disable web 2-8, 3-9, 19-2
download bootrom 2-13, 20-8
download configuration 2-13, 6-24, 20-5, 20-9
download configuration cancel 20-6, 20-9
download configuration every 20-6, 20-9
download configuration incremental 20-6
download image 2-13, 20-2, 20-9

E

enable access-list counter 16-10
enable bgp 13-14

enable bgp aggregation 13-8, 13-14
enable bgp always-compare-med 13-14
enable bgp compare-as-path 13-14
enable bgp compare-med-within-as-only 13-14
enable bgp neighbor 13-14
enable bgp synchronization 13-15
enable bootp 2-8, 11-20
enable bootp vlan 3-4
enable bootprelay 11-16, 11-21
enable cli-config-logging 2-8, 18-10, 18-11
enable clipaging 2-8
enable diffserv examination ports 9-17
enable diffserv replacement ports 9-17, 9-18
enable dlcs 9-29
enable dlcs ports 9-29
enable dot1p replacement ports 9-14, 9-18
enable dvmrp 14-5, 14-7
enable dvmrp rxmode 14-7
enable dvmrp txmode 14-7
enable edp ports 4-13, 5-12
enable esrp vlan 10-16
enable gvrp 6-21
enable icmp address-mask 11-23
enable icmp parameter-problem 11-23
enable icmp port-unreachables 11-24
enable icmp redirects 11-24
enable icmp time-exceeded 11-24
enable icmp timestamp 11-24
enable icmp unreachables 11-24
enable icmp useredirects 11-24
enable idletimeout 2-8
enable igmp 14-8
enable igmp snooping 14-8
enable ignore-stp vlan 8-8
enable ipforwarding 11-11, 11-21
enable ipforwarding broadcast 11-21
enable ipmcforwarding 14-4, 14-7
enable ip-option loose-source-route 11-24
enable ip-option record-route 11-24
enable ip-option record-timestamp 11-24
enable ip-option strict-source-route 11-24
enable ip-option use-router-alert 11-25
enable iproute sharing 11-22
enable ipxrip 15-10
enable ipxsap 15-11
enable ipxsap gns-reply 15-11
enable irdp 11-25
enable isq 9-25
enable jumbo-frame 4-5, 5-4
enable jumbo-frame ports 4-7, 5-6
enable learning port 7-4
enable learning ports 4-5, 5-4

enable license 2-8
enable log display 18-9, 18-12
enable loopback-mode vlan 11-21
enable mac-vlan 6-23
enable mirroring 4-12, 5-11
enable multinetting 11-21
enable ospf 11-11, 12-24
enable ospf export 11-4, 12-24
enable ospf export direct 12-24
enable ospf export rip 12-12, 12-24
enable ospf export static 12-12, 12-24
enable ospf export vip 12-12, 12-24
enable pim 14-5, 14-7
enable ports 4-3, 4-5, 5-1, 5-4
enable qosmonitor 9-22
enable radius 3-16
enable radius-accounting 3-16
enable red port 9-8
enable rip 11-11, 12-15
enable rip aggregation 12-15
enable rip export 11-4, 12-13, 12-16
enable rip originate-default 12-16
enable rip poisonreverse 12-16
enable rip splithorizon 12-16
enable rip triggerupdates 12-16
enable rmon 18-14
enable route sharing 11-5
enable sharing 4-6, 4-10, 5-4, 5-10
enable slb 17-17
enable slb 3dns iquery-client 17-44
enable slb failover 17-31, 17-44
enable slb failover manual-fallback 17-35, 17-44
enable slb failover ping-check 17-32, 17-44
enable slb global synguard 17-44
enable slb gogo-mode 17-17
enable slb l4-port 17-17
enable slb node 17-17
enable slb node ping-check 17-23, 17-45
enable slb node tcp-port-check 17-23, 17-45
enable slb vip 17-16, 17-17
enable slb vip client-persistence 17-26, 17-45
enable slb vip service-check 17-25, 17-45
enable slb vip sticky-persistence 17-27, 17-45
enable slb vip svcdown-reset 17-46
enable smartredundancy 5-4
enable snmp access 3-12
enable snmp traps 3-12
enable sntp-client 3-22, 3-25
enable ssh2 2-9, 3-7
enable stpd 8-6, 8-8
enable stpd port 8-8
enable subvlan-proxy-arp 11-14

enable syslog 18-9, 18-12
enable tacacs 3-21
enable tacacs-accounting 3-21
enable tacacs-authorization 3-21
enable telnet 2-9, 3-6, 3-7
enable type20 forwarding 15-9
enable web 2-9, 3-9, 19-2
enable web access-profile 3-9

H

history 2-6, 2-9

L

logout 3-6

N

nslookup 2-13

P

ping 2-13, 2-14

Q

quit 3-6

R

reboot 20-2, 20-9
restart ports 4-6, 5-5
rtlookup 11-22
run diag 18-3

S

save 20-3, 20-9
show access-list 16-6, 16-10
show access-list-fdb 16-10
show access-list-monitor 16-6, 16-10
show access-profile 16-29
show accounts 2-12
show banner 2-9
show bgp 13-15
show bgp bgp-policy 13-15
show bgp neighbor 13-15
show configuration 20-9
show debug-tracing B-8
show diag 18-2
show dlc 9-29

show dns-client 2-13
show dot1p 9-14
show dvmrp 14-13
show edp 4-13, 5-12
show esrp 10-5, 10-13, 10-16, 10-20
show esrp vlan 10-16
show fdb 7-5, 9-12, 9-23
show flow-redirection 17-47
show gvrp 6-21
show igmp snooping 14-13
show iparp 11-11, 11-15, 11-27
show iparp proxy 11-27
show ipconfig 11-11, 11-16, 11-27
show ipfdb 11-11, 11-28
show ipmc cache 14-13
show iproute 11-11, 11-28
show ipstats 11-28
show ipxconfig 15-6, 15-13
show ipxrip 15-6, 15-13
show ipxroute 15-6, 15-13
show ipxsap 15-6, 15-13
show ipxservice 15-6, 15-13
show ipxstats 15-13
show log 18-2, 18-9, 18-12
show log config 18-2, 18-12
show mac-vlan 6-23
show management 3-7, 3-9, 3-13, 19-2
show memory 18-2
show mirroring 4-12, 5-11
show ospf 12-12, 12-28
show ospf area 12-28
show ospf ase-summary 12-28
show ospf interfaces 12-28
show ospf lsdb 12-28
show ospf virtual-link 12-28
show pim 14-13
show pim rp-set 14-13
show ports collisions 4-6, 5-5
show ports configuration 4-6, 4-11, 5-5, 5-10
show ports info 4-6, 5-5, 9-19, 9-21, 9-23, 10-3
show ports packet 4-6, 5-5
show ports qosmonitor 4-6, 5-5, 9-22
show ports rxerrors 4-6, 5-5, 18-5
show ports stats 4-6, 5-5, 18-4
show ports txerrors 4-6, 5-5, 18-5
show ports utilization 4-6, 5-5
show protocol 6-17
show qosprofile 9-12, 9-21, 9-23
show radius 3-16
show rip 12-19
show rip stat 12-19
show rip vlan 12-19

show session 3-6
show slb 17-17
show slb 3dns members 17-46
show slb failover 17-46
show slb l4-port 17-18
show slb node 17-17
show slb pool 17-18
show slb vip 17-18
show slot 4-2, 4-6
show sntp client 3-22
show sntp-client 3-25
show stpd 8-8
show stpd port 8-9
show switch 3-23, 6-24, 9-23, 18-2, 20-6
show tacacs 3-21
show tacacs-accounting 3-21
show udp-profile 11-19
show version 18-3
show vlan 6-16, 9-21, 9-23, 10-13, 10-20, 11-15, 15-6
synchronize 20-7, 20-9

T

telnet 2-13, 3-3
traceroute 2-13, 2-15

U

unconfig diffserv examination ports 9-17
unconfig diffserv replacement ports 9-17
unconfig dvmrp 14-14
unconfig icmp 11-25, 11-29
unconfig igmp 14-14
unconfig ipxrip 15-14
unconfig ipxsap 15-14
unconfig irdp 11-25, 11-29
unconfig management 3-12
unconfig ospf 12-29
unconfig pim 14-14
unconfig ports display-string 4-6, 5-5
unconfig rip 12-20
unconfig slb all 17-18
unconfig slb vip service-check 17-46
unconfig slot 4-6
unconfig stpd 8-9
unconfig switch 2-9, 20-3
unconfig switch all 20-3
unconfig tacacs 3-21
unconfig tacacs-accounting 3-21
unconfig udp-profile 11-19
unconfig vlan ipaddress 6-15
unconfig vlan xnetid 15-14

upload configuration 2-13, 20-4, 20-9
upload configuration cancel 20-4, 20-9
use configuration 20-3, 20-10
use image 20-2, 20-10

X

xping 15-9